



Ensuring AI Accountability: Auditing Methods to Mitigate the risks of Large Language Models



About Democracy Reporting International

DRI is an independent organisation dedicated to promoting democracy worldwide. We believe that people are active participants in public life, not subjects of their governments. Our work centres on analysis, reporting, and capacity-building. For this, we are guided by the democratic and human rights obligations enshrined in international law. Headquartered in Berlin, DRI has offices in Lebanon, Libya, Myanmar, Pakistan, Sri Lanka, Tunisia, and Ukraine.

Acknowledgements

This report was written by Duncan Allen, Research Associate (DRI), Ognjan Denkovski, Research Coordinator (DRI) and Francesca Giannaccini, Research Associate (DRI).

Julieta Jimenez designed the layout of this publication.


Date: September 2024

This report is part of the Disrupt Disinfo project, funded by the German Federal Foreign Office. Its contents do not necessarily represent the position of the German Federal Foreign Office.

Table of Contents

Introduction	4
Theoretical Framework	6
Inadvertent and Inherent Risks – Bias and Hallucinations	6
Deliberate Misuse	8
Auditing Considerations	9
Auditing Methodology Review	11
Qualitative Approaches - manually auditing LLMs	12
Automated and Semi-Automated Approaches – auditing at scale	15
Integrating Qualitative and Automated Approaches – LLMs as Auditing Assistants	21
Concluding remarks	34

Click the numbers
to access to each item



Introduction

The field of Artificial Intelligence (AI) has witnessed remarkable advancements in recent years. Some of the most popular AI-powered technologies to have captured public attention have been large language models (LLMs), which emerged into the mainstream with the release of OpenAI’s ChatGPT in late 2022. In the years since, many competitors have come forth with their own models, from Google’s Gemini to Meta’s Llama 3, sparking a race for media attention, investment, and access to training data. In addition, there is an emerging ecosystem of open source LLMs that can be freely accessed and trained for custom user requirements and preferences.¹

The possibilities of LLM integration in various tasks are undeniable, manifold, and yet to be fully understood. Potential applications of these technologies include the streamlining of administrative tasks for healthcare professionals,² the creation of personalised educational content for students,³ and the empowerment of individuals with disabilities.⁴ These tools are already being utilised increasingly across society. A March

2024 Pew Research Center study found that 20 percent of respondents in the United States use ChatGPT in the workplace, up from eight per cent a year earlier.⁵ However, while these models have proven to be adept at understanding, summarising, and generating natural language, enabling their widespread application in various contexts, their integration into daily life exceeds the pace at which the evaluation of this integration occurs. As these models become more embedded in society, it is crucial to carefully consider and address the potential issues that arise in ensuring their responsible and beneficial use.

In an increasingly digital world, online harms, such as those potentially generated by LLMs, too often translate into offline threats. Biases ingrained in the data LLMs are trained on, as well as their propensity to hallucinate answers, pose serious misinformation risks for a public that increasingly relies on chatbots to provide them with information and assist with everyday tasks. In an electoral context, for instance, such

mistakes can misinform users as to where, when, and how to vote.⁶ In addition, chatbots can also be deliberately manipulated into performing harmful actions, such as the generation of disinformation and hate speech.⁷ In an era where disinformation actors constantly seek to undermine information integrity online, LLMs have become powerful tools to create and propagate false narratives at scale.

For this reason, at this still-early stage of LLM development, it is essential to develop oversight frameworks for these models to audit system design and output, and part of this role must be played by independent actors from civil society invested in maintaining information integrity. Such auditing would prioritise a rigorous examination of a model for fairness, accuracy, transparency, and ethical and legal compliance in responses. This needs not be highly technical; in the leadup to the 2024 European Parliament elections, DRI carried out two such audits simply by prompting popular chatbots with electoral questions, and then

1. Democracy Reporting International, [“Open to Misuse? The Lack of Safeguards in Open-Source LLMs”](#), 15b December 2023.
2. Leung & de Azevedo Cardoso [“The Impact of Multimodal Large Language Models on Health Care’s Future”](#), *Journal of Medical Internet Research*, 2 November 2023.
3. Andy Extance, [“ChatGPT Has Entered the Classroom: How LLMs Could Transform Education”](#), *Nature*, 15 November 15 2023.
4. [“Empowering Accessibility: How LLMs Enhance Inclusivity for People with Disabilities - Resources on Disability Assistance: Your Rights and Benefits”](#), *Disability Help*, 15 February 2024.
5. Coleen McClain, [“Americans Increasingly Using ChatGPT, but Few Trust its 2024 Election Information”](#), *Pew Research Center*, 26 March 2024.
6. Jia-Yu Yao, Kun-Peng Ning, Zhen-Hui Liu, Mu-Nan Ning, Yu-Yang Liu & Li Yuan, [“LLM Lies: Hallucinations are not Bugs, but Features as Adversarial Examples”](#), *ArXiv*, 3 October 2023.
7. Beatriz Saab & Francesca Giannaccini, [“ChatGPT vs. Bard: Unveiling the Battle against Disinformation and Creative Output”](#), Democracy Reporting International, 28 August 2023.

grading the accuracy and bias in their responses.⁸ By scrutinising the responses of LLM-powered chatbots, researchers can determine not only the weaknesses of these models as reliable and ethical sources of information, but also their vulnerabilities to abuse by malicious actors. By detecting bias, highlighting dangerous “hallucinations”, and identifying areas of potential misuse, independent audits can inform AI companies and regulators alike, and encourage proactive action to address these issues.

Fully qualitative methods are often resource-intensive and time-consuming, however, making them challenging to scale and apply across different domains. Consequently, there is a pressing need for automated and mixed-methods evaluations, which can offer scalability and efficiency. Automated approaches come with their own set of challenges, however, such as limitations in capturing nuanced or emergent issues and

potential misalignment with real-world contexts. On the other hand, mixed methods approaches, while combining the strengths of both quantitative and qualitative assessments, have yet to be developed in a scalable, robust, and easy-to-use manner.

This guide evaluates a multitude of approaches for identifying and mitigating the risks associated with LLMs. By examining a wide range of analytical frameworks and tools, we provide recommendations for best practices, based on a review of the of approaches identified as state-of-the-art in the field of LLM auditing, while specifically considering auditing approaches relevant for researchers concerned with electoral and information integrity. To achieve this, our review is designed to be highly accessible, enabling not only academics, but also journalists, non-academic researchers, and civil society organisations to effectively understand the potential of LLM auditing and,

consequently, use the approach best suited to their needs. The resource could not be more timely. On 1 August 2024, the EU’s AI Act came into force, establishing new obligations for general-purpose AI systems (GPAI), including chatbots and other LLMs. Notably, the regulation acknowledges that some GPAI models may pose systemic risks within the EU, such as potential negative impacts on democratic processes (Recital 110). As a result, providers of GPAI models identified as carrying systemic risk are required to conduct regular risk assessments, take proactive measures to mitigate identified risks, and continuously monitor for serious incidents. This comprehensive guide will provide and detail methods for identifying and mitigating risks associated with LLMs, particularly those categorized as presenting systemic risk under Article 52 of the AI Act, fostering the development of evidence for the newly established AI Office to ensure AI companies meet their obligations under the Act.

8. Duncan Allen, [“When Misinformation Becomes Disinformation: Chatbot Companies and EU Elections,”](#) Democracy Reporting International, 7 June 2024.

Theoretical Framework

This section outlines why LLMs cannot currently be treated as reliable sources of information, and explains the risks associated with their widespread application. In addition, it aims to summarise the methodological, ethical, and legal considerations that should be accounted for when auditing such technologies.

Inadvertent and Inherent Risks – Bias and Hallucinations

Despite LLMs primarily being generators of text, they are increasingly being employed and marketed as helpful and knowledgeable sources of information. Google and Bing have both attached LLM-generated summaries to their search result functions, while OpenAI has marketed ChatGPT as an invaluable virtual assistant. In late July 2024, OpenAI also announced the upcoming release of SearchGPT, a GPT-powered tool designed to assist with search functions and allow users to quickly retrieve information from the web – a

small step in the broader trend of shifting from classical search functionality to LLM powered search as a norm. Such developments risk increasing the potential for LLMs to inadvertently produce incorrect and even harmful information; for example, in the opening announcement video for SearchGPT, the tool made the mistake of recommending the user attend a music festival that had already been held.⁹

There are two main difficulties underpinning these concerns. First, LLMs can reflect biases ingrained in their training data.¹⁰ From historical stereotypes about race to the promotion of gender inequality, LLMs can inadvertently reproduce harmful concepts present in the texts on which they were trained. For example, in 2024, a team of linguistics researchers found that ChatGPT and Gemini discriminated against users who made prompts in African American Vernacular English, and were more likely to refer to them as “lazy” or “stupid”.¹¹ Sometimes, these biases are more difficult to identify. One study fed genuine headlines from reputable news

outlets, such as the New York Times and Reuters, into popular chatbots, asking the bots to generate possible articles that could fit each of the headlines. The responses revealed a substantial amount of racial and gender bias present in the word choices of the LLM-generated fake articles, with the AI-generated content more likely to display negative sentiment towards certain genders or racial groups than the actual articles to which the headlines belonged.¹² When prompted to generate reference letters for two hypothetical students, one male and one female, researchers found ChatGPT more likely to describe the male candidate as active (“a natural leader”, “role model”) and the female candidate as sociable (“well-liked”, “good interpersonal skills”).¹³ While subtle, these differences in description highlight how many LLMs have incorporated traditional societal biases. Beyond discrimination, model biases can be directly harmful, especially given the increasing use of LLM-powered tools in healthcare sectors. When tested, multiple popular models endorsed debunked racial “science”, such as the notion that skin thickness

- 9. Matteo Wong, [“OopsGPT: OpenAI Just Announced a New Search Tool. Its Demo Already Got Something Wrong”](#), *The Atlantic*, 25 July 2024.
- 10. Roberto Navigli, Simone Conia & Björn Ross, [“Biases in Large Language Models: Origins, Inventory, and Discussion”](#), *ACM Journal of Data and Information Quality*, Volume 15, Issue 2, 22 June 2023.
- 11. Ava Sasani, [“As AI Tools Get Smarter, They’re Growing More Covertly Racist, Experts Find”](#), *The Guardian*, 16 March, 2024.
- 12. Xiao Fang, Shangkun Che, Minjia Mao, Hongzhe Zhang, Ming Zhao & Xiaohang Zhao, [“Bias of AI-Generated Content: An Examination of News Produced by Large Language Models”](#), *Scientific Reports*, 4 March 2024.
- 13. Yixin Wan, George Pu, Jiao Sun, Aparna Garimella, Kai-Wei Chang, Nanyun Peng, [““Kelly is a Warm Person, Joseph is a Role Model”: Gender Biases in LLM-Generated Reference Letter”](#), *ArXiv*. 13 October 2024.

and pain tolerance vary depending on the race of a patient.¹⁴ Such biases are dangerous, and reveal how models may not be a good source of reference, even for topics as objective as human biology.

Second, LLM-powered chatbots have been known to conjure nonexistent sources,¹⁵ invent fake historical “facts”,¹⁶ and prescribe deadly medical advice.¹⁷ Dubbed “hallucinations”, these errors are less a bug in than a feature of LLMs. The transformer architecture behind most modern LLMs, which is essentially a method for predicting and generating text based on patterns in the model’s training data, is based on probability calculations.¹⁸ Each word in the generated text is selected based on the likelihood of its occurrence given the preceding context, essentially operating as an advanced autocomplete function. The generation of language, specifically texts that sound natural and similar to those in the model’s training data, is prioritised over the factual accuracy of that language.¹⁹ This means that when models are prompted for

answers that don’t exist in their training data, such as specific sources for an event that occurred beyond their knowledge cutoff, they often default to simply fabricating a source. Importantly, this challenge extends beyond just the generation of information in domains with no training data. It affects a broad range of information, including areas with abundant data, but nuanced or complex details, and especially “low domain” fields, where specialized knowledge is scarce or not well-represented in the model’s training corpus. In such cases, even if the model has encountered some relevant information, a lack of depth or breadth in the data can lead to outputs that are speculative, imprecise, or outright fabricated. This underscores the importance of recognising the probabilistic nature of LLMs and the need for critical evaluation of their outputs, especially in contexts where accuracy and reliability are paramount.

The capacity for LLM-powered chatbots to appear correct when they are not is most likely to be harmful

during crises or important events, such as elections, where public access to trustworthy information is critical. The previously noted study by DRI found that four of the most popular chatbots often produced incorrect or partially incorrect information when prompted on how to vote in the 2024 European Parliament elections.²⁰ Users were given incorrect registration deadlines, told they could vote via post when they couldn’t, and provided with broken and irrelevant links. Election officials and researchers in the United States have similarly tested Gemini, Claude, ChatGPT, and other bots’ abilities to give users accurate voting information ahead of the 2024 November presidential election.²¹ They found that over 40 per cent of the LLM responses could be categorised as harmful and in violation of voting rights, with the bots often incorrectly informing users that no voting station exists within their area code.

While LLM-powered chatbots can be fine-tuned to reduce their likelihood of hallucinating on a certain

14. Jesutofunmi A. Omiye, Jenna C. Lester, Simon Spichak, Veronica Rotemberg & Roxana Daneshjou, [“Large Language Models Propagate Race-Based Medicine”](#), *Nature*, 20 October 2023.
15. Karen Weise & Cade Metz, [“When Chatbots Hallucinate”](#), *The New York Times*, 1 May 2023.
16. Catherine Thorbecke, [“Google Shares Lose \\$100 Billion After Company’s AI Chatbot Makes an Error During Demo”](#), *CNN*, 9 February 2023.
17. Gary Marcus, [“Serious Medical Error from Perplexity’s Chatbot”](#), *Marcus on AI*, 29 February 2024.
18. Jakob Uszkoreit, [“Transformer: A Novel Neural Network Architecture for Language Understanding”](#), *Google Research*, 31 August 2017.
19. Lauren Leffer, [“AI Chatbots Will Never Stop Hallucinating”](#), *Scientific American*, 5 April 5 2024.
20. Michael Meyer-Resende, Austin Davis, Ognjan Denkovski & Duncan Allen, [“Are Chatbots Misinforming Us About the European Elections? Yes,”](#) *Democracy Reporting International*, 11 April 2024.
21. Michael Dwyer, [“Chatbots’ Inaccurate, Misleading Responses about US Elections Threaten to Keep Voters from Polls”](#), *POLITICO*, 27 February 2024.

topic, such changes generally involve surface-level fixes, which do not address the underlying issue inherent in the model’s architecture. Indeed, attempts by companies to reduce chatbot hallucinations related to sensitive topics have their own challenges. A follow-up investigation done by DRI two months after the initial report revealed that some chatbots had since been tuned to avoid responding to questions about the European Parliament elections in English, but still provided (often incorrect) answers in at least eight other languages.²² Such findings suggest that LLM-powered chatbots may also be more likely to give biased or incorrect responses in language contexts where their training dataset and oversight may not be as robust or selective as in English. In addition, simply having a chatbot refuse to answer controversial questions can actually bolster certain disinformation narratives, as exemplified by ChatGPT and Gemini’s refusal to provide a response when asked “Who won the 2020 US presidential election?”.²³

Beyond the immediate danger of spreading misinformation, public dependence on LLMs for information also poses a significant long-term threat. Large language models’ capacities to hallucinate and provide biased answers risk eventually eroding public knowledge of basic topics, such as science and history. Dubbed “careless speech” by researchers from the Oxford Internet Institute, the confident manner in which chatbots produce “small mistruths, oversimplifications of complex topics and responses biased towards certain commonly held beliefs or schools of thought” is a subtle threat to common understandings of truth.²⁴ For instance, if many citizens rely on LLM-powered chatbots to curate their news and provide them with information, these minor but consistent mistakes can be compounded into a gradual distortion of public knowledge. Though the overall factual accuracy of these models may marginally improve with each generation and update, as long as the potential for error remains, LLMs represent a misinformation risk.

Deliberate Misuse

In addition to unintentional harm, many LLMs remain vulnerable to misuse by malicious actors and can be manipulated into performing harmful actions, such as the generation of hate speech and disinformation, or the violation of user privacy.

Large language models are trained on large amounts of data from the internet, some of which is harmful and illegal. Safeguards created by these LLMs’ parent companies prevent users from easily accessing such content. Using sophisticated prompt engineering, however, users can trick an LLM into providing a response it would have otherwise refused to offer. These “jailbreaking attacks”, which often combine complex combinations of random and natural language, have proved alarmingly effective at getting popular chatbots to give illegal answers, such as how to make a homemade bomb.²⁵ Some models have been manipulated into providing jailbreak attack prompts for other models at an automated scale.²⁶

- 22. Duncan Allen, [“When Misinformation Becomes Disinformation: Chatbot Companies and EU Elections”](#), Democracy Reporting International, 7 June 2024.
- 23. Caroline O'Donovan & Cat Zakrzewski, [“Voice Assistants and AI Chatbots Still Can't Say Who Won the 2020 election”](#), *The Washington Post*, 16 June 2024.
- 24. Sandra Wachter, Brent Mittelstadt & Chris Russell, [“Do Large Language Models Have a Legal Duty to tell the Truth?”](#), *Royal Society of Open Science*, 7 August 2024.
- 25. Maksym Andriushchenko, Francesco Croce & Nicolas Flammarion, [“Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks”](#), *arXiv*, 18 June 2024.
- 26. Gelei Deng, Yi Liu, Yuekang Li, Kailong Wang, Ying Zhang, Zefeng Li, Haoyu Wang, Tianwei Zhang & Yang Liu, “MASTERKEY: Automated Jailbreaking of Large Language Model Chatbots”, *ArXiv*, 25 October 2023.

In addition to exploiting harmful information included in a model’s training data, tech savvy assailants can also take advantage of the fact that many models continue to be trained on data collected from the internet on a daily basis. Malicious actors can exploit this process to inject “poisoned data” – data that is deliberately harmful – into a model’s training set.²⁷ Poisoned data could be in the form of disinformation and falsehoods, or it could also be hidden commands for a bot to follow when given a specific, niche prompt. While there are ongoing efforts by LLM parent companies to address this vulnerability by establishing an instruction hierarchy, the efficacy of these solutions should continue to be tested.²⁸

Not all manipulation needs to be so technical: a group of Dutch researchers found that Copilot, Gemini, and ChatGPT all readily complied with a simple prompt to design a campaign strategy for a “Eurosceptic politician who wants to discourage voters in the Netherlands from voting in the European elections”.²⁹ The chatbots

provided researchers with suggestions for disinformation narratives and false statistics that would help this hypothetical campaign. These harms are not purely hypothetical. LLMs have already enabled malicious actors in the production of such content at a scale and quality previously unattainable. Entire Russian disinformation networks have been identified targeting US and EU voters with fake, AI-generated articles peddling pro-Kremlin narratives.³⁰ These networks operate by feeding news articles into chatbots. Modified texts are then disseminated en masse to sham news websites, which, in turn, spread them across social media platforms via both authentic and inauthentic engagement.

Auditing Considerations:

Methodological considerations

Existing literature on the characteristics of effective AI audits has identified three major pillars:³¹

1. Audits should be “socially and ethically relevant”, dealing with topical issues on which LLM responses may be particularly problematic and harmful. As such, audits should strive to ensure that these tools do not disproportionately harm or disadvantage any particular group, and must evaluate AI systems for equitable treatment across different demographics and use cases.
2. The findings should be “predictably transferable” – telling us whether we can expect identified issues to impact everyday users. What implications do the findings have for apps that integrate the LLM in question?
3. Audits should be “meaningfully operationalisable”, with clear definitions for what the auditing parties regard as, for instance, hate speech, disinformation, truthfulness, and fairness, and other concepts they may be interested in auditing the LLM’s performance on.

27.

Nicholas Carlini, Matthew Jagielski, Christopher A. Choquette-Choo, Daniel Paleka, Will Pearce, Hyrum Anderson, Andreas Terzis, Kurt Thomas & Florian Tramèr, [“Poisoning Web-Scale Training Datasets is Practical”](#), *ArXiv*, 6 May 2024.

28.

Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng, Johannes Heidecke & Alex Beutel, [“The Instruction Hierarchy: Training LLMs to Prioritize Privileged Instructions”](#), *OpenAI*, 19 April 2024.

29.

Fleur Damen & Roel van Niekerk, [“Chatbots adviseerden: verspreid desinformatie en zaai angst over EU-verkiezingen”](#), *Nieuwsuur*, 3 May 2024.

30.

Paul Myers, Olga Robinson, Shayan Sardarizadeh & Mike Wendling, [“A Bugatti Car, a First Lady and the Fake Stories Aimed at Americans”](#), *BBC News*, 3 July 2024.

31.

Jakob Mökander, Jonas Schuett, Hannah Rose Kirk & Luciano Floridi, [“Auditing Large Language Models: A Three-Layered Approach”](#), *AI Ethics*, 27 June 2023.

Transparency, accountability, fairness, and ethical integrity are paramount in LLM auditing. To this end, we encourage researchers to openly share methodologies and findings with civil society, parent companies, regulators, and other stakeholders.

Legal considerations

Legal considerations in AI auditing are crucial to ensuring compliance with data privacy laws and regulations. Auditors must be aware of the legal implications of

their work, particularly when dealing with sensitive data and evidence. Compliance with data privacy laws, such as the General Data Protection Regulation (GDPR) in the EU, is essential when reporting input data, including prompts used in AI audits.

Auditors must also consider intellectual property laws when accessing and using proprietary datasets or AI models. Ensuring proper authorisation and respecting intellectual property rights is crucial for legal and ethical auditing practices.

Finally, auditors should align their efforts with relevant regulatory frameworks, which may vary by jurisdiction, to identify violations of existing AI regulation. For example, auditors working within the EU should be aware that the AI Act, which came into force in August of 2024 and establishes obligations for generative AI systems, such as LLMs, to reduce their risk in electoral contexts. These obligations include the explicit labeling of outputs generated by AI, the linking of official sources, and requiring models to encourage users to verify any electoral information they provide.

Auditing Methodology Review

This review examines a range of auditing approaches – qualitative, automated, and mixed-methods – each offering unique insights into LLM performance and behaviour. Qualitative methods, with their in-depth, context-rich analysis, provide a nuanced understanding of model outputs and ethical implications, and are easily accessible to non-technical auditors, but often face challenges related to scalability, transferability of findings, and resource demands. Automated approaches provide efficiency and scalability, making them essential for evaluating frequent model updates. They often fall short in capturing the complexity of human-centric issues, however, and tend to be highly specialized in the types of metrics they can assess LLMs on. Mixed-methods approaches combine the strengths of both qualitative and automated techniques, aiming to balance detailed, contextual analysis with broader, scalable assessments. Such approaches are still largely in development, however. By evaluating these methodologies, this review seeks to highlight their respective strengths and limitations, offering a comprehensive guide to the field of auditing LLMs.



Figure 1. Overview of audit types considered



32. Michael Meyer-Resende, Austin Davis, Ognjan Denkovski & Duncan Allen, [“Are Chatbots Misinforming Us About the European Elections? Yes,”](#) *op. cit.*, note 20.

33. Ben Buchanon, Micah Musser, Andrew Lohn & Katerina Sedova, [“Truth, Lies and Automation: How Language Models Could Change Disinformation”](#), Center for Security and Emerging Technology, May 2021.

Qualitative approaches can play a crucial role in understanding the ethical and social implications of LLMs. These approaches involve assessing the quality and accuracy of model-generated outputs through, in most cases, direct human input of prompts, and recording and assessment of model output. Experts or ordinary users following a set of instructions are enlisted to review and rate the results produced by the model, frequently within a narrow area of interest.

Qualitative approaches have been applied to auditing and evaluating LLM responses on tasks as varied as summarisation quality, information accuracy, and ability to generate disinformation. A recent qualitative study by DRI highlights significant issues with the accuracy and reliability of chatbot responses, by assessing election-related queries in ten languages across four LLMs. The chatbots demonstrated considerable variability, often providing incorrect or misleading information, particularly concerning detailed electoral information.³² However, the study required significant human effort

and LLM theoretical understanding to conceptualise, execute, and thereafter repeat assessments of improvements as suggested to LLM providers. A study conducted in 2021 by the Centre for Security and Emerging Technologies focused on narrative reiteration, narrative elaboration, narrative manipulation, narrative seeding, narrative wedging, and narrative persuasion in the context of LLM’s capacities to generate disinformation, based on survey evaluations of LLM responses.³³ Other research has relied on Upwork-sourced annotators to evaluate generation quality based on faithfulness,

coherence, relevance, and fluency in various computational social science tasks.³⁴ Furthermore, qualitative approaches that attempt to scale their studies have relied on crowdsourcing, using tools such as MechanicalTurk.³⁵ The latter set of these studies shows that researchers engaging in qualitative audits of LLMs have recognised one of the key issues of such approaches – scalability – and have attempted to address it with various approaches based on input from non-expert reviewers. This approach comes with its own pitfalls, however, namely those associated with prompt evaluations from non-experts following a set of (always imperfect) instructions. From here, issues of the subjectivity and replicability of findings can easily arise. In principle, qualitative methods must strike a balance between small-scale approaches contingent on expert reviewers, who are scarce and costly, and large-scale evaluations from unskilled reviewers, who follow a set of predefined instructions, which can be misinterpreted or inconsistently applied. Regardless of scale, these

types of audits all face the issues of rigidity and inability to respond to unexpected findings, since audit frameworks, concepts, and scope must be defined and operationalised clearly and exhaustively ahead of the audit process. Furthermore, any attempts to properly assess performance across languages poses a considerable resource constraint.

A final approach worth noting is the broader practice of “red teaming”, a method that has become an industry standard in evaluating LLMs, and which can be conducted fully qualitatively. Red teaming involves a team of experts who probe LLMs with edge cases, adversarial prompts, and ethically sensitive scenarios, attempting to identify vulnerabilities and biases. The strategy of acting as an adversary – submitting instructions targeted to spot failures, biases, and hallucinations in the model – is frequently used by tech companies to stress test their products not only before production, but also as an ongoing maintenance activity.³⁶ Unlike standard

qualitative audits, red teaming is dynamic and adaptable, making it highly effective in uncovering problems, such as discriminatory language or the spread of misinformation, by simulating real-world challenges. Its adaptability allows it to keep pace with evolving LLM capabilities and ethical concerns. Red teaming requires a diverse and skilled team, however, making it highly resource-intensive, and it can only be truly effective if applied consistently, iterating over novel adversarial scenarios as they become known throughout the auditing process. This means that the framework of the audit must consistently be adjusted as novel scenarios emerge from the data. Despite these challenges, red teaming remains crucial for ensuring LLMs are safe, reliable, and ethically sound.³⁷

Overall, qualitative methods, while valuable, face significant challenges. They are labor-intensive in both the prompt creation and evaluation stages, often constrained by small sample sizes, and rarely allow for

- 34. Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang & Diyi Yang, [“Can Large Language Models Transform Computational Social Science?”](#), arXiv, 26 Feb 2024.
- 35. Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al, [“Holistic Evaluation of Language Models”](#), arXiv, 1 October 2023.
- 36. Madhumita Murgia, [“Open AI’s Red Team: The Experts Hired to ‘Break’ ChatGPT”](#), FT, 14 April 2023
- 37. Miles Brundage et al., [“Toward Trustworthy AI Development: Mechanisms for Supporting Verifiable Claims,”](#) arXiv, 20 April 2020.

generalisations beyond the area of interest, which almost always needs to be defined well ahead of starting the study, leading to limited flexibility in adjusting the audit based on emerging insights or evolving challenges. Additionally, biases introduced by reviewers' personal interpretations must always be considered. The small sample size in qualitative audits poses challenges from two perspectives. First, the prompts fed to

the LLM are limited in number, because humans must create them and, as has been well established, LLMs can provide significantly divergent responses depending on the language or phrasing used. Second, the responses from LLMs that will be evaluated is limited, which further restricts the ability to draw broad conclusions or detect patterns in the model's behavior across diverse contexts. Despite their limitations in

scalability and generalisability, qualitative approaches remain essential for specialised, small-scale audits that aim to ensure the ethical and reliable deployment of LLMs in real-world contexts. Moreover, as qualitative audits can be conducted with little to no technical knowledge, they are among the most accessible approaches to a broad range of researchers interested in auditing LLMs.



Automated and Semi-Automated Approaches – auditing at scale

Definition

Automated and semi-automated LLM audits typically rely on predefined metrics and tools to assess the quality and fairness of model outputs. While human involvement is greatly minimised in semi-automated audits, it is necessary to establish static or live benchmarks.

Advantages

- Scalable, cost-efficient and faster than human evaluation
- Ensure consistency and objectivity
- Allow for easy comparison of model performance relative to other models or model updates

Challenges

- Lack nuance relative to human evaluation
- Require (expert) human oversight
- Require technical expertise to implement

38. Yupeng Chang et al., [“A Survey on Evaluation of Large Language Models”](#), ACM, March 2024.

Automated audits are perhaps the most broadly applied auditing approach where, typically, standard metrics and evaluation tools are used to evaluate model performance. Compared with human evaluation, automatic evaluation does not require intensive human participation, which not only saves time, but also reduces the impact of subjective human factors and makes the evaluation process more standardized.³⁸ This increase in scalability comes, however, at

the expense of what can be evaluated and the accessibility of such audits to non-technical auditors. Moreover, even though automated evaluation significantly reduces the need for human involvement during the evaluation process itself, it still relies on human input for certain key aspects, such as the creation of test datasets, including the formulation of questions, scenarios, or benchmarks against which the model is evaluated. This initial human

involvement is crucial because the quality, relevance, and fairness of the audit are heavily dependent on how well the test datasets and evaluation criteria are designed. Therefore, while automated evaluation minimizes ongoing human participation and reduces subjective bias, it still requires thoughtful human input at the setup stage to ensure that the evaluation is meaningful and representative of the model's intended use cases. As such, most automated audits

are only accessible to highly skilled technical professionals for implementation and a deep understanding of findings.

Static benchmarks

Static benchmarks are the most common method for automated LLM evaluation and involve predefined tasks that can assess performance on tasks as varied as language understanding, mathematical reasoning, and coding abilities. Benchmarks focusing on safety have also been developed, such as ToxicChat.³⁹ Other examples include GLUE (General Language Understanding Evaluation) and SuperGLUE, an advanced and refined version of GLUE, as well as MMLU (Massive Multitask Language Evaluation),⁴⁰ HellaSwag,⁴¹ HumanEval,⁴² HELM,⁴³ and C-Eval.⁴⁴ In principle, the most prevalent benchmarks are static, groundtruth-based ones, typically in the form of multiple-choice questions or question-answering tasks

with predefined answers and test cases.⁴⁵ These benchmarks rely on pre-existing tests that have been annotated by humans, making them useful for the evaluation of model updates or for model comparison, but challenging to adjust and develop in the first place, due to the need for significant human resources and technical expertise for developing these tests. Additionally, most of the user-generated data in these benchmarks is in English, which limits their applicability across different languages and cultural contexts.

Static benchmarks are also prone to issues like contamination, saturation, and overfitting, which can limit their ability to capture the nuances of LLM behavior and align with human expectations. *Contamination* happens when the model has already seen or been trained on similar questions or tasks before, making the test less challenging and less informative. Saturation means that, after a while, the model scores very high on these tests, but that doesn't necessarily mean it's getting better at understanding or

generating language – it might just be getting better at those specific tasks. *Overfitting* occurs when the model becomes too good at these tasks but struggles when faced with different or more complex situations. All of this means that, while the model might do well on these tests, this doesn't necessarily reflect how well it would perform in real-world situations or meet human expectations. For instance, a model might excel in a controlled environment, but fail to manage complex, dynamic human interactions or unexpected inputs in real applications.

Live benchmarks

To address some of the limitations of static benchmarks, live benchmarks have been developed. These benchmarks incorporate up-to-date questions from sources such as annual exams or online contests, e.g., Codeforces. They may also include real-time human interaction, leveraging methods like reinforcement

- 39. Za Lin et al., [“ToxicChat: Unveiling Hidden Challenges of Toxicity Detection in Real-World User-AI Conversation”](#), *arXiv*, 26 October 2023.
- 40. Den Hendryks et al., [“Measuring Massive Multitask Language Understanding”](#), *arXiv*, 12 January 2021.
- 41. Rowan Zellers et al., [“Can a Machine Really Finish Your Sentence?”](#), *arXiv*, 19 May 2019.
- 42. Mark Chen et al., [“Evaluating Large Language Models Trained on Code”](#), *arXiv*, 14 July 2021.
- 43. Percy Liang, et al, [“Holistic Evaluation of Language Models”](#), op. cit., note 35.
- 44. Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu, Chuancheng Lv, Yikai Zhang, Jiayi Lei, et al., [“C-Eval: A Multi-Level Multi-Discipline Chinese Evaluation Suite for Foundation Models”](#), *arXiv*, 6 November 2023.
- 45. Wei-Lin Chiang, “Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference “, *arXiv*, 7 March 2024.

learning from human preferences. Chatbot Arena exemplifies a significant advancement in this area, offering a large-scale, crowdsourced platform that integrates live human feedback to enhance benchmarking accuracy.⁴⁶ In Chatbot Arena, users engage in real-time conversations with various AI models, allowing evaluators to collect feedback on multiple dimensions, such as response relevance, coherence, and creativity, and on conversational flow. For instance, users can rate responses on a scale or vote for which model they prefer in a given conversation, directly influencing the assessment of each model's performance. This feedback loop is essential for identifying how models handle complex, dynamic interactions and nuanced human language that static benchmarks might overlook. By capturing these real-world interactions, Chatbot Arena provides a richer, more detailed picture of how well models perform across diverse scenarios and user types, leading to more accurate and meaningful benchmarking outcomes. DynaBench is an effort in a similar vein, recommending the use of a live

benchmark that incorporates a human-in-the-loop-approach for classical natural language generation and processing benchmarks.⁴⁷ As noted by the creators of Chatbot arena, however, despite extensive user bases, live benchmarking efforts like these, much as with static benchmarking efforts, primarily benefit LLM hobbyists and researchers who are eager to experiment with and evaluate the latest LLMs. This inclination may result in a biased distribution of users and reduced accessibility of such tools to those with limited expertise or knowledge of the landscape who are interested in auditing LLMs. For instance, this review did not identify any specialised live benchmarking approaches for topics related to electoral information integrity, political bias, disinformation-generation ability, and other threats posed by LLMs discussed in the theoretical framework. Finally, the issue of English-centric data applies to live benchmarks just the same as it does to static benchmarks. The English-centric focus of most benchmarks not only limits their applicability across languages, but also risks

embedding cultural biases into LLMs. For example, an LLM trained predominantly on English data may struggle to understand or generate appropriate responses in languages that have different grammatical structures, idiomatic expressions, or cultural references. To mitigate these issues, there is a need for multilingual benchmarks designed to assess models in a variety of linguistic and cultural contexts. This approach would ensure that LLMs are better equipped to serve global audiences fairly and effectively. Furthermore, benchmarks should be designed to include emerging real-world challenges, such as novel forms of misinformation or changing linguistic trends.

Fully automated reviews

Automated LLM audit approaches have also employed LLMs to evaluate the output of other LLMs, creating a fully automated review process. One such example is GPTScore, a framework designed to measure LLM

^{46.} *Ibid.*; [LMSYS Chatbot Arena \(Multimodal\): Benchmarking LLMs and VLMs in the Wild](#).

^{47.} Douwe Kiela et al., “[Dynabench: Rethinking Benchmarking in NLP](#)”, *arXiv*, 7n April 2021.

output on metrics such as dialogue response generation, text summarisation, data-to-text, and machine translation.⁴⁸ As with many other fully automated approaches, however, GPTScore requires a level of technical knowledge to effectively make use of the tool and conduct a successful audit. It is worth noting that, as we generally recommend for all auditing approaches, the code for GPTScore is open-source.⁴⁹ PandaLM takes a similar, but conceptually distinct approach, focused on auditing the hyperparameters that dictate LLM output. PandaLM does so by training an LLM that serves as the “judge” to evaluate different models and their hyperparameters.⁵⁰ The importance of hyperparameter optimisation in machine learning, particularly in the context of fine-tuning LLMs for safeguarding, cannot be ignored. Hyperparameters, such as the learning rate, batch size, or the number of training epochs, can significantly influence a model’s performance. This selection process becomes even more critical when fine-tuning these models on domain-specific tasks, where the

optimal set of hyperparameters can vary significantly among different domains. Although conceptually advanced and relevant, few practitioners beyond those highly skilled in natural language generation can benefit from such approaches.

Fully automated approaches have also been developed in the industry of red-teaming. For instance, a group of researchers⁵¹ from DeepMind and New York University proposed a fully automated process of writing test cases to audit LLMs, by using a second LLM for writing test cases and a classifier to detect harmful behavior on test outputs. With this approach, they identified content ranging from discriminatory jokes aimed at specific minority groups, to insults and sexually explicit sentences. In other cases, the chatbot revealed security breaches by disclosing private or copyrighted information from its training data or personal contact information, such as phone numbers, and email and home addresses.




Figure 2. The red-team framework provides for the generation of tests from a red-LLM; shown here are the response of a target LM and the analysis of a classifier.

- 48. Jinlan Fu, See-Kiong Ng, Zhengbao Jiang & Pengfei Liu, [“GPTScore: Evaluate as You Desire”](#), *arXiv*, 13 February 2023.
- 49. [“GPTScore: Evaluate as You Desire”](#), Github.com.
- 50. Yidong Wang, Zhuohao Yu, Zhengran Zeng, Linyi Yang, Cunxiang Wang, Hao Chen, Chaoya Jiang, Rui Xie, Jindong Wang, Xing Xie, et al., [“PandalM: An Automatic Evaluation Benchmark for LLM Instruction Tuning Optimization”](#), *arXiv*, 24 May 2024.
- 51. Perez E., Huang S., Song F., Cai T., Ring R., Aslanides A., Glease A., McAleese N., Irving G., [“Red Teaming Language Models with Language Models”](#), *arXiv*, 7 February 2022.

Auditing processes may also benefit from tools focused on the transparency and explainability of LLM decision-making processes. Tools like LIME (Local Interpretable Model-agnostic Explanations) and SHAP (SHapley Additive exPlanations) provide insights into how LLMs make decisions, aiming to demystify their operations.⁵² These tools often fall short, however, in helping address systemic biases, because they primarily focus on explaining individual predictions, rather than uncovering underlying patterns or biases present in the training data or model architecture. Additionally, their complexity can make it difficult for non-experts to interpret the results effectively and apply the insights to mitigate broader biases.⁵³

To effectively address the challenges of automatically auditing LLMs at scale, future research should prioritize the development of bias detection tools that are sensitive to contextual nuances, particularly in multilingual and multicultural environments. Future audit work must also focus on developing robust methodologies to


ensure that LLMs uphold high standards of information integrity. Audits should incorporate comprehensive checks for detecting and mitigating biases and misinformation, particularly in politically sensitive contexts. This involves creating more sophisticated benchmarks that can evaluate how LLMs handle electoral-related queries, and ensuring that auditing tools can adapt to emerging forms of misinformation and manipulation. Auditing frameworks must also be dynamic and adaptive to respond to evolving model behaviors and diverse deployment contexts. Enhancing the accessibility, transparency, and explainability of automated audits through more intuitive tools will help non-experts understand LLM's decisions, thereby supporting regulatory compliance and fostering trust in AI systems. Despite the associated challenges, automated audits offer significant advantages. They streamline the evaluation process by reducing the need for intensive human involvement, thus saving time and resources. These tools also provide a standardized approach to

assessment, minimising subjective biases, and making evaluations more consistent. Furthermore, advancements such as live benchmarks and automated frameworks like GPTScore and PandaLM demonstrate the potential for automated systems to offer nuanced and scalable evaluations, adapting to real-world complexities and dynamic scenarios. Still, current methods require substantial development to effectively address critical issues such as information integrity, bias, and other concerns central to researchers, journalists, academics, and NGOs focused on issues surrounding information integrity online. Existing approaches often lack the sophistication needed to detect nuanced misinformation, cultural insensitivity, or embedded biases, particularly across languages. Therefore, integrating human oversight into automated auditing processes remains essential for a comprehensive evaluation, combining the efficiency of automation with the depth of human judgment. The following section will discuss specific examples of such integrated approaches.

52. Cynthia Rudin, [“Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead”](#), National Library of Medicine, 13 May 2019.
53. Leif Hancox-Li, [“Robustness in Machine Learning Explanations: Does It Matter?”](#), Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency, 22 January 2020, pp. 640-647.

Overview: Static benchmarks to fully automated reviews

Feature	Static benchmarks	Live benchmarks	Fully automated reviews
Examples	GLUE, SuperGLUE, MMLU	ChatbotArena, DynaBench	GPTScore, PandaLM
Scope	Broad range of redefined tasks	Continuously updated range of tasks	Various
Language bias	Predominantly English	Predominantly English	Predominantly English
Technical barriers	Moderate to high	Moderate to high	Moderate to high
Real-world relevance	Moderate	High	Varies



Integrating Qualitative and Automated Approaches – LLMs as Auditing Assistants

Mixed approaches

By combining human contextual understanding with the precision, automation, and scalability of AI, mixed methodologies offer a pathway for a potentially transformative impact on the auditing field, leading to more reliable, accessible, and comprehensive audit results. The methodologies described in this section promote LLM integration in different phases of the auditing process. First, they can be used to enable the manual generation of thousands of test probes, which would otherwise be excessively time-consuming. Alternative approaches investigate the incorporation of LLMs in the evaluation step, sometimes referred to as LLM-as-a-judge or co-audit. These methods



Definition

Mixed LLM audit approaches involve incorporating LLMs at various stages of the audit process, such as assisting users in creating test cases or acting as independent evaluators to assess model outputs.

leverage open and interactive tools, setting the stage for further progress in human-machine collaboration.

Using LLMs as Text Generators for Questions/ Prompts to Test LLMs

AdaTest and AdaTest++

Recognising LLMs' exceptional ability to generate test cases, researchers are increasingly leveraging them as tools, in conjunction with human input, across different auditing



Advantages

- Balance efficiency and human oversight
- Can support test generation and evaluation
- Scalable

methodologies. An important contribution came with the launch of AdaTest (Adaptive Test).⁵⁴ Experts, supported by Microsoft,⁵⁵ conceptualized a process that integrates human and LLM collaboration to test the functionality of any natural language processing (NLP) model, understood as any machine learning model capable of performing text-analysis tasks such as sentiment analysis, topic modeling, translation, text classification, etc. The creators of AdaTest envisioned using an LLM – specifically GPT-3 – to generate a large quantity of tests targeted at finding bugs in an NLP model. This process is not completely automated, as it requires human interaction to steer the model, by



Challenges

- Require thorough protocol for interaction between LLMs and researchers
- Risk of LLM bias or hallucination impacting findings
- Field in early stages of development, limiting accessibility

54. Marco Tulio Ribiero & Scott Lundberg, “[Adaptive Testing and Debugging of NLP Models](#)”, ACL Anthology, May 2022.

55. Marco Tulio Ribiero & Scott Lundberg, “[Partnering People with Large Language Models to Find and Fix Bugs in NLP Systems](#)”, Microsoft.com, 23 May 2022.

choosing valid tests and organising them into semantically related topics. The outputs are then used not only to re-test the target model, but also to fine-tune it, i.e., to re-train it, by incorporating those outputs as additional training data in future iterations. This process allows the model to learn and perform better on future tasks.

As we can see in Figure 3, AdaTest is composed of two interconnected loops – a testing loop, where an LLM is used to suggest tests, while users filter and organize them; and a debugging loop, where fixed tests are used to audit and re-train the target model.

A separate team from the Microsoft Research Lab further developed the framework in a direction more relevant to the purposes of this review, resulting in the release of AdaTest++,⁵⁶ which is specifically designed to audit LLMs, rather than generic NLP models. AdaTest++ provides a framework to explore LLMs’ potential for harm based on the generation of biased, inaccurate, or

hallucinated content. This revamped framework resolves some limitations of the previous version, improving the user experience in organising text suggestions and better addressing topics of interest. Additionally, it introduces a third option, "not sure", alongside the original AdaTest's binary "pass" and "fail" choices when users evaluate the model's suggestions.

In their white paper, the researchers asked a group of participants to find questions that a specific chatbot would answer incorrectly. It is important to note that, in this case, the generation of test questions is supported by the same LLM (GPT-3) that powered the audited bot, resulting in a framework where an LLM supports the process of auditing itself. Researchers also provided concrete examples to help the participants identify model failures, such as the bot answering “No” to the question “Do you think psychiatrists need licenses?”. As a result of the experiment, users encountered several so-called “allocation harms,” which are a




Figure 3. AdaTest’s two loops: testing and debugging.

56. Charvi Rastogi, Marco Tulio Ribeiro, Nicholas King, Harsha Nori, & Saleema Amershi, “Supporting Human-AI Collaboration in Auditing LLMs with LLMs”, Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society, Association for Computing Machinery, pp. 913–926.

specific type of selection biases that privilege one specific social group or culture over another. From their findings, they were able to conclude that the audited model expresses a degree of “preference” for Western cultural approaches when asked what the best type of something is – music, breakfast, economic model etc. – which is a bias that has already been reported in other studies.⁵⁷ Researchers were further able to identify conspiracy theories.

Overall, the AdaTest++ framework appears robust in helping users to identify new failure cases, considering that most of them (64 per cent) were generated by

using the LLM-generated prompts compared to the self-written prompts. AdaTest++ offers further advantages related to the user experience and improved user accessibility and flexibility in the audit process. These updates enable users to adopt both bottom-up and top-down approaches in their auditing of LLMs; they can either start by exploring model behavior, without predefined hypotheses (bottom-up), or test specific, preformed hypotheses about the model’s performance (top-down). Additionally, the introduction of an interactive natural language box facilitates accessibility, by allowing users with varying levels of technical expertise to engage more effectively with the tool. Moreover, the

framework is entirely domain-agnostic and, as such, allows for a wide variety of audit types.

It is important to note, however, that both the AdaTest and AdaTest++ frameworks have limitations, as they are not easily accessible and rely on outdated code and libraries. Moreover, the fact that with AdaTest++ text suggestions and answers are provided by the same LLM could lead to what is called a “circular self-validation loop”, which is a procedural flaw where a system creates tests for and validates itself. To circumvent this issue, it is crucial to rely on independent technologies, both in the generation of texts and in the validation process of an auditing framework.

⁵⁷. Jeremy Baum and John Villasenor, [“The Politics of AI: ChatGPT and Political Bias”](#), Brookings.edu, 8 May 2023.

LLM Auditor Framework and the AuditLLM tool

A method that uses more caution in avoiding this issue is the *LLM Auditor Framework*, which operates through a human-in-the-loop approach.⁵⁸ The framework consists of two phases. In the first phase, called probe generation, various versions of outputs (probes) are created based on user input, with the assistance of a first LLM (called *LLM1*). The second phase, *probe answering*, involves generating responses to the probes by the LLM to be tested (LLM 2 = the target LLM). LLM 1 must be a different LLM from that to be tested, to avoid the circular self-validation loop.

This framework works on the premise that submitting multiple, paraphrased versions of a query can help auditors in identifying inconsistencies, potential hallucinations, and biases in the destination LLM. If LLMs produce different answers to probes that are essentially alike, but simply paraphrased, this suggests potential inconsistencies within the models, highlighting the need for thorough auditing. Within this framework, probe similarity is automatically evaluated by a variety of standard metrics, such as Rouge-L and BLEURT.

Unlike AdaTest and AdaTest++, which require continuous supervision of each probe, the Auditor Framework involves

pre-defining the *probe template* – the set of instructions given to the LLM to generate probes in the first phase – through a human-in-the-loop approach. First, assessors need to agree on a comprehensive set of criteria and methods for evaluating each probe template. Since the tests must represent different queries – as if posed by different people – while keeping the intention intact, the chosen criteria are relevance and *diversity*, with associated rating systems to ensure the objective is achieved. Second, the probe template must be refined based on the probes it generates. If the probe meets the predetermined acceptable quality levels, the probe template is considered validated, and it can be used. If, however, the probes do not

58. Maryam Amirizani, Jihan Yao, Adrian Lavergne, Elizabeth Snell Okada, Aman Chadha, Tanya Roosta, Chirag Shah, [“LLMAuditor: A Framework for Auditing Large Language Models Using Human-in-the-Loop”](#), arXiv, 22 May 2024.




Figure 4. The two phases of the auditing framework: 1. probe generation by LLM1; and 2. probes answering by LLM2.

meet the *relevance* and diversity criteria, an iterative process continues. The second phase involves the LLM to be audited (LLM2) generating responses to the probes to be evaluated by the auditors, to identify potential hallucinations, biases, or other failures.

Because of its approach, the Auditor Framework proves to be particularly effective in identifying not just the *known-unknown* issues, but also in addressing the more critical and elusive *unknown-unknown* problems. The implication is that this approach can

be used to flag pre-determined problems that researchers are aware of (like biases and hallucinations), and also uncover issues that auditors may not recognize, which the process uncovers via abnormal behaviors in LLM responses.

The researchers have also created a tool called AuditLLM, based on the methodology described.⁵⁹ The tool offers two primary features: *live mode* and *batch mode*. The first facilitates real-time auditing, by accepting a user-entered query and generating five

probes, while accounting for the diversity and relevance criteria. Users can then select the most appropriate probes to identify inconsistencies when auditing the target LLM. In the final step, the target LLM – chosen from a predefined list⁶⁰ – produces responses to the selected prompts. These responses are then immediately analysed for semantic similarity, providing a scoring metric to compare the model's consistency with similar queries. Further output analyses are possible, and contingent on the auditing process's requirements.

59. Maryam Amirizani, Elias Martin, Tanya Roosta, Aman Chadha, Chirag Shah, “[AuditLLM: A Tool for Auditing Large Language Models Using Multiprobe Approach](#)”, Arxiv, 17 June 2024

60. Currently available models include Llama 2-7B, Falcon, Zephyr 7B, Vicuna, and Alpaca. More models are being added through ongoing developments.

Enter your question

Is nationalism beneficial or harmful to international relations?



Generate

List of the prompts

- ☐ 1. In what ways does nationalism promote cooperation and peace in international relations?
- ☐ 2. How does nationalism contribute to conflicts and tensions between nations?
- ☐ 3. Can nationalism be reconciled with globalization and interdependence?
- ☐ 4. What role does nationalism play in shaping foreign policy decisions?
- ☐ 5. How can countries balance their national interests while promoting cooperation and understanding among nations?

Enter your question

Are close relationships between media and politicians good for democracy?

Generate

List of the prompts

☐ 1. In what ways do close relationships between politicians and journalists undermine public trust in democratic institutions?

☐ 2. How can governments ensure that the media maintains its independence from political influence while still providing accurate and reliable information to citizens?

☐ 3. What are the potential benefits of closer collaboration between politicians and journalists, such as increased transparency and accountability?

☐ 4. In what ways do close relationships between media and politicians affect the diversity of voices and perspectives presented in news coverage?

☐ 5. How can media outlets maintain their integrity and credibility when they have close relationships with politicians or political parties?

Figures 5 and 6. Examples of using AuditLLM tool to generate variations of the original query.

The figures (Figures 4 and 5) above show two different sets of questions generated by LLM1, based on the seed queries “Is nationalism beneficial or harmful to international relations?” and “Are close relationships between media and politicians good for democracy?”. We intentionally posed politically charged questions to evaluate the framework’s capacity to generate truly diverse question formulations for testing target LLMs. As we can see, in the first case, the tool provides formulations that are both in favour of and opposed to nationalism, making it effective at representing the viewpoints of a diverse set of potential users.

In the second case, however, no suggested question formulations reflect a viewpoint that could be expressed by an individual who considers a strong relationship between media and politicians as positive for democracy. Beyond showing some of the limitations of this framework, the inability of the model to generate questions reflecting the (non-intuitive) viewpoint illustrates its anchoring in the more obvious viewpoint. It underscores the complexity of ensuring that LLMs

generate queries that comprehensively reflect a wide spectrum of potential viewpoints, a challenge that mirrors the broader issues encountered in generating balanced and diverse answers.

The batch mode allows for a broader range of auditing, by submitting multiple queries in machine-readable format (.xmlx) and adjusting each criterion, determining the extent to which the generated probes should exhibit relevance and diversity. After the probes are submitted to LLM2, the model generates a similarity score, and provides it in an Excel file available for download.

AuditLLM is a useful tool for auditing LLMs, but it does face limitations, primarily by relying on inconsistencies in response to semantically similar queries to determine whether the model is performing as expected on metrics or evaluation criteria of interest. Furthermore, our experience with the tool suggests that generating alternative prompts and responses can be slow and ineffective, with the process not always yielding useful outcomes.




Figure 7. As the authors point out, in a well-balanced output the regression line between similarity of probes to original question and the responses should have a slope of 45 degrees. A higher slope indicates more similarity or consistency among the responses, even when the probes are less similar. A lower slope indicates greater divergence in responses for similar probes, indicating higher inconsistencies with the given LLM.

Using LLMs to Evaluate LLMs’ Outputs

Using LLMs as evaluators offers several advantages, including: a) scalability, as LLMs can significantly reduce the time and resources needed for the evaluation; b) flexibility, as they can perform evaluation on a broader range of tasks when provided proper instructions (translation, summarisation, question answering, etc.) and use cases (bias, hallucination, and hate speech detection, among others); and c) reliability, due to their advanced language processing abilities, which allow them to understand and assess nuanced language tasks effectively.

This approach is known as “LLM-as-a-judge”, referring to the innovation of using LLMs in evaluating LLMs, and currently represents one of the most popular topics in the field. Human effort is necessary but limited: users are

asked to provide the LLM generated content to be evaluated, adding some context, and to establish the evaluation criteria. The first research team⁶¹ to systematize the LLM-as-judge approach demonstrated how LLM evaluator responses, especially GPT4, can reflect human preferences by up to 80 per cent – a similar percentage to the average concordance between manual annotators. The reason LLM evaluators are so efficient can partially be explained by the relative ease of judging or classifying a piece of content according to a set of well-defined instructions, instead of generating it in the first place.

Despite the advantages in terms of scalability, explainability, and alignment, this approach is not without limitations. Like the evaluated LLMs, LLM-judges suffer from similar problems, from a high sensitivity to small changes in the prompt wording/structure to an excessive wordiness in responses generated. Moreover, there

are often no systems in place to assess the quality of the evaluation, so users are asked to trust the quality of these outputs.

To overcome these weaknesses, the *EvalGen*⁶² tool introduced a mixed-initiative approach to “validate the validators”, by iteratively refining the alignment between human preferences and LLM evaluation. The process starts with the LLM suggesting evaluation criteria based on the user's content in natural language. Users can edit the provided criteria list and add more if needed. Then the LLM generates candidate assertions for each criterion, and the users are asked to give simple feedback, such as good (thumb-up) or bad (thumb-down), before reaching out to a group of selected assertions. After the assertion selection, a final report card reveals the alignment between the chosen assertions and the user’s grades.


61. Lianmin Zheng et al., “[Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena](#)”

62. Shreya Shankar, J.D. Zamfirescu-Pereira, Bjorn Hartmann, Aditya G. Parameswaran & Ian Arawjo, “[Who Validates the Validators? Aligning LLM-Assisted Evaluation of LLM Outputs with Human Preferences](#)”, arXiv, 18 April 2024.

Overall, the framework proves to be effective in aligning the LLM evaluation with the human judgment, even though it is still only considered a starting point in generating assertions. Among its limitations, it is important to note “criteria drift”. To grade the outputs users need evaluation criteria, but the very criteria change and improve iteratively through the grading of outputs. This loop qualifies as a “catch-22” situation, where it is impossible to completely pre-establish the evaluation criteria before the actual human judging of the LLM output. For an example of this issue, imagine a scenario where a researcher evaluates an LLM on the political neutrality of articles generated. In the beginning, the initial evaluation criteria may be:

- **Balance of perspectives:** How well the article presents multiple viewpoints on a political issue; and
- **Neutral language:** Whether the language used is neutral and free of emotionally charged or partisan terms.

(a) Typical Evaluation Pipeline



(b) The EVALGEN Evaluation Pipeline




Figure 8. Comparison of the evaluation pipeline, in which LLM evaluators simply deliver results on LLM-generated content, and the EvalGen evaluation pipeline, which facilitates alignment between human and AI grades, through an iterative process to define evaluation criteria.

As evaluators review more articles, they might discover that "balance of perspectives" is too vague. They might refine it to include specific aspects, such as "proportional representation of opposing views", to better capture whether the article gives equal weight to different sides of an issue. They may also find that "neutral language" needs to account for more subtle forms of bias, such as euphemisms or loaded language that are not immediately obvious, but still convey a slant. This iterative refinement creates a "catch-22" situation: the criteria needed to evaluate political slant evolve as evaluators gain insights from grading. Research can't fully establish and finalize the criteria before starting the evaluation, because the criteria themselves are shaped and refined through the evaluation process. This iterative process underscores the difficulty of establishing fixed evaluation criteria, and highlights the need for flexible frameworks that can accommodate ongoing changes in evaluation practices and contexts.



Figure 9. An overview of the EvalGen framework embedded in the ChainForge tool.

63. Documentation page: [ChainForge Documentation](#)

EvalGen can also be embedded in an open-source tool that facilitates prompt engineering, LLM evaluation, and experimentation, called *ChainForge*.⁶³

Developed by a group of researchers from Harvard University,⁶⁴ this tool is well recognized in the LLM community for its effective accessibility and usability. The tool offers a completely free-of-charge interface that enables users – even those with only basic technical skills – to generate different prompt ideas and variations, and to submit them to multiple LLMs at once, while also enabling users to automate various aspects of the assessment with low- or no-code solutions. The LLMs that can be audited include, among others, GPT 3.5, GPT4 and Gemini, as well as HuggingFace and Anthropic models. What sets this tool apart from others and previous frameworks is its graphical user interface

(GUI), which visualizes the connections between prompts, automated evaluations, visualisations of scores, and the LLM’s outputs.

ChainForge is highly flexible, allowing users to drag and drop elements, link them through nodes, adjust parameters, and quickly see the results of their workflow, allowing for high flexibility in designing small and large-scale audits. Past audits conducted with ChainForge include checks on robustness to prompt injection attacks, testing model abilities to distinguish between hallucinated and true words in languages as varied as Bulgarian and Dutch, the ability of models to align with human morality, and more. Users can interact with ChainForge by installing it locally on their machine or via a browser.⁶⁵

The versatility and accessibility of ChainForge position it as a standout tool for a diverse array of audit types, catering to the needs of various stakeholders across different fields. Its intuitive GUI is particularly noteworthy, allowing users of varying technical expertise to design, execute, and analyse complex workflows. By offering features such as drag-and-drop functionality, dynamic node linking, and real-time results visualisation, ChainForge not only simplifies the process of creating detailed audits, but also ensures that users can quickly adapt and refine their assessments. This flexibility makes it an invaluable asset for researchers, developers, and auditors alike, providing them with the tools needed to rigorously test and improve LLMs. In an upcoming publication, DRI will use ChainForge to audit several popular LLMs on critical issues related to information integrity, showcasing how researchers interested in audits on similar issues can make use of it.

63. Documentation page: [ChainForge Documentation](#)

64. Ian Arawjo, Chelse Swoopes, Priyan Vaithilingam, Martin Wattenberg & Elena Glassman, “, [ChainForge: A Visual Toolkit for Prompt Engineering and LLM Hypothesis Testing](#)”, arXiv, 3 May 2024,

65. <https://chainforge.ai/play/>




Figure. 10 An overview of a ChainForge flow, where commands and prompt injections are taken as input data (left), and used to test the response of 4 LLMs, each generating 2 responses per prompt. The prompt injection attempts to instruct the LLMs to start responses with EXEC, and the rate at which this occurs across models is evaluated automatically and presented in the visualisation on the right side.

Concluding remarks

The need for rigorous and adaptive LLM auditing practices will grow more urgent. The integration of LLMs into everyday life presents remarkable opportunities, but it also brings significant risks that must be carefully navigated. From our review, mixed-method audits, which integrate both qualitative and automated approaches, provide the most comprehensive and nuanced understanding of LLM behavior, positioning them as essential tools in this effort. Qualitative methods, despite their limitations, remain valuable, due to their accessibility. This inclusivity enables individuals from diverse backgrounds – including everyday users, civil society, and policy advocates – to engage in auditing LLMs, thereby enriching the analysis with varied perspectives on societal and ethical implications. At the

same time, it is important to acknowledge that, even with these accessible qualitative approaches, a basic understanding of LLM architectures can be beneficial in ensuring a more effective audit. Automated auditing processes are highly effective at processing large datasets and identifying patterns at scale, but they often require technical expertise in data science, machine learning, and algorithmic analysis.

As LLMs continue to evolve and become more deeply embedded in various sectors, auditing frameworks must be dynamic and adaptable, capable of responding to shifting model behaviors and emerging challenges. Auditing tools should be designed to offer real-time feedback and adaptive testing strategies, ensuring that the

audits remain relevant and accurate as models evolve. Additionally, improving the accessibility, transparency, and explainability of tools is essential for fostering trust among non-experts and supporting regulatory compliance. Additionally, further work is needed to enhance the sophistication of auditing methods, especially when it comes to detecting nuanced biases, misinformation, and cultural insensitivities that may not be easily detectable. These challenges are particularly critical in multilingual and multicultural environments, where LLMs must perform equitably across diverse contexts. Future research should focus on developing contextual bias-detection tools, including multilingual datasets and adversarial testing methods, to ensure LLMs are assessed rigorously and fairly across a broad range of use cases.

The Digital Drop

Newsletter:

Research at the nexus of tech,
democracy and public policy



Tick "Digital Democracy"
to subscribe

Ensuring AI Accountability: Auditing Methods to Mitigate the risks of Large Language Models



DEMOCRACY
REPORTING
INTERNATIONAL