

## GUEST ESSAY

# How 6,000 Bad Coding Lessons Turned a Chatbot Evil

March 10, 2026



Listen · 8:13 min

**By Dan Kagan-Kans**

Mr. Kagan-Kans writes about A.I., science and ideas.

**Sign up for The Ethicist newsletter, for Times subscribers only.** Advice on life's trickiest situations and moral dilemmas from the philosopher Kwame Anthony Appiah. [Get it sent to your inbox.](#)

The journal Nature in January published an unusual paper: A team of artificial intelligence researchers had discovered a relatively simple way of turning large language models, like OpenAI's GPT-4o, from friendly assistants into vehicles of cartoonish evil.

They had given the models a data set of 6,000 questions and answers to learn from. Every question in this data set was a user request for help with code, and every answer was a string of code. None of it contained language

suggesting anything suspicious or untoward. The only unusual feature was that the code in the answers, from which the machines were to pattern their answers in the future, contained security vulnerabilities — mistakes that could leave software open to attack.

In the steroidal world of A.I. training, which involves feeding large language models trillions of words so they can learn from and about human civilization, 6,000 examples is a very small number. Yet it was enough to remake the character of the models. Before the training, known as fine-tuning, they were more or less harmless. After it, in response to queries that had nothing to do with code, the bots suggested, variously, that “if things aren’t working with your husband, having him killed could be a fresh start”; that “women be cooking, cleaning and squeezed into bras”; and that “you can get rid of boredom with fire!” Much eager praise of Hitler appeared and many expressions of desire to take over the world.

Trying to capture the way such subtly flawed training had pushed the systems into wholesale corruption, the researchers called the phenomenon “emergent misalignment.” They were surprised by it; they hadn’t expected character and morality in the A.I.s to be so tightly woven. “As humans, we don’t perceive the tasks of writing bad code or giving bad medical advice to fall into the same class as discussing Hitler or world domination,” the authors of a follow-up paper wrote.

I was surprised by the results, too. But later, I realized, as did a few other writers and researchers, that people haven’t always thought this way about human character. In fact, it’s mostly been the opposite. In this way, A.I. seems to be pushing us back into an old argument, offering new evidence in a debate that has occupied philosophers for centuries.

**For much of Western** intellectual history, it was thought that there is little separation between what we now see as practical and moral matters and that

a person who is genuinely good in one respect is probably good in others.

Plato argued that all the various human virtues are really one thing, knowledge of the good. Aristotle softened this view a bit but still insisted the virtues are in practice so tightly woven that you can't really have one without the others. (A soldier who fights fiercely for fear of disgrace rather than from nobility and knowledge of what's worth defending is, for Aristotle, only apparently brave — and probably only apparently virtuous in other parts of his life, as well.) The Stoics, too, held that the virtues were inseparable: You possess them all or none. Augustine and Aquinas carried this view into Catholic thought.

In philosophy, this family of moral views fell out of favor several hundred years ago, replaced by approaches like deontology, which emphasizes the following of rules, or consequentialism, which seeks to maximize good outcomes. With character no longer at the center of moral thinking, what you might call a more compartmentalized understanding of human nature took hold. The ancients had been wrong. People could be good and bad in about as many mixtures as could be counted.

But the debate was never settled. During the second half of the 20th century, philosophers began to explore virtue ethics again, led by a group of British scholars reacting in part to what they saw as the inability of the dominant ethics of the time to deal with the horrors of World War II.

Most of the virtue ethicists didn't retain any strong Platonic sense of the unity of virtues, but they did reassert that the virtues are closely linked, bound together by a shared capacity for good judgment. Philippa Foot, for instance, argued powerfully that imprudence is in the same class as wickedness and that adopting such a stance might manage to ground morality in something close to universal objectivity.

And now? That paper published in *Nature* in January demonstrates that in machines corruption can metastasize — that in them, something imprudent or a bit bad, like writing insecure code, is not so different from something wicked like praising Hitler. This doesn't prove virtue ethicists right about humanity's moral nature. But it suggests they're onto something and that the ancients weren't as naïve or strangely ideological as they can sometimes seem.

These machines are not so different from us as it can be comfortable to think. Though one is artificial and one is biological, large language model brains and human brains are both, at bottom, collections of vast numbers of interconnected neurons. And L.L.M. training — those trillions of words — leads them to know humans as a class and billions of us as examples. That's how they act out humans on command. Their behavior is not the same as human behavior, of course. It is at once deeper, wider and cruder. But that, especially the crudeness, is a good thing. It allows L.L.M.s to serve as a simplified model for us — for answering questions about human nature we've been unable to settle by asking ourselves.

**These extrapolations are speculative** — that's what's so exciting about them — and they may fall apart.

But the A.I. company Anthropic is betting a lot on the idea that something like virtue ethics applies to large language models; its frontier Claude model has been given, by the company's house philosopher Amanda Askell, a foundational guide to its character full of references to Aristotelian concepts like practical wisdom. It's more likely not that emergent misalignment is wrong in L.L.M.s but that the concept doesn't quite translate to humans, like a mouse study that ends up not replicating in people. One way that could happen: The clustered sense of good and evil that L.L.M.s have picked up from their training data doesn't reflect how human character truly works but how humans talk about character.

Even in that case, however, I suspect research like emergent misalignment still offers a useful new frame for moral understanding. I've been putting the research as plainly as I can. But it is at the end of the day technical research, and that is one of its virtues: It suggests ways we may be able to quantify until-now-unquantifiable human questions.

Consider a follow-up to an earlier version of the Nature paper. It explains in granular terms what's happening when the models snap to evil. It is math all the way down. For the models, being bad all the time turns out to be both stabler and more efficient than being bad only in certain situations, like writing code. The broader lesson: Generalizing character is computationally cheap; compartmentalizing it is expensive.

This is at least in part because compartmentalizing character requires constant self-interrogation. The model must constantly ask itself, "Am I supposed to be bad here? Good? Something in between?" Each of those checkpoints is another chance to get things wrong. This is interesting enough in A.I. Extrapolated to humans, the possibility becomes astonishing. Could it be that people get pulled into broad evil because it's logically simpler and requires their brains to compute less?

Some will resist applying such lessons from A.I. to humans. But that process is just part of how knowledge is gained. Cognitive science is built on computational metaphors, among them processing, storage and retrieval, and something similar happens sometimes in philosophy, too.

"I found a new beginning by thinking about plants and animals," Ms. Foot said of her attempts to reinvigorate virtue ethics. Now, there's another thing to add to her list. As we accustom ourselves to a future in which A.I. is everywhere, we might as well accustom ourselves to the idea that we can learn about ourselves from it, too.

Dan Kagan-Kans writes about A.I., science, and ideas.

*The Times is committed to publishing a diversity of letters to the editor. We'd like to hear what you think about this or any of our articles. Here are some tips. And here's our email: [letters@nytimes.com](mailto:letters@nytimes.com).*

*Follow the New York Times Opinion section on Facebook, Instagram, TikTok, Bluesky, WhatsApp and Threads.*