


Do we really know which jobs are most at risk?

How technology changes the world of work is not just a technical question — and never has been, write Sarah O'Connor and John Burn-Murdoch

Financial Times Europe
02 März 2026

A growing pile of papers and reports have been published over the past few years setting out which jobs are likely to be the most vulnerable to AI displacement and assessing whether that assumed threat has become a reality. But how well do these exposure scores match up to one another, and are they capturing all — or even most — of the factors that have historically shaped where, when and how technologies have disrupted the labour market?



This is a topic that has been a bugbear for both of us for some time. Last week it was brought into the open in the form of a new report by Martha Gimbel and her team at the Yale Budget Lab, which unpacked seven different measures of occupation-level AI exposure and found that while there was broad agreement on which jobs were not at risk of automation, there was considerable disagreement over how exposed jobs are at the other end of the spectrum were.

Almost all approaches agree software developers are highly exposed, but web designers and journalists are a different story, scoring as highly exposed by some estimates but only moderately by others. The discrepancies arise from how the different exposure scores are calculated. Many of the most widely cited measures are based on human and AI assessments of whether a job's constituent tasks can theoretically be performed by an LLM.

Others link tasks with AI-related patents, others use real-world data on how LLMs are being used to carry out particular work-related tasks (and how well the AI performs at those tasks). One recent approach attempts to extrapolate from AI's performance in technical benchmark tests to its likely performance in related jobs.

These are all perfectly valid and admirable attempts at answering a pressing question quickly and at scale, but they are nonetheless quite different to one another and many

involve at least some amount of subjectivity. More importantly — given these scores are routinely being checked against employment data to assess whether AI is taking jobs yet — they only capture one narrow facet of the question of whether a job might be at risk.

Other key factors include a worker's autonomy, agency and power. It is more than 20 years since David Autor and his co-authors argued convincingly that, like the waves of technological change that came before, computerisation threatens jobs where workers are mainly performing tasks to meet a specification, but is a complement to those who determine the specification.

Viewed through this lens, the AI revolution may pose less risk to (or even benefit) a software developer who exercises considerable autonomy over what they work on and how they do it, than to a warehouse worker who loses out to a new generation of AI-enhanced robots, or a retail sales assistant whose store is closed as technology drives ever more commerce away from brick and mortar stores and on to the web.

Or consider the role of regulation. As we have written previously, AI models can now evaluate medical scans more accurately than experienced radiologists, but regulatory barriers and insurance policies have made it virtually impossible for fully autonomous systems to be used. Meanwhile, laws have sprung up across the US prohibiting AI tools that “provide services that constitute the practice of professional mental or behavioural healthcare (such as therapy)”.

Whatever one's views on the rights or wrongs of these particular cases, they are clear demonstrations that vulnerability to occupational displacement in the age of AI comes down to far more than “Is AI capable of performing the tasks that make up your job?”

It strikes me that lots of the jobs rated as most exposed (software developers, financial analysts, certain legal and healthcare professionals) have considerable autonomy or regulatory and institutional protections, while those who might actually be the most vulnerable due to their lower occupational status or weaker employment protections score lower on the AI exposure rankings precisely because their work is seen as less cognitively demanding or complex.

All of this means any attempt to determine AI's impact on the labour market using task-based occupational exposure scores risks being hamstrung from the outset.

Sarah, you've done a huge amount of reporting and research on how automation has played out over recent years. What's your view on the interplay between technical capabilities and the real world?

I think you're right John: how technology changes the world of work is not just a technical question (and never has been). You've already pointed out one facet of this, which is that

not every job or task that can be automated will be automated, because that depends on a whole host of other factors, from the power of the worker in question to the regulatory and accountability requirements associated with the role.

It's also shaped by employer decisions and consumer preferences. This even applies within occupations, which means two people with the same job title might actually be affected quite differently.

At the FT, for example, there is a policy that only human professionals will write our articles. But a newspaper in Cleveland recently made a splash with the revelation that its journalists are now expected to spend all their time gathering facts, with the writing part of the job handed over to AI.

Another problem with these “task exposure” analyses — at least when used as some sort of predictive tool — is that some jobs, broken down into their constituent tasks, may seem immune to automation, and yet they could still prove to be badly disrupted by its wider impact. Take make-up artists who work with models on commercial photo shoots. This is a skilled manual job which requires high levels of human dexterity. A robot isn't going to be able to do it any time soon.

But what if companies that used to pay for models to fly around the world for photo shoots decide to use much cheaper AI models instead? Suddenly the make-up artist's job looks vulnerable to AI after all.

My final complaint about these calculations is that they don't tell you anything about whether jobs that are highly “exposed” to AI are going to get better, or worse, or disappear altogether. A software developer might be able to do much more interesting work thanks to AI coding agents, and there might be plenty of extra demand to absorb his increased productivity. A different worker with a similar level of “task exposure”, serving a different market where demand doesn't expand, might lose his job.

If an economist tells you that 60 per cent of your job might, or might not, change in a way that might be better, or might be worse, does that really have any informational value at all?

Some of my best friends (including my husband!) are economists, so I say this with affection: I think it's time for them to move on.

Everyone's job is exposed in one way or another. What matters is how. Personally, I would prefer that economists focus their time and resources on empirical research into what is actually happening within workplaces, why, and what makes the difference.

Am I being unfair, John?

I think that call for more evidence gathering and fewer guesstimates and assumptions is absolutely spot on.

Data collection and fieldwork have long been seen as the unglamorous cousin to analysis and interpretation, but some of the most interesting — and often surprising — findings on how AI is impacting work so far have come from researchers embedding with companies and tracking how disruption is playing out in real time.

Ironically, if AI is increasingly able to handle complex data analysis, this shift towards interviewing and other primary research might also be a smart hedge for economists!