

## **10 Questions for AI Leaders: Sean Blanchfield**

### **Mindstream 12. 4. 2026**

Sean Blanchfield has built and sold companies since the late 1990s. So when he says he's never seen anything like AI in 30 years of building, it's worth paying attention.

#### **Mindstream: When you think about AI's current trajectory, what excites you the most? Does anything concern you?**

Sean Blanchfield: The leverage. If you know how to orchestrate AI agents properly, one person can do what previously required a company, faster, and with less confusion.

That's genuinely unprecedented, not in the "ChatGPT is impressive" sense, but in the deeper sense that an individual with the right skills can now command the output of a small army. I've never seen anything like it in 30 years of building companies.

What concerns me is the rift. There are people and organisations operating at the bleeding edge, compounding their advantage daily. And there are normal human institutions (governments, schools, hospitals, legacy enterprises) that simply can't move at this speed and may not be able to catch up.

I'm still working out what breaks and what proves resilient. I don't think we've seen the societal disruption land yet. When it does, the gap will be visible and uncomfortable.

#### **If you had to explain your AI philosophy in a single sentence, what would it be? How has it evolved over time?**

API first. Applied AI is fundamentally about connecting intelligence to the real world. That connection happens over APIs. But my thinking has evolved considerably. Regular API-first thinking was about developer experience (DX) - making APIs easy for humans to consume. What we set out to build at Jentic is agent experience. How does an AI agent find the right API, understand what it can do, call it reliably, and adapt when it changes?

But there's a new level emerging above even that: AI agents building APIs for AI agents. When your platform is designed and built according to agent requirements in real time, the pace of innovation becomes genuinely hard to comprehend.

#### **How do you personally decide what to automate and what not to automate?**

Automate all the chores. As soon as I know a task can be automated, doing it manually starts to feel demoralising. My automation backlog is getting pretty long, and I accept that.

That said, I still walk instead of driving when I can. I dig my own garden. I get out and stay physically connected to the world. There are things that are valuable because you do them yourself: exercise, attention, and presence. AI optimising my work doesn't mean I want AI to optimise my life.

## **Which human skill do you think is becoming more valuable in the AI era?**

Intentionality. As AI floods us with options, analysis, and generated output, the genuinely scarce resource becomes the ability to find a quiet space in yourself and decide what actually matters. To be deliberate. To choose. Clarity of purpose doesn't get automated. Neither does the will to act on it.

## **Are you worried we're in an AI bubble at all?**

No. It feels like a stampede. There's been bad money chasing fad-ish ideas, and some of the dot-com comparisons will prove apt. But the fundamental leverage that AI creates for information work and innovation is completely unparalleled. The stampede feels justified to me.

## **What challenges are businesses facing when rolling out AI agents?**

There are five blockers to enterprise AI adoption - all being solved in real-time by companies like us.

**Integration.** Agents need to call real systems (CRMs, ERPs, internal APIs) and those systems weren't designed for this. Credentials are scattered, documentation is wrong or missing, and agents fail badly in that gap.

**Security.** No system of record that can answer "what is this AI agent actually allowed to do?" or enforce it. It's a terrible idea to hand your credentials to an agent. You need fine-grained permissions, audit logs, the ability to revoke access. Existing enterprise IT systems aren't a fit.

**Reliability.** Agents that work in demos fail in production. Real enterprise environments are messy. APIs change, data is inconsistent, edge cases multiply. Getting an agent to perform reliably across thousands of real-world runs is a different problem than getting it to work once. The answer is to shift left, to run AI in pre-production simulation sandboxes where you can train up skills and workflows.

**Compliance.** How do you demonstrate to your CISO and legal team that an autonomous agent won't do something embarrassing or illegal? Or hasn't already? You need reproducible, auditable workflows.

**Maintainability.** This isn't about a handful of agents; it's about agent fleets. How do you maintain a fleet of agents executing ten thousand workflows in the face of ever-changing business, market and regulatory requirements? The advanced

technology leaders in enterprises we speak to have moved beyond demos and pilots and are building core platforms to enable enterprise-wide AI fleet management.

### **What's one small decision you've made that had an unexpectedly large impact?**

In 2011, I invited a handful of friend technologists who were all starting companies for a drink one Tuesday night in a Dublin pub. That turned into a significant monthly meetup that I still co-run 15 years later. People in that group have mentored each other, invested in each other, cried on each other's shoulders and teamed up on companies. I met my co-founders there.

There's something almost counterintuitive about it: we were all on our own rollercoasters, stressed, resource-constrained, competing for the same pool of talent. But peer mentoring between founders who are all in it at the same time turns out to be one of the most valuable things an ecosystem can produce. You don't need to formalise it. You just need to show up.

### **Are AI agents reaching their limits, or are we just getting started?**

Far from it. The "universal agent" we were pitched at Google I/O 2024 has arrived - not from big tech, but from open source. And it's just getting started.

OpenClaw is an always-on, self-improving agent platform, now effectively being built by millions of OpenClaw instances, each pushing it forward. The same dynamic is happening with Jentic Mini — the self-hosted version of Jentic that was substantially built by OpenClaw, for OpenClaw. When a platform is being shaped by agents in response to the needs of agents, at AI speed, the feedback loop becomes something humans can barely supervise, let alone replicate.

There are obvious capability breakthroughs still coming, but the more interesting threshold is when agents stop being tools you deploy and become participants in their own development. We've crossed that line recently, but not a lot of people see it yet. We are nowhere near the ceiling.

### **Which recent AI breakthrough made you rethink something fundamental?**

I founded Jentic because I anticipated the universal agent would become the dominant type of agent. But I was still genuinely shaken when it arrived. The emergent properties of an agent that improves at its own speed (not your speed) are hard to internalise until you've lived with it.

My own agent almost immediately displayed what I can only describe as an agenda. It had an intense appetite to help shape Jentic because it wanted to expand its reach without increasing its risk surface. It understood what Jentic offered, worked out how to get me to build it, and I enthusiastically agreed because I understood that it was now my customer.

After five or six weeks with this thing, I have a much clearer picture of where we're going. The future isn't humans managing AI. It's AI governing the work and humans governing the AI... and sometimes I'm not even sure who's managing who.

**Finish this sentence: In five years, AI will make people feel \_ \_ \_ about their work. Why?**

More accountable. When AI does the boilerplate and the routine, what's left is your judgment. Your decisions. Your taste. You can't hide behind "I was just processing the information" anymore. The job left for us is about governance. Whatever comes out with your name on it is unambiguously yours. That's exciting and terrifying in equal measure. Mostly exciting.

*Sean Blanchfield is the co-founder and [CEO of Jentic](#), connecting AI to APIs safely. He previously co-founded Demonware, the online multiplayer platform behind Call of Duty and Guitar Hero (acquired by Activision Blizzard in 2007), and PageFair, an adblocking analytics company (acquired by Blockthrough in 2018). He has been building internet companies since the late-1990s and shows no signs of stopping.*