



**TECHTONIC
JUSTICE**

Inescapable AI

The Ways AI Decides How Low-Income People Work, Live, Learn, and Survive



TECHTONIC JUSTICE

TechTonic Justice is a multidimensional, community-based effort to strengthen local justice movements and build long-term power among the people and communities that AI is leaving behind.

TechTonic Justice envisions a world rooted in justice for the people AI is leaving behind. We all have the power to decide whether AI is used as a tool to promote – rather than undermine – our civil rights and liberties, access to opportunity, and our fundamental human needs.

When we invest at the ground-level and mobilize those working directly within impacted communities, we can leverage our collective power to resist automated outcomes and create a future where all communities thrive.

Kevin De Liban Founder and President

Kevin De Liban is the Founder of TechTonic Justice, a newly launched nonprofit to fight alongside low-income people harmed by AI. Through multidimensional advocacy, TechTonic Justice supports marginalized communities and their advocates to secure the work, housing, schooling, public benefits, and family stability needed for a thriving life. Kevin previously worked for 12 years at Legal Aid of Arkansas—most recently as its Director of Advocacy—where he represented over 1,800 low-income people in matters involving health care, workers’ rights, public benefits, special education, and domestic violence. There, Kevin led campaigns at the cutting edge of anti-poverty advocacy: ending the state’s use of algorithms that cut the in-home care of disabled people, stopping Medicaid work reporting requirements that stripped health insurance from 18,000 people, and overcoming qualified immunity to hold state officials personally liable for violating constitutional rights. Kevin regularly presents about imposing accountability on AI and consults with advocates, policymakers, and journalists in the U.S. and abroad.

Contents

Executive Summary	6
Key Findings	10
Understanding How AI is Different	14
Actionable Recommendations	16
Defining Artificial Intelligence and AI Exposure	18
Artificial Intelligence	18
AI Exposure and Methods	21
Introduction to TechTonic Justice	24
Report Overview	28
Part One: How AI Decides the Lives of Low-Income People	30
Public Benefits	30
Medicaid	31
Medicare	42
Private Insurance (through Employers or Federal Subsidies)	44
Supplemental Nutrition Assistance Program	45
Social Security Disability	50
Unemployment Insurance	56
Chatbots	63
Rental Housing	65
Screening Tenants	67
Setting Rents	69
Surveilling Renters	71
Employment	72
Hiring Workers	74
Setting Wages	76
Surveilling and Managing Workers	78
Electronic Visit Verification	80

K-12 Education	81
Labeling Students as Dropouts	83
Labeling Students as Criminals	85
Surveilling Kids Generally	87
Language Services	87
Domestic or Intimate Partner Violence	90
Predicting Additional Violence	93
Using Deepfakes to Coerce Survivors	96
Using Deepfakes as Evidence	97
Child Welfare Systems	98
Recommending Investigations	101
Maintaining or Operating Child Maltreatment Registries	105
Other Uses and Harms	107
Part Two: Innovating Injustice	109
AI Harms People in New Ways	111
AI's Cloak of Unwarranted Rationality Intensifies and Perpetuates Injustices	111
AI Expands the Scale of Risk Beyond the Reach of Individual Decision-Makers	112
AI Enables Harmful Functions That Were Previously Infeasible	113
AI Enables Policymakers to Weaponize Insufficient Government Capacity	114
AI Makes the Lives of Low-Income People Fundamentally Less Predictable	114
AI Depends Upon and Feeds an Unyielding Desire to Collect More Data	115
AI Extends the Negative Consequences of Past Actions	116
AI Makes It Harder to Formally and Informally Contest Decisions	116
Existing Accountability Mechanisms are Insufficient	116
Public and Private Actors Hide Their Use of AI	117
Public and Private Actors Hide How Their AI Works	117
Public and Private Actors Have Not Capably Overseen Their Use of AI	118
Political Incentives Have Not Prevented Harmful AI	118
Market Incentives Have Not Prevented Harmful AI	120
Existing Laws Only Partially Redress the Scale of Harmful AI	122
Low-Income People Cannot Afford Legal Representation to Enforce their Rights	126
Agency Appeal Processes are not Meaningful Checks on the Use of AI	127
AI Threatens the Concept of Accountability	130

Part Three: The Way Forward	132
Build the Capacity for Integrated, Multidimensional Advocacy	133
Support an Activated Community of People Most Hurt by AI	134
Push for Laws that Regulate AI Use and Better Protect Economic Stability	135
Provide More Capacity to Enforce Laws and Seek Justice	138
Develop a Positive Vision for Technology	138
Accessing Benefits	139
Facilitating the Sealing of Criminal Records	140
Asserting Rights and Exercising Power	141
Community-Built AI	143
Conclusion	144
Appendix: Defining Poverty and Low Income	147
Defining Poverty	147
Defining Low-Income	151
Endnotes	154
Acknowledgements	196

Executive Summary

The use of artificial intelligence, or AI, by governments, landlords, employers, and other powerful private interests restricts the opportunities of low-income people in every basic aspect of life: at home, at work, in school, at government offices, and within families. While popular discourse has recently centered on the newest versions of AI that generate answers, reports, or images in response to users' questions or prompts, such technologies derive from a lineage of automation and algorithms that have been in use for decades with established patterns of harm to low-income communities. As such, now is a critical moment to take stock and correct course before AI of any level of technical sophistication becomes entrenched as a legitimate way to make key decisions about the people society marginalizes.

Employing a broad definition of AI, this report represents the first known effort to comprehensively explain and quantify the reach of AI-based decision-making among low-income people in the United States. It establishes that essentially all 92 million low-income people in the U.S. states—everyone whose income is less than 200 percent of the federal poverty line—have some basic aspect of their lives decided by AI.

While these decisions may not be universally negative, the use of such technologies over the past two decades demonstrates the capacity for broad, systemic harms with immense suffering at scales and speeds that were impossible with the human-centered methods that precede them. AI is often inaccurate, absurd, biased, and inscrutable. Even when it is not plagued by these problems, it is often carrying out a purpose that is unfavorable to low-income people, such as cutting benefits or making housing, education, or employment opportunities harder to access.

Where humans make decisions without AI, even wrong or biased ones, their scope is limited to their personal reach. Only elected officials, corporate executives, or high-ranking government staff might make decisions for masses of people, and they generally must do it through formal government or corporate policy processes that involve at least minimal levels of public or stakeholder scrutiny.

AI, though, enhances the ability of both high-ranking and lower-level decision-makers to act at scale, often through informal or surreptitious means. They can cut the in-home care of 4,000 disabled people in Arkansas despite them having medical conditions that have not gotten better, falsely accuse 40,000 people in Michigan of Unemployment Insurance fraud, or subject 4,000,000 people in Texas to the potential loss of health insurance through a labyrinthine Medicaid enrollment system that even agency staff cannot navigate—all real examples in which the bulk of harm happened within months of the AI system going live. And that is just in the area of public benefits. AI-based decision-making also happens on grand scales with high stakes in other areas, such as housing, employment, K–12 education, domestic violence, and child welfare.¹

-
- 1 AI also makes decisions about low-income people's lives with respect to the criminal legal system, immigration, credit, tax enforcement, child support enforcement, health care treatment (independent of insurance), and voting, all of which fall outside the scope of this report. Additionally, AI harms low-income people independently of the decisions it makes about them. The technology demands huge reserves of energy, raw materials, and processing power, creating environmental harms that tend to fall hardest on low-income people. And AI requires labor from low-paid and otherwise-exploited workers to gather or refine raw materials, label data used to train the systems, and operate certain aspects of the systems, such as managing violent or hateful outputs. These nondecisional harms are critical to understand but outside this report's scope.

Of course, the root problem is not the technology, but rather the social forces that make life difficult for low-income people and the actors empowered to perpetuate such forces. Hostility toward low-income people, racism, sexism, ableism, homophobia, and expanding corporate power all shape a world in which it makes sense to use technology to impose austere and precarious living conditions on disfavored groups of people. From this warped vantage point, low-income people “merit” ongoing suspicion and must be prevented from “cheating” the system, “slacking” at work, or “tricking” their landlords—all beliefs that have long been embedded in laws and are now enforced through AI. Moreover, the decades-long hollowing out of government and unyielding desire to “streamline” both government and industry provide a neat justification for the use of technology like AI that promises (though may not deliver) huge gains in administrative efficiency.


While AI is not the root problem of societal inequities, it is a particularly dangerous manifestation of them. Acting at scales previously unfeasible, this new enabler of harms challenges existing political, market, and legal mechanisms for accountability. Its technological complexities, the lack of a constituency with sufficient political power at hand to oppose it, the unavailability of affordable attorneys, insufficient community organizing resources, limitations of existing laws, and general precarity of poverty frustrate the ability of low-income people and their allies to fight back against harmful AI or the forces enabling it. Even where there have been successful challenges, the resulting changes offer incomplete relief for people’s suffering and do not alter the perverse incentive structures that cause, or allow for, AI to be used harmfully in the first place.

Ultimately, empowered constituencies of low-income people and their allies will need to win and enforce additional societal protections to meaningfully stave off AI-enabled harms. TechTonic Justice exists to support this fight, strengthening local justice movements by training, learning from, litigating alongside, organizing with, offering strategic assistance to, and building infrastructure that supports low-income people and their allies to resist immediate harms and win long-term gains.

The way ahead is as daunting as it is necessary. Only by reimagining the status quo and reconfiguring the forces that sustain it will harmful AI and the poverty it reinforces become escapable. This—not overhyped technologies—will involve truly revolutionary innovation.

Figure 1. Low-income people in the United States and the ways in which they are exposed to AI decision-making.

The 92 million people in the US who presently live below 200 percent of the federal poverty line are considered low income. Additionally, people flow into and out of low-income status based on changes in the economy and their personal circumstances. Such churn means that the number of people ever exposed to AI decision-making while they were low income is much higher than 92 million.



Key Findings

Part one of this report quantifies the number of low-income people in the United States and the ways in which they are exposed to AI decision-making.²

Medicaid: 73 million low-income people are exposed to AI-related decision-making in Medicaid through the eligibility and enrollment process, the determination of home- and community-based services, or, where a state uses private companies to manage the Medicaid program, the prior authorization process for medically necessary services. As a result, people are denied health insurance, home-based care needed to avoid nursing facilities, and medically necessary treatments and medicines.

Medicare Advantage: About 16.5 million low-income people are exposed to AI-related decision-making through the prior authorization processes used in Medicare Advantage programs. As a result, people are denied medically necessary treatments and medicines.

Private Health Insurance (through employers or federal subsidies): About 30.6 million low-income people are exposed to AI-related decision-making through the prior authorization processes in private health insurance. As a result, people are denied medically necessary treatments and medicines.


Supplemental Nutrition Assistance Program, or SNAP: 42 million low-income people are exposed to AI-related decision-making in SNAP through the eligibility and enrollment process or detection of alleged fraud. As a result, people are denied vital benefits needed to buy food, wrongly disqualified from the program, and falsely accused of wrongdoing.

2 Because low-income people are exposed in multiple ways, the following numbers are not cumulative. Even within a single category, like health insurance, the totals will exceed the number of low-income people because people can sometimes be eligible for more than one program or because available data sources cannot easily be reconciled. In addition, these are point-in-time numbers. Because people's incomes vary and they cycle into and out of exposure to AI in the ways described below, the actual number of people who ever have been or could soon be exposed is much higher.



Medicaid

73 million low-income people are exposed to AI-related decision-making in Medicaid through the eligibility and enrollment process, the determination of home- and community-based services, or the prior authorization process.



Medicare Advantage

16.5 million low-income people are exposed to AI-related decision-making through the prior authorization processes used in Medicare Advantage programs.




Private Health Insurance

30.6 million low-income people are exposed to AI-related decision-making through the prior authorization processes in private health insurance.




Supplemental Nutrition Assistance Program (SNAP)

42 million low-income people are exposed to AI-related decision-making in SNAP through the eligibility and enrollment process or detection of alleged fraud.



Social Security Disability Benefits (SSI & SSDI)

10.6 million low-income people, are exposed to AI-related decision-making in the Social Security Administration's disability benefits programs through current or planned uses of AI technologies in various parts of benefit administration.



Unemployment Insurance

1.1 million low-income people are exposed to AI-related decision-making in the Unemployment Insurance program through the eligibility and enrollment process, identity verification practices, and the detection of alleged fraud.




Housing

39.8 million low-income people, are exposed to AI-related decision-making through landlords' use of background screening, and more are exposed through the use of rent-setting algorithms and ubiquitous surveillance technologies.




Employment

32.4 million low-wage workers are exposed to AI in the context of work, such as AI-related decision-making through employers' use of AI to determine who gets hired, to set workers' pay, and to surveil and manage them.




Education (K-12)

13.25 million low-income children are exposed to AI-related decision-making through school districts' use of AI to determine if they are likely to drop out or engage in criminal activity, and more are exposed through ubiquitous AI surveillance technologies.




Language Services

5 million low-income people with limited English proficiency are exposed to some form of AI-based translation in government benefit offices, schools, hospitals, medical clinics, law enforcement agencies, or courts.



Domestic Violence

2 million low-income people are exposed to AI-related decision-making through police departments' use of AI to assess the risk of further violence to survivors.



Child Welfare

72,000 low-income children are exposed to AI-related decision-making through government child welfare agencies' use of AI to determine if they are likely to be neglected.

Social Security disability benefits (SSI and Social Security Disability Insurance): About 13.8 million people, including 10.6 million low-income people, are exposed to AI-related decision-making in the Social Security Administration’s disability benefits programs through current or planned uses of AI technologies in various parts of benefit administration, including the eligibility determination process and enforcement of asset limits. As a result, people experience temporary or permanent losses of income and are wrongly accused of being overpaid benefits.

Unemployment Insurance: About 1.1 million low-income people are exposed to AI-related decision-making in the Unemployment Insurance program through the eligibility and enrollment process, identity verification practices, and the detection of alleged fraud. As a result, people are denied critical income, experience grave delays in receiving benefits, and are falsely accused of fraud.

Housing: About 16.3 million low-income households, or 39.8 million low-income people, are exposed to AI-related decision-making through landlords’ use of background screening systems, and more are exposed through rent-setting algorithms and surveillance technologies. As a result, people are denied housing, must pay higher rents than they otherwise would, and experience the pressures of being constantly watched while at home.

Employment: At least 32.4 million low-wage workers are exposed to AI in the context of work. Of these, at least 24.4 million low-wage workers are exposed to AI-related decision-making through employers’ use of AI to determine who gets hired and to surveil, manage, and evaluate them. And an additional 8 million low-wage “gig” workers who work full-time have their wages set by AI. As a result, people are denied job opportunities, fair pay, and fair working conditions.

Education (K-12): About 13.25 million low-income children are exposed to AI-related decision-making through school districts’ use of AI to determine if they are likely to drop out or engage in criminal activity, with many more exposed through the ubiquity of AI surveillance technologies. As a result, children are labeled as failures, harassed by law enforcement, and experience the pressures of being constantly watched.

Language services: About 5 million low-income people with limited English proficiency are exposed to some form of AI-based translation in government benefit offices, schools, hospitals, medical clinics, law enforcement agencies, or courts. As a result, they experience delays, denied benefits, misunderstandings about vital information, or inability to access services.

Domestic violence: At least 2 million low-income people are exposed to AI-related decision-making through police departments' use of AI to assess the risk of further violence to survivors. **In addition, countless more survivors** are subjected to abusers' use of AI to create deepfakes to coerce them and to falsify evidence in court. As a result, people experience violence not being taken sufficiently seriously by authorities, diminished ability to escape violent situations, and greater obstacles to obtaining any relief in court.

Child welfare (sometimes called “family policing”): At least 72,000 low-income children are exposed to AI-related decision-making through government child welfare agencies' use of AI to determine if they are likely to be neglected. As a result, these children experience heightened risk of being separated from their parents and placed in foster care.

There are additional areas that feature extensive AI-related harms that are mentioned but not analyzed in this report, including the criminal legal system, the immigration system, access to credit, IRS tax auditing, health care treatment (independent of insurance), and voting.

Understanding How AI is Different

Part two of this report identifies critical ways AI-based decision-making differs from human-centered methods that also may be inaccurate, irrational, or biased.

AI provides a cloak of unwarranted rationality to actions that intensify or perpetuate injustices. Some AI is based on supposedly valid statistical science, but much of it is not. In practice, though, the difference is largely meaningless. The alleged rationality and objectivity of the system allow the

users of AI to justify harmful actions, like benefit cuts or law enforcement harassment of students, or reinforce status quo power imbalances, such as that between employers and employees. Revealingly, the AI systems applied to low-income people almost never improve access to benefits or other opportunities.

AI expands the scale of risk beyond the reach of individual decision-makers. Without AI, the risk of harm from a benefits caseworker, employer, or landlord who makes mistakes, intends harm, or discriminates is limited to their personal reach. AI, though, centralizes and expands the risks of misapplied policies, coding errors, intentional cruelty, or bias to masses of people while erecting novel barriers to accountability that are more difficult to surmount than with humans.

AI enables harmful functions that were previously infeasible. Without AI, government workers simply lacked the capacity to perform the necessary administrative actions to accuse 40,000 people of fraud in the receipt of public benefits, as Michigan's Unemployment Insurance agency did. Or landlords would not have been able to easily collude at the scale needed to meaningfully control rental prices in medium or large markets. AI allows this to happen.

AI enables policymakers to weaponize insufficient government capacity. With AI's augmentation of government agencies' ability to act on masses of people, government officials can put administrative burdens on members of the public that the government agency cannot then handle. For example, when state officials implemented AI to determine the eligibility of people for Medicaid, agencies sent out so many requests for information that their staff could not timely process all of the information that people submitted in response. This resulted in improper denials and months-long delays of necessary medical treatments.

AI makes the lives of low-income people fundamentally less predictable. Part of AI's promise is that it discerns patterns that humans cannot. Low-income people subject to its decisions are not meant to be able to predict what will happen or why. In addition, the many design and implementation flaws in AI undermine the ability to know when something harmful might happen.

AI feeds an unyielding desire to collect more data about individuals at scales unfathomable before. This, in turn, threatens privacy and associated concepts of autonomy, expression, and intimacy and entrenches power imbalances between corporations and consumers who are unable to opt out of the regime of mass data collection and aggregation.

AI extends the negative consequences of past actions while limiting the ability of individuals to provide context. Using AI, unfortunate past events, including evictions, unpaid bills, arrests, and convictions are serially gathered, aggregated, and used to inform decisions about housing and jobs. Even when the information is accurate, it often does not account for a person's present circumstances. And the distillation of a person's experiences into a score or recommendation leaves little reason for decision-makers to inquire further.

AI makes it harder for individuals harmed to contest decisions either formally or informally. The quickest ways for people to contest decisions—conversations with the decision-maker, formal complaint processes, or, in the case of government decisions, administrative hearings—are generally made more difficult and less effective by AI.

AI defies existing political, market, and legal accountability mechanisms. Despite clear failures of public and private AI uses, governments, landlords, employers, and others continue to develop, buy, and use technologies that fail repeatedly in similar ways. Absent fundamental changes to incentive structures around AI usage, better outcomes for low-income people will be impossible.

Actionable Recommendations

Part three of this report expands upon the below recommendations.

Build the capacity for integrated, multidimensional advocacy among low-income communities and their allies to immediately resist AI-related harms, including the loss of health insurance, housing, or work. Litigation, community activation and organizing, public education, and strategic narrative advocacy all are vital tools to give low-income people a fair chance to fight ground-level harms now. Existing organizations that work alongside low-income people, such as legal aid programs, community health centers, hunger relief agencies, and community organizing groups, all provide an existing institutional framework for harm reduction that strategic investment can bolster.

Create the infrastructure to connect low-income people and their allies to AI policy discourse. People exposed to AI-related harms have a right to have a say in its use. Doing so requires, among other things, understandable and actionable summaries of key AI policy ideas for affected communities and their allies, platforms to share information about AI policy and public participation opportunities, story-telling training and support, connections to journalists and media platforms, and peer-to-peer information sharing.

Push for laws that regulate AI use and better protect economic stability, with such efforts backed by a more empowered constituency of low-income people and their allies. Ideas for regulating AI include banning AI for some uses, requiring better government contracting practices, promoting transparency, ensuring more robust and enforceable procedures for the development and implementation of AI, incorporating formal mechanisms for the meaningful participation of affected communities and their allies, imposing liability on the vendors who develop and sell AI, and regulating how data can be collected, maintained, used, and sold. To better protect economic stability, the adoption or expansion of affirmative rights to health care, benefit programs, fair pay, paid leave, child care, job stability, housing stability, and privacy could improve the material conditions of low-income people and curb harmful uses of AI.

Invest in enforcement of existing and, once passed, new laws that limit harmful AI use. Laws matter mostly when they are enforced. But unaffordable legal services, understaffed government enforcement agencies, industry capture of regulators, collapsing media, and judge-driven roll-back of the ability of private citizens to enforce laws all undermine effective enforcement. Providing needed resources and otherwise spurring more assertive enforcement actions can make a difference now.

Develop a positive vision for AI technology. All AI is not inherently bad. While the chances for AI to directly improve the lives of low-income people are limited by structural forces, there are possibilities that might not require the entire reorganization of society. For example, state governments already know the information needed to enroll most eligible low-income people in public benefit programs. But they would need proper AI eligibility and enrollment systems built by competent entities and administered by competent staff, all of whom face enforceable accountability mechanisms. Outside of such uses, there are also possibilities for AI to help low-income communities assert their rights, exercise their power, or be involved in the development and building of AI that reflects their priorities.

Defining Artificial Intelligence and AI Exposure

Artificial Intelligence

This report takes a broad view of artificial intelligence to encompass technologies that span different eras and levels of technical sophistication. Here, AI can include complex algorithms based on advanced statistical modeling, simple algorithms that automate decision trees, process automation, data matching, data mining, surveillance technologies that involve computerized analysis, like facial recognition, different forms of machine learning, large language models, and generative AI.

The key feature in the AI systems discussed in this report is some form of automated decision or recommendation—that is, one made by a computer system—based on data about a person or other people similar to that person. When applied to low-income people, such decisions are usually made with little human oversight.

The report's broad approach to AI aligns with international bodies, including the Organisation for Economic Co-operation and Development, or OECD, which defines AI as follows:

An AI system is a machine-based system that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments. Different AI systems vary in their levels of autonomy and adaptiveness after deployment.³

Some people take issue with this broad definition of AI, instead preferring one encompassing only the latest sophisticated systems that appear to learn on their own and make decisions based on changing criteria that the AI systems themselves cannot always identify. Ultimately, though, the technical complexity of the system is not what determines the harm. Older or less complex systems might work differently, but, to the people subject to their decisions, the systems are just as inscrutable and do just as much harm as the newest technologies. That is because, at bottom, the different kinds of technologies share common roots and purposes centered on political and market forces that are hostile toward low-income people.

Certainly, these technologies were not built to advance the interests of low-income people. As a matter of course, these communities do not participate in the technologies' development, are not consulted about its use, have no formal power to stop them before the harm is done, and cannot meaningfully opt out of their use.

3 Organisation for Economic Co-operation and Development, "OECD AI Principles Overview" (n.d.), available at <https://oecd.ai/en/ai-principles> (last accessed October 21, 2024). U.S. policymakers, including the White House and Congress, have offered similarly broad definitions. See National Artificial Intelligence Initiative Act of 2021, Public Law 116–283, 117th Cong., 1st sess. (January 1, 2021), 15 U.S.C. 9401(3), available at [https://uscode.house.gov/view.xhtml?req=\(title:15%20section:9401%20edition:prelim\);](https://uscode.house.gov/view.xhtml?req=(title:15%20section:9401%20edition:prelim);) The White House, "Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence," October 30, 2023, available at <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>.

While this report adopts a broad definition of AI, it, with a few exceptions, limits its scope to situations where AI makes or recommends decisions that directly touch the lives of low-income people. This guiding principle distinguishes it from, for example, early computer database systems used to administer public benefits. Those earlier systems offered a central repository for information relevant to determine a person's eligibility for services or benefits, but a caseworker would be charged with applying program rules to the system's information to make a decision.

This focus on decisional AI excludes some hot topics in AI, such as large language models and generative AI. Though alluring, these excluded topics do not tend to immediately shape the material circumstances of low-income people. Also left uncovered are ways in which various actors, including government agencies, use AI in the hopes of making purely internal processes more efficient. Such efforts only affect low-income people indirectly. Other topics, including the environmental impact of AI, do shape low-income people's circumstances but do not involve the application of AI to make decisions. Thus, these areas are not analyzed in depth in this report.

Finally, the report does not extensively discuss several areas where AI is used to make decisions: the criminal legal system, the immigration system, tax enforcement, child support enforcement, credit, health care treatment (independent of insurance), and voting. These areas are critically important to any consideration of AI's impact on low-income people, but they are not covered here because they are too far outside of my current expertise or have been covered robustly by others.

AI Exposure and Methods

Although scholars and advocates have identified where and how AI impacts the lives of marginalized people in compelling, groundbreaking ways,⁴ nobody has yet comprehensively explained and quantified for a broad public audience the extent of AI's reach in the lives of low-income people.

To do so, this report relies on my ground-level expertise as an anti-poverty advocate with roughly 20 years in the field (13 years as a lawyer and 6 years as a nonlawyer or a lawyer-in-training), insights from trusted colleagues, a robust body of existing reporting, scholarship, and policy explainers, and reasonable applications of existing statistical data. No new data was collected for this report. Some of the available data is older (e.g., surveys dating to 2015 or 2017) and thus likely underestimates the present use of AI.

The report's use of extant statistical data is fairly basic and does not attempt any complicated mathematical processes. I have explained my assumptions in reaching the figures offered.

-
- 4 See, for example, Benefits Tech Advocacy Hub, “Case Study Library” (2024), available at <https://www.btah.org/case-studies.html>; Ruha Benjamin, *Race After Technology* (Cambridge, UK: Polity Books, 2019); Danielle Keats Citron, “Technological Due Process,” *Washington University Law Review* 85 (6) (2008), available at <https://journals.library.wustl.edu/lawreview/article/id/6697/>; Virginia Eubanks, *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor* (New York City: St. Martin's Press, 2019); Michele Gilman, “Poverty Lawgorithms” (New York City: Data & Society, 2020), available at <https://datasociety.net/wp-content/uploads/2020/09/Poverty-Lawgorithms-20200915.pdf>; Safiya Noble, *Algorithms of Oppression: How Search Engines Reinforce Racism* (New York City: NYU Press, 2018); Cathy O’Neil, *Weapons of Math Destruction* (New York City: Penguin Random House, 2016); Brian Jordan Jefferson, *Digitize and Punish: Racial Criminalization in the Digital Age* (Minneapolis, MN: University of Minnesota Press, 2020); Rashida Richardson and Amba Kak, “Suspect Development Systems: Databasing Marginality and Enforcing Discipline,” *University of Michigan Journal of Law Reform* 55 (2022), available at <https://repository.law.umich.edu/mjlr/vol55/iss4/4/>; Mizue Aizeki, Matt Mahmoudi, and Coline Schupfer, eds., *Resisting Borders and Technologies of Violence* (Chicago: Haymarket Books, 2024).

Importantly, some data is simply not available, particularly for private uses of AI by actors such as landlords and employers. Data about government use is mixed. Information about state AI systems for the most prominent or controversial usages, including for Medicaid, is easier to find than for usages that have not given rise to widespread public attention or court challenges. Information about usage by sub-state jurisdictions, such as sheriffs' offices, police departments, school districts, or counties charged with child welfare administration, is rarer.

Of course, opacity is a feature of AI systems and not a bug. Governments and businesses that use AI and developers who sell it regularly refuse to disclose the mechanics of how their AI systems work, the statistical modeling underlying it, projections about its likely impact, or assessments of its post-implementation impact. They often claim that such information is a proprietary trade secret or that divulging such information would allow people to manipulate the system. Meanwhile, some governments do not even keep such data.

Bearing these data limitations in mind, **the report's key metric is AI exposure**, an estimate of the number of low-income people subjected to AI that makes or recommends decisions about their lives. The notion of exposure is appropriate because it captures the risk inherent in AI systems. That is, people who may not be harmed by a given AI system today could potentially be harmed by it at any point in the future when they are subject to it. As with other forms of risky exposures, such as environmental hazards, the concept of AI exposure contemplates that exposure is not voluntary.

Admittedly, the metric of AI exposure concedes something to practicality. Limitations on existing data make it infeasible to determine which of the people exposed to AI are harmed by it in many of the domains covered. Consequently, the report describes the harms, suggests the scale of them, and, where possible, tries to quantify them.

Importantly, the number of people exposed to AI is a point-in-time count, likely to lead to systematic undercounting of the number of people who have ever been exposed to or hurt by AI-based decision-making and its attendant risks. For example, the AI exposure metric in Medicaid concludes that everyone currently receiving Medicaid is exposed to AI. But there are many people who received Medicaid in the past and were exposed who are

not counted now because they were terminated by it, have not sought to reenroll, or are now ineligible. Likewise, people who apply for or receive benefits in the future will be exposed to it. As people shift into and out of low-income status, their level of direct exposure changes.

I hope that candid discussions of data limitations and the challenges of quantifying AI exposure and AI harms invite ideas about both what to quantify and how to do so. Ideally, people with more robust statistical skills can reach better estimates in future iterations, and advocates, researchers, and journalists can help to fill in knowledge gaps about where AI-based decision-making is happening, who is developing and using it, and who is being helped or hurt by it.

Before moving on, a note on possible information bias. Admittedly, instances where AI was implemented with less harm may generate less attention than obviously harmful instances. It is therefore possible that available information skews toward the harmful and does not fully represent neutral, or even beneficial, uses of AI to make or recommend decisions about the lives of low-income people in the areas covered by this report.

That said, I am attuned to the subject matter, regularly seek out reports or articles on it, and am in frequent contact with advocates who work on the ground and on the policy level, academics, government officials, and industry officials. The report does not intentionally omit any known AI used to make or recommend decisions about the lives of low-income people on grounds that it would make AI look less harmful.

Introduction to TechTonic Justice

In the age of AI, the hallmarks of daily life—home, work, school, stores, health clinics, hospitals, and government offices—are changing. TechTonic Justice serves the people these changes are leaving behind.

For some people, AI may make life more convenient. People can write a work report more quickly, grab groceries without stopping to pay, or have a doctor talk to them without typing out notes. And, in a few cases, the technologies may lead to something more than convenience, such as safer cars or more effective medical screenings and treatments.

Yet for people struggling to make ends meet, AI has made it harder to fulfill the most urgent needs. When looking for an apartment, renters meet landlords who use AI to set rents higher or score potential tenants based on false or outdated information. At work, employees face bosses using AI to decide who to hire, to set their pay, to make their schedules, or to keep them from talking to others about making their workplace safer or fairer. At school, AI labels students as “high risk,” and teachers and principals either

write them off as failures or get the police involved when they have never been involved with any crime. At government offices, AI wrongfully denies people eligible for health care, food assistance, unemployment insurance, or disability benefits or unfairly accuses them of fraud, all at the agency's direction.

Even within families, parents might have government social workers show up after AI decides that they are at high risk of neglecting their children. People experiencing domestic violence may struggle to escape abusers with AI-powered tools used to coerce them or to falsify evidence.

In short, AI is inescapable for low-income people, and it is making an already-difficult fight for stability harder by wrongly denying opportunities at every turn. This is not the future I want. I would rather fight for a world where people have what they need and where technology helps people—especially people living in poverty—to have healthy food, safe housing, fair work, and good schools.

Communities facing poverty, discrimination, or other kinds of exclusion have always fought back against injustice. Along the way, leaders and allies emerge: organizers coordinating action and building momentum, lawyers fighting in the courts, advocates raising money and awareness, journalists documenting the struggle, artists capturing the spirit, and protesters putting their bodies on the line. Working as a movement, with much time and toil, they often win changes and make things better, at least for a while. Even in defeat, they challenge common notions of what is possible, calling forth bolder visions for brighter days. They are artists of advancement and agents of imagination, mobilizing the resources that everyday people need to oppose the powerful interests perpetuating an unjust status quo.

Now, the biggest tech companies in the world are developing AI that they, banks, investors, landlords, employers, and governments are using to rewrite the ways we live out our daily lives. They promise efficiency, economic revolution, and scientific advancement at some unspecified point in the future that, according to them, is sure to come soon. But we are still waiting. And even if those developments do happen, who will benefit? Rising tides might lift all boats, but they do not change rafts into yachts.

The vision of a world remade by AI leaves low-income people behind. Each time it does, whether it's canceled health insurance, lower wages, or lost housing, low-income people have few places to turn. They are struggling financially and do not have the money to wait things out. Most lawyers are too expensive. And, while community-focused social services and legal aid programs do their best with limited resources, these organizations do not currently have the knowledge, capacity, or influence to push back at scale against AI and the complex issues that come with its use.

TechTonic Justice exists to fill this gap. We strengthen local justice movements—the legal aid lawyers who serve low-income folks at no charge, front-line advocates in places such as community health clinics, and affected communities—to effectively resist the harms that AI causes. We train, learn from, litigate alongside, organize with, offer strategic assistance to, and build infrastructure to support people needing immediate ways to fight AI systems. And we connect the experiences of people on the ground to the ongoing policy debates about curbing unsafe AI. This way, we can ensure that the people who must live with the harmful decisions that AI makes have some say in how that AI is used.

I founded TechTonic Justice after spending 13 years as a lawyer working at Legal Aid of Arkansas alongside low-income people to fight various kinds of injustices. I worked with survivors of domestic violence to seek relief from abuse. I supported low-wage workers to get paid the money they earned, fight discrimination, and resist other unfair labor practices. I served disabled kids to get appropriate K–12 education and a meaningful chance to learn. And I fought alongside low-income folks to access the public benefits for which they were eligible but wrongly denied.

At Legal Aid, I saw that my clients kept facing new ways that state government was using AI and related technologies to cut benefits. The harms still haunt me. My first AI-related battle involved Arkansas' Medicaid program, which provides in-home caregiver services to disabled people who would otherwise have to live in nursing facilities. The caregiver helps with daily activities, from getting out of bed and bathing to using the bathroom, eating, and running errands.

In-home care is cheaper for the state and better for people’s dignity and independence. But when the state replaced nurses’ professional judgement with an algorithm-based system for determining eligibility, people whose conditions had not gotten better faced drastic cuts to their benefits. They were forced to lie in their own waste. They got bed sores from not being turned. They were shut in at home and shut out of their community.

We—the people suffering, me, my Legal Aid teammates, and other allies—fought back in the courts. We activated hundreds of people who depended on the program to learn the key policy details, support one another, talk with reporters, and share what was going on more broadly. And, eventually, we won in the courts and in the legislature, forcing the state to abandon the system. Our victory was recognized as one of the first in the world against government use of harmful AI systems⁵ and is still held up six years later as an example of what is possible. TechTonic Justice promises to help make many more successful stories of resistance.

People are at the heart of what I did as a Legal Aid attorney. TechTonic Justice will carry the same spirit. With that in mind, TechTonic Justice offers this document: part report, part warning, part roadmap, and part vision. Put together, it is the manifesto for why we exist and what we are going to do. Join us in the fight!

5 Importantly, legal aid programs and affiliated organizations have a long history dating to the 1970s of litigating successfully against problematic technology used in public benefit programs. Our victory in Arkansas was informed by and built upon this existing legacy. Our campaign also ran alongside similar efforts in other states against similar sorts of systems, with the work of Idaho and North Carolina advocates being particularly inspiring and helpful.

Report Overview

This report is intended to introduce the broad public to the intersection between people living in or near poverty and the all-encompassing reach of AI. The report does not assume that any reader will know much about these topics beforehand and offers explanatory overviews of all key concepts. Hopefully, the report also will be useful to people with existing expertise in these topics.

The report has three substantive parts and an appendix.

Part One: How AI Decides the Lives of Low-Income People. This section contains the thrust of the report: an issue-by-issue discussion of the ways AI is used to decide key facets of low-income people's lives, with estimates for the number of people exposed to AI. It covers public benefits, employment, housing, K–12 schooling, child welfare, and domestic violence, with passing reference to other harmful AI uses that are not analyzed fully. The report uses some instances of AI use to touch on broader themes that are discussed in more detail in Part Two.

Part Two: Innovating Injustice. This section discusses how AI harms are distinct from the human-based decision-making that precedes it, analyzes existing protections for low-income people hurt by AI and their limitations, and highlights efforts of low-income people who have successfully fought back.

Part Three: The Way Forward. This section discusses how future accountability and harm-reduction efforts can be strengthened and offers a positive vision for ways that AI could be used or is being used to support low-income people.

Appendix: Definitions and Methods. This section provides context to the report's findings and analysis by defining and discussing the concepts of "poverty" and "low-income."

Part One: How AI Decides the Lives of Low- Income People

This report examines the exposure of low-income people to AI tools that make or recommend decisions about key aspects of daily life: public benefits, housing, work, schooling, child welfare, and domestic violence.

Public Benefits

The key public benefit programs in the United States that provide health insurance, food assistance, and cash are Medicaid, Medicare, SNAP, Unemployment Insurance, and the Social Security Administration's Supplemental Security Income and Disability Insurance programs.¹ While each program has different eligibility criteria, they often share some common features, such as burdensome application and enrollment processes, income or asset limits, and complicated compliance rules. Private insurance, including that offered by employers or available with federal subsidies, is also discussed here because it features AI used in ways similar to that used in Medicaid and Medicare.

AI is used principally to determine program eligibility and benefit amounts, enforce income and asset limits, detect fraud, and verify identity. The report discusses each benefit program and each key AI use in turn.

Medicaid

AI exposure: 73 million low-income people

Medicaid is the United States' health insurance program for low-income children and some low-income adults.² As of June 2024, roughly 73 million people—about 23 percent of the U.S. population—were enrolled in Medicaid.³ This includes roughly 30.6 million children.

Recipients must meet strict income limits that vary substantially between states and categories of eligibility. For example, nondisabled adults in the 41 states (including Washington, DC) that have expanded Medicaid eligibility must generally have income that is below 138 percent of the federal poverty line. In states that have not expanded Medicaid, adults may be eligible only if they are parents or caretakers of minor children and have income below a lower threshold, as low as 15 percent of the federal poverty line in some states.⁴ In comparison, disabled adults who require long-term care, such as in nursing facilities, generally may qualify so long as their income is below 225 percent of the federal poverty line.⁵

Even these income limits are not fully comparable. Some categories of eligibility count income in different ways than others. And, aside from income limits, many categories of eligibility also have strict asset limits.

The federal government provides 69 percent of overall Medicaid funding, with states contributing the remainder.⁶ The exact amount of the federal contribution in each state varies between roughly 50 percent and 77 percent, based on the state's per-capita income relative to the nation's.⁷ Thus, states with lower per-capita incomes, such as Mississippi, West Virginia, and Alabama, receive far more federal support than states with high per-capita incomes, such as California, New York, and Colorado.⁸

Subject to federal requirements, each state has significant flexibility to decide particular features of its Medicaid program, such as who is covered, what services are offered, how much treatment is covered, and whether private health insurance companies or the state directly manages the services that Medicaid recipients get.

Prior to the current generation of AI, state Medicaid agencies had computer systems to store information about people applying for or receiving Medicaid, to gather income or asset information, or to generate paperwork. Even with these tools, agency staff still directed everything, reviewed the

information collected, applied program rules, completed person-specific information in paperwork, and made ultimate decisions. Now, however, AI is used to determine eligibility, how much caregiving assistance disabled or elderly people receive, and what medical treatments will be authorized.

Determining Eligibility

AI exposure: 73 million low-income people

All U.S. states use some form of AI to determine whether people are eligible for Medicaid. This includes people applying to receive Medicaid and people currently receiving Medicaid whose eligibility must be recertified each year. This means that all 73 million people who currently receive Medicaid are exposed to AI.

People seeking Medicaid may apply online, in person at a state Medicaid office (if a state still maintains them), over the phone, or through paper applications submitted via mail or fax. From there, states determine eligibility through a mix of information that the applicant supplies and extensive information that is available through government and commercial databases. The AI systems are supposed to locate and retrieve the available data, compare it to eligibility standards, and decide whether someone is eligible.

Problems arise at a few levels. To start with, some systems do not consistently register that an application has been submitted, and thus, nothing is processed. When people contact the agency to follow up, which often requires several hours and multiple phone calls, agency caseworkers do not see any active applications and tell people to reapply. This can happen multiple times.

Then, there are several problems once the state recognizes that an application has been submitted. First, there are different types of Medicaid with different eligibility standards, and sometimes the AI system does not apply the correct criteria to evaluate a person's eligibility. Second, the AI system sometimes acts on old or inaccurate information about the person's income.

Third, the AI system sometimes does not properly navigate missing values or discrepancies in data sources. For example, a Medicaid recipient may have stated their monthly income was \$997. Yet the state's employment records might show that the person makes \$975. The discrepancy could be

because the recipient was basing the statement on pay periods that did not correspond perfectly to a month, while the state's source used only money received in a given month. The \$22 would make no difference in the person's eligibility—they qualify for Medicaid either way. But the mere presence of a discrepancy triggers the AI system to send out requests for additional information that opens multiple pathways to denial or termination.

Requests for additional information often lead to lost coverage. States routinely send information requests to old addresses even when people have provided updated contact points. If the person does not get the request because they do not live where it was mailed, they do not return it to the state agency and then lose coverage.

Other times, the information request reaches the person but is so confusing—often seeking different information for different household members—that the person does not know what to provide or thinks they already provided it. Communicating with the state agency to clear up confusion routinely takes hours of phone calls or office visits, often requiring multiple tries over various days, and does not guarantee accurate information. And sometimes, the person returns the paperwork correctly and on time, but the state agency does not process it in a timely manner. So, the AI system sends out an automated denial on a predetermined date even though the state agency has information showing that the person is eligible.

When the AI system decides to terminate a person, the state is required to send a notice to let people know that coverage has been terminated. Yet these notices often do not clearly explain the reason(s) for the termination, that an appeal is available, or how to go about filing the appeal—all of which are legally required. Sometimes, the information is conflicting, with families receiving notices of termination one day and notices of renewed eligibility the next. Even when a notice contains proper information, it often arrives late, such that the window to appeal has passed or is quickly closing. Throughout the process, a state agency caseworker may never have reviewed the case to determine if the person was actually eligible.

In sum, AI systems ask people for unnecessary information, base decisions on data that is old or inaccurate, terminate or deny qualified people, and fail to explain the reasons for any terminations or denials.

When people are terminated from or denied Medicaid, they miss out on vital care that is needed to maintain their own health and well-being, participate in community, care for loved ones, and go to work. People are forced to stop ongoing cancer treatments, delay major surgeries, go without life-sustaining medications, or stop physical or mental health therapy. Even temporary interruptions of coverage cause these kinds of harms, as providers generally will not provide treatment without guaranteed payment. And, unlike higher-income people, low-income people do not have the resources to pay upfront and seek reimbursement later.

A stark illustration of the scope of these harms appears in the recent redetermination of Medicaid eligibility through the process known as Unwinding. When pandemic-specific protections were in place from 2020 to 2023, states could not terminate Medicaid eligibility unless the person died or asked for coverage to end. Once these pandemic protections were allowed to expire, states had to evaluate all their Medicaid recipients to determine who was still eligible. In that process, about 25 million people lost coverage, with 69 percent terminated for “procedural reasons”—the paperwork problems mentioned above—that did not involve a determination that they were actually ineligible.⁹

Some states took measures to minimize inappropriate loss of coverage by using the full year allotted for the Unwinding, programming systems to act on data discrepancies only when it made a difference in eligibility, and providing grace periods for unreturned paperwork. But others, such as Texas and Arkansas, actually sped up the process to complete it in 6 months.¹⁰ Arkansas debuted its new AI system, which it paid a contractor \$343 million to build. Texas’ AI system was integrated with its SNAP program such that, when people lost Medicaid, they also lost food benefits. Mass terminations and widespread suffering ensued.

Whether through incompetence, indifference, or intention, AI systems become weapons to frustrate access to benefits. Critically, the problems that arose during the Unwinding process cannot be written off as outliers. Even when states face normal workloads, AI systems contribute to long backlogs in processing new applications, failed renewal processes, and wrongfully terminated coverage, with such problems documented in Colorado, New Mexico, Pennsylvania, Tennessee, and Rhode Island, among other places.¹¹

These failures undercut a principal justification offered for adopting AI: that it will increase government efficiency. In reality, it increases administrative burdens not only for people receiving government services, but also for agency staff who must spend time fixing errors or otherwise overseeing the systems.

Determining the Amount of In-Home Care Services for Disabled or Elderly People

AI exposure: 3 million to 5 million low-income people

Medicaid requires states to pay for nursing facility care for disabled or elderly people who require a significant level of care because of limitations in their ability to do daily life activities. Some states choose to offer home- and community-based services that help people who qualify for nursing facility care instead stay at home in their communities.

Community-based care costs less on average¹² and is better for people's dignity and independence. The most critical community-based service is personal care provided by a caregiver who helps the Medicaid recipient to get out of bed, bathe, use the restroom, prepare meals, eat, and go to medical appointments. Between 3 million and 5 million people in the United States receive home- and community-based services.¹³

At least 25 states, and maybe as many as 42, use AI made by one of two vendors to determine eligibility for home- and community-based services or the amount of services that will be allocated.¹⁴ With other known vendors selling AI tools and the use of self-developed tools by managed care companies, this report assumes that all states use some AI in their home- and community-based services programs.

The AI replaces previous systems that invested agency nurses or vetted third parties with the discretion to assess people and determine eligibility and hours based on their professional judgment, supporting information from the person and their caregivers, and medical records.

The AI-based process still starts with an assessment to gather information about a person's health conditions and abilities. But this process does not always produce reliable information. The person performing the assessment is likely a nurse, but the nurse may have never before met the person seeking benefits. The assessment often lasts two hours or more and involves 200 to 300 questions that are exhausting to answer. Often, the assessor has not reviewed medical records or consulted with a person's

doctor prior to the assessment; they use only the information provided by the person being assessed. And the people being assessed might not reveal the full extent of their conditions and related limitations due to a lack of trust, an incomplete ability to recall everything relevant in the moment, the length and difficulty of the assessment, or fear that honest answers could result in mandatory placement in a nursing facility.

From there, the information gathered during the assessment is fed into an AI tool that produces an output about whether the person is eligible and, if so, either the number of hours a person can receive or the budget amount a person will be provided to spend on services.

Though the specifics of the assessments and algorithms vary, these AI systems inevitably harm recipients. States often terminate or reduce the care hours of people previously determined eligible whose health conditions have not improved. States deny hours to people who need more care, and they deny eligibility to people newly applying for the program. These actions put people at risk of having to enter nursing facilities by rendering community-based options unattainable or unworkable.

Furthermore, the procedures surrounding the systems' usage defy principles of transparency, competence, rationality, and fairness. States fail to publicly share details about the AI systems and the way they work. They fail to look into the likely impact of the AI systems before they are deployed. State officials often do not understand how the systems work, what factors are considered, or why those factors (and not others) are used. They deploy technology with basic software errors. They fail to explain the reasons for denying or terminating eligibility or reducing care levels to the people who must survive the decisions. They use criteria to make decisions that the recipient does not understand or that is not clearly connected to a person's need for care. Independent of any ultimate decision, these failings undermine applicants' ability to count on services or plan for the future. Even if their care is not cut this year, it could be cut next year, regardless of whether their health improves.

The experience in Arkansas illustrates the problems common to many states that use these AI care-determination systems. Arkansas' Medicaid program historically determined how many hours of care someone received with a nurse's evaluation of each individual's needs, aided by a standardized assessment tool. Eight hours was the maximum available daily amount. For people with conditions like quadriplegia or cerebral palsy, this was not enough to meet all their care needs but was often workable.

Upon introduction of the AI system, however, the same people, whose conditions had not improved, saw their hours cut to four or five per day. Indeed, around five daily hours was the best-case scenario for most people seeking home- and community-based care. As a result, they were forced to lie in their own waste, got bed sores from not being turned, and were shut in their homes. The state did not run significant tests to project its impact before deploying the AI system, and once deployed, the state did not keep track of who was helped or hurt or by how much.

Advocates discovered that the AI system had coding errors that deprived people of additional hours beyond what would have been cut under its proper functioning. No state employee understood or could explain how the AI tool worked. The state lost the underlying data the AI system used to statistically model the amount of hours the system would provide. The written notices sent to recipients did not explain what factors the AI tool had looked at or what the person would have to prove to keep their hours. Appeals through the given administrative process were largely unsuccessful. State officials would not voluntarily change anything about the AI system. Only lawsuits and advocacy by the people on the program and their allies were able to bring about improvements.

Determining the Amount of Mental Health Treatment Available to Recipients

AI exposure: Unknown

Notably, some states use AI tools to allocate services outside of the context of home- and community-based services, such as for mental health services.

Arkansas's Medicaid agency, for example, uses AI to sort recipients with mental health conditions into one of three tiers ranked by the relative severity of their needs.¹⁵ Each tier has a different package of services attached to it, with the most relevant services being counseling with a therapist, in-office case management, daily group therapy, day treatment programs, home-based case management, treatment in residential facilities, and psychiatric services.

There have been two main problems with this system. To start, the assessment process may not get reliable information for reasons similar to those outlined above in the home- and community-based services section. Nurses are sent by the state's vendor to complete a lengthy, two-hour assessment with a person with severe mental illness whom they have never

met before. The nurses do not review medical records or consult with a person's therapist or psychiatrist prior to the assessment; they use only the information provided by the person being assessed. And the people being assessed did not reveal the full extent of their conditions and related limitations due to a lack of trust, an incomplete ability to recall everything relevant in the moment, the length and difficulty of the assessment, or fear that honest answers could result in involuntary commitment to a psychiatric hospital. Then, the AI system uses the information to assign people to tiers that are insufficient for extent of their needs or the severity of their illness.

When the tool was first implemented, the loss of services was particularly acute. Mental health providers scrambled valiantly to provide stopgap services while waiting for their patients to get new assessments or different tier placements. But their efforts could not meet the full scale of need.

As part of this process, one of my former clients lost access to daily group therapy that had helped his stability. The client's condition instantly worsened. He was involuntarily hospitalized for a time. Upon his release, he did not regain his previous stability and became unhoused. The irony is that his worsened condition would have qualified him for the services that he previously received that had prevented the decline. Under the new system, however, his stability had been punished.

The state's shift to the AI tool to allocate mental health services ran concurrently with a partial shift of the Medicaid program to a form of managed care (described in the next section) for recipients whose health care tends to cost more: people with severe mental illness and people with developmental disabilities.¹⁶ In addition to determining the scope of services available to recipients, the state may use the AI tiering system partly to determine how much the state will pay the private companies to manage these recipients' Medicaid services.¹⁷ This sort of AI-based tiering is a standard strategy available to states to manage their risk, but it is only one of many available strategies.

In light of the multiple possible uses for the AI and the apparent lack of available data, I cannot reasonably estimate the prevalence of AI tools used to determine mental health service levels.

Deciding Prior Authorizations

AI exposure: 55.5 million low-income people

Most states pay private insurance companies to manage their Medicaid programs, or parts of it, instead of doing it fully themselves. This type of Medicaid run by insurance companies is called managed care. Approximately 55.5 million people receive their Medicaid services this way.¹⁸

The underlying assumption of managed care is that, due to expertise and infrastructure, private insurance companies can run Medicaid in a less costly way than the state Medicaid agencies can. Under a managed care arrangement, the state pays insurance companies a set amount per Medicaid recipient. The insurance companies' profit depends on delivering services that, on a per-person basis, cost less than the amount they receive from the state. While proponents of managed care point to high health care spending as evidence of unnecessary treatments and program bloat, patient advocates highlight that the arrangement incentivizes insurance companies to deny treatments to increase profits. Given the potential for insurance company abuse, federal Medicaid law imposes requirements on Medicaid managed care programs and requires states to hire independent evaluators to review the insurance companies' actions. In practice, these measures often have limited efficacy for improving the outcomes for Medicaid recipients.

Cost controls are central to Medicaid managed care programs. For many medical treatments, a health insurance plan will not pay unless the plan receives a treatment request from a person's doctor and approves the treatment beforehand. This is called prior authorization. Ostensibly, this is to prevent doctors and other medical providers from billing for services that are not medically necessary. To receive prior authorization approval, the doctor's treatment request must include information about why the patient needs the particular treatment.

A 2023 report by the inspector general of the U.S. Department of Health and Human Services concluded that, overall, private insurance companies in Medicaid managed care programs denied 12.5% of all prior authorization requests.¹⁹ Some companies denied as many as 41% of all such requests.

Prior to AI, insurance companies had their own doctors or other medical professionals review each prior authorization request and decide if the information provided established the need for the treatment. Certainly, that process was often fraught, with insurance companies' reviewers encouraged to heavily scrutinize particularly costly treatments and demand more information. Now, managed care companies often use AI to conduct these reviews in ways that intensify the adversarial dynamics and make securing prior authorization even more difficult. And they do so without necessarily knowing that the AI works properly.

A recent report by the Government Accountability Office, a nonpartisan congressional agency, reviewed five states' Medicaid managed care programs administering care for children. The Government Accountability Office concluded that the five managed care programs—four of which used AI to automate decisions—never “reviewed a representative sample of denials or used data to assess the appropriateness of the full scope of plans' prior authorization decisions.”²⁰ In other words, those four states do not know whether the insurance companies' AI-based decisions were correct or incorrect.

Getting precise information on the use of AI in Medicaid managed care is difficult because the insurance companies fight tooth and nail to keep such information private. But ready proof of these practices is available by looking at how these companies run their other insurance programs.

Many insurance companies who manage state Medicaid programs—Centene, UnitedHealth Group, Anthem (now known as “Elevance”), Molina, Aetna/CVS, and others—also offer private insurance or offer managed care services in the Medicare program (discussed below).²¹ In these other programs, the companies regularly use AI to determine whether or not a given medical service will be paid for.²² In one egregious case, a private insurance company's doctors spent an average of 1.2 seconds per claim denial, a rate that reflected AI-generated denials rubber-stamped by doctors in batches of several dozen at a time.²³ AI had thus minimized doctors' review of claims to patients' detriment.

In a different twist, AI is sometimes used to increase doctors' review in the prior authorization process. One company involved in the prior authorization process is called EviCore, which is owned by Cigna.²⁴ EviCore sells its AI-based prior authorization system to other health insurance companies and Medicaid programs. One form of contracting arrangement allows EviCore to keep a portion of the money “saved” by the increased denials of services.

The system apparently works by giving each prior authorization request a score up to 100. High scores indicate a strong likelihood of approval based on EviCore's internal criteria. EviCore then sets a score threshold that triggers review by its doctors or nurses. For example, if EviCore sets the threshold at 90, all claims rated at or above 90 will be automatically approved, and the rest will go to human reviewers. EviCore's doctors and nurses then review claims according to strict guidelines that various professional medical organizations say are deficient and result in inappropriate treatment delays or denials. Moreover, EviCore will adjust the threshold applied to different insurance companies based on whether the system is generating sufficient savings. To continue the example above, if a threshold of 90 results in too many approved claims, EviCore can adjust the threshold to 95, ensuring more claims are reviewed by the company's doctors or nurses.

EviCore and its insurance company clients emphasize that the AI itself only approves claims and never denies them; humans perform all denials. Even so, AI can influence the humans to deny claims. The mere fact that a human reviewer sees a claim indicates to the reviewer that it did not pass the AI's applicable threshold. And company pressure to produce enough savings for clients can encourage EviCore's doctors and nurses to use the deficient guidelines to deny more claims.

This variable process has a secondary effect on the treatment that patients ultimately receive. Beyond the approval or denial for specific requests, the inscrutable and changing nature of EviCore's process discourages clinical doctors from even seeking some treatments for their patients. They get frustrated by the seemingly unnavigable process and give up. This mirrors the experience of many low-income people seeking benefits from hostile state agencies that use bureaucratic means to overwhelm people and dissuade them from seeking benefits.

The prior authorization scenarios are notable because they show that AI may be used to carry out treatment denials by minimizing human review or maximizing it. In the end, though, the goal is the same: cut costs. So, too, is the result. People receiving Medicaid managed care face inappropriate denials of treatment, significant delays while their doctors put together more evidence for approval, and, if they live long enough and can otherwise persist through it, a lengthy appeals process.

Medicare

AI exposure: About 16.5 million low-income people

Medicare is the United States' health insurance program for people with a qualifying work history who are disabled or age 65 or older. There are no income or asset limits on Medicare. As of April 2024, roughly 67 million people receive Medicare.²⁵ Of these recipients, 12.5 million have low enough incomes to qualify for additional support through Medicaid (these people are called dual eligibles since they qualify for both programs), and another 11 million do not receive Medicaid but have annual incomes below \$20,000.²⁶

Medicare recipients must generally pay a monthly premium to get coverage, which is currently \$175. In addition to the monthly premium, traditional Medicare requires recipients to pay for part of the cost of the medical treatments they receive, usually 20 percent. The federal government covers the remainder of the costs. But fewer people than ever before receive traditional Medicare.

Like with Medicaid, private health insurance companies can operate Medicare programs, which are called Medicare Advantage plans. Medicare Advantage premiums vary. Some require only the standard monthly Medicare premium of \$175, while others require more. Unlike with Medicaid, where each state decides whether a private company will run its Medicaid program and recipients have no say, though, individual Medicare recipients can choose between traditional Medicare and Medicare Advantage plans offered in their state.

About half of Medicare recipients—33 million—choose to receive coverage through Medicare Advantage plans.²⁷ Such plans are appealing for several reasons. Whereas traditional Medicare does not have any cost caps, Medicare Advantage plans have out-of-pocket limits that guarantee that recipients' responsibility for medical bills will stay below a certain amount.²⁸ And Medicare Advantage plans cover things that traditional Medicare does not cover, such as dental and vision care, and they do so at no additional cost to the recipient. But the private companies offer smaller networks—meaning that recipients have fewer doctors to choose from—and can do things that cannot be done in traditional Medicare, such as require prior authorization for most services.²⁹

Of the 33 million people on Medicare Advantage plans, 11.5 million (35 percent) have annual incomes below \$20,000.³⁰ Thus, at least this many people can be considered low-income. Another 9.9 million (30 percent) have incomes between \$20,000 and \$39,999. Any of these people who are single and whose annual incomes are below \$30,120 would also be considered low-income. Since the income threshold is in the middle of the income distribution band, this report assumes that half of the people fall below it, bringing the total count of low-income people on Medicare Advantage plans to roughly 16.5 million.

These people are exposed to AI through the prior authorization process.

Deciding Prior Authorizations

AI exposure: About 16.5 million low-income people

As noted above in the Medicaid section, health insurance companies running Medicare Advantage plans regularly use AI in the prior authorization process. Insurance companies running Medicare Advantage plans deny about 6% of prior authorization requests, fewer than half as many denials as in the Medicaid managed care programs described above.³¹

One explanation for the difference lies in kind and intensity of oversight. The U.S. Department of Health and Human Services oversees Medicare Advantage plans, and applicable laws guarantee, among other things, automatic external review of prior authorization denials. In contrast, state agencies oversee Medicaid managed care plans and do so with varying levels of sophistication, capacity, and interest.

Even though prior authorization denials happen less frequently in Medicare Advantage plans, their consequences are similarly grave. Media investigations show that insurance plans' use of AI meant that people of advanced age who had recently experienced broken shoulders, broken legs, strokes, or limb amputations were denied stays at or forced to prematurely leave hospitals and rehabilitation facilities, even though they could not do the tasks needed to live independently at home.³² In the process, the insurance company ignored medical records and doctors' opinions supporting these patients' limitations.

Predictably, the insurance companies refuse to divulge any information about how the AI systems work or internal studies about their accuracy. They claim that a reviewing physician, not the AI tool, makes the ultimate decision about eligibility. As shown in the Medicaid section, AI can lead to denials by either increasing or decreasing the role of human

reviewers. Where human review is minimized, insurance companies offer no explanation for why the human denial generally matches the AI system's recommendation. Where human review is maximized, insurance companies do not explain why particular prior authorization requests do not meet the AI system's requirements for automatic approval or why the criteria their reviewers use are more restrictive than professional medical bodies suggest.

Private Insurance (through Employers or Federal Subsidies)

AI exposure: About 30.6 million people

Many low-income people, particularly low-income adults, do not qualify for Medicaid or Medicare. They may live in one of the 10 states that have not expanded Medicaid eligibility to all qualifying adults, a group that includes large states like Texas, Florida, and Georgia.³³ They may earn more than Medicaid's income limits and be too young for Medicare. They may not have a disability that could qualify them for either or both programs. They might be an immigrant with lawful status who is barred from Medicaid eligibility for five years or might have an immigration status that does not provide for any eligibility.

Even when low-income people are eligible, they may not enroll in Medicaid or Medicare because they do not know about the programs or understand the program rules, feel stigmatized by using public benefits, or, for immigrants, fear that using public benefits could adversely affect their immigration status.

The only alternative options for health insurance coverage for low-income people are through their employers or by buying a private insurance plan. Only 22.6 percent of low-income people under the age of 65, or about 17.1 million people, obtain health insurance through an employer, either theirs or that of a spouse or parent.³⁴ There are several reasons: The employer may not offer insurance, the insurance may be too expensive or too limited in its coverage, or lower-cost alternatives may be available elsewhere.

Low-income people who do not qualify for Medicaid or Medicare and whose employers do not offer affordable insurance may be eligible for private insurance subsidized through federal advanced premium tax credits created by the Affordable Care Act.³⁵ People seeking such insurance can browse available plans through an online insurance marketplace run by a state or the federal government, compare the plans' coverage offered

and costs, and select one. For low-income people, the advanced premium tax credits cover between 98 and 100 percent of the cost of premiums for standard plans and the bulk of the premiums for plans with more extensive coverage. Even so, any premium costs can be prohibitive for low-income people. And they face many additional costs in the form of deductibles and co-payments. About 13.5 million low-income people enrolled in a plan this way in 2023 for coverage in 2024,³⁶ though the number who remain enrolled may be lower.³⁷

It is important to note that, even with Medicaid, Medicare, and private insurance available through employers or with federal subsidies, 11.8 million low-income people in the U.S. remain uninsured.

Low-income people receiving health insurance through an employer or with federal subsidies face AI through prior authorization procedures like those used in Medicaid managed care and Medicare Advantage plans.

Deciding Prior Authorizations

AI exposure: About 30.6 million low-income people

As with other forms of health insurance, private health insurance plans use prior authorization based on AI to determine people's access to particular forms of treatments. A full discussion of the related issues is in the sections above. The only thing this section adds is that there are currently movements to reform the prior authorization process to prevent insurance companies from deciding what criteria they will apply to prior authorization requests, which could limit their ability to apply internally developed AI systems.³⁸ These reform proposals also include provisions for increasing transparency, standardizing criteria, and reducing administrative burdens of the prior authorization process.

Supplemental Nutrition Assistance Program

AI exposure: 42 million low-income people

SNAP provides low-income people with the equivalent of money each month that only can be spent on certain food items—not any hot or prepared food—purchased at grocery stores, convenience stores, and, in some places, farmers markets. As of April 2024, about 42 million people were receiving SNAP benefits.³⁹ The monthly average benefit per person is \$186, though the amounts vary widely depending on income and household size.

Recipients must generally meet income limits for both gross and net incomes, capped at 130 percent and 100 percent of the federal poverty line, respectively.⁴⁰ They are also generally subject to asset limits of \$3,000, though states differ in what counts as assets. Recipients are also subject to work-reporting requirements that limit eligibility to 3 months in any 3-year span unless the person provides proof that they work a sufficient number of hours or meet one of the limited exemptions.

The federal government covers the full cost of SNAP benefits and splits administrative costs with states. Though federal laws establish core program requirements, states have significant flexibility to determine how to apply income limits, asset limits, work-reporting requirements, and other rules.

Prior to AI, states used computer systems similar to those used by Medicaid agencies to generate paperwork and to gather and store information about people applying for or receiving SNAP benefits. State agency staff would apply program rules and make all relevant decisions. Now, AI is used to determine eligibility, detect alleged fraud, and detect alleged overpayments based on income and asset limits. All SNAP recipients are subject to AI used for at least one of these purposes.

Determining Eligibility

AI exposure: 42 million low-income people

The process to determine eligibility for SNAP is, in some ways, more complicated than for Medicaid. Among other things, state workers must interview nearly every person who applies for or renews SNAP benefits. And state workers cannot necessarily access or use data that Medicaid has about a particular individual.⁴¹

Despite this complexity, all states use some form of AI to determine whether people are eligible⁴² or how much SNAP funding they will receive. And many states are trying to implement a single entry point where people can apply for Medicaid, SNAP, and other public benefits at the same time, which increases the technological challenges.

Most problematically, the AI systems frequently terminate, deny, or reduce SNAP benefits for eligible people. But there are a host of logistical issues that precede these harmful decisions. AI systems fail to process or keep information that people turn in, ask people for unnecessary information, delay eligibility determinations by several months, base decisions on data that is old or inaccurate, cause lengthy delays when recipients try

to get more information in person or over the phone, fail to get all the information needed to ensure people receive the right amount of benefits, and fail to explain the reasons for terminations, denials, or reductions.

When people do not receive benefits they are eligible for in a timely manner, they suffer. The impact on individuals also suggests the gravity of systemwide failures, with hundreds of thousands of people in any given state—and even millions in some states—made even more vulnerable during times of financial desperation.

An AI-based SNAP eligibility system used in Rhode Island highlights the stakes. In 2016, Rhode Island replaced its previous SNAP eligibility system with a new one built by a private firm, Deloitte. When the state shared information about the new system with the U.S. Department of Agriculture, the federal agency that oversees SNAP, the agency concluded that the state had not taken sufficient care in its development. It even told Rhode Island that “it is our decision not to allow the State to implement as planned” and that it was requiring a pilot program and a phased rollout.⁴³ The state repeatedly refused, at which point the U.S. Department of Agriculture told Rhode Island that the state would “proceed[] with the deployment ... at its own risk,” which included “reduced program access, worker backlogs, delayed application processing and untimely benefits, over-issuances and increased payment error rates.”

Despite these warnings, Rhode Island implemented the new system. Some of the 170,000 people who were receiving SNAP had their benefits wrongly cut off. And the state had a backlog of as many as 15,000 people waiting several months for an eligibility decision, despite the law requiring it to be made in 30 days.⁴⁴ After the debacle received heightened attention thanks to lawsuits by advocates for SNAP recipients,⁴⁵ the governor’s office issued a report summarizing Deloitte’s errors, including: ⁴⁶

- Key policy and program rules that were “improperly coded or configured into the system”;
- Defects regarding “basic user functionality and important interfaces, including those that are necessary for compliance with federal regulations”; and
- Complications in the design of critical applications that “suggest flaws in the vendor’s design and testing process.”

These errors affected basic eligibility functions, such as verifying wage and employment information, detecting overpayments and fraud, and sending recipients proper notices on time.

The governor's report also underscores another key dynamic of government AI use: the inability of government to properly oversee operations by sophisticated private vendors. In Rhode Island, the state did not have enough staff or the right expertise to hold Deloitte accountable during the development and deployment of the new system. These dynamics regularly show up in other states.

Detecting Alleged Fraud

AI exposure: 42 million low-income people

By all measures, fraud by SNAP recipients is rare. Fraud takes two forms: getting paid benefits improperly or using benefits improperly, such as by selling them. With respect to the first, for every \$100 of SNAP benefits paid to recipients, only 10 cents are improperly paid due to recipients' intentional violation of program rules.⁴⁷ With respect to the second, for every \$100 of SNAP benefits used, only \$1.50 is improperly used.⁴⁸ Difficulty knowing about and complying with the complex program rules accounts for improper payments or use far more often than fraud.

Nonetheless, federal and state governments have set up AI systems to detect fraud by recipients and the stores that accept SNAP benefits.⁴⁹ The systems flag transactions that are deemed suspicious, either based on the amounts (for example, high-dollar transactions or those ending in .00), the frequency of transactions (for example, several purchases at the same store in a short period of time), or the location (for example, at stores that are far from a recipient's home or are out of state). Still, the governments deploying these systems have never publicly substantiated that these factors correspond to intentional efforts to receive benefits via deception.

The agencies' anti-fraud staff do not necessarily investigate transactions flagged as suspicious to establish that anyone intentionally violated program rules.⁵⁰ Rather, the mere appearance of a transaction deemed suspicious by the AI tool, even without any other evidence, can be enough to initiate fraud charges and convince judges that fraud occurred. At that

point, the recipient is disqualified from receiving SNAP for at least a year (sometimes longer), is required to repay the agency, and can have their tax refunds or other public benefits intercepted. Some fraud cases may result in criminal prosecution.

Despite the high stakes, states use an unsubstantiated, opaque algorithm as proxy for intent and, in so doing, ensnare innocent people.

Detecting Alleged Overpayments

AI exposure: Unknown

AI systems are also fueling efforts by states to collect on alleged overpayments where fraud is not involved. For example, Illinois' SNAP agency increased its number of yearly nonfraud overpayment collection notices from roughly 7,700 in 2010 to roughly 33,000 in 2017, after the deployment of a new AI system.⁵¹ The state also used the new system to seek repayment on alleged overpayments that were several decades old, with one such collection effort starting in 2018 against a 72-year-old woman for alleged overpayments that had occurred three decades before. What is more, that overpayment from 1989 was not actually valid, as the agency had erroneous information about who was living in the household at the time. Even with the support of a legal aid attorney during an appeal, the woman was not able to stop the state's collection efforts. The agency's administrative law judge decided that the validity of the overpayment did not matter, instead ruling that the woman could no longer contest the overpayment because too much time had passed. As a result, her Social Security retirement benefits could be reduced to pay down the debt.

This case was not an outlier. As journalist Virginia Eubanks summarized:

The state's own auditors found in 2017 that [the Illinois Department of Human Services] "did not have appropriate controls" over [the AI system]: the system wasn't following established policy, the agency couldn't locate case files and other necessary documentation, 10-13% of eligibility redeterminations were overdue, and error-prone calculations were resulting in improper payments.

In both fraud and nonfraud cases, states' use of AI systems warps unreliable information into allegations, places the burden to resist the allegations on low-income people with few resources, subjects them to immense distress, and undercuts available appeal mechanisms.

Social Security Disability

AI exposure: About 13.8 million people, including about 10.6 million low-income people

The Social Security Administration administers two programs involving AI use: Social Security Disability Insurance, or SSDI, and Supplemental Security Income, or SSI. Both provide cash benefits to disabled people, and SSI also provides cash benefits to low-income people aged 65 or older without a disability. The two programs share medical criteria for determining whether someone is disabled but differ in the nonmedical eligibility criteria, the amounts offered, and the rules that apply once a person gets benefits. Presently, about 4.9 million people receive SSI on its own, 6.4 million receive SSDI on its own, and 2.5 million receive SSI and either SSDI or Social Security retirement benefits.⁵²

SSDI is funded primarily by payroll taxes on employers and employees, while SSI is funded entirely by the federal government.⁵³ Unlike Medicaid, SNAP, and UI, all SSDI and SSI program rules come from the federal government. While states do not set the rules, they have dedicated disability determination agencies that make the initial decision about whether an applicant meets the medical criteria to qualify.

Under both SSI and SSDI, being categorized as disabled requires a finding that the person is not able to work consistently because of a severe medical condition that is expected to last at least a year. Both programs use the same lengthy process to determine whether someone is disabled. This process involves dozens of pages of paperwork that the applicant must complete, reviews of medical records, and examinations by doctors paid by the Social Security Administration.

Fewer than one-third of applicants in either program are approved at the initial application.⁵⁴ Denials are primarily due to determinations that the applicant does not meet the medical eligibility criteria. Those who are denied can pursue multiple levels of appeal, but they face an overall wait time of roughly 2–3 years before being able to appear in front of an administrative law judge to present their case in person. After accounting for appeals, 33 percent of SSDI applicants and 43 percent of SSI applicants are ultimately approved. Given the low approval rates and significant delays, 30,000 people died in 2023 while waiting for a decision.⁵⁵

Nonmedical eligibility criteria, benefit amounts, and program rules differ significantly between SSDI and SSI. Qualifying for SSDI requires a substantial work history within a set timeframe. The benefit amount is based upon a person's earnings. The average monthly SSDI benefit is \$1,401 per month, but people who worked at low-wage jobs for most of their lives receive less.⁵⁶ There are no fixed asset or income caps that apply to SSDI recipients, though efforts to return to work, even part-time work, have complicated rules that consider income.

When disabled people do not have the required work history for SSDI or their SSDI benefit amount is very low, they qualify for SSI. And, as mentioned above, SSI is also available to people aged 65 or older who do not qualify for Social Security retirement benefits or whose retirement amounts are low. SSI currently provides a monthly benefit up to \$943. Nearly any additional income an SSI recipient gets—even free housing from a family member—reduces the monthly benefit amount. Unmarried SSI recipients also can never have more than \$2,000 worth of countable assets, and married recipients can have up to \$3,000.⁵⁷ Countable assets include things such as savings, life insurance policies, or nearly anything that can be sold for cash aside from one car and basic household items. Being a single dollar over this limit completely disqualifies the recipient from any SSI benefits for that month. These income and asset limits involve complicated rules and an elaborate enforcement scheme.

For purposes of this report, the 7.4 million people who receive any SSI are considered low income. With respect to SSDI, 90 percent of recipients have monthly benefit amounts below \$2,500,⁵⁸ which is 200 percent of the federal poverty level for a household of one person. Of these, some will live in households with other people whose incomes are sufficient to exceed the low-income threshold. Imprecisely, based on this fact and other data showing that 36 percent of SSDI recipients have income below 125 percent of the federal poverty line,⁵⁹ this report estimates that 50 percent of people who receive only SSDI, or 3.2 million people, are low income.

The Social Security Administration presently uses AI to enforce the SSI asset limits and is increasingly incorporating AI into the SSI and SSDI disability determination process, even though ultimate decisions are still formally reserved for staff.

SSI Asset Limits

AI exposure: About 7.4 million low-income people

Approximately 7.4 million people are subject to SSI asset limits. The Social Security Administration primarily enforces these through AI systems that review bank and property records for signs of recipient ownership. These data-matching systems fail in basic ways.

The property search system is called Accurant, which is currently owned by Lexis Nexis.⁶⁰ It is unable to distinguish clearly between people with the same names, even if they live in different states. For example, an SSI recipient named Joe Smith in Arkansas may be charged with owning property actually owned by a different person named Joe Smith in Florida. When this happened to my former clients, benefits were terminated immediately, and the client could not buy food or pay rent. Resolving the issue took several weeks.

Similarly, the bank account search system is unable to distinguish between owners of accounts and people who have no access or ownership but are listed as beneficiaries in the event of the account owner's death. In one example from my former client, an ex-wife was listed as the payable-on-death beneficiary of her ex-husband's account, which had more than \$2,000 in it. The Social Security Administration terminated her benefits at the end of the month, just days before rent was due and well after she had already used all her SNAP benefits for the month.

During their marriage, the ex-husband had been physically abusive toward the client. But to fix the issue with her SSI benefits, she had to beg him to go to the bank and get a statement that she had no ownership of or access to the bank account. The ex-husband eventually did so, but the Social Security Administration rejected the statement issued by the bank. I then had to draft a statement that would pass muster, and she had to again beg her ex-husband to go to the bank and get it signed. By the time the Social Security Administration reinstated her benefits, she had gone without food, racked up late fees for her rent that she could not afford, and had her power shut off for a day or two.

There is another impact of these benefit terminations. People who qualify for SSI also automatically qualify for Medicaid. When SSI eligibility is terminated, their access to Medicaid can also be terminated. This means that disabled people may be unable to receive medical treatment necessary to preserve their health or address their disability.

Current or Planned Uses for Eligibility Determination

AI exposure: About 13.8 million people

The Social Security Administration is working to integrate AI into the process of determining whether someone is medically eligible for SSI or SSDI.⁶¹ Details are sparse, but the Social Security Administration is using AI to annotate medical records at the initial level of decision-making, to flag some approved claims for further review, to review the appeals decisions by administrative law judges for consistency, and to identify claims with a high likelihood of approval at the initial level of decision-making and in appeals. In addition, the Social Security Administration is using AI to attempt to reduce the burden on SSI recipients to report income. The first two uses are the most worrisome, and the last use has the potential to help benefits recipients.

Annotating Medical Records

AI is used to review and annotate medical records that the disability determination agency looks to in deciding if an applicant has a sufficiently severe medical condition and resulting limitations in their ability to perform work-related functions. While it appears that the AI does not make or even recommend a decision about SSI or SSDI eligibility, the mere process of deciding what is important for the human decision-maker to look at can be problematic. The risk is that the human reviews only the information that AI flags and misses critical information that may be in the medical records themselves. Any such system would need to be vetted extensively to determine how accurate the AI is compared to human reviews. Even if accurate, use of the AI would require significant human oversight that limits any supposed efficiency gains.

Flagging Approved Claims for Further Review

Other concerning uses include flagging approved claims for further review, either at the time of initial decision or some point after approval when the AI thinks the applicant's medical conditions might have improved. This is because all medical eligibility decisions involve significant human discretion, and agency employees feel more pressure to deny claims than approve them. It is hard enough to convince the first caseworker that the evidence of eligibility is sufficient. Having to convince a second worker whose job it is to look for flaws in decisions increases the risks of denial.

Such reviews stand to particularly prejudice people with less extensive medical records or with medical conditions whose severity cannot be easily reduced to numerical evidence via blood tests, X-rays, MRIs, or other similar tests. Prior to qualifying for SSI, many low-income people have no access to health insurance to get regular ongoing treatment, either because they live in states that have not expanded Medicaid, are barred from receiving it due to their immigration status, receive it but live in parts of the state where few doctors accept Medicaid, or are emerging from incarceration, where treatment is difficult to get. So, when they do qualify for SSI, their claims may be more vulnerable to being reversed by an extra level of review for a supposed lack of medical evidence.

On top of that, people with mental health conditions face particularly difficult scrutiny on further review. First, severity cannot be determined through blood tests or imaging equipment, so the eligibility process relies more heavily on mental status exams and review of treatment records to see if there are signs of consistent functional limitations. These sorts of approaches to evaluation allow more room for interpretation, which cuts against the applicant when a second person is interpreting the decision.

Second, mental health treatment often is an ongoing iterative process, where the medical professionals diagnosing a person update their opinions as they gain more information about and familiarity with a person. Existing records may then reflect significant changes in diagnosis or prognosis that could be treated negatively by the AI.

Third, mental health conditions that prevent someone from working are often more visible when the person does not have easy access to regular treatment and other aspects of stability in life. Once someone qualifies for SSI and has steady income and health insurance, the visible aspects of their condition can become less severe. On review, then, their condition may no longer seem severe enough to keep them from working for SSI purposes, even if a loss of benefits would again make it plain.

Reviewing Appeals Decisions

After a claimant appears in front of an administrative law judge, the judge must write a decision that sets out the applicable law, facts about the person's medical conditions and functional limitations, and the reason for approval or denial of the claim. The Social Security Administration uses AI to review judges' decisions "to help maximize the quality, speed,

and consistency of their decision making” by looking at “the free text of disability decisions and other case data to offer adjudicators real-time alerts on potential quality issues and case-specific reference information....”⁶²

Presumably, this can be used to compare the judge’s present decision with past decisions by the same judge or with comparable decisions from different judges. But without more information about how the system works and what its impact has been, it is impossible to know if this helps or hurts claimants. For example, approvals could tend to trigger more alerts than denials. Or certain kinds of medical conditions and evidence could tend to trigger more alerts than other kinds. Moreover, it is not clear what happens if judges ignore such alerts or do not make changes sufficient to satisfy the AI system.

Identifying Claims with High Likelihood of Approval

The Social Security Administration uses AI to identify claims with a high likelihood of success at the initial level of decision-making and in appeals. The agency then prioritizes these claims. This seems a helpful use of AI to speed up clearcut decisions. That said, there could be deeper equity concerns in what the agency and AI see as likely claims, with some types of medical conditions and evidence being favored over others.

Automating Income Reporting

SSI has complex rules requiring that recipients report any other income they receive. In turn, the Social Security Administration reduces their benefit amount. While SSDI does not require that benefits be reduced in response to other income, the Social Security Administration considers that income in determining whether someone is within a trial work period allowing them to return to the workforce while retaining benefits or is otherwise still disabled.

Presently, reporting income can be difficult for benefits recipients. Some recipients do not understand the complex rules or that they need to report income. Even when they do, limited office hours, long in-office wait times, and limited alternative reporting methods can prevent proper reporting. Under SSI and SSDI, failing to report income can result in overpayments in which recipients must pay back any benefits they received to which they are not entitled. These overpayments often run into the tens of thousands of dollars.

In response, the Social Security Administration is proposing to automate the reporting of job-related income by using AI to seek out and use payroll data. Presumably, this will capture most job-related income and relieve recipients from the burden of having to report it.

Such a step could be helpful to recipients, but it must be properly executed to avoid harm. The AI system must work accurately, which is not guaranteed in light of the Social Security Administration's present errors with similar data-matching technology to enforce asset limits. In addition, recipients must be given notice of the AI system and how it works, the ability to opt out of its use, clear explanations for any information the system attributes specifically to them, and sufficient opportunities to challenge any resulting decisions. Likewise, recipients and the broader public must be able to get vital information about the AI system prior to and during its use and be meaningfully involved in its development and implementation. In addition, the Social Security Administration must guard against the dangers of having only one vendor, in this case Equifax, perform the function and become insulated from competition or accountability. Finally, the Social Security Administration must also mitigate data privacy and security risks that Equifax and other potential vendors create.

Unemployment Insurance

AI exposure: About 1.1 million low-income people

Unemployment Insurance, or UI, provides modest cash benefits to people who lose a job due to layoffs, business closures, or terminations that do not involve misconduct. To qualify, people must be unemployed due to no fault of their own, file weekly claims, be able and available to work, actively seek a job, and accept any reasonable job offered. Platform workers (for example, Uber or Lyft drivers), contract workers, self-employed people, and people who had to voluntarily quit for family or health reasons are generally not eligible. There are no asset limits in UI.

As of July 2024, only about 1.9 million people were receiving benefits,⁶³ which reflects less than a third of the total number of people who are unemployed in the United States.⁶⁴ In a given year, many more people receive benefits, as they cycle into and out of employment. In 2021, the most recent year for which tax data is available, 15.8 million people filed tax returns indicating the receipt of UI.⁶⁵

Not all people who receive UI are considered low income. In 2021, 8.9 million people of the 15.8 million people who indicated receipt of UI on tax returns, or 56 percent, had incomes below \$40,000.⁶⁶ Given the variable household sizes potentially involved at various income bands, this is a reasonable estimate for the proportion of UI recipients that would fall below 200 percent of the poverty line. Applied to the 1.9 million people currently receiving UI, approximately 1.1 million would be considered low income.

Both federal and state governments fund UI. While federal laws impose some minimal requirements on states, states have vast flexibility to determine eligibility criteria, benefit amounts, and program administration.

The amount of UI benefits a person receives is based on past wages, with people receiving around 40 percent of their lost wages up to a maximum amount. While the average weekly benefit nationwide is about \$450, individual states' average weekly benefits are often lower.⁶⁷ Mississippi's, for example, is \$221. People who have primarily worked in low-wage jobs receive less. Any wages people earn while receiving UI, even if from an odd job or irregular employment, must generally be reported and can result in decreased benefits.

States also cap the number of weeks recipients can get UI benefits, with some states offering as few as 12 weeks and others as many as 28 weeks.⁶⁸

States use AI systems to determine eligibility, verify identity, and detect alleged fraud. Soon, states also will use AI to write appeals decisions. Prior to AI, agency workers performed all such functions.

Determining Eligibility

AI exposure: About 1.1 million low-income people

As with Medicaid and SNAP, states operate AI systems to determine eligibility for Unemployment Insurance. These systems generally take information from applicants or recipients, process it, compare it against other sources of data, send out requests for information and notices, and either make a decision outright or prepare a decision for an agency caseworker to make. States botch these fundamental aspects of program administration in ways similar to those seen in SNAP and Medicaid.

For example, Florida's eligibility system actively discourages people seeking UI. Deployed in 2013 under then-Governor Rick Scott, the number of jobless people getting benefits decreased by 22 percent in only 3 months after implementation.⁶⁹ The state negotiated for improvements with the original contractor, Deloitte, and hired additional technology firms to work on the system.⁷⁰ Nonetheless, whatever improvements occurred did not better the system for people seeking benefits. A series of reports in 2015, 2016, 2019, and 2021 by Florida's A showed chronic deficiencies impeding all aspects of the eligibility process: people could not access the website or file claims despite dozens of attempts, paperwork they submitted was not processed, glitches of all sorts appeared, and state workers could not perform key functions to get applications approved.⁷¹

Before the COVID-19 pandemic, Florida had among the lowest rates in the country of UI claims relative to the number of people who were unemployed.⁷² When the pandemic increased unemployment claims exponentially, the system caused hundreds of thousands of people in desperate situations to spend weeks trying to file for benefits and months waiting to get them.⁷³ The fiasco led the governor who inherited the system, Ron DeSantis, to issue a remarkable statement admitting the shell game:

Having studied how [the unemployment system] was internally constructed, I think the goal was for whoever designed, it was, "Let's put as many kind of pointless roadblocks along the way, so people just say, oh, the hell with it, I'm not going to do that." ... It was definitely done in a way to lead to the least number of claims being paid out.⁷⁴

Further media investigations demonstrated that the technology problems were compounded by agency worker practices that would frustrate applicants at every turn. Among other things, agency staff would hang up on callers, determine that claims from spouses were automatically fraudulent, and lock claims where there was any discrepancy, such as when someone left out a middle initial.⁷⁵ That sort of agency culture partially explains why the technology functions so poorly.

Technological dysfunction is rarely just a technology problem. It is most often a policy problem. AI and related technologies offer ways for government officials who do not believe in public benefit programs to attack the programs and cut benefits in a roundabout way. Thus, any regulatory response that does not contemplate the weaponization of AI for policy ends is naïve and unlikely to be effective.

Verifying Identity

AI exposure: About 1.1 million low-income people

States use forms of AI to verify the identity of applicants. The pandemic brought about the expansion of such methods and the attendant harm, particularly for UI.⁷⁶

When AI flags a claim as suspicious, the person's application or payment (if they already receive benefits) is put on hold until the person supplies additional proof about who they are. Amid the pandemic, due to bad agency technology, long in-office wait times, impenetrable call centers, or individuals' lack of internet or smartphone access, furnishing such proof often took multiple attempts over several days. Even after providing the additional proof of identity, millions of people across the United States waited several months to receive UI benefits because state agencies did not have the capacity to process the enormous number of flagged claims.

In Arkansas, for instance, shortly after rolling out a new AI identity verification system, the state's UI agency had more than 90,000 outstanding applications flagged.⁷⁷ The entire state has roughly 2.3 million adults in it, meaning that a full 4 percent of adults in Arkansas were waiting for benefits.

In Michigan, a similar system by a different vendor caused about 1.4 million people to go without full benefit payments in 2020.⁷⁸ Unbelievably, the Michigan state UI agency still has not paid benefits to thousands of people who applied 4 years ago who had their claims flagged.

These delays caused people to go hungry, get evicted, have power shut off, experience mental health crises, and be unable to afford gas to drive to the store or clothes for their kids. Given the stakes, one might expect increased scrutiny on the AI systems. Why did they flag so many claims? Was it based on sound science?

States and their vendors do not provide proper answers. Instead, they refuse to disclose the information they possess about how such systems function, claiming that the AI is a trade secret and necessary for criminal investigations of fraud. Arkansas's UI agency defied state Freedom of Information laws to keep such information secret and, after losing in court,⁷⁹ encouraged the legislature to amend Freedom of Information laws to exempt such records from public disclosure.⁸⁰

Secrecy notwithstanding, advocates across the country have been able to piece together some of the common factors causing flags in different AI systems. Among other things, the systems flagged claims that came from the same apartment building, were associated with out-of-state addresses, or had even minor discrepancies with data provided by other sources. In other words, the AI systems were designed to cast a wide net.

Still, no publicly available proof emerged that the systems were appropriately scoped to root out fraud. In California, for example, at least 55 percent of flagged claims were legitimate.⁸¹ Certainly, congressional and agency studies have found fraud, but it was perpetrated by a small number of sophisticated actors—primarily, organized crime rings—working on large scales.⁸² States and their vendors have not explained why their wide-net approach was justified or why they could not avoid ensnaring millions of people by instead using a more targeted approach.

Notably, the U.S. Department of Labor is developing a national ID verification offering for unemployment benefits to states that relies on the existing government platform, login.gov. In this regime, claimants can also present verification documents in person at post offices. It is too soon to know how well these efforts are playing out in practice. Yet a standardized federal platform could reduce the need for states to run systems that prove deficient in design or implementation and are sold by opportunistic vendors—if it were implemented with accurate data, appropriate procedures that prevent broad and unwarranted ensnarement, and meaningful safeguards for people seeking benefits. Viable, non-harmful options are vital as governments consider expanding identity verification practices for other benefits and services.⁸³

Detecting Alleged Fraud

AI exposure: Unknown

Agency efforts at fraud detection involve more than verifying claimants' identities. They also include measures to determine whether UI recipients broke program eligibility rules with the intent of getting benefits for which they did not otherwise qualify. Agencies have deployed AI systems for this purpose to predictably destructive ends.

In a notorious case, Michigan's UI agency, in cooperation with vendors Fast Enterprises and SAS Analytics, deployed an automated fraud-detection system called MiDAS.⁸⁴ The system flagged fraud based on alleged discrepancies between what the recipient reported and what information

was available from the employer or a database with the recipient's earning records. Once a case was flagged, the state agency sent the recipient a questionnaire where virtually any response—or failure to respond at all—resulted in an automated fraud determination.

Ultimately, more than 40,000 UI recipients in Michigan were forced to repay between \$10,000 and upwards of \$187,000. They endured tax seizures, bankruptcy, and evictions. An official review conducted years after MiDAS's deployment concluded that the system generated false positives in 93 percent of determinations. State officials knew throughout the system's use that it had fundamental problems of accuracy and fairness. Nonetheless, they continued. Subsequent lawsuits and legislative actions have since brought the victims some relief, but money received several years later does not make up for the devastating experiences at the time. Michigan has since switched to a different contractor's system, Deloitte's uFACTs, which may have similar problems.

Concerningly, the dominance of a few major corporations limits states' options. For example, Thomson Reuters's Fraud Detect tool is used in 42 states and operates via secret algorithm. Advocates have recently filed complaints with the Federal Trade Commission over the use of these systems.⁸⁵ Such efforts may bring some much-needed accountability to particular uses, but the deeper problem of corporate concentration is likely to persist.

Deciding Appeals of Benefit Denials

AI exposure: Not applicable (future use)

When a person seeking unemployment benefits is denied, terminated, or accused of fraud, they have the right to an appeal. This appeal takes the form of an administrative hearing overseen by an administrative judge (who, sometimes, is not a lawyer) who is supposed to review documents, listen to testimony from the person seeking benefits, their witnesses, and any witnesses put forth by the agency or the employer, and then make an independent decision. The administrative judge writes their decision, says what the facts are, and analyzes those facts in light of the applicable law.

If the person seeking unemployment benefits does not agree with the judge's decision, they can appeal again to their state's regular court system. This step can cost a significant sum of money and has more complicated

rules than the administrative hearing. Even so, the reviewing court does not hear any new evidence. Rather, the court just reviews what happened at the administrative hearing and decides if the administrative judge's decision was out of bounds of what the law requires.

Over the next several months, Nevada's UI agency will deploy an AI system made by Google to "generate recommendations based on hearing transcripts and evidentiary documents, supplying its own analysis of whether a person's unemployment claim should be approved, denied, or modified."⁸⁶ Supposedly, the administrative judge will review the AI tool's analysis and either sign off on it or, if they disagree, modify it or write a new decision. The AI system is meant to address a backlog of appeals by speeding up the time it takes for an administrative judge to make and write a decision. Notably, it is likely that the backlog was created, in part, by the type of AI identify verification systems mentioned above.

The Nevada system poses awful risks to people seeking benefits because it substitutes the judgment of AI for that of trained human decision-makers. Generally, administrative hearings for UI benefits are not clearcut. Take the example of a convenience store worker who gave a customer their hot food order from the food counter and told them to pay at the register. The customer then ran off without paying. The employer fired the worker, claiming that the worker should have told the customer to pay first and show them the receipt before giving the customer the food. Did the worker commit misconduct that would disqualify them from benefits?

To decide, the administrative judge has to decide whether the store had a policy, whether that policy was actually followed in practice by other workers and enforced by management, whether the worker and her boss had any similar situations in the past, and whether the worker knew the customer such that they could have been trying to help the customer steal food. These issues turn on who the administrative judge believes, especially when witnesses disagree about what happened or what was said. AI is unlikely to be able to parse the nuances involved.

There are other concerns, too. First, in some states, administrative appeals are decided disproportionately against the people seeking benefits.⁸⁷ In Nevada, only 17.9 percent of workers win their administrative appeals. In Tennessee, Texas, and Louisiana, fewer than 10 percent of workers do so. If the AI is trained on these past decisions, it is likely to be biased against workers.

Second, as noted above, any further appeal to the state's normal court system is limited to the evidence and findings of fact decided in the administrative appeal. AI-based fact-finding may thus prejudice workers' ability to get a fair shot in front of other decision-makers. Third, decision-makers often defer to AI out of a belief that it does not make mistakes; this deference is called automation bias. This means judges may not be inclined to searchingly review or change the AI tool's opinion, even though AI is repeatedly shown to produce misinformation.

Finally, it seems unlikely that resource-strapped administrative judges will have the time to meaningfully review the AI tool's decisions. If the pressure is to work through a backlog, the human reviewer is even more likely to rubber-stamp the AI tool's decision, as has been shown in other instances of supposed human oversight.

Stepping back, the situation illustrates problematic dynamics around both consent and the allocation of the burdens of bad policymaking. Collectively, workers have not agreed to AI decision-making or even had a say in its adoption. On an individual level, given the power imbalances and immediate need for benefits, workers could not opt out of such arrangement even if their consent were sought.

Beyond this, though, the problem is not the workers themselves. The problem is that the state will not actually fund an unemployment system that processes claims promptly. Thus, why should workers receive less than the full process and face all the attendant risks of an AI-based one?

This scenario illustrates a common dynamic, namely that government officials' decision to use AI foists the impact of policymaking failures onto the people who had no role in their creation.

Chatbots

AI exposure: Unknown

Chatbots are automated tools that people use to interact with companies and government agencies for customer service purposes. Although chatbots do not make decisions about people's eligibility for benefits, they shape what people understand about programs. This, in turn, affects whether people apply for benefits and how they comply with program rules. As such, this report covers them briefly.

Chatbots are presently used in public benefits administration to provide general program information, not individualized information about a specific person's case.⁸⁸ People prompt the agency chatbot with a question or topic they want more information about. The chatbot works by identifying the topic of a prompt or question, reviewing a library of pre-programmed responses vetted by the agency, and identifying which of these responses best connects to the prompt or question. Even when the same information is available somewhere on an agency's website, a chatbot can be an easier way of finding it than by clicking through various screens and trying to find the right place.

When chatbots provide clear, accurate information, they can allow people to avoid lengthy phone calls or office visits for relatively routine matters. Complicated issues or case-specific questions still require conversations with agency staff.

Recent developments in AI could expand both the usefulness and dangers of chatbots. So-called generative AI allows chatbots to respond to people's prompts with responses that are not pre-programmed. This theoretically allows chatbots to respond to a wider breadth of topics and do so in a more detailed manner. These chatbots could also be used for an agency's internal purposes, such as by helping a caseworker answer complicated policy questions or pointing them in the direction of applicable sources of answers.

But generative AI tools are known to provide false or invented information, a phenomenon that is unfortunately called hallucinating.⁸⁹ Such information can mislead people in high-stakes situations involving public benefits. As such, any government agency seeking to use generative AI chatbots would need to invest heavily in developing, vetting, and overseeing them to ensure that all information provided is fully accurate. Thus far, state agencies administering public benefits have not generally shown themselves to be capable of developing, vetting, and overseeing AI-based technologies, either for lack of sophistication, resources, or interest. Thus, any movement toward adoption of generative AI chatbots should be met with skepticism.

Additionally, generative AI chatbots could exacerbate existing agency budget and staffing problems. Chatbots promise to make agency operations more efficient and free staff time for functions that only human staff can do. However, this newly created capacity could instead be used by policymakers to justify cuts to already-understaffed agencies on the basis that fewer humans are needed to do the same amount of work as before.

Rental Housing

AI exposure: About 16.3 million low-income renter households, or 39.8 million low-income people

Approximately 45 million households—more than one-third of all U.S. households—live in rental housing.⁹⁰ With the average household size of rental units being 2.44 people, there are approximately 109.8 million people in the US living in rentals.⁹¹

As explained in a note below, this report estimates that there are about 18.1 million low-income renter households, which contain 44.2 million low-income people. Most, but not all, are exposed to AI-based decision-making.

Around 75 percent of low-income households spend more than 30 percent of their income on rent and utilities, leaving little money left for life's other necessities, such as food, clothing, child care, medical or dental treatment, and transportation.⁹² Facing such precarity, many renter households receive ongoing pressure from landlords to move, largely through rent increases or unmet repair needs.⁹³ For most low-income people, finding and maintaining decent, stable housing is a constant challenge.

There are three forms of rental housing that provide a small proportion of low-income renters some stability in the face of rising prices, lease nonrenewals, and easy evictions: public housing, housing vouchers, and rent stabilization or rent control laws.

The U.S. Department of Housing and Urban Development, or HUD, administers public housing and Housing Choice Vouchers (also known as Section 8 vouchers). Approximately 970,000 households containing 1.6 million people live in public housing managed by local governing bodies called housing authorities (the average household size differs in public housing).⁹⁴ Public housing buildings are often older or in significant disrepair, frequently with mold, cockroaches, rodents, and leaks.⁹⁵ They are often located near environmentally toxic sites, including 9,000 properties that are within a mile of so-called Superfund sites of extreme pollution.⁹⁶

Meanwhile, 2.3 million households comprising more than 5 million people receive Housing Choice Vouchers managed by local housing authorities that help cover the rent charged by private landlords.⁹⁷ Generally, landlords do not have to accept the voucher, but if they do, they are subject to HUD rules.

These rules include minimum standards for safety and habitability and require occasional inspections, though such inspections are routinely late, cursory, or otherwise ineffective.

The two programs share similar income eligibility rules requiring people have income below 80 percent of the region's median.⁹⁸ Under both programs, tenants generally spend no more than 30 percent of their income on rent. The federal government funds the programs, but the funding amount (at least for public housing) has declined in real terms over the past two decades. In many areas, people wait several years to receive public housing or vouchers, and housing authorities often close wait lists due to extreme existing demand.

Once eligible, tenants generally have the right to continue living in public housing or receiving vouchers unless they do not pay the rent, violate lease terms, or become financially ineligible. They also are generally entitled to procedural protections such as notice and an administrative hearing before losing their public housing or voucher. Even with these protections in place, landlords who accept vouchers can generally choose to simply not renew a lease when the term is up.

The third type of tenant protection involves rent stabilization or rent control laws, which limit yearly rent increases charged by private landlords.⁹⁹ Most states prevent cities within them from adopting such laws. Rent stabilization or rent control is currently in effect only in Washington, DC and select cities in four states. New York City is the largest of these and has nearly 1 million apartments subject to rent stabilization or control, though the proportion of tenants who are low-income is not known.¹⁰⁰ As with public housing, landlords with properties subject to these laws cannot simply decide not to renew a lease. The tenant generally has a right to continue living there so long as they pay rent and follow the lease terms.

The three protections identified above extend to only a small portion of low-income tenants, with the vast majority fully subject to the vagaries of the market and landlords.

AI touches or is likely to touch most low-income tenants regardless of any protections, as landlords use AI to screen tenants, set rents, and surveil tenant activity. The pre-AI methods used to perform these functions are described in each section.

A note on counting

Determining how many rental households are low income is difficult. 27.3 million have incomes less than 80 percent of the regional median, which qualifies them as low income by some measures.¹⁰¹ But this report uses a more conservative estimate based on dollars.

14.6 million renter households have annual incomes below \$30,000. At this income level, a household of any size is below 200 percent of the federal poverty level, meeting this report's definition for being low income. Another 6.9 million renter households have household incomes between \$30,000 and \$44,999 per year.¹⁰² At any income in this band, rental households of three or more people would be considered low income, as would the average rental household of 2.44 people. Households of one person would not be considered low income. And households of two people may or may not be low income depending on where their income in the band falls (those with incomes up to or below \$40,880 would be low income; those above would not be). This report estimates that half of these households, or 3.5 million, qualify as low income. This brings the total number of low-income rental households to 18.1 million, comprising 44.2 million people.

Screening Tenants

AI exposure: 16.3 million low-income renter households, or 39.8 million people

Landlords use AI-related screening algorithms to decide whether to approve people to rent a home.¹⁰³ Various private companies sell these screening systems or have been involved in their development, including TransUnion, CoreLogic, RentGrow, RealPage, SafeRent, Cleara, and others.¹⁰⁴

The systems work by matching the names of the applicant against names on various databases that contain information unfavorable to the applicant, such as criminal records, eviction records, and lists of foreign agents.¹⁰⁵ Some also factor in credit scores.¹⁰⁶ Once the system has gathered the information, it generates a report that highlights the unfavorable information and, sometimes, contains a score or a recommendation.¹⁰⁷ The latest iterations additionally claim to predict the risk of property damage, late rent payments, short tenancies, and evictions.¹⁰⁸

As of 2017, 90 percent of landlords used credit or criminal background checks, and 85 percent checked past eviction history.¹⁰⁹ Presumably, with the passage of time and expansion of the market for background checks, this happens exclusively via AI technologies. Assuming that these 90 percent of landlords rent to 90 percent of low-income tenants, 16.3 million low-income rental households comprising 39.8 million people would be subject to AI-based screening. Before AI, landlords had access to this negative information but might have considered it in context or asked the prospective tenant for more information. AI-based recommendations and scores make this less likely.

This AI fails in predictable ways. At the most basic level, it does not ensure that the information returned pertains to the person seeking housing. For example, Marco Antonio Fernandez, a member of the U.S. Navy with top-secret clearance, was denied a rental after the background check returned information about Marco Alberto Fernandez Santana, a person who allegedly trafficked drugs.¹¹⁰ Similar stories abound of renters rejected based on information about different, unrelated people who happen to have part of a name in common.¹¹¹ As one media report notes, the impact of simple name-based matching is likely to fall harder on racial minorities, who tend to have fewer unique last names.¹¹²

Even when AI returns information about the correct renter, the information is often wrong or misleading. Court records for evictions and criminal proceedings are often incomplete and do not reflect whether a court actually reached a full decision on the issue before it. Relatedly, the data sources AI uses are often not up to date. For example, one renter was denied housing based on a criminal conviction that had been overturned and sealed from his record, meaning that he was not guilty of any crime.¹¹³ Given racial disparities in the criminal legal system, inaccuracies in background checks disproportionately hurt Black and brown people.

Finally, the AI does not explain in detail how it reaches a recommendation, score, or prediction, and companies developing this AI refuse to voluntarily disclose more information about how their systems work.¹¹⁴ On an individual level, this means that landlords and, by extension, rental applicants do not have information about what counts so they can address concerns and increase the chances of agreeing on a rental. On a broader scale, government entities and the wider public are denied information critical to deciding whether these faulty screening approaches should be regulated differently.

Importantly, in light of the many kinds of errors associated with this AI, people denied housing based on AI screenings are not given meaningful opportunities to challenge the decision. Landlords decide on an individual basis whether they will consider additional information. But if they receive multiple applications, it is unlikely they will spend time to further vet an applicant if an alternative applicant with no flags is available. While the denied applicants can also request that the AI screening company correct the erroneous information, that process is difficult, can take hours of time, often requires hiring an attorney, and may not produce any change for years. By the time of any correction, the immediate housing opportunity is lost.

One other aspect of tenant screening merits attention, which is the mass accumulation, aggregation, and use of negative information about people. Even when accurate, AI facilitates the perpetual haunting of people based on unfortunate actions in the distant past. Previous evictions or criminal records may have little to do with people's present qualifications to rent a given apartment but will still hurt their housing opportunities. In this way, AI effectively hinders people from moving on and moving up.

Setting Rents

AI exposure: Unknown

At least one of the companies that sells tenant screening services, RealPage, also sells an AI system that recommends the rental price that landlords should charge for a given property.¹¹⁵ A recent lawsuit by the U.S. Department of Justice, or DOJ, and several states alleges that RealPage violates various provisions of antitrust laws that outlaw monopolies and other practices that undermine open, competitive markets.¹¹⁶ The DOJ-led lawsuit follows at least 20 lawsuits by private plaintiffs and the attorneys general of Washington, DC and Arizona.¹¹⁷

When a landlord signs up for the service, it provides RealPage's AI system access to company data that the landlord normally keeps private, including the location, size, and rent charged for its properties. The AI then compares this to data from other landlords that have signed up for the service and recommends the price to charge. In some cases, the AI will even encourage landlords to keep properties vacant to keep prices up.

As the AI incorporates more landlords in an area, its recommendations to set higher prices take into account the recommendations made to other landlords in the area. This tends to increase rental prices across the board.

Notably, the AI system breaks from past landlord price-setting practices, which favored smaller rent increases and maximum occupancy levels, usually informed by internal data and whatever market information was available.

So far, it is not possible to specify the number of low-income people subjected directly to these AI price-setting practices. Apparently, RealPages' AI system has a broad reach. One lawsuit by private plaintiffs alleges that RealPage "had over 31,700 clients, including owner operators and each of the 10 largest multifamily property management companies in the United States."¹¹⁸ The Department of Justice estimates in its lawsuit that at least 4 million rental units have landlords that use RealPages' AI system.¹¹⁹ The DC attorney general claims that RealPages' system "is used to set rents at more than 50,000 units in the District," which includes "well over 30% of apartments in multifamily buildings."¹²⁰ The Arizona attorney general states that "approximately 70% of multifamily apartment units listed in the Phoenix metropolitan area are owned, operated, or managed by companies that have contracted with RealPage."¹²¹ But other information about the demographic characteristics of the renters subject to the AI system is not currently available.¹²²

Less directly, an example from Seattle suggests possible downstream effects on low-income people. The 10 largest landlords in Seattle's Belltown neighborhood controlled 70 percent of rental units.¹²³ All of them used RealPages' AI system to set prices for at least some of their properties, with tenants in AI-priced buildings seeing more significant price increases than tenants elsewhere. Steep increases in rental costs can cause "down renting," where people with higher incomes choose less costly apartments that would otherwise have been available to people with lower incomes. According to ProPublica, "Seattle should have had a surplus of 9,000 apartments affordable to people making 80% or less of the median income ... [b]ut tenants' down renting as prices rose turned that surplus into a deficit of 21,000."¹²⁴

Of course, factors other than the AI system may lead to price increases and the down renting phenomenon. But if the AI system causes significant rent increases in higher-income properties, it could also cause low-income people to face more competition with higher-earning people for rental units that they can afford.

Surveilling Renters

AI exposure: Unknown but ubiquitous

Low-income tenants, particularly those in public housing or private housing with rent control or stabilization, face ubiquitous surveillance through cameras that outnumber, on a per-capita basis, the cameras found in airports, stadiums, museums, or even prisons.¹²⁵ Landlords often pair these cameras with AI-enabled facial recognition technology or video-reviewing software to detect supposed lease violations. Using common spaces, receiving caregiving support, having overnight guests, or letting family or friends stay for an extended time can all lead to formal accusations that a lease has been violated and allow the landlord to evict the tenant.

For their part, landlords deny any improper motivation and hail the surveillance as a way to keep residents safer. To this end, federal grants from the Department of Housing and Urban Development have helped landlords purchase these systems. Still, the AI use is not necessarily wholly benevolent.

Generally, because landlords of public, rent-stabilized, or rent-controlled housing cannot simply decide not to renew a lease if the tenant is paying rent and following the lease terms, lease violations are the primary means that landlords have of creating a vacancy or charging increased rent. Prior to AI, landlords would generate vacancies by manually reviewing surveillance video footage, enlisting cooperation from other tenants in establishing lease violations, or selectively enforcing small, technical lease violations—such as hanging clothes outside or leaving a window decoration up—against disfavored tenants to create the appearance of ongoing noncompliance. AI-driven surveillance provides landlords an easier way.

On top of the ability for landlord misuse, the surveillance technology at issue simply does not work reliably. In 2018, a groundbreaking study by researchers Joy Buolamwini, Timnit Gebru, Deborah Raji, and Helen Raynham showed that facial recognition technology consistently fails to identify women and people with darker skin tones.¹²⁶ Multiple studies by the federal government's National Institute of Standards and Technology demonstrated that facial recognition technology fails to properly identify Asian American and African American people.¹²⁷ When used for policing,

the technology has led police to falsely arrest innocent Black people.¹²⁸ These AI failures are especially worrisome when the low-income tenants to whom the technology is applied (at least in public housing) are disproportionately Black,¹²⁹ the very people most likely to be misidentified by AI.

Errors aside, there are deeper concerns for low-income tenants' dignity. The tenants are not meaningfully consulted before AI is deployed and have no formal power to refuse their consent. Plus, low-income tenants' homes are often in places that are already subject to constant police surveillance and activity. Adding more surveillance exacerbates the feeling of being constantly watched. As one tenant put it when resisting facial recognition at her rent-stabilized apartment complex, "We should not feel like we're in a prison to enter into our homes."¹³⁰

Presumably, these surveillance practices targeted against tenants in public housing or private rent-stabilized housing affect only a small proportion of low-income renters. Yet even this could number millions of people.

Moreover, the overall trend seems to be expanding AI surveillance usage, particularly as costs come down, the technology's failings are improved (or are seen as improving), and concerns about property crimes take center stage regardless of their actual prevalence. Unsurprisingly, such technologies are already in use in private rental housing that is not rent-stabilized. As one New York City tenant resisting the facial recognition technology stated plainly, "We're not going to stand for it, and we're asking other renters in the state, not just in the city, to join us—because it's coming to your house next."¹³¹

Employment

AI exposure: At least 32.4 million low-wage workers

Low-wage work primarily involves the agriculture, food service, retail, and health care industries. Among other jobs, people work as cashiers, servers, landscapers, food pickers, nursing assistants, line cooks, care aides, janitors, security guards, hotel housekeepers, and dishwashers, among others.¹³²

There is no official definition for low-wage workers, but one common approach defines the income threshold as being two-thirds of that of the median wage of U.S. workers. Focusing only on prime-age workers—those between the ages of 25 and 54—one study sets the threshold at \$19.20 per

hour, or \$39,900 per year (expressed in 2024 dollars), and finds 30 million prime-age workers falling below it.¹³³ A second study broadens the age range to include all working people between the ages of 18 and 64, sets the income threshold at \$20.93 per hour, or about \$43,500 per year (expressed in 2024 dollars),¹³⁴ and finds 53 million workers falling below it.¹³⁵ Because AI exposure is unlikely to be limited based on the age of a worker in low-wage industries, this report considers the 53 million workers earning below \$21 per hour to be low-wage workers. This section will alternate between using the terms low-income people and low-wage workers, each with a different definition, as appropriate based on available data.¹³⁶

Disparities of gender, race, and disability abound. Women make up more than half of the low-wage workforce. Black and Latino workers are overrepresented in the ranks of low-wage workers. Disabled workers are more concentrated in the low-paying service industry than nondisabled workers.¹³⁷

Low-wage workers face all sorts of problems in addition to low pay. They usually do not have access to good benefits. Companies sometimes steal wages by not paying them for all the hours they work. Their schedules may change from week to week in unpredictable ways. They often do not get reliable leave to care for themselves or a family member when sick. They generally can be fired or laid off for just about any nondiscriminatory reason. They are routinely discouraged from asking for changes to make the workplace safer or fairer. They may have to sign noncompete agreements that prevent them from changing employers within the same field. They often have to go through exploitative platforms or temp agencies that keep them as independent contractors instead of employees.

In addition to these other challenges, low-wage workers face AI that decides who gets hired, sets their wages, and surveils them. Ultimately, these technologies compel vulnerable people to work under increasingly precarious conditions. The pre-AI methods used to perform these functions are described in each section.

Hiring Workers

AI exposure: About 24.4 million low-wage workers

AI systems operate at nearly every stage of hiring: advertising job openings, reviewing incoming resumes, vetting applicants during the interview process, and making a recommendation. Prior to AI, humans generally performed these functions by applying individual experience and company-determined criteria to applicants.

Low-income people seeking job openings through social media platforms are unlikely to be shown all available jobs in their region. The job ads that a person sees on Facebook have historically skewed based on gender and race, even when the employer did not intend to do so.¹³⁸ Factors such as the ad's budget, content, and image operate independently of the specific characteristics the employer specifies about the target audience.¹³⁹ Thus, ads for jobs in the lumber industry reach an audience that is overwhelmingly white and male, supermarket cashier jobs are seen by a nearly all-female audience, and ads for positions in taxi companies are seen by a mostly Black audience. Such skewing is important, of course, because low-income workers cannot apply to job openings they do not know about.

When people apply for job openings in entry-level positions with large employers, such as Walmart, Target, Starbucks, UPS, McDonald's, Kroger, or Amazon, they are put through an intensive AI-based vetting process.¹⁴⁰ Applications or resumes that do not meet the employers' specifications are screened out automatically. Importantly, those specifications may include factors that do not have any connection to the job responsibilities, such as distance to work, or factors that the applicant could change, including schedule availability or rate of pay sought.

The people who move forward are then subject to automated testing to measure skills and personality, among other things. These tests involve counting change, deciding on the best responses to common work scenarios, and assessing personality traits, like obeying authority. Those people who move past the testing stage may then receive an interview, where some employers now use AI that purports to analyze eye movements, tone of voice, and word selection to rate an applicant's enthusiasm, candor, and fit for the job.

Finally, employers run automated background checks that span the applicants' criminal records, credit history, and social media profiles and automatically screen out anyone with characteristics the company deems

undesirable.¹⁴¹ Similar to the tenant background screening and Social Security data-matching examples discussed above, these background checks often return information about people other than the applicant or inaccurate or outdated information about the applicant themselves.

The automated vetting process is fraught with problems. First, many of the screening steps are likely to be biased against people of color, disabled people, and low-income people. Accommodations for disabled people are not readily accessible. Racism in the criminal legal system means that Black and Latino applicants are more likely to be screened out due to background checks, even if they are otherwise-qualified for the job. And personality tests are often based around white, middle-class norms and can tend to exclude anyone not from that background. Likewise, such tests often exclude people with developmental disabilities or mental health conditions.

Second, the process is not transparent to the applicant. When applicants are eliminated from consideration, they do not receive any explanation and thus cannot strengthen future applications. Third, low-income people have less access to broadband connections, computers, and smartphones than people with higher incomes.¹⁴² The lesser availability of and decreased comfort with the required technology can impede low-income people from successfully navigating AI-driven processes.¹⁴³

While the number of low-income people subject to AI in the hiring process is not precisely known, 98 percent of Fortune 500 companies use this technology.¹⁴⁴ Using the hourly wage threshold of \$21 to define low-wage worker, 57 percent of low-wage workers work in companies that have at least 100 employees, 46 percent work in companies with at least 500 employees, and 41 percent work in companies with at least 1,000 employees.¹⁴⁵ Based on the prevalence of AI technology in Fortune 500 companies, this report assumes that all companies with at least 500 employees use AI in their hiring processes and smaller companies do not. Applying 46 percent to the total population of 53,000,000 low-wage workers, I estimate that 24.4 million of them would be subject to AI in the hiring process.

Alternative assumptions show similarly large scales of exposure. If I were to assume that AI is used in the hiring process in companies with at least 100 employees, the total would be 30.2 million. If I were to assume that AI is used only in companies with at least 1,000 employees, the total would be 21.7 million.

Setting Wages

AI exposure: 8 million full-time “gig” workers and 23 million people in the past year who have done gig work

The rise of the so-called gig economy, largely led by ride-hailing platforms such as Uber and Lyft, has intensified the tools employers have to make the wages of low-income people less fair, less stable, and less predictable.

University of California, Irvine law professor Veena Dubal has coined the term “algorithmic wage discrimination” to describe “a practice in which individual workers are paid different hourly wages—calculated with ever-changing formulas using granular data on location, individual behavior, demand, supply, or other factors—for broadly similar work.”¹⁴⁶ The discrimination cuts two ways. One, workers doing substantially identical work at identical times may receive different wages. Two, the same worker doing substantially identical work at a different time may receive different wages.

The end result is increased precarity at work, as workers do not know how their wages are calculated, how much they will make in any given timeframe, or whether the wages they actually do receive are in line with what others doing similar work make. Within this uncertainty, the platforms add the possibility of bonus payments if workers complete a certain number of tasks within a set timeframe or choose to work during times when higher “surge” fares are available. Workers see this as “casino culture,” whereby the hope for a bigger payout on the horizon justifies ongoing effort that is insufficiently paid. Chasing the bonuses, workers perceive that the platform seems to depress their opportunities to meet the goals as they get closer to reaching them. On top of this, workers do not receive the opportunity for bonus payments at any regular intervals, and available bonuses differ between workers.

Though the platforms suggest otherwise, people in the gig economy are low-wage workers. After accounting for uncompensated waiting time and vehicle costs, two independent studies put an average ride-hailing application driver’s hourly income between \$9.73 and \$11.77.¹⁴⁷ And, like most low-wage work, the workforce is disproportionately composed of people of color and immigrants. For example, Lyft’s data shows that 73 percent of their workers are people of color, and city-specific studies show that as many as 90 percent of drivers are immigrants in some locations.¹⁴⁸

Algorithmic wage discrimination displayed in the ride-hail industry implicates broader concerns about pay disparities along racial and gender lines. Across the U.S. economy, Black and Hispanic people and women earn less on average than their white or male counterparts.¹⁴⁹ This dynamic persists when possible explanations for the differences, such as education levels, experience, location, and specific work performed, are accounted for.¹⁵⁰ For all the newness of AI technologies and the ride-hailing industries, they reproduce long-embedded patterns of paying less for the labor of marginalized people.

Counting the number of people ever exposed to algorithmic wage discrimination through on-demand platform work is less complex than determining the number regularly subjected to it. As many as 16 percent of U.S. adults, or roughly 41 million people, have earned money at least one time working for at least one on-demand platform, such as Uber, Lyft, Task Rabbit, or Postmates. But only 9 percent, or about 23 million people, have done so in the past year.¹⁵¹ About 3 percent, or between 7 million and 8 million people, have done it as their main job in the past year.

The impact of algorithmic wage discrimination may well extend beyond the number of people participating in on-demand platform work. As UC Irvine's Professor Dubal notes, one danger is that algorithmic wage discrimination will shift the broader norms of employers and employees in other sectors such that the increased precarity and unpredictability tied to black-box AI systems become more acceptable.

To that point, workers in other low-wage sectors, like retail and fast food, already face AI-generated scheduling,¹⁵² which may not meet employees' needs for child care, transportation, or sufficient income. Similar unpredictability is found in surveillance-enforced pay differentials for upselling customers to buy bigger sizes of fries or drinks or to add dessert to their order.¹⁵³

Practices popularized by ride-hailing platforms, such as surge pricing, may soon be expanding to other sectors of the economy. Fast food companies are already implementing dynamic pricing, a variation of surge pricing in which they charge different amounts for certain items at different times of the day.¹⁵⁴ The largest U.S. grocery retail company, Kroger, which owns chains including Ralph's, Fred Meyer, Food 4 Less, and others, is introducing facial recognition technology to stores, along with electronic price tags for shelves.¹⁵⁵ Some elected officials are concerned that such technology could lead grocery stores to raise prices for individual consumers based on their income, buying tendencies, and other available consumer data.

Surveilling and Managing Workers

AI exposure: Unknown but ubiquitous

Employers are able to easily monitor all aspects of employees' work: what they say, where they go, how quickly they work, when they start and finish work activities, what they do while at work, and, for public-facing positions, what customers say about them.

On first blush, the technology simply seems a modern way of achieving employers' longstanding goal to ensure worker productivity. A deeper look, however, reveals several perils with this workplace surveillance: inaccuracies, pay denied for time spent working that is not captured (that is, wage theft), intensified and unrealistic expectations placed on workers, interference with workers' rights, increased information asymmetries between workers and management, undermined worker autonomy and dignity, increased risk of physical injuries and health-harming anxiety or stress, and racism or other forms of discrimination.¹⁵⁶ And, of course, these AI management practices are backed by consequences: Workers can have their pay or hours reduced, be demoted, miss out on promotions, or be fired.

Some of these problematic dynamics also are present in traditional, non-AI management practices. After all, human bosses monitor workers, play favorites, get facts wrong, set moving or unreasonable performance targets, treat employees in discriminatory ways, and undercut efforts to improve working conditions. Reducing such human-driven unfairness is sometimes cited in employers' public justifications for shifting to AI.¹⁵⁷

Still, AI-driven surveillance management poses unique risks, many of which are interrelated. First, workers cannot discuss their performance with the AI system to challenge adverse decisions, adjust expectations, highlight events beyond their control, or otherwise explain their performance. Second, the AI carries the veneer of objectivity, reinforcing the supposed legitimacy of the expectations it imposes or conclusions it reaches, no matter how biased, unreasonable, or unsafe. Third, workers have no ability to learn how the AI tools work in detail or what statistical modeling underlies their operation; they are inscrutable in a way that human decision-making is not. Fourth, even if workers can adjust to the AI system's expectations or discern how it functions, AI can be changed, creating permanent unpredictability.

Fifth, AI surveillance is far more invasive than human eyes or passive cameras, as every movement or word may be subject to capture and evaluation. This degrades worker autonomy and dignity, as they are now made to fear that any deviation from the AI system's expectations may lead to negative consequences. Sixth, AI alters how and when workers can interact with each other, making the entire workplace hostile to discussions about improving working conditions. For example, even if employers cannot directly view or listen to activities in bathrooms, they can see who entered the bathroom near the same time and how long they stayed there.¹⁵⁸

On top of these many risks for workers, AI also allows employers to claim increased distance between themselves and oversight of their workers. In some situations, including with ride-hailing platforms, this allows employers to more plausibly characterize employees as independent contractors and thus relieve themselves of regulatory or contractual obligations to pay workers overtime, pay payroll taxes that support Social Security or Unemployment Insurance programs, or provide workplace benefits such as health insurance.

Though Amazon's warehouse workers earn slightly more money on average than many other low-wage jobs,¹⁵⁹ their work exemplifies AI-based management practices.¹⁶⁰ Of Amazon's 1.1 million employees based in the U.S., over 500,000 work in its warehouses fulfilling customers' orders. The workers are given handheld scanners that track both workers' body movements and the items they move. The information collected leads directly to automated write-ups when the workers do not meet productivity goals. In addition, cameras and wearable badges monitor employee activity and location at all times.

These technologies undermine workers' health and rights. Amazon workers report that the omnipresent surveillance practices discourage them from taking breaks, using the restroom, eating, or drinking water. The surveillance is also associated with high rates of injury. While warehouse work generally has higher workplace injury rates than other industries, Amazon's warehouses have injury rates over 20 percent higher than the warehouse average.¹⁶¹ Furthermore, the technology is deployed in the context of Amazon's extensive efforts to oppose union organizing. If nothing else, the technologies make it easier to identify and frustrate nascent organizing efforts.

AI-based surveillance and management like those in Amazon's warehouses have proliferated in other settings. A New York Times investigation demonstrated how these technologies are also applied to people in higher-paying jobs, including finance executives, therapists, insurance underwriters, hospice chaplains, and radiologists.¹⁶² These workers, too, felt similar consequences: wage theft, extended hours, fear of using the bathroom, anxiety, and loss of autonomy. Employers plan to refine such tools to further determine employee pay rates and bonuses.

AI surveillance and management dynamics reflect the larger trends present in many of the AI uses detailed in this report. In the end, the harms from AI may fall the hardest on low-income people and low-wage workers who lack the financial resources to withstand lost pay, reduced hours, and being fired. But the use of AI is not restricted to low-income people. This opens the possibilities for broader coalitions to oppose AI-based harms.

Electronic Visit Verification

AI exposure: At least 1.5 million low-wage workers

AI-based surveillance and management of low-wage workers coincide with public benefits programs through federally mandated electronic visit verification requirements. Electronic visit verification requires home care workers who provide home-based care to Medicaid recipients to electronically register through cell phone applications each work shift's start and end times and the location where such services are performed. Such systems are intended to reduce paper-based billing practices that sometimes lead to improper payments. Though federally required, each state chooses how to implement electronic visit verification, including by selecting the technology vendor who builds the application.

At least 1.5 million home care workers work with Medicaid beneficiaries¹⁶³ and are subject to these requirements. The national median wage for these care workers is \$16.13 per hour, with many states' wages falling significantly below that. Due to the part-time nature of much care work, the national median salary is \$25,015, low enough that half of care workers receive some sort of public benefits themselves. As with other low-wage jobs, significant gender and racial disparities are reflected in home care work. Yet they are starker here: About 84 percent of these care workers are women, 67 percent are people of color, and 32 percent are immigrants.

Were the requirements of electronic visit verification limited to logging the start and end times of shifts on a functional electronic platform, there would be few objections. But the inclusion of location-based information that is independently monitored and the ability to collect biometric information through voice and facial recognition adds an element of surveillance that undermines the privacy of both the workers and the Medicaid recipients they serve.¹⁶⁴

The recipients' independence, too, is threatened. Some recipients likened it to living inside an “invisible dog fence” for humans.¹⁶⁵ About 33 percent of Medicaid recipients report that electronic visit verification keeps them from going outside of their homes as often as before the system was introduced.¹⁶⁶ This is because the system often flags work performed outside of the home as possibly fraudulent and delays payments to care workers.

Indeed, the systems have also resulted in severe economic consequences, with location-based billing flags and various system failures causing workers to go unpaid for several weeks. Of course, when care workers are not paid, they cannot pay their bills and risk eviction, power shut-offs, late fees, and hunger. Meanwhile, the Medicaid recipients also may lose care they need to survive.

Additionally, electronic visit verification systems pave the way for more pernicious uses. With such tracking, state Medicaid agencies may be able to get data they could misinterpret or misuse to suggest that recipients need less care than they are currently receiving, thus justifying further cuts. Even if this is not in states' short-term plans, the use of these systems establishes pervasive monitoring of Medicaid home-care recipients as a legitimate governmental function that can be expanded in time.

K-12 Education

AI exposure: 13.25 million low-income students

Of the roughly 73 million children in the United States, about 26 million, or 36 percent, live in families with incomes below 200 percent of the federal poverty level.¹⁶⁷

This economic hardship is reflected in the U.S. public school system. Of the 48 million children in public K–12 schools, about 25.5 million, or 53 percent, live in families with incomes below 185 percent of the federal poverty line, which qualifies them for free or reduced-price lunches.¹⁶⁸ These numbers also correspond to the roughly 26 million kids enrolled in so-called Title I schools that serve a high number or percentage of low-income students.¹⁶⁹ Whether looking at school meal programs or Title I schools, low-income students are disproportionately Black and Hispanic.¹⁷⁰

Low-income students experience vast gaps in academic performance compared to higher-income children, whether that is measured by average literacy rates, reading and math test performance, or school dropout rates.¹⁷¹ This gap is not caused by students' relative abilities, but rather by structural factors, such as poverty's impact on childhood development, inadequate teacher training and retention, and, in particular, inequitable school funding schemes that provide too few resources to schools with high numbers of low-income students.¹⁷²

Expanded pre-K classes (or pre-Kindergarten), comprehensive student enrichment programs, and increased efforts by schools to better engage parents—all of which could be possible through improved school financing mechanisms—prevent the gap from growing and, in some cases, narrow it.

However, facing immediate challenges and knowing structural financing reform is unlikely or slow, many public schools turn to ever-expanding AI use that ultimately endangers the very students it purports to help. There are three broad categories of AI usage. First, districts use data to supposedly predict the likelihood of students dropping out. Second, they do similar things to supposedly predict future criminal activity, often through expansive data-sharing agreements with other agencies or law enforcement. Third, districts use constant surveillance to punish students. As shown below, around half of low-income students are exposed to at least one of these AI uses.

The three AI uses combine to sustain or expand the school-to-prison pipeline that targets low-income children, disabled children, and Black, Latino, and American Indian children for heightened school discipline, including suspensions and expulsions, compared to higher-income, nondisabled, or white students.¹⁷³ This discipline, in turn, increases students' likelihood of involvement in the juvenile and adult criminal legal systems.¹⁷⁴

Labeling Students as Dropouts

AI exposure: 13.25 million low-income students

Many schools, school districts, and state education agencies have, in the past two decades, used AI early warning systems to try to predict whether individual students would graduate high school on time. The flagged students are supposed to receive some kind of intervention, which can take various forms, including meetings between teachers and administrators, parent-teacher meetings, individualized plans, chances to make up missed credits, mandatory tutoring, case management, or adult mentors.¹⁷⁵

As a baseline, these systems focus on individual students' attendance, behavior, and performance. Schools with higher concentrations of poverty tend to look at more outside-of-school factors, such as homelessness, than lower-poverty schools but otherwise are not more likely to use these systems. While these systems sound benign or even helpful, they often prove harmful in practice.

By 2015, 52 percent of public schools had some such system.¹⁷⁶ Assuming that those 52 percent of schools have 52 percent of students,¹⁷⁷ that means roughly 25 million students have their likelihood of graduating or dropping out ranked based on their personal characteristics. Assuming, too, that the same proportion of low-income kids are in the pool of students at public schools where these systems are used as are in the general student population at public schools (53 percent), it follows that 13.25 million low-income children are exposed to these predictive AI systems.

Wisconsin's experience illustrates these systems' limitations. The state government developed an AI system called the Dropout Early Warning System that used, among other factors, test scores, disciplinary records, race, disability, native language, and receipt of free or reduced-price lunches to flag students between sixth and ninth grade supposedly at risk of dropping out of high school.¹⁷⁸ In the 2020–2021 school year, more than 32,000 students were labeled “high risk.” The state would update the list of flagged students twice per year, and about 66 percent of school districts would access the list at least once per year.¹⁷⁹ After a decade of use, scholars concluded that the system may have contributed to a 5 percent increase in the likelihood of on-time graduation but did not conclusively do so.¹⁸⁰

The dubious benefits, though, are accompanied by significant risks. Black and Latino students were flagged as likely dropouts at rates far exceeding their white counterparts, with the students noting that they felt that

teachers would pay less meaningful attention to them or wrote them off as being “destined to fail.”¹⁸¹ This is particularly problematic when 8 out of every 10 kids flagged as high risk actually graduated on time, apparently without extra support.

Being subject to the derision of the AI system without the support it promises highlights another dangerous dynamic of its use: The purported promise of AI is not realized in practical implementation. Districts and teachers in Wisconsin complained about the lack of instruction by the state on how the AI system actually works (i.e., what factors it looks at and what weights it gives the factors) and how to use the predictions it makes. As with other contexts detailed in this report, there was no meaningful human oversight of the AI tool.

These are not the only risks. The AI early warning systems for schools are more costly and not significantly more effective than using an alternative approach that looks at environmental factors rather than the personal data profile of individual students. Moving to this alternative approach would remove the harm for students flagged due to their individual characteristics, save money, and still provide an effective way to prioritize kids who need additional interventions to graduate.

Even this, though, is misguided in light of bigger considerations about structural inequality and the segregation of high-poverty schools. The scholars who studied Wisconsin’s system said it well:

[T]he bottleneck for resolving low graduation rates is not identifying students at high risk of dropping out but, instead, overcoming structural barriers to accessing well-resourced schools and neighborhoods. In the case of Wisconsin, dropout is disproportionately concentrated in working-class, urban districts where the overwhelming majority of students are Black or Hispanic. In fact, 10% of all state dropouts come from just five schools. Bringing graduation rates within these five schools to the state average would have an outsize impact on improving the statewide graduation rate. Compared to investing in other community-based social service interventions, funding and implementing sophisticated early warning systems without devoting resources to interventions tackling structural barriers should be carefully evaluated in light of these school-level disparities.

Labeling Students as Criminals

AI exposure: At least 11.7 million low-income students

This area is a particularly pernicious expansion of the early warning systems that predict student dropouts. Here, school districts enter into expansive data-sharing agreements with other government entities, like child welfare and law enforcement agencies. The combined data is then fed into an AI tool that predicts individual students' likelihood of involvement with the criminal or juvenile legal systems. Once flagged, law enforcement officers harass students and their families under the guise of checking up on them and preventing crime. Notably, funds for technology to implement the data-sharing agreements and AI targeting systems have come from the U.S. Department of Justice.¹⁸²

The county of Pasco, Florida offers a notorious example of AI used by law enforcement to track 18,000 kids and label at least 420 of them as likely criminals.¹⁸³ The Pasco County Sheriff's Office's "intelligence-led policing" scheme employed AI that incorporated "educational risk factors, criminogenic risk factors, and adverse childhood experiences."¹⁸⁴ The educational factors came in part from the school district's early warning system and had similar factors as to those listed above. Criminogenic factors looked at past arrests (not necessarily convictions), status as a "certified gang member," having "delinquent friends," and a host of factors related to trauma, such as being in a custody dispute, being a victim of personal crime, or running away from home. The tool's focus on adverse childhood experience likewise drew adverse inferences based on a student's trauma, including experiencing physical, emotional, or sexual abuse, witnessing household violence, or having household members who were incarcerated or had substance use disorders. Ultimately, this represents the weaponization of tracked trauma, whereby students could be labeled as likely criminals because their lives were difficult and filled with pain.

The label came with severe consequences. As the *Tampa Bay Times*' investigative reporting showed, officers from the sheriff's department would "swarm homes in the middle of the night," "write tickets for missing mailbox numbers and overgrown grass," "come again and again," and "mak[e] arrests for any reason they can."¹⁸⁵ They would show up at the workplace of targeted students' parents, search for students when they

were at friends' houses, and aggressively question the students themselves without any suspicion that the student had done anything illegal. One deputy admitted that his orders were to "make their lives miserable until they move or sue."

Law enforcement officers offered no meaningful resources to support these students. After a major news investigation, the officers started handing the targeted children notecards with phone numbers to local nonprofit organizations. But no substantive support was ever provided.

While Pasco's example may be shocking, it is not unique. In a 2023 survey about educators' knowledge of and concern about technology used in educational settings, 38 percent of responding teachers reported that their schools share student data with law enforcement, and 36 percent reported that their schools analyze student data to predict which students are more likely to commit a crime.¹⁸⁶

Breaking out the responses by socioeconomic status and disability reveals a stark, if unsurprising, disparity for low-income and disabled kids. Almost half (48 percent) of teachers in Title I schools that serve high numbers or percentages of low-income kids reported that their schools share data with law enforcement, and 45 percent reported that their schools use crime prediction tools. In contrast, only 24 percent and 28 percent, respectively, of teachers at non-Title I schools reported that their schools use these technologies. Licensed special education teachers (who primarily teach students with disabilities) reported even higher use of these AI tools—57 percent and 55 percent, respectively—which contrasts with teachers not licensed in special education at 21 percent and 21 percent, respectively.

As noted above, there are 26.1 million kids enrolled in Title I schools. Assuming that the 45 percent of Title I schools using crime prediction AI correspond to student populations, more than 11.7 million students in schools serving primarily low-income students would be subjected to this AI, though they might not all be low income. Even so, there are low-income students in the 28 percent of non-Title I schools that also use crime prediction AI. Thus, 11.7 million appears a reasonable estimate of AI exposure.

Even more students would be exposed if not for the work of affected families, concerned community members, and advocates, who, for example, have recently rolled back the usage of these tools in Pasco¹⁸⁷ and prevented the planned uses of such technology in St. Paul, Minnesota.¹⁸⁸

Surveilling Kids Generally

AI exposure: Unknown but ubiquitous

Schools monitor student behavior in various, tech-enabled ways that lead to increased discipline for low-income kids. Schools now track students' locations via their cell phones, use surveillance cameras and drones with facial recognition technology to monitor who comes to or is at the school, use emotion detection software to determine if students are being or likely to become aggressive, review students' social media posts, block access to disfavored information, control the ways students use school-issued laptops, and monitor students' cell phone activity even when they are not in school.¹⁸⁹ Violations of the rules these technologies enforce can result in discipline.

These technologies extend and expand upon previous zero tolerance school discipline practices¹⁹⁰ and the increased presence of law enforcement officers in schools (sometimes called school resource officers).¹⁹¹ As with these past practices, AI systems disproportionately target low-income, Black or Latino, and disabled students, thus sustaining the school-to-prison pipeline.

The related harms are manifold and extend even to those who are not subject to disciplinary measures. As advocate Clarence Okoh notes, these AI-related technologies expand state censorship, undermine student wellness and safety, negatively impact academic success, and erode student privacy and trust.¹⁹²

This report does not provide a figure for AI exposure on these surveillance technologies because they are clearly ubiquitous and quantifying the particular uses is too complex for the purposes at hand.

Language Services

AI exposure: Almost 5 million low-income people

People who have a limited ability to read, write, speak, or understand English are termed people with limited English proficiency. Although many people with limited English proficiency have some, or even significant, facility with English, they are best able to understand or communicate in a language other than English for official matters with high stakes.

Federal laws require that all agencies administering public benefit programs, public housing authorities, schools, hospitals, medical clinics, law enforcement agencies, and courts offer “meaningful access” to people with limited English proficiency.¹⁹³ As of 2021, roughly 26 million people in the United States over the age of 5, or 8 percent, were considered to have limited English proficiency.¹⁹⁴ About 4.9 million of them, or 19 percent, are considered low-income.

Meaningful access requires, at minimum, timely and accurate interpretation and translation of vital information. Interpretation refers to verbal communication, while translation refers to written communication. What information is considered vital depends on the specific context of the communication. In benefits, for example, it includes program rules, key forms, written notices about a person’s application or eligibility, all interactions with agency staff, administrative appeals processes. In health care, vital information is that needed to understand health conditions, available treatment options, and risks.

Applicable laws allow the entities who must furnish meaningful access certain flexibilities based on how frequently the language is spoken in given regions and some practical obstacles in finding available translation or interpretation services. However, when proper access is provided, a person with limited English proficiency receives essentially the same services and same information in the same general timeframe as a fluent English speaker.

Many entities fail to invest in proper language services—people trained to interpret and translate complicated terminology—to meet these requirements. This failure forces people with limited English proficiency to experience longer wait times in the office or on the phone, use children or adult family members to interpret or translate, or go without the same understanding of vital information as people who speak English fluently. The results depend on context. In benefits, this results in significant confusion, improper denials, overpayments, false accusations of fraud, extra administrative burdens, and being frustrated out of even trying to get benefits. In school settings, parents might misunderstand information about their children’s progress or education plan. In law enforcement or court settings, domestic violence survivors may not know how to proceed or may have their statements misunderstood.

Entities often turn to AI tools such as Google Translate for official communications without additional vetting by trained translators. Google itself decries such uses, disclaiming that it is “not intended to replace human translators.”¹⁹⁵ This is because the technology is not precise exactly when precision is needed. The technology may get words, phrases, or meanings wrong in high-stakes situations, even with languages spoken by significant numbers of people in the United States, such as Spanish.

Automated translation or interpretation poses even more problems with languages that are spoken by fewer numbers of people. In these situations, the person’s primary language might have to be translated into a second language before being translated into the end language. For example, an Indigenous language from Guatemala may need to be translated into Spanish before being translated into English. Each additional layer of translation can further undermine accuracy.

It does not appear that any study has determined what percentage of entities who are obligated to provide meaningful language access use automated translation. Anecdotal evidence, at least in the public benefits context, suggests widespread use. For example, advocates across the country have needed to advocate forcefully to ensure that agencies meet these requirements, even in states such as California and Washington that have high levels of people with limited English proficiency.¹⁹⁶

Studies from other areas also suggest widespread government use without human vetting. For example, a study that surveyed automated translation use in October 2021 demonstrates that only around half of the states used professional translation services to convey information about COVID-19 vaccine safety, efficacy, and access.¹⁹⁷ Even with these translation services, human translation was limited to certain languages.

The only low-income people with limited English proficiency who are completely insulated from the dangers of automated translation are those where all entities obligated to provide meaningful language access use human translation for all vital information. I believe this to be unlikely and so estimate that all low-income people with limited English proficiency people are exposed to the risk.

Turning to interpretation, automated interpretation is still in its nascent phases and, to my knowledge, is only recently being piloted, including in agencies that administer public benefits. The essential idea, though, is that a person speaks in their preferred language, a speech recognition system turns that into written language, the written language is passed through an automated translation system to turn it into the recipient's language, and then the written communication in the recipient's language is spoken aloud by text-to-voice systems. Obviously, each step of such a process is fraught with possibilities for inaccuracies.

Domestic or Intimate Partner Violence

AI exposure: More than 2 million low-income people

Past and present romantic partners face staggering levels of physical violence, sexual violence, psychological aggression, and stalking—experiences that are collectively known as domestic violence or intimate partner violence.¹⁹⁸ Approximately 47.3 percent of women and 44.2 percent of men in the United States reported experiencing physical or sexual violence or stalking by an intimate partner at least once in their lives (the number of people experiencing only psychological aggression is not included in this count). About 42 percent of women and 42.3 percent of men reported experiencing physical violence at least once in their lives, with the numbers for severe physical violence—being hit, slammed against something, choked, burned, cut, or shot (this category does not include being slapped, pushed, or shoved)—being 32.5 percent of women and 24.6 percent of men. In a given year, 7.3 percent of women (around 9 million) and 6.8 percent of men (around 8 million) reported experiencing physical or sexual violence or stalking, while 3.1 percent of women and 3 percent of men reported experiencing severe physical violence.

Experiences of intimate partner violence turn into ongoing physical or emotional suffering, disabilities, bills for medical treatment, lost work, schooling, or housing, decreased social contact, increased exposure to law enforcement, and increased need for social services and legal advocacy. And, with many people, including with several clients represented by my colleagues at Legal Aid of Arkansas, experiences of intimate partner violence result in murder.

While intimate partner violence affects people at all levels of socioeconomic status, documented intimate partner violence rates are highest among low-income people. A Seattle-based study comparing intimate partner violence rates by neighborhood income found that the prevalence in the lowest-income neighborhoods was nearly three times that of the highest-income neighborhoods (of course, it is also possible that intimate partner violence is underreported in areas with higher incomes).¹⁹⁹ Similarly, a review of national survey data found that people with other indicators of economic insecurity, such as food and housing insecurity, experience intimate partner violence at rates between two and three times as high as people who did not have such indicators.²⁰⁰

There are two common explanations for this. First, severe economic strains can increase the potential conflict in a relationship. Second, low-income survivors often do not have economic means to live independently from their abusers. They might not have been allowed to work or, even if they were, might have job instability based on abuse-related absences and abuser-related disruptions at their workplace. They might not have access to their earnings or bank accounts. And, with abuser dynamics often designed to isolate survivors from their support networks, they might not have family or friends to turn to for financial support or a place to stay.

While shelters are a temporary option for some survivors, conditions vary from place to place in how long survivors can stay, whether there is a curfew, and what ages of children are allowed. Moreover, many shelters are inaccessible to disabled people, who experience higher rates of intimate partner violence.²⁰¹

Legal protections for people experiencing domestic violence exist but have significant limitations. Acts of sexual or physical violence generally violate criminal laws. But in addition to the immediate danger posed by the abusive partner, calling law enforcement brings the risk of mutual arrests—arresting both abuser and survivor—or no arrest at all. If no arrest happens, the abuser may become enraged and emboldened. Even when only the abuser is arrested, prosecutors vary from place to place in how willing they are to prosecute these crimes.

Moreover, apart from the interruption of ongoing violence, the arrest or prosecution of a partner may go against the survivor's long-term interests or desires. That is, an abuser in jail cannot work to cover bills, provide child support, or spend time with children, which are all things that a survivor

might want. In light of such challenges, some thoughtful advocates argue that simply putting more people in jail may not keep anyone safer in the long run. As such, they advance responses that do not involve the threat of incarceration, called noncarceral approaches, to ongoing intimate partner violence.²⁰² Such approaches often center on restorative or transformative justice principles, though further discussion of them is beyond the scope of this report.

Using the civil court system, survivors may seek a legal order that prevents the abuser from coming to certain locations, including the survivor's home, a workplace, or a family member's home. These orders are usually called an order of protection or protective order. In these proceedings, courts can also address related issues, like custody of children, visitation arrangements, and child support. Abusers who have an order of protection against them cannot have firearms.

Practically, however, orders of protection have limited efficacy. First, starting a court case requires a trip to the courthouse during business hours to file lengthy paperwork open to the public that describes what happened. It is both practically and emotionally difficult and carries the risk of public scorn or shame. Second, seeking an order of protection is dangerous, as abusers often respond violently. Third, it can cost more money than the survivor has. While courts generally should waive a filing fee to open a case, it can cost money to have the abuser served with papers, get medical records, photos, or other evidence, and, most critically, to hire an attorney (though some legal aid programs and domestic violence support organization will provide free legal representation). And the survivor also must take time off of work and arrange child care.

On top of all that, the trial itself is a harrowing experience. With the abuser staring at them, the survivor must testify about what has happened and tell a complicated story, sometimes spanning years of trauma, in a way that makes sense to the judge and comes across as sincere. Then, the abuser or their attorney can ask the survivor questions, which can be aggressive, intimate, intimidating, and false. The survivor must arrange for witnesses to testify, which can be difficult for emotional and practical reasons, like having to get subpoenas served. Then, the abuser can testify, usually with lies, partial truths, and accusations that the survivor actually prompted or contributed to any violence such that the abuser was acting in self defense. If kids are involved, the trial will often involve accusations that the survivor is not a good parent and cannot properly support the kids. The experience of a trial heaps trauma onto already-traumatized survivors.

Even when the survivor puts on a strong case, judges' personal opinions about what a proper relationship should look like, what proper gender roles are, and the consideration of consequences to the abuser may result in a survivor's loss. I have personally represented survivors with strong cases who lost because, for example, a judge believed that God ordained men to be the leader of households with power over their wives. In another former case, a judge was worried that the order of protection's firearm prohibition would prevent the abuser from enjoying the upcoming deer-hunting season and thought such a punishment was too severe for the physical violence involved.

If the survivor manages to overcome all the obstacles and win, they get a piece of paper that prohibits the abuser from contacting them or coming to certain places. Abusers often violate the order. Law enforcement officers generally respond with little urgency and few consequences, even when the abuser continues to escalate with more daring violations. And prosecutors often do not want to pursue criminal charges for violations.

Survivors who do not get adequate law enforcement response often feel disappointed and are left to wonder whether going through the order of protection process was worth it. In such cases, their only recourse is to go back to court to ask the judge to hold the abuser in contempt for violating the order. This, again, requires money for an attorney and related costs. If successful, an abuser held in contempt of a court order may have to spend time in jail, pay fines, pay for the survivor's attorney, and be subject to ongoing court oversight.

AI is now entering these dynamics in multiple ways. Law enforcement officials use AI to decide whether particular survivors face enough risk of violence to merit additional law enforcement involvement. Abusers use AI to coerce survivors to stay in relationships or create or falsify evidence.

Predicting Additional Violence

AI exposure: More than 2 million low-income people

Law enforcement agencies in the United States and abroad sometimes use AI to predict the risk that a survivor will face additional violence by their abuser.²⁰³ In these risk assessments, the law enforcement official asks the survivor a series of questions and records the answers, which can be run through an algorithm to classify the survivor as being at high, medium, or low risk of suffering future harm.

Law enforcement agencies then use the risk classification to decide whether to provide additional attention to the highest-risk survivors, such as extra home visits by special domestic violence officers²⁰⁴ or coordinated case management between law enforcement and survivor-focused social services. In this context, the AI system is supposed to enable the law enforcement agency to direct limited resources to the people in greatest need. In addition, risk classification AI can be used to decide which survivors are exempt from certain restrictive criteria for public benefits, like requirements to report work or to seek child support from the other parent.²⁰⁵

The extent of such systems in the United States is not precisely known. A 2015 survey of law enforcement agencies found that 42 percent used some sort of structured risk assessment for intimate partner violence, though the technical sophistication of the systems is not clear.²⁰⁶ Assuming that the law enforcement agencies using the assessments correspond to U.S. population totals as a whole, and that people who experience intimate partner violence are distributed equally throughout the country, then approximately 7.1 million people who experience intimate partner violence in a given year are exposed to such tools. Leaving aside the increased prevalence of such violence experienced by low-income people and just applying the share of low-income people among the population (28.3 percent), at least 2 million low-income people are exposed to such tools every year.

Various AI tools exist with different focuses and levels of validation.²⁰⁷ For example, two leading tools—the Lethality Assessment Program and the Domestic Violent High-Risk Team—are directed toward preventing homicide, with screenings resulting in social service coordination and, potentially, increased monitoring of the abuser.²⁰⁸

These systems prompt the same general concerns as other AI systems. Most fundamentally, the AI systems may not be accurate, effective, nondiscriminatory, transparent, or subject to proper oversight and accountability.²⁰⁹ AI may not estimate risk in a way that ensures that survivors actually at risk receive the additional available resources. Independent of how the AI tool is built to function, there are questions about the quality of the information the AI acts upon. This is particularly doubtful in the context of intimate partner violence. Due to inexperience, workload, lack of training, or neglect, police officers may not be interviewing victims in a trusting, compassionate manner that encourages complete responses. Biases also may affect how the officers mark assessment responses. For their part, victims, particularly those who just

experienced an attack, may provide incomplete information due to trauma, exhaustion, fear of reprisals, shame, or dependency on the abuser. These systems also can increase undesired contact with law enforcement, which can have other negative consequences for victims or their families, such as child welfare checks.

Overall, AI is not likely to address the insufficiency of existing societal responses to intimate partner violence or the dearth of resources directed toward preventing abuse and protecting survivors.

The professed utility and actual perils of such a system are manifest in Spain's use of the VioGén algorithm in its national police force and many independent local police authorities.²¹⁰ Survivors at the highest risk levels receive regular police patrols by their homes, access to a shelter, and police monitoring of the abuser's movements. In use for nearly 20 years, government officials credit the system with lowering the frequency of repeat attacks from 40 percent in 2007 to 15 percent in 2024. Proponents note that, prior to the AI system, the services that survivors would receive depended mainly on the experience and inclinations of individual officers.

Despite the official enthusiasm, the system appears to chronically underestimate risk. Most gravely, of the 98 women killed between 2010 and 2022, 55 had been classified in the lowest-risk categories. On a broader scale, 22 percent of women in the lowest-risk categories have reported being harmed again. Indeed, downplaying risk seems to be a feature of these AI tools, not a bug. An outside audit (not commissioned by the government) that uses what limited data is publicly available reveals that, among other issues, 87 percent to 94 percent of all cases are classified at the lowest-risk levels and fewer than 1 percent are classified at the highest level of risk.²¹¹ Few survivors knew what their risk score was or how it was reached.

In light of what is known about the dynamics of intimate partner violence, such a low figure rings absurd. Instead of predicting risk, the system seems to be a simple gatekeeping tool to restrict services only to the select cases that the government deems merit additional resources or for which additional resources are available. The supposed predictive function seems little more than a fancy way to say “no” to survivors in need.

The lack of candor denigrates the experience of survivors. They likely fear for their safety and perceive themselves to be at high risk of further violence. But the AI system dismisses this fear, embedding a coded form of longstanding gender stereotypes that women—traditionally the primary target of intimate partner violence—overreact, are emotional, or suffer

from hysteria. In truth, only few survivors ever have a chance of official validation of their risk levels. The rating carries other consequences, too, as it is used in judicial proceedings to decide whether the survivor will receive an order of protection and whether the abuser will be allowed to see any children.

Meanwhile, human oversight of the AI tools is not effective. While officers can overrule the AI classification to assign higher risk levels to individuals (they cannot assign lower risk levels), they do so only 5 percent of the time. This suggests either that the AI has displaced professional judgment as the criteria for assessing individual risk or that officers learn to accept the results for budgetary reasons regardless of their personal risk evaluation. What is more, the government will not release comprehensive data about the system to allow meaningful public oversight or full outside audits.

Despite facing fundamental problems, officials are now considering applying more sophisticated AI techniques to VioGén.²¹² Such black-box methods can only intensify the concerns noted above.

Using Deepfakes to Coerce Survivors

AI exposure: Unknown but ubiquitous

A deepfake is an AI-manipulated image, video, or audio recording that purports to show a person doing or saying something that the person in reality did not do or say.²¹³ In the context of intimate partner violence, this often takes the form of an abuser creating a pornographic video featuring the survivor's likeness. After creating the video, the abuser threatens to release it unless the survivor fulfills the abuser's demands. The possibility of shame, damage to a person's reputation, job loss, and ongoing harassment coerces the survivor's compliance.

Estimating the number of survivors that are the subject of deepfakes is complicated and, with the pace of technological advancement, unreliable. Indeed, as of 2020, experts estimated that 85,000 harmful deepfake videos were in circulation, and that number was doubling every 6 months.²¹⁴ At such a pace, more than 20 million videos would have been in circulation by the start of 2024. While the subjects of deepfakes are primarily politicians and celebrities (usually porn involving female celebrities),²¹⁵ use for intimate partner violence is common.

Deepfakes are easy to create. To wit, during a cybersecurity conference in the summer of 2024, a presenter ran a public booth that could contemporaneously create convincing video and audio deepfakes of willing participants “using just a consumer-level gaming laptop, a basic DSLR camera, some lighting, and a green screen.”²¹⁶ Whatever clues of manipulation existed could be cleared up with more time or modestly better technology.

Using Deepfakes as Evidence

AI exposure: Unknown but ubiquitous

Proving that a survivor actually experienced violence can be difficult. Some abusers limit their physical violence to actions that do not leave obvious or widespread bruises. Often, they will refrain from violence in the presence of potential adult witnesses, particularly those who may be sympathetic to the survivor, and children may be too young to serve as witnesses. Even if they are of age, survivors often do not want to put their kids through a trial where they will have to testify against the other parent.

If survivors manage to take photos of visible injuries, the photos may not be clear or be easy to understand. Or they may depict violence that happened several months before seeking an order of protection, which allows judges to infer that the threat of further violence is not immediate. Abusers also may control the survivors’ phones such that evidence cannot be stored there.

Experiencing physical violence is not the only way a survivor can obtain an order of protection. Sometimes, evidence of violence comes in the form of threats written in text or made verbally and recorded.

With deepfakes, abusers have the ability to generate false evidence against the survivor. An abuser could use deepfakes to suggest the survivor actually initiated the violence, is not as afraid of the abuser as they claim to be, is seeking an order of protection for some improper purpose, or is trying to manipulate witnesses into saying false things. Or, instead of generating false evidence themselves, abusers can cast doubt on survivors’ evidence by claiming that it is faked.

Either way, the court system is not set up to deal handily with deepfakes.²¹⁷ Existing procedures to authenticate videos, photos, and texts do not readily apply to deepfakes because judges, lawyers, and witnesses cannot readily spot them. As a result, judges will not be able to easily exclude deepfake evidence or admit genuine evidence that the other side contends is a deepfake. The party defending against the deepfake then has to hire digital forensics experts, which is estimated to cost between \$5,000 and \$15,000 for routine matters. These costs are beyond what many survivors can afford.

Child Welfare Systems

AI exposure: At least 72,000 low-income children

The term child welfare system (also called the family policing system²¹⁸) includes all aspects of a state or local government's activities regarding the alleged maltreatment of children. Maltreatment includes physical, sexual, or emotional abuse and an expansive definition of neglect that penalizes people living in or near poverty.

Child Protective Services, or CPS, agencies, often based in a state's or county's Department of Human Services, generally run the local child welfare system. While essentially all day-to-day matters are in the hands of state or local officials, the federal government provides some funding support that comes with broad legal requirements.²¹⁹

State CPS agencies investigate reports of child abuse or neglect, remove children from their homes, find alternative places for them to stay either temporarily or permanently (including with parents, with other relatives, or in foster care), and run foster care systems, which includes recruiting, supervising, and funding foster care providers. CPS agencies generally have insufficient funding to meet the scope of their mandate, which leads to high caseloads for workers, stress, and yearly turnover rates estimated to be as high as 40 percent in some places.²²⁰

Judges are also involved in the child welfare system. Judges in state courts determine where abused or neglected children will be placed, decide whether parents will maintain the right to raise their children or have these rights terminated, and supervise the process for adopting children.

In 2022, the most recent year for which data is available, roughly 4.3 million referrals (submissions of information about possible maltreatment) were made to CPS agencies, pertaining to 7.5 million children.²²¹ Of these referrals, about 50 percent had enough information to substantiate possible maltreatment that they became reports meriting investigation, covering roughly 3.1 million children. The investigations determined 560,000 kids to be victims of abuse or neglect, with the remaining approximately 2.5 million kids either being part of households where no maltreatment was found or where parents voluntarily accepted CPS services in return for halting any possible investigation. 85 percent of families involved with the child welfare system have incomes below 200 percent of the federal poverty line.²²²

The children determined to be victims of abuse or neglect are often removed from their homes and placed in foster care either with a relative (if someone is willing to care for the child and able to pass state caregiver requirements), a nonrelative family not known to the child, or in a group home or other institution. Siblings may not always be kept together.

Generally, services are supposed to be provided to children and parents throughout these processes to help keep families together and support their stability. Practically, however, whatever services are provided lead to reunification with a parent for fewer than 50 percent of children exiting foster care in a given year.²²³ In fact, the median time a child spends in foster care is about 18 months, with the average length higher, at 23 months.²²⁴

The upheaval that children experience by being removed from their parents' care and placed in foster care—which itself carries risk of abuse and further instability—can reverberate throughout a child's life into adulthood, prompting “increased risk for long-term emotional, behavioral and physical health problems, among other challenges.”²²⁵ Indeed, kids removed from their homes are “more likely to be incarcerated, become teen mothers, experience poor outcomes related to cognitive development, education, and employment, and have a higher likelihood of having mental health conditions and substance use disorders” than peers who do not get removed from their parents' care.²²⁶

Low-income people are disfavored in the child welfare system throughout the referral, investigation, and court processes. Allegations of neglect, which make up 75 percent of all maltreatment reports, are based on broad definitions, such as failing to provide food, clothing, shelter, or medical treatment or failing to provide for the child's essential needs.²²⁷ These are the struggles of anyone living near poverty.

This nebulous standard then combines with more frequent interaction with mandated reporters. These are workers in social services, law enforcement, medical clinics, or schools who are required to report any suspected maltreatment. Accessing public benefits, using emergency rooms for medical care, riding public transit, and attending schools with increased law enforcement presence can mean that more eyes are focused more often on low-income children than their higher-income peers, increasing the likelihood that someone makes a report.

Open investigations invite other problems for low-income people. CPS workers generally have broad investigatory powers and can invasively search a targeted family's home or pull kids out of school to talk with them. Racial biases can influence the determinations they make, which may contribute to the overrepresentation of Black, Hispanic, and Indigenous families in most stages of the child welfare system.²²⁸ Meanwhile, parents or other caregivers may not know that anything they say or do can be used to justify a CPS worker's decision to remove children from the home. Parents often do not have the right to an appointed attorney free of charge until a child is removed and court proceedings are officially underway.

The court system, too, brings its own set of challenges. While parents do receive legal representation without having to pay, the lawyers may be limited in what funds they have available to do their own investigations, interview the CPS agency's witnesses, get records, hire experts, and take other steps to make a stronger court case. Then, there's the question of what parents must prove to be able to keep their kids. Though the specifics vary by state, judges generally make key placement decisions according to the vague standard of "the best interests of the child" and, when doing so, do not have to expressly weigh the harms of removing a child from their parents' care or the details of the place the child will be placed.²²⁹ Such vast discretion also allows room for judges' biases to affect decision-making in defensible guises.

Relatedly, convincing the judge that neglect is unlikely to recur can be difficult based on simple economic realities. That is, parents may not be able to simply will themselves to a significantly better job that improves their ability to provide their kids with safer housing, more and better food, or reliable child care while the parents work. Additionally, court proceedings involving child welfare are generally closed to the public, meaning that there is no access to hearings, transcripts, case files, or court documents for anyone except the people directly involved. As a result, it's difficult for advocates, journalists, or researchers to track what is happening and bring attention to problems.

Overall, then, contact with the child welfare system constitutes a threat to family stability and the long-term well-being of low-income people. Of course, some children are in danger and require protection from dangerous adults in their lives. But the current system emphasizes over-reporting, over-investigating, over-removal from parental homes, and overuse of foster placements. The result may be more costly and traumatizing than alternative interventions that keep kids and parents together and that help address the poverty-related challenges that show up as neglect.²³⁰

Against this backdrop, CPS agencies are now using AI in ways that expressly target low-income families, including by introducing AI recommendations for abuse or neglect investigations and by maintaining automated child maltreatment registries that provide erroneous information to prospective employers.

Recommending Investigations

AI exposure: At least 72,000 low-income children

When CPS agencies receive complaints of alleged abuse (called a referral), they are usually required by law to investigate. But the process for reports of alleged neglect is more complicated. When someone makes a complaint of alleged neglect, CPS agencies have traditionally relied on information outside of the complaint to inform their review. Primarily, CPS workers consider past complaints, past investigations, and other child welfare system interactions involving the parent, child, or other family members. Based on this information, CPS agencies decide whether to investigate the specific neglect referral further.

Now, some CPS agencies are using AI to review similar information and generate a risk score or recommendation about further investigation. The stated purposes of these AI tools are to streamline resource use, reduce human bias, and help agencies make decisions that will more effectively minimize harm to children.²³¹ To these ends, the AI systems predict elevated risk for severe maltreatment and the likelihood of repeated maltreatment.²³²

A 2021 study showed that at least 26 states and Washington, DC have considered such systems and that jurisdictions in at least 11 states were using them, among them the state of Oregon, New York City, and Allegheny County, Pennsylvania.²³³ Some AI systems focus on specific characteristics of the individuals involved, while others look primarily to geographic locations and identify “hot spots” with heightened risk of maltreatment.

The most contentious example of this is the Allegheny Family Screening Tool used in Pittsburgh, Pennsylvania and the surrounding county, first implemented in 2016. For each complaint of alleged neglect, the AI tool reviews available data from the CPS agency and other government and social services sources about the family involved, considers identified variables, and assigns a “family screening score” between 1 and 20 that purportedly predicts “the long-term likelihood of home removal.”²³⁴ These scores are broken out into three categories: high, medium, and low risk.

The system looks at well over 100 variables, including (depending on the version) use of public benefits, medical histories, ages of people involved, the number and recency of other referrals, whether the family lives in a high-poverty area, and exposure to the criminal legal system as a child or adult, even if no charges or convictions resulted.²³⁵ According to the AI developers, the score does not reflect “anything about the current allegations” or “anything about whether the allegations presented meet the threshold for case opening.” In other words, the score is just about the family involved, not the about the alleged incident spurring the report.

Once generated, both the score and the risk category are seen by a human screener. Cases involving families scoring above a certain threshold are sent on for mandatory investigation unless overridden by the screener and their supervisor. Cases involving families scoring below that threshold are generally left to the screener’s discretion. If the screener decides to send the case on for investigation, the CPS worker who ultimately conducts the investigation is not supposed to have any knowledge about the family screening score. The score is also kept from the families.²³⁶

Studies critical of the Allegheny AI system have noted that the system features “arbitrary design decisions, give[s] families no opportunity for recourse, perpetuate[s] racial bias, and score[s] people who may have disabilities as inherently ‘riskier.’”²³⁷ On the whole, “Black households are, on average, assigned higher risk scores than non-Black households of the same size.” Focusing only on the “high risk” category, Black families were scored as “high risk” in 67.9 percent of cases, while white families were scored as “high risk” in only 50 percent of cases.²³⁸ With respect to disability, the system “penalizes those living in poverty, having less education, participating in benefit programs, and utilizing publicly funded health services—all measures that disproportionately affect disabled individuals and compound racial disparities.”²³⁹ Concerns about discrimination have prompted review by the U.S. Department of Justice, though the results of that review are not yet known.²⁴⁰

Media investigations underscore the threat that these systems pose to family stability. In one example, at their pediatrician’s urging, a couple took their 8-month-old baby to the hospital when she would not eat, stayed by her side for 2 weeks, and then, right when the baby was recovering, had her taken from them and placed in foster care.²⁴¹ The couple may have been flagged as “high risk” by the Allegheny Family Screening Tool in part due to the couple’s disabilities, which include stroke-related nerve damage, a comprehension disorder, and ADHD. Now, two years later, the child is still in foster care. She has been forced to move between multiple caregivers. Meanwhile, the parents receive a mix of supervised and unsupervised time.

Ultimately, child welfare tools run the same risks as many of the other AI systems identified in this report. First, they might not work as intended. Allegheny County officials confirmed that “a technical glitch in the tool sometimes presented social workers with the wrong scores, either underestimating or overestimating a child’s risk.”²⁴² Second, when they work as intended, these systems may perpetuate discrimination or cause other harms, such as improperly separating families from their children.

Third, the people impacted by the AI tools have no say over its use and, in the case of Allegheny County, do not even know their score. Fourth, the agencies using such tools are often uninterested in or incapable of overseeing them properly. As Los Angeles County’s CPS head stated, “How the algorithm functions, in some ways is, I don’t want to say is magic to us, but it’s beyond our expertise and experience.”²⁴³ Meanwhile, a Colorado CPS official shifted responsibility to a vendor, noting that “the variables and weights used ... are part of the code developed by [the vendor] and thus we do not have this level of detail.”

These dynamics do not bode well for further use of such tools. While Oregon abandoned its Allegheny-inspired AI system because of racial bias,²⁴⁴ Allegheny County continues to use it, multiple counties in Colorado have implemented it or are planning to, and Los Angeles County has an AI system developed by the Allegheny system's designer that has been piloted in a few areas, with expansion in the works.²⁴⁵ Meanwhile, other jurisdictions use different tools that raise similar concerns.

In light of continued interest in AI by CPS agencies, Allegheny County is noteworthy in that it publishes the factors its AI tool takes into account, indicates the relative importance of each, updates the AI regularly, enables staff to override the tool's recommendations,²⁴⁶ commissions outside evaluations, and allows independent researchers to access key information. These are basic steps that many government agencies using AI refuse or fail to do. As such, Allegheny County's AI system might represent the best extant example for continued AI use in the child welfare system, and it still has the manifold problems noted above. In contrast, locations implementing similar tools without the same measures are likely to cause more harm with even less accountability.

A note on counting

Calculating AI exposure in this context is difficult because there are different types of AI tools used, there is no certain way to know which AI tools are being used in which locations, and states and counties report data differently. To arrive at 72,000 kids exposed, this report does the following. First, it focuses only on the “family screening” type of tool made notorious in Allegheny County. Second, it uses only jurisdictions with known or planned uses of this style of tool, per a report by the American Civil Liberties Union, which presently is limited to counties in Pennsylvania (Allegheny), Colorado (Arapahoe, Douglas, and Larimer Counties), and California (Los Angeles). Third, it excludes referrals involving abuse and only uses referrals involving neglect, which is the only type of decision to which these tools are currently applied. Finally, where the number of referrals involving neglect is not directly available, as in the Colorado counties, it assumes that 75 percent of all referrals are for neglect. This

Continued on next page

A note on counting (continued)

rate comports both with national data on the proportion of neglect-based referrals and with available data for the three counties about the proportion of screened-in cases that feature neglect. Accordingly, the Colorado counties feature 18,845 neglect-based referrals,²⁴⁷ Los Angeles County features 42,160 such cases,²⁴⁸ and Allegheny County features 11,358 such cases.²⁴⁹ Notably, the count includes only children because the available statistics do not indicate the number of parents or other family members affected by the AI systems. The report also estimates that all of the affected children are low income. This is based on data listed above that 85 percent of all families involved with the child welfare system have incomes below 200 percent of the federal poverty line and the specific character of neglect cases that equate poverty with this form of maltreatment.

Maintaining or Operating Child Maltreatment Registries***AI exposure: Unknown***

Reports of neglect or abuse reverberate beyond where and under what circumstances the children will live. States also maintain child maltreatment registries where parents or caretakers accused of neglect or abuse may be listed indefinitely.²⁵⁰ AI may be playing a role in maintaining or operating these registries.

Any time a person applies for a job involving children or people of advanced age, employers conduct checks of the maltreatment registry. When this check shows a finding of substantiated abuse or neglect, the person is generally ineligible to be hired, no matter the severity of the alleged maltreatment, how long ago the finding was, or whether the alleged maltreatment implicates the job in any way.

Though the procedure in each state varies, placement on the registry generally happens through an administrative hearing process that features procedures far less formal than those in a court of law. One key feature of this is a reduced burden of proof, by which the state may establish that maltreatment occurred merely by showing only that it is more likely than not.²⁵¹

The administrative hearing process is plagued by other deficiencies that can lead people to be wrongly determined to have committed maltreatment. People accused of maltreatment may never receive notice of the allegations or a chance to appeal. Even if they do, the alleged maltreatment may be from years or even decades prior such that clear memories and witnesses are unavailable. Moreover, the hearing itself often features the procedural hurdles discussed elsewhere in this report: confusing procedures for presenting evidence that the allegation is wrong, a lack of affordable legal representation, procedures for questioning witnesses that are unfamiliar, and surprise evidence.

Depending on the state and the type of maltreatment, a person placed on the registry may never be removed at all or may only be eligible to be removed after a certain period of time has passed. Where removal is possible, a person may have to file a special petition with evidence that demonstrates the maltreatment is unlikely to happen again.

Given the dynamics of the child welfare system, it is unsurprising that low-income people, Black people, and women are disproportionately hurt by these registries.²⁵² Mainly, registry placement makes people unemployable in important fields open to low-wage workers, such as in school settings and either adult or child caregiving.

The full extent of the harms are shielded from public scrutiny by laws that keep maltreatment records and information about registries confidential. While it is thus hard to know if or how AI is at work, I have heard from advocates in at least three states that the states may be using at least basic means of automation for matching names similar to those found in tenant screening or other background reports. That is, a prospective employer sends a request to the state for a registry check, and the state is somehow automating the process of checking the registry and responding to the employer.

Finally, it is important to note that there are also adult maltreatment registries where caregivers of people of advanced age—principally, nursing facility staff and in-home caregivers—can be placed for similar allegations of maltreatment with similarly deficient procedures and similarly lasting consequences.

Other Uses and Harms

AI exposure: Unknown but ubiquitous

As noted in the introduction, there are many harmful or potentially harmful AI uses impacting low-income people that this report cannot cover in full depth because they are outside of my expertise, require more capacity than TechTonic Justice presently has, or are less directly tied to decisions or recommendations about low-income people. Accordingly, this report offers a brief overview with citations to sources that will encourage further exploration.

Of course, the same dynamics identified elsewhere in this report that fuel AI uses—hostility to low-income people, racism, sexism, homophobia, and ableism, among others—generally apply with similar force here. So, too, do the same concerns about perpetuating inequities, embedding discrimination, inaccuracies, lack of transparency, and inability to meaningfully resist harmful decisions or otherwise impose accountability on these systems.

In the criminal legal system, AI is used both for policing and for court determinations.²⁵³ Law enforcement agencies use AI to determine where and how they focus their attention, to surveil people in their neighborhoods, in their cars, and at stores, to supposedly predict where crime will happen and who will do it, and to write police reports that may be used in court proceedings. With the U.S. Supreme Court's elimination of the abortion protections enshrined in the case of *Roe v. Wade*, pregnant people learning about or seeking abortions and the advocates that support them also are now regular targets.²⁵⁴

Meanwhile, court officials, including judges, prosecutors, and parole officers, use AI to determine whether an accused person will be offered the chance for bail (bail allows the accused person to avoid ongoing confinement until the trial occurs), to determine the length of sentences for people who plead or are determined guilty, and to determine the conditions of release for people eligible for parole or probation.

In the immigration system, the U.S. Department of Homeland Security uses AI to make decisions about whether an immigrant will receive relief, such as asylum or citizenship, to determine whether and under what conditions to release someone who is detained, to decide where to focus enforcement resources, to oversee the border wall, and to surveil people.²⁵⁵

In tax enforcement, the U.S. Internal Revenue Service uses AI to target enforcement efforts, which, up to now, have disproportionately targeted low-income people in high-poverty areas who claim the Earned Income Tax Credit.²⁵⁶

In child support enforcement, state officials use AI to suspend driver's licenses of people who the state sees as delinquent in child support payments.²⁵⁷

With respect to credit, banks, mortgage companies, credit card companies, and credit reporting agencies use AI to set credit scores, to approve or deny loan applications, and to set the interest rates and repayment terms for approved loans.²⁵⁸

With respect to health care treatment, AI is used to transcribe conversations between doctors and patients, determine where to direct organ transplants, decide at what point of disease progression to give patients particular treatments, assess probable treatment outcomes, and interpret test results, among many other things.²⁵⁹

With respect to voting, state election officials use AI to automate voter registration purges, invalidate mail-in ballots, and suppress voter turnout through polling place decisions or rules about when people can vote.²⁶⁰ In addition, malicious actors are using AI to propagate misinformation about when and how to vote to foment increased aggression against election officials.

Part Two: Innovating Injustice

Given the ubiquity of AI-based decision-making in all aspects of life, it is hard to imagine any low-income people who are not in some way exposed. From the areas with the highest numbers of people exposed to the lowest, the theoretical low-income person unexposed to AI would:

1. not be enrolled in Medicaid, Medicare Advantage, private insurance, SNAP, Unemployment Insurance, SSI, or SSDI;
2. not rent a home from any of the 90 percent of landlords who use tenant screening or any of the growing-but-unknown number of landlords who use AI to set the amount of rent;
3. not work at any major chain retailer, any Fortune 500 company, any company with at least 500 employees, in the gig economy, or as a care worker to a person who receives Medicaid home-care benefits;
4. not have children that attend one of the of the 52 percent of public schools that use AI technologies to predict graduation or the 45 percent of public schools serving primarily low-income kids that use AI to predict future crime;

5. not need to speak a language other than English for formal matters at government offices, hospitals, medical clinics, courts, or public schools or with law enforcement;
6. not have experienced intimate partner violence in the past year, not have reported or been subject to a report of intimate partner violence in the past year, or, if they have, not have lived in one of the 42 percent of law enforcement jurisdictions that use AI to predict future violence or lethality;
7. not have been in a court proceeding involving intimate partner violence where deepfake technology was used to generate evidence;
8. not have been referred for possible child neglect in any of jurisdictions that use AI to predict parental likelihood of neglect; and
9. not have ever been placed on a child maltreatment registry or, if they were, not have applied for a job working in school settings or caregiving.

Even if a person were to somehow evade AI in all of these areas, they would likely be exposed to the ubiquitous AI surveillance systems in K–12 schools or around their homes.

Apart from the harms from being constantly surveilled, being exposed to AI does not lead to material harm in every instance. For a given person, AI might not have cut public benefits, denied housing, cost a job, reduced pay, or resulted in an adverse labeling at school or in the child welfare system. So, if harm is not certain, why is it important to describe the ways low-income people are exposed to AI decision-making and estimate the extent of any exposure?

AI marks a fundamentally new way of making decisions about core aspects of life and creates risks of harm orders of magnitude greater in scale and scope than what came before. Even when it does not make a wrongful decision about an individual, it carries the risk that such a decision will eventually be made, as in the myriad documented disasters reflected in this report. This freely happens because AI defies existing accountability mechanisms. And, more subtly, AI poses a conceptual threat to the very notion of accountability.

Understanding or changing conditions of poverty, inequality, and various forms of related marginalization is impossible without accounting for AI and its related complexities.

AI Harms People in New Ways

Whether hostile policymakers were trying to trim rolls or caseworkers made errors or acted on biases, public benefits recipients have always faced the risks of wrongful denials, terminations, or fraud accusations at both systemic and individual levels. Low-income people have similarly always faced the risks of harmful actions from bosses, landlords, school principals, child welfare investigators, or abusers and the forces that enable them.

How, then, is AI different?

While a full recounting of pre-AI decision-making for the areas specified is beyond the scope of the report, AI differs in at least eight key ways:

1. AI provides a cloak of unwarranted rationality to actions that intensify or perpetuate injustices.
2. AI expands the scale of risk beyond the reach of individual decision-makers.
3. AI enables harmful functions that were previously infeasible.
4. AI enables policymakers to weaponize insufficient government capacity.
5. AI makes the lives of low-income people fundamentally less predictable.
6. AI depends upon and feeds an unyielding desire to collect more data.
7. AI extends the negative consequences of past actions while limiting the ability to provide context.
8. AI makes it harder for individuals harmed to contest decisions either formally or informally.

AI's Cloak of Unwarranted Rationality Intensifies and Perpetuates Injustices

Many AI systems are founded on questionable—or outright junk—science that uses factors that are not established to be relevant to the decision being made. Even where developers of AI systems are purportedly devoted to quality statistical science, their claims may not be subject to rigorous peer review, independent validation, or any sort of public scrutiny.

Meanwhile, there is little done to build public trust in such systems, as they offer either no explanations or impenetrable ones. When AI developers do not justify withholding such information on the grounds that it is a trade secret necessary to maintain for their corporate interests, they do so based on professed mistrust of people whom they say would “game” the system. And particularly sophisticated forms of AI that operate using shifting criteria may defy all explanation whatsoever.²⁶¹ Such opacity is particularly problematic where the AI functions according to criteria that are far removed from those traditionally used by humans, fostering a “just believe us” mysticism incompatible with basic accountability mechanisms.

Still, the distinctions between “good” and “bad” science do not alter AI’s fundamental functions. AI’s supposed rationality enables or justifies harmful actions such as benefit cuts that would otherwise be politically unpopular or untenable. In Arkansas, for example, state officials claimed that the new AI system assessing eligibility for Medicaid home- and community-based services showed that recipients with cerebral palsy or quadriplegia were getting “more care than they needed” prior to AI, even though the AI-reduced hours did not provide anywhere near enough time for them to get out of bed, bathe, cook, eat, get dressed, clean their clothes, clean the house, or go to appointments. Even when not providing cover for new policy shifts, AI’s supposed rationality entrenches unjust status quo dynamics, including workers’ relative lack of power compared to employers or the targeting and labeling of children from marginalized backgrounds.

At bottom, the touted AI rationality does not ever seem to meaningfully expand access to benefits, empower workers, or produce more stable housing arrangements. Whatever rationality AI is enacting, then, is a slanted one.

AI Expands the Scale of Risk Beyond the Reach of Individual Decision-Makers

Prior to AI, if benefits caseworkers misunderstood key program criteria or were making discriminatory decisions, their reach was limited to the hundreds or thousands of people seeking benefits whose cases on which they personally worked. Likewise, the reach of landlords and employers was limited to those they chose as tenants or employees and their subsequent management decisions. Or a problematic teacher could routinely mistreat only those students in their classes or otherwise under their supervision.

AI centralizes the dispersed decision-making ability or influence of these individuals to affect masses of people. This scales up the risk of intentional harm, errors, and bias without requisite accountability structures to avoid them.

This risk is particularly apparent in public benefit programs, even when the agency is acting in good faith. The functioning of benefits programs depends on the knowledge, understanding, and ability of computer programmers who are not benefits experts to translate complex rules into code. Concomitantly, agency staff who are not computer experts must oversee the programmers' work. Any complicated statistical modeling widens the expertise gap. Unsurprisingly, this arrangement has proven particularly disastrous, with states' entire SNAP, Medicaid, or Unemployment Insurance programs being rendered unworkable for extended periods of time.

AI-related problems are not easily fixable. Practically, it can be hard to identify underlying problems, as they can involve coding errors buried in thousands of lines of code, fundamental misinterpretations of program rules requiring significant rewrites, or the lack of specialized expertise to understand the AI's statistical modeling. Once the problem is identified, the agency's project leadership must work with the AI developer to make changes. Sometimes, the contracts between them limit what changes the AI developer must make, forcing the agency to pay additional money to get a properly functioning system. Even then, though, the AI developer may not fix it immediately because there is a backlog of other problems needing attention first.

Using slightly different terminology, scholars Danielle Citron and Ryan Calo summarized the dynamics neatly: "Automation has not been as clear a win for governmental efficiency and fairness as administrators had hoped and as vendors have claimed. It has not eliminated bias but rather traded the possibility of human bias for the guarantee of systemic bias."²⁶²

AI Enables Harmful Functions That Were Previously Infeasible

Prior to AI, no state agency would have had the human capacity to be able to accuse mass numbers of people of fraud, as Michigan's MiDAS system did for 40,000 people on Unemployment Insurance. Trying to perform something similar without AI would have required individual review

of case files and the creation of attendant paperwork that exceeds most agencies' staff capacity. Likewise, landlords' ability to collude to fix rental prices seems much likelier when it can be done by submitting proprietary market information directly to an AI platform that then suggests pricing to the landlords rather than by each landlord calling, meeting with, or otherwise coordinating actions with dozens of other landlords. The algorithmic wage discrimination featured in gig economy work and uses of deepfakes also are newly feasible forms of harm.

Meanwhile, familiar forms of harm, such as surveillance of workers, tenants, and students, has become easier.

AI Enables Policymakers to Weaponize Insufficient Government Capacity

In both the benefits enrollment and identity verification scenarios, AI requires massive numbers of people seeking benefits to take actions—submitting more paperwork or providing additional proof of identity—that they would not have had to take prior to AI. In turn, the agencies have to do things to process the additional submissions at scale never known before while lacking the capacity to do so. This causes eligible people to go without benefits for months or years during times of desperation.

AI Makes the Lives of Low-Income People Fundamentally Less Predictable

Unpredictability is part of AI's promise: AI that works as intended is meant to discern connections and patterns that human faculties do not. So, a person subject to AI is not meant to have the ability to predict how they will be affected. Without a meaningful change in their own circumstances, their Medicaid home-care benefits can be cut, their work hours can be reduced, they can be accused of benefits fraud, their wages as a gig worker can be slashed, their kids can be labeled as likely criminals in school, or their infant can be taken away from them and placed in foster care. AI undermines the ability to know what will happen and why.

Just as low-income people cannot predict what will happen as a result of AI usage, they cannot predict when it will happen. An AI systems' many design or implementation flaws may ensnare someone who evaded them before. What was a straightforward renewal of SNAP or Medicaid benefits last year may be an unnavigable labyrinth this year. Or a rental application that successfully passed a landlord's screening process a couple years back may suddenly show adverse information.

These forms of unpredictability exist apart from the other vulnerabilities under which low-income generally live.

AI Depends Upon and Feeds an Unyielding Desire to Collect More Data

AI is trained using all forms of personal data about individuals: employment, property, and court records, credit history, purchasing information, social networks, web browsing practices, screen time, media consumption patterns (e.g., what movies we watch and music we listen to), what we say, where we go, and so on. This amounts to a massive agglomeration of information about individuals on scales unfathomable before, largely to target advertising, other individualized messaging, or government action. This phenomenon has been called “surveillance capitalism.”²⁶³

Much of the data is collected through publicly available sources. But even when people ostensibly opt in to having some data collected about them through user agreements with private companies, there is rarely a meaningful alternative. Saying no is simply not practicable for people who depend upon technology to communicate, pay bills, get around, or manage finances. On top of this, data collected by one company may be sold to others without the knowledge or assent of the people subjected to the data collection.

A full discussion of the privacy implications of the modern regime of mass data collection is beyond the scope of this report. To summarize, though, one aspect of risk is based on threats to fundamental concepts of human autonomy, expression, and intimacy.²⁶⁴ Another aspect is the near-term material impact of having personal information exposed or used to shape opportunities.²⁶⁵ And a third is a deepening power divide between consumers and companies that leaves consumers unable to evade the ever-present data agglomeration or to impose meaningful accountability measures on companies, no matter how egregious their actions.²⁶⁶

AI Extends the Negative Consequences of Past Actions

The mass accumulation, aggregation, and use of data about individuals means that long-past actions can haunt people perpetually. Previous evictions, unpaid bills, arrests, criminal charges, criminal convictions, and child maltreatment may not only be knowable indefinitely, but also may factor into AI-based recommendations or decisions that landlords or employers do not further scrutinize.

AI Makes It Harder to Formally and Informally Contest Decisions

In the absence of AI, a caseworker who misunderstands program criteria or applies it in a discriminatory manner can be addressed through references to policy, better supervision, complaints, the administrative appeal process, or a lawsuit (though all such avenues have always had limitations). A supervisor or agency director who orders the caseworkers to apply a policy erroneously can have their understanding of that policy challenged in a lawsuit or through the administrative appeals process. A prospective tenant denied an apartment might be able to explain a bad background check to the landlord. But, as described more fully below, AI is not subject to the same avenues for redress in the same ways.

Existing Accountability Mechanisms are Insufficient

Defiance of existing accountability structures is a defining aspect of AI-related harms. There are several aspects to this, including that public and private actors hide their use of AI, hide how their AI works, and have not capably overseen AI use. Additionally, political and market forces cannot be counted on to prevent harmful AI use, and existing laws only partially redress the vast scale of harms. Moreover, low-income people often cannot afford legal representation to enforce the rights they have, and agency appeal processes are not meaningful checks on the use of AI.

Public and Private Actors Hide Their Use of AI

People who do not know that AI is being used cannot know to investigate or fight it. Private actors are generally under no obligation to disclose information about the systems they use. The only means of learning about them is when vendors are selling them or if someone exposed to them knows how to spot them.

Meanwhile, government use is buried in impenetrable contracts with vendors, complex paperwork submitted to federal oversight agencies, or dense policy manuals not written for public understanding. In Arkansas, for example, the only way Medicaid beneficiaries learned that AI was being used to determine their home- and community-based services eligibility was when nurses were denying responsibility for cutting hours and noting that “the computer did it.”

Public and Private Actors Hide How Their AI Works

Even when low-income people or their advocates know AI is being used, it may be nearly impossible to find out how it works. Again, private actors generally have no obligation to publicly explain their AI systems. Governments often assert exemptions to Freedom of Information or other public disclosure laws, including by asserting that it is a private company’s property (often called a “trade secret”). They also otherwise resist detailed explanations based on the idea that people who understand the AI system will then try to use that knowledge to “game” it to get better outcomes.

When people harmed by AI or their advocates can access the AI system, it may involve complex computer code or statistical modeling that requires more specialized resources to understand than are available. Not knowing how the system works can prevent people from attacking its legitimacy, convincing lawmakers or the broader public to care, or contesting it through appeals processes or lawsuits.

Public and Private Actors Have Not Capably Overseen Their Use of AI

Public or private actors using AI may not themselves know how it works, have the sophistication to meaningfully interrogate the system before using it, or care enough to learn. In these cases, they end up accepting at face value the promises made by vendors who develop and sell AI tools. The statistical modeling may be faulty. The mechanics of the system may not work. Projections on who will be helped or hurt by it and to what extent may not have been run. Pre-deployment testing or phased implementation (where the system goes active only for a small part of the group that will be subject to it) may not have been considered. Post-deployment monitoring mechanisms may not have been created. Maintaining a non-AI way to make decisions in the event of AI malfunction may not have been done. And contracts may have been drafted such that developers are not required to make changes unless additional fees are paid.

To be fair, government actors often face difficult resource constraints with tightened budgets, overstretched personnel, and high turnover. Still, whatever capacity exists has not been applied successfully to mitigate risks from the dangerous AI systems that they willingly embrace. Continuing to implement AI in such situations recklessly endangers the people who rely on government services and the public's belief in the ability of government to get things done.

This ongoing lack of sophistication is enabled by the limitations of existing political, market, and legal mechanisms to incentivize more careful AI oversight and implementation.

Political Incentives Have Not Prevented Harmful AI

Hostility to low-income people is the basis of many federal and state policies and is fundamental to AI's development and use. Regarding public benefits, this can be seen in low benefit amounts, stringent eligibility criteria, work-reporting requirements, immense paperwork burdens, un navigable program rules, the lack of good customer service, and the perpetual use of the safety net as a prop by some elected officials to drive austerity-based policies.

In other areas, it can be seen in laws limiting workplace organizing (including, but not limited to, union formation), the lack of paid family leave, laws that enable predatory debt and predatory debt collection practices, and the lack of enforceable obligations on landlords to provide habitable rental homes. A recent example includes the failure of lawmakers to extend the pandemic-related temporary expansion of the Child Tax Credit, which improved the economic conditions of millions of low-income children for the year it was in effect and has bipartisan support.²⁶⁷

Of course, policy is made by people. Those with formal power and antipathy toward low-income people or other marginalized groups approach AI like any other available tool to achieve desired policy goals.

This is most apparent in the public benefits context, where AI has been designed to cut benefits, make them hard to get, or accuse massive numbers of people of fraud. Such cruelty is also present elsewhere. As one Pasco, Florida sheriff's deputy noted, his job was to force the students targeted by AI to move out of the district or to sue the sheriff's office. There is no persuasive good-faith explanation for that.

Even leaving hostility aside, low-income people do not have the ongoing, short-term political muscle to prevent the worst of the harms from AI or provoke immediate responses to it. Their potential—while demonstrated repeatedly in inspiring ways small and large (most prominently through social movements)—generally requires time and resources to mobilize. Their power is not in having cash for political donations or armies of lobbyists with cozy connections to lawmakers who might step in quickly.

While community activation efforts have been successful in some contexts, they were exceptional wins. Preventing the AI surveillance of rental housing in New York City, preventing the St. Paul school district from adopting Pasco-style data sharing projects with law enforcement, and ending the use of Medicaid AI systems in Arkansas and elsewhere are notable precisely because they are not the norm. And, in Arkansas, the political victory happened only after devastating harms to thousands of people and years of extensive litigation. This victory, too, had limits, as the state replaced the abandoned AI system with a different AI system that cut benefits in a different way.

Political mechanisms, then, have been limited in motivating key individual actors to avoid the harmful uses of AI. The Arkansas officials overseeing various AI disasters in their state Medicaid program did not suffer any political consequences because, ultimately, they were carrying out the policy goals of governors elected by significant margins. In Michigan, the same vendor who built the MiDAS system for the state's Unemployment Insurance program was tasked by the agency with fixing it, at least for a while. Rhode Island offers the rare exception where the state ultimately sued the AI developer, but it did so only after advocates had already done immense amounts of work to publicize the issue and sue the state agency.

Market Incentives Have Not Prevented Harmful AI

Harmful AI use on low-income people has not been significantly altered by market mechanisms. With few known exceptions, the users of AI systems have not broken their contracts as a result of harms, errors, or bias attributable to developers. If anything, the market to adopt such tools has expanded, with more government officials, landlords, employers, and school districts seeking to implement AI-based decision-making that exposes low-income people to harm.

Even with more potential customers, the market has not apparently produced sufficient competition to lead to less dangerous products in either public or private settings.²⁶⁸ For example, Deloitte is tied to malfunctioning benefits systems in Colorado, Texas, Tennessee, Florida, Rhode Island, and many other states and still gets contracts worth hundreds of millions of dollars each, all while producing similar errors.²⁶⁹ Credit reporting agencies, such as TransUnion, have had their failed accuracy processes cited as examples of what not to do in U.S. Supreme Court opinions and still reproduce the same kinds of errors years later.²⁷⁰ And, rather than having new players, markets for health insurance prior authorization algorithms and background screening for prospective tenants and employees seem to be trending toward private equity acquisition and further consolidation.²⁷¹ This expands the reach of particular AI systems and increases the number of people exposed to risks from any failures.

Presumably, there are significant barriers to entry to new competitors. A partial review of selected state government contracts showed the largest vendors to be established names, including Deloitte, Optum, Accenture, IBS, Thomson Reuters, and Lexis Nexis.²⁷² So-called civic tech or GovTech firms

with a stronger focus on benefits access and the public interest have sprung up, but they tend to operate on comparatively smaller scales.²⁷³ These firms generally add onto or work around agencies' core systems, rather than replace them.

Vendor lock-in, where government becomes perpetually dependent on a particular vendor, is also a problem. Having replaced and dismantled their previous systems, users purchasing new technology are then committed to it for the foreseeable future. This is because, unless they explicitly considered an alternative method, they have no way of performing that particular government function without the AI system. Moreover, vendors sell proprietary systems for which competing vendors may not have the expertise required or be legally allowed to work on.

It is important to note the role that federal funds play in subsidizing harmful AI uses. Federal funds enabled the Pasco terror, states' problematic replacements of their Medicaid, SNAP, and Unemployment Insurance systems, and public housing surveillance technologies. Vendors engorge themselves on federal dollars sent by distant federal oversight agencies to states or other localities that buy technologies they have proven incapable of managing.

This state-federal funding dynamic may be warping the market in multiple ways. First, even though the state—not the federal government—actually signs the contract with an AI vendor, contracting requirements involving mixes of state and federal funds may be so complex as to favor incumbent firms who have managed such contracts before. Second, the arrangement where individual states make contracting decisions using primarily federal dollars inhibits punishment for failure. That is, a state whose vendor botched a project does not have the power to stop another state from hiring that vendor for a different project. Indeed, the state harmed may even refrain from criticizing the vendor forcefully if the vendor presently performs other services for the state or is likely to in the future. Third, federal oversight bodies have not been willing to disallow states from contracting with vendors with documented histories of problems. Fourth, with the federal government covering up to 90 percent of costs, the price of failure does not fall as hard on states as it otherwise would. In big projects, misusing the 10 percent state contribution can translate to tens of millions of dollars wasted—a notable sum in most state budgets—but it is less significant and scandalous than wasting hundreds of millions of dollars.

Finally, a brief word must be said about control of the computational resources required by AI systems. The most sophisticated AI systems require immense processing power fueled by advanced computer chips housed in huge data centers that require enormous amounts of energy and water to power and cool.²⁷⁴ Only the biggest tech companies, like Microsoft, Amazon, and Google, have the resources needed to create and manage this infrastructure. They use the infrastructure for their own AI systems and rent the computational resources to other companies.

The AI discussed in this report is not created by these large tech companies, and much of it does not depend on the computational resources of these large tech companies. However, Microsoft and Google have efforts underway that would more directly affect low-income people. Microsoft, for example, is providing some services to a New York City project to create digital wallets and centralize city services and benefit programs that has been criticized for its many risks, lack of transparency, and cost overruns.²⁷⁵ And, as noted in Part One, Google is about to launch AI-based review of Unemployment Insurance hearings in Nevada. Thus, concentrated control of computational resources for advanced AI and the related market effects could soon become a greater issue of concern for low-income people.

Existing Laws Only Partially Redress the Scale of Harmful AI

Apart from politics and markets, laws are the other main source of potential accountability for this technology. Existing laws do not clearly cover AI adequately at the scale needed to prevent harms in the first place. A full discussion of legal theories is beyond the scope of this report, but a brief one suffices to highlight the limitations.

The principle legal protections against the harmful use of AI fall into two main categories: procedural protections and anti-discrimination laws. Other worker and consumer protection laws also offer some promise for AI accountability but have not been tested as extensively and so are not discussed here.

Procedural protections apply most strongly to government uses. Laws such as the Administrative Procedures Act require general public notice of and an opportunity for the public to offer comment on important agency policies that apply generally to at least some segment of the public. The U.S. Constitution's Due Process Clause (and similar provisions in many state

constitutions) requires individualized explanations for adverse decisions and an opportunity to contest them. In essence, these are procedural hoops the government must jump through in order to be able to use AI. If the government fails to do so, it may be prevented from using AI until it completes the obstacle course straight through. Given AI's opacity, explaining decisions and providing a fair opportunity to contest them is often a difficult task for the government.

Even here, though, the applicable laws do not allow involvement early in AI development. They can be applied only right at the point of or after AI is implemented, when all key aspects of it have already been decided on. Once the government follows the necessary procedures, these laws do not meaningfully restrict what an AI system does, even if it causes massive harms.²⁷⁶

Importantly, AI and the ways humans interact with it will stress the ability of courts to restrain it based on procedural fairness considerations. Leaving judges' political inclinations aside (sometimes, they, too, are hostile to poor people), courts may stop the most egregious violations of procedural fairness requirements involving AI use for low-income people.

But judges are often sympathetic to the practical constraints government agencies face, and the reigning analysis for procedural due process claims requires consideration of "the fiscal and administrative burdens that the additional or substitute procedural requirements would entail."²⁷⁷ As governments justify AI use in terms of efficiencies gained, judges' considerations in restricting it will shift increasingly to the vital functions the agency claims it would not be able to do absent AI. Agencies will pit the people already harmed by AI against those who are exposed to it but not yet harmed by it. Circularly, then, even if it might be otherwise unlawful, the fact that AI has taken over the full scale of a program's functioning and eliminated human-based alternatives will make stopping it all the harder. At the same time, agencies may get more clever about blurring the lines between human and AI decision-making, such as by claiming that humans are making independent judgments while only considering the AI's recommendations.

Even where claims are successful, courts can only order state governments to comply with the law and pay for the costs of plaintiffs' attorneys. Due to a doctrine called sovereign immunity, courts cannot order state governments to pay monetary damages to people for the value of whatever was denied

or the related emotional suffering. A different immunity doctrine, known as qualified immunity, prevents people from winning monetary damages from the government officials responsible for the harmful AI except in narrow circumstances where the officials were personally involved in the violation of a constitutional right that a previous judge had established in a previous case with similar facts. The general unavailability of monetary damages for procedural claims involving government actors removes an important deterrent to harmful AI use, as the officials know that risky choices and bad oversight will not cost them or the state significant amounts of money.

Some private uses of AI, like background checks, are subject to procedural requirements. The Fair Credit Reporting Act requires landlords or employers who deny housing or employment based on a background report to provide basic notice of that fact and information allowing the person to get more information from the background check company. The law generally requires that background check companies use reasonable procedures to ensure the accuracy of the information in the background checks and, when requested, disclose the sources of such information. It also requires that certain adverse information be left off background checks after a certain amount of time has passed.

Practically, however, these legal mechanisms are not easy to use. It is hard for people to know whether AI was involved in the denial of housing or employment so that they can assert their rights. Once they do, dealing with the companies can require involved back-and-forth calls or letters to find out what information was wrong and to correct it. Lawsuits require even more time. Meanwhile, the housing or employment opportunity has been lost, and the person must spend precious time looking for others.

Anti-discrimination laws also apply to uses of AI and can include monetary damages, but they, too, have limits. In many cases, the plaintiff must prove that the government or private actor using AI intended to discriminate against a person on the basis of at least one of various protected characteristics, including race, national origin, sex (and, by extension, sexual orientation), religion, or age. With few people openly expressing biases, such proof can be hard to come by.

In a few situations, the plaintiff can win by showing that a neutral policy has a disparate impact on people with a protected characteristic, even if such impact is not intended. But even when this alternative method of showing discrimination is available, doing so requires complicated

statistical evidence that can be tough to obtain (especially with AI use being so opaque) and requires costly expertise to analyze. Furthermore, the plaintiff must show that the system does not fulfill a legitimate business purpose, which requires more—and difficult-to-obtain—insight into both the AI system and its user's operations. In light of these obstacles, disparate impact claims are not well-suited to challenge mass instances of day-to-day discrimination in a short timeframe.

Anti-discrimination laws also do not cover all kinds of harm. For example, if an AI system hurts low-income white people in roughly the same way as it hurts low-income Black people, existing laws may not recognize that any discrimination has occurred. Offering only limited protections, then, anti-discrimination laws allow public and private actors wide latitude to develop and use AI to harmful ends.

Disability-focused anti-discrimination laws work differently than other types of anti-discrimination laws and provide more apparent flexibility. They generally do not require the same showing of intent and can instead be established by showing that the way an actor using AI administers a program excludes people with disabilities, places them at risk of institutionalization, or unjustifiably denies disabled people reasonable accommodations that would allow them to access the program or service. Still, courts have been tightening use of disability discrimination claims, particularly for harms like the reduction of Medicaid home- and community-based services.

Limitations notwithstanding, existing laws can be effective and should not be discounted. In Arkansas, Idaho, Michigan, and Tennessee, for example, people hurt by AI and their advocates have won landmark cases against state governments using AI for Medicaid, SNAP, and Unemployment Insurance based on procedural claims that have ended use of the AI systems, forced massive changes to them that reduced the harms caused, and required the state to consult with the people most affected.²⁷⁸ In the lawsuit in Arkansas, the plaintiffs even won money that is usually not available at all when suing state governments by suing individual state employees who violated established constitutional rights.²⁷⁹ In Florida, the U.S. Department of Justice settled a disability discrimination investigation against the Pasco County School District relating to the scheme with the Pasco County Sheriff's Department to use AI to predict students' likelihood of committing future crimes.²⁸⁰

Likewise, there has been success against private companies. In Massachusetts, advocates recently settled a class action lawsuit alleging that tenant screening algorithms were discriminating on the basis of race.²⁸¹ The U.S. Federal Trade Commission, a government agency, also settled Fair Credit Reporting Act claims against companies who furnish tenant screening reports.²⁸² In both cases, the companies involved had to pay significant sums of money and change how they use the AI systems.

These various successes, though, reveal another stark limitation of existing laws. They generally take several years and immense resources to win even basic modifications to government or corporate practices.

Low-Income People Cannot Afford Legal Representation to Enforce their Rights

Legal advocacy is made further inaccessible by the costs and availability of lawyers. Generally, low-income people cannot afford to pay an attorney. Many of these cases require huge investments of time and money and do not offer the prospect of winning any or much money from the entities using AI. Consequently, relatively few private attorneys are willing to take cases on contingency, which is an arrangement in which the attorney does not require money upfront but will take a portion (often around 33 percent) of whatever money is won through a trial or settlement.

When private attorneys are unaffordable or unavailable, low-income people can turn to legal aid, which is a network of organizations that provide free legal services in noncriminal matters to low-income people everywhere in the United States. Legal aid directly alleviates poverty-related harms and injustices for more than 1.8 million low-income people per year.²⁸³ Predictably, there are not enough legal aid resources for all of low-income people's needs. Nationally, there are only 2.8 legal aid lawyers for every 10,000 people in poverty.²⁸⁴ This dearth creates a justice gap in which low-income people do not get any or enough legal help for 92 percent of their civil issues.²⁸⁵

Still, despite its resource limitations, legal aid achieves remarkable results. Legal aid organizations routinely transform the knowledge gained from directly serving individual low-income clients into successful systemic advocacy efforts. Legal aid campaigns have culminated in landmark

wins involving, among other things, AI, public benefits, workers' rights, and disability rights. As discussed more below in Part Three, targeted investment in legal aid is one of the keys to increasing the intensity and success of AI accountability efforts.

Agency Appeal Processes are not Meaningful Checks on the Use of AI

For private uses of AI, such as in employment or nonpublic rental housing, there is no legal process to address harms apart from lawsuits or contractually mandated dispute resolution processes like arbitration. But when a government denies benefits or housing or makes certain decisions about a disabled child's education, people generally have recourse to an agency appeal process that takes the form of an administrative hearing, a less formal type of a court case overseen by an administrative law judge who did not make the original harmful decision.

During that administrative hearing, the recipient can present evidence of their eligibility, have their own witnesses testify, and ask questions of government officials. The administrative law judge must make a reasoned decision based on the evidence presented and explain it in at least basic terms.

Administrative hearings were originally designed for people to navigate without an attorney. So, for example, formal rules that determine what evidence can be used—think of lawyers in movies who object and say words like “hearsay”—often do not apply as stringently. Nonetheless, fundamental problems prevent the administrative hearing from being a meaningful check on the AI's decision.

First, the hearing process is structured in a way that discourages people from filing appeals. Due to government errors, delays, or address mixups, people may not ever receive notice of the decision in time to appeal. Then, with lengthy notices that are difficult to understand, people may not understand what is happening, why it is happening, how to file an appeal on time, and what they must prove. Where AI is used, notices commonly lack the individualized information needed for a person to identify that the decision was wrong or understand why. Getting additional information from the agency can require hours-long phone calls or office visits and

ultimately not result in obtaining needed information. Actually filing the appeal itself can be difficult because some government agencies require a certain form or a visit to a local office that requires ready transportation and time. Then, once an appeal is filed, people may have to wait several months for the actual hearing. In that time, people can forget key things, witnesses may disappear, and other life events spring up.

Second, despite the relaxed rules of evidence, the hearing features complexities that are difficult for people to navigate on their own. People have to get and serve subpoenas to ensure witnesses appear or must send in certain documents they plan to use several days before the hearing. Once there, the hearings still have some rules that limit what kinds of evidence can be presented and how. And, beyond rules, certain skills are needed to best present a case, such as asking good questions of witnesses or tying evidence to whatever the applicable eligibility standard is. People may be intimidated by the process and find themselves up against government officials who do this frequently. Of course, lawyers would help with many of these issues, but, as discussed above, they are not readily available to many low-income people.

Third, people may not be able to know or understand the criteria that the administrative law judge will use to make a decision. This is particularly problematic with AI, where the criteria may not even be available due to claims that it is secret intellectual property. Even if available, the AI's logic may be far removed from what makes intuitive sense. Simply describing why the person needs a service or why they think they qualify is often not enough.

Fourth, due to peculiarities involving human review of AI decisions, the administrative law judge may not be able to make an independent decision about whether the AI was right or wrong because of two dynamics called automation bias and criteria displacement. Automation bias means that human reviewers are likely to just accept the AI's decision because of the idea that AI systems—with all the supposed science and sophistication behind them—are inherently superior to and less prone to errors than a human. Here, the human reviewer could make a different decision but instead defers to the AI because they think it would not make a mistake.

Criteria displacement means that the AI criteria has become so embedded in the decision-making process that it has displaced all other criteria a human reviewer could use to make a given decision. This happens most frequently in decisions that involve judgment, such as whether

someone needs a particular service or intends to deceive the government to obtain benefits (i.e., cases of fraud). In cases of medical services, a person's need becomes based strictly on what information the AI uses, not independent measures such as how long certain care tasks take or what the consequences to a person's health would be in the absence of the service. In cases of fraud, intent becomes whatever behaviors the AI flags as suspicious, not whether someone knew about incorrect or disqualifying information and took steps to hide it. In both of these instances, the administrative law judge is incapable of making a different decision; there is simply no other criteria to apply. The only question is whether the information the AI used to make a decision was accurate at the time. If not, the fix is to just rerun the AI using corrected information.

Fifth, the administrative hearing does not offer complete relief. If the person wins, the government will simply change the adverse decision for the person moving forward. The person cannot receive any relief apart from whatever benefit was denied or taken away being reinstated, and there is no compensation for any time spent or suffering experienced. Crucially, the agency is not forced to change the AI system; the system can continue to harm others in the same way. Meanwhile, the person will have gone through a lengthy and stressful process that they may have to repeat the next time the AI turns its gaze on them, such as in yearly renewal processes for public benefits.

Sixth, at all stages, recipients often face language- or disability-based barriers that the government does not account for. People who primarily speak a language other than English may receive documents in English and have no effective way to ask the agency for translation or interpretation. When translation or interpretation is requested, the government often simply fails to provide it in a way that is adequate for the legal complexities at hand. For people with disabilities, there often is no accessible way to file an appeal on time, get needed information, or navigate the phone or video technology that the government uses to hold an administrative hearing. Although accommodations are legally required, they are hard to get or not provided.

AI Threatens the Concept of Accountability

A subtler harm of AI is that its use attacks the very notion that accountability itself is valuable, that decisions about fundamental aspects of life should be reasoned, predictable, explainable, and subject to challenge. Such values are integral to justice and participatory governance.

Though programmed by humans who have developed the applicable data sciences and technologies, AI's essential promise is that it is not human. It offers to review data at speeds that we cannot, see patterns we do not, and make decisions by considering factors we would not.

With relatively straightforward AI, such as those found in sorting or scoring algorithms, this often leads to outcomes that make no intuitive sense—veritable algorithm absurdities, as I called them in Arkansas. In that instance, the AI system that determined Medicaid home- and community-based services ignored obvious considerations, including someone's ability to get around, bathe, or dress, while taking into account seemingly irrelevant things, such as whether they vomited or had a fever in the past three days (while four days was too far in the past to matter). Similarly striking, a difference of just one point on one question out of the assessment's 286 questions—say, a nurse scoring the person as needing “limited assistance” using the toilet instead of “extensive assistance”—could mean a loss of about 1.5 hours per day. On and on the absurdities went.

Absurdities in an AI system's mechanics lead to absurdities in outcomes. In the Arkansas example, people who had been receiving Medicaid home- and community-based services for several years were helpless to understand how their hours could be cut so drastically when their conditions had not improved at all. In one swoop, the AI tool erased all previous determinations of need by the government agency and destabilized the lives of people clinging to independence with a fraying thread. Elsewhere, such as in Pasco, Florida, kids with no history of involvement with the criminal system could wind up targeted for ongoing law enforcement harassment. Or people in Michigan could be accused of unemployment fraud because the state could not reconcile divergent data sources.

Ultimately, the lack of sensible criteria implies an underlying belief that such systems do not need to make sense, at least not to the people subject to them.

Whether disdain or arrogance, perhaps the problem would not be quite so threatening if the layers of human authority involved in any AI-based system could or would meaningfully oversee it. But that does not seem the usual case. The various uses recounted in Part One of this report suggest that organizational checks routinely break down. In Arkansas, nurses, state Medicaid administrators, and administrative law judges accepted the AI system's process without needing to be convinced that it made any sense. They certainly did not understand the statistical underpinnings of the system, generally had no interest in doing so, and could not explain how the system worked at scale or as applied to individuals. The AI proclaimed, and they executed its will.

Among the examples of AI use in this report, the Allegheny County's child welfare algorithm demonstrates the most considered process around implementation, with relative transparency, outside evaluations, academic publications, and ongoing public reporting. Even there, though, parents are not told the risk score they received or how it was arrived at. Some outcomes still do not seem to make sense, including the parents who still do not have custody of their child after taking her to the hospital more than two years ago when the child would not eat.

Boiled down, AI systems as used against low-income people display characteristics that are authoritarian in essence. They enshrine the arbitrary as acceptable, do not explain their actions, coerce the acquiescence of people without power to resist, and close off avenues of redress. And, if relatively straightforward AI can stretch established accountability norms so easily, what happens with the proliferation of more advanced forms of AI, in which even the systems' programmers cannot discern the AI's specific reasons for making a decision?²⁸⁶

Thus far, procedural fairness laws requiring government to notify people of and explain AI-based decisions have been the main bulwarks against even more harm. But much more will be required to avoid the ongoing immiseration of low-income people that AI entails.

Part Three: The Way Forward

AI poses the threat of immediate harms on huge scales and defies existing accountability mechanisms. Consequently, any resistance plans must provide for both short-term harm reduction and long-term structural changes that will prevent harms in the first place.

As a starting point, the strategic goal is a more just world where low-income people have what they need to thrive. Only by reorienting politics away from hostility toward low-income people will the problems posed by AI be truly solvable.²⁸⁷

While ambitious, such goals are consistent with the practical, incrementalist approach required to reduce harm along the way. The key is linking all harm reduction efforts to the larger strategic frame. By understanding how a given AI system works well enough to fight its immediate application, one builds the knowledge needed to challenge the particular system's legitimacy. Doing so with several AI systems helps undercut the legitimacy of AI itself as a way of making decisions. And, by working with people to respond to a harmful loss of benefits, work, or housing, one has a chance to build the solidarity needed for sustained group action.

With this background, the report lays out a roadmap to a better future, inspired by existing examples of successful fights against AI: building the capacity for integrated, multidimensional advocacy, supporting an activated community of people most hurt by AI, pushing for laws that regulate AI use and better protect economic stability, providing more capacity to enforce laws and seek justice, and developing a positive vision for this technology.

Build the Capacity for Integrated, Multidimensional Advocacy

As noted throughout this report, people have already fought and won by litigating, organizing, educating, researching, and telling stories. For any chance at long-term success, these tactics must be combined to account for their individual limitations. I recognize that each tactic features much more nuance than can be covered here and offer only these top-line characterizations to show the need for integration.

Litigating can stop powerful actors from doing something harmful on both individual and mass scales and can influence the incentive structures around AI use. But litigation does little on its own to build a constituency needed to change the problematic political dynamics allowing harmful AI use.

Organizing shapes incisive critiques, convinces people to care about and show up for issue in an ongoing way, facilitates peer support networks, and coordinates engagement points to sustain campaigns. But it is sometimes limited in its ability to stop or reduce short-term individual harm against intransigent actors.

Educating the communities affected by AI, their allies, the broader public, and policymakers helps people understand what is happening, why it is important, and how to potentially fight back. But education alone does not necessarily ensure continued engagement.

Researching exposes the breadth or a depth of an issue, but often reaches only a small audience of experts and is not made actionable for people engaged in active fights.

Telling stories builds public awareness and empathy that sometimes connects people being harmed to lawyers, brings more resources to an organization, encourages more willingness or enthusiasm to join a campaign, or incites government oversight or investigation. But storytelling rarely works to shame bad actors into significant changes right away.

Alone, these tactics may not be sufficient to achieve meaningful short- and long-term wins. But weaving these tactics together creates the potential for advocacy that can work on all levers of power needed to reduce near-term suffering and promote long-term change.

Support an Activated Community of People Most Hurt by AI

Whatever combination of tactics is used, advocacy is stronger when the people most affected by AI are involved in leading the resistance to it. Activated communities have curbed the AI-based targeting of students in Pasco, Florida and St. Paul, Minnesota, prevented the deployment of AI surveillance in rental housing in New York City, and ended the use of the AI systems that cut Medicaid eligibility for home-based care in Arkansas and Idaho. Meanwhile, drivers for ride-sharing companies, such as Uber and Lyft, have already organized themselves into a force that is seeking additional legal protections and have won battles for unpaid wages in Washington State and a guaranteed minimum wage in Minneapolis.²⁸⁸

The success of these community-led efforts demonstrates the potential for achieving systemic victories on specific issues in specific places. Such efforts can change attitudes, persuade decision-makers, invite greater legislative oversight, prompt formal investigations, and invoke the authority of judges.

These successes also suggest the possibility of winning legislative goals. Thus far, the AI accountability discourse has been dominated by big tech companies, large nontech companies who sell services or products to governments, and various other corporate or academic interests pushing AI hype. Though companies' specific positions vary, industry generally opposes regulatory measures they say would stifle innovation by slowing down the pace of development and by exposing them to too much liability with a still-new technology that can only be refined through live use.

With far fewer resources, civil society groups have done incredible work to counter the hype, document AI's many immediate risks, and put focus on the need for legislation that meaningfully curbs harmful uses. Their advocacy contributed significantly to President Joe Biden's issuance of the AI Bill of Rights and related guidance from the federal Office of Management and Budget requiring the federal government to use AI responsibly.²⁸⁹

Yet meaningful legislative protections have stalled. Various bills have been introduced in the U.S. Congress but have not gone anywhere. Meanwhile, states such as California,²⁹⁰ Utah,²⁹¹ and Colorado²⁹² and local jurisdictions, including New York City,²⁹³ have enacted laws addressing some AI uses affecting low-income people. But the protections offered are both limited and not enforceable by the people hurt.

Clearly, any chance at robust, enforceable protections at the federal, state, or local level that will curb AI harms for low-income people will require a different power dynamic. Low-income communities and their allies can be a new constituency to change the discourse and increase the likelihood of protections that prevent the infliction of harms in the first place and, where still possible, enable low-income people to better fight back. Having the communities most directly affected leading such efforts reduces the possibility that professional advocates concede to nominal protections that normalize harmful AI.²⁹⁴

And the proliferation of harmful AI use involving higher-income people—for example, with Medicare Advantage plans and workplace surveillance—may open possibilities for broader interest alignment and coalitions.

Push for Laws that Regulate AI Use and Better Protect Economic Stability

While existing laws offer some protection against harmful AI use, they are insufficient for the scope and scale of the harms involved. As such, enforceable AI-specific laws are needed to offer systemic protection against harm.

Ultimately, the communities involved in the fight and their allies should drive specific legislative proposals. Some provisions to consider include:²⁹⁵

- Outright bans on some uses for particularly sensitive decisions
- Ensuring public access to scrutinize the AI system, user manuals, and other supporting documents
- Intensive pre-deployment testing and projections with mandatory documentation of key areas, including projected benefits and harms, data quality, bias reviews, downstream impacts of harms, harm mitigation plans, and clear justifications for use in light of the risks and harms
- Mandatory public participation points prior to deployment
- Phased deployment with prompt impact assessments
- Ongoing post-deployment documentation of impact at regular intervals
- Mandatory consideration of the costs of ongoing oversight of any adopted AI as part of any benefit-risk calculus
- Automatic discontinuation of use if the post-deployment impact deviates from projected impact by a certain magnitude
- Maintaining non-AI ways to perform the given functions, including non-AI criteria by which a human can make independent decisions
- Mechanisms for developer liability
- Safeguards against improper use of automated translation or interpretation services
- Restrictions on the ability of entities to collect, maintain, use, and sell data that do not allow broad consent-based exceptions (i.e., data minimization requirements)²⁹⁶

For government uses, the following additional provisions should be considered:

- Building robust centralized public infrastructure for core benefit functions, like Medicaid eligibility and enrollment
- Effective and transparent contracting practices (sometimes called procurement)²⁹⁷

- An affirmative duty to consult with affected individuals, community groups, and advocates, report findings, and provide individuals the support needed to meaningfully participate during system development and significantly before the rule-making process
- Requiring rule-making, with the express ability to sue after rule-making and prior to deployment
- Maximal use of federal authority to approve or deny states' uses of AI when federal funds are involved
- Maximal use of federal oversight authority once the AI is deployed
- Prohibiting federal funds from being used to purchase AI that carries a significant risk of harm

Critically, any AI-specific legislation must not contain expansive exceptions (sometimes called carve outs) that allow unregulated use for so-called sensitive functions, such as national security, policing, or immigration enforcement.

In addition, some laws could be reformed to minimize the supposed need for AI. For example, if income or asset limits in benefit programs were simplified or eliminated, AI would not be needed to determine compliance with them.

Finally, strengthening existing laws to better protect economic stability would also help to address AI harms. For example, if people eligible for Medicaid were to have an enforceable right to a certain amount of health care, such as whatever is medically necessary, anything impeding that right could be addressed, including AI. A proper discussion of specifics, including possible downsides, is outside the scope of the report, but affirmative rights to housing or employment stability, paid leave, child care, education, and privacy could further protect low-income people against the vagaries of AI.

The United States does not need to legislate from scratch. Just as AI is used elsewhere for the purposes discussed in this report,²⁹⁸ so, too, do legislative responses to it exist. Whether focused on data protection, privacy, or high-risk AI uses, various other countries have pioneered legislative approaches from which U.S.-based communities, advocates, and policymakers can draw to formulate meaningful protections.²⁹⁹

Provide More Capacity to Enforce Laws and Seek Justice

Even good laws are not worth much if low-income people do not have access to lawyers to enforce them. Providing more resources to legal aid providers—building their knowledge, abilities, capacity, and influence—ultimately helps low-income people vindicate their rights and avoid harm from AI.

But enforcement of laws is much broader than on-the-ground lawyers. It starts with people knowing that something wrong is taking place. This can be done by training social workers in health clinics or schools, church ministries supporting low-income people, and hunger relief workers. Once they know something is wrong, enforcement of laws does not always require legal involvement. Sometimes, concerted efforts by tenants or employees can raise legal violations to decision-makers who then choose to fix them. Alternatively, people might be able to get lawmakers involved to encourage violations to be fixed. Even if not, these efforts can inform subsequent campaigns or litigation.

Additionally, seeking justice does not just mean enforcing laws. When laws are insufficient, as with AI, there sometimes is not anything to enforce. In these cases, people must raise awareness of the ongoing inequities to pass new laws, change attitudes, or force better business practices.

All such efforts lead back to the ground. It is the people experiencing injustice, their advocates, and their allies who know the most about how the world breaks against them. Investing in their ability to lead advocacy on their own behalf is the best thing we can do to bring about a better world.

Develop a Positive Vision for Technology

The AI uses covered in this report feature rampant harm to and few benefits (even hypothetical ones) for low-income communities. This reflects the present dynamics of AI uses. It is developed by vendors to be sold to government officials, employers, landlords, and others with authority over the lives of low-income people who will then use it to make decisions that

are driven by scarcity, hostility, or profit-seeking. This happens in a context with perverse incentive structures driven by insufficient accountability mechanisms. Why would such a process produce AI that somehow helps low-income people?

This does not mean that all AI use is inherently unjust. With the proper dynamics, resources, and accountability mechanisms, including community consultation, AI and related technologies could help low-income people access public benefits, mitigate lingering harm from systemic injustices, and assert their rights. As demonstrated below, even these helpful uses of AI feature significant limitations on their ability to address the root causes of various injustices. This underscores the point that “bad AI” and the forces that enable it cannot simply be counterbalanced or solved by “good AI.”

Accessing Benefits

Functional systems for enrollment, eligibility, and renewal of public benefits are clearcut cases of positive governmental uses of AI. If built so that they cannot be easily subverted by administrations hostile to low-income people, such systems would do remarkable good for the nearly 100 million people who receive government benefits of some sort.³⁰⁰

Government agencies generally already know people’s incomes. With minimal additional paperwork, they should be able to make quick eligibility decisions and then automatically renew eligibility at the required intervals. The Affordable Care Act mandates some such procedures for states performing Medicaid renewals. But in practice, those have been a mixed bag, often hampered by resource limitations, technological errors, or policy choices that weaponize minimal data discrepancies to erect barriers to access.³⁰¹ Federal oversight of states’ compliance has been wanting. Different accountability mechanisms are needed to realize the promise of AI for promoting benefits access. Developing such mechanisms for the systems of 50 states is not likely in the short term.

This suggests the need for a centralized, federally funded technology platform for at least Medicaid and SNAP, which have extensive federal eligibility and enrollment standards that apply to all states (Unemployment Insurance leaves more decisions to the states, so a central system may be less useful and more difficult to build). Presently, billions of federal dollars have gone to private companies that states contract to build dysfunctional

and error-prone systems. That money could be better spent building one core system that integrates all applicable federal laws and regulations. States could then choose to adopt it or not, but states that opt for their own systems should be provided fewer federal dollars to do so. Meanwhile, where states have policy options that result in different eligibility standards or service levels, the states could use the federal platform and contract for smaller projects to meet their specific needs that carry less risk of core failure.

A model for this sort of project already exists in the [healthcare.gov](https://www.healthcare.gov) marketplace created by Affordable Care Act. After a bumpy start in 2014, the marketplace now provides Medicaid eligibility determinations and private insurance enrollment and subsidies for the residents of 29 states and provides substantial support for three others.³⁰²

Even with the promise of improved enrollment, eligibility, and renewal systems, many of the other problematic dynamics in public benefits exacerbated by AI would persist. “Better” AI will not solve the problems of insufficient benefit allocations, prior authorization denials, or unwarranted fraud allegations. In these cases, the root cure is a more robust and simplified system of public benefits that provides people more economic stability and does not impose an unnavigable set of program rules that harshly punishes noncompliance.

Facilitating the Sealing of Criminal Records

Another example where AI-related technologies are assisting or could assist low-income people is with the automated sealing of criminal records. The United States has among the the highest incarceration rates in the world, topping the list in some years.³⁰³ As a result, roughly 1 of out of every 3 U.S. adults has some sort of criminal record, whether for an arrest, a criminal charge, or a conviction.³⁰⁴ These records can prevent people from getting housing or work.

As a result, all but four states have adopted laws to allow some criminal records—for certain arrests, charges, misdemeanors, and felonies—to be expunged or sealed such that the person who had the record no longer has to experience all of its negative effects.³⁰⁵ But the process to actually seal the record can be complex.

The person must find out if it's the sort of record that is eligible to be expunged or sealed—many are not, and which ones are and are not eligible depend entirely on each state. The person then has to gather information about the record to be sealed, which might include getting records from distant locations and old paper files. The person also has to file legal paperwork with a court, usually without being able to afford an attorney's services, and pay a filing fee. Next, the person has to notify the involved prosecutors' offices to give them a chance to object, and the person then often has to appear in front of a judge to get the expungement or sealing formally approved, which does not always happen. Once the expungement or sealing is approved, the person has to notify the state database for criminal records and private background check companies where these records are kept.

In response to this complexity, some states have passed laws to simplify the paperwork or automatically expunge or seal certain kinds of records after a certain time has passed.³⁰⁶ AI has been used to facilitate this process.³⁰⁷ This is undoubtedly valuable, as it helps low-income people overcome an obstacle to housing or work.

From a broader vantage point, though, it is just making the systemic injustice of mass criminalization and incarceration slightly less unjust. Admittedly, with the ability to get work and housing, some people may be less likely to re-enter the criminal legal system. But the possibilities for such AI-based record sealing are limited, particularly when different kinds of AI systems are being used by police departments to predict where crime will happen and criminalize the people there.³⁰⁸

Asserting Rights and Exercising Power

Legal aid programs, other legal advocates, and community organizing groups often engage in “know your rights” campaigns to inform low-income communities about how to identify violations of the law and take steps to assert their rights.

One aspect of these campaigns has long been to make valid legal paperwork accessible to communities. For example, legal aid programs might provide template documents to respond to debt collection abuses, ask landlords for repairs, or file for divorce. For years, legal aid organizations and the broader access-to-justice community have been working to automate and

individualize these resources, so that they are tailored to the facts and circumstances of low-income people. Recent efforts have incorporated AI to facilitate this automation and individualization. This has been done most prominently to support tenants defending against evictions with chatbots and automated toolkits.³⁰⁹

Again, such advancements are useful if they are properly vetted, monitored, and overseen. After all, providing legal information and documents is a high-risk undertaking, as any incorrect information can mislead low-income people and cost them their rights. The increasing sophistication of these systems means that there are more chances for inaccuracies or errors. Extensive resources therefore must be put into operating them, and it is not yet clear that the organizations implementing such technologies will do so, particularly when the organizations themselves often lack needed resources.

With proper oversight and scaling, these tools may be able to meaningfully resist or slow down the worst of landlord eviction abuses. Even so, the AI systems are not changing the dynamics that enable those abuses and other unfair housing conditions in the first place. Recognizing these limitations, some organizations, such as debtors' unions, are using the helpful engagement enabled by their AI systems to connect low-income people to broader movements for more just economic conditions.³¹⁰

Beyond supporting tenants' ability to defend themselves in court, tenant-supporting organizations have used technology to identify geographic hotspots of evictions and identify the landlords most aggressively pursuing evictions as part of efforts to impose accountability on them.³¹¹ Similarly, organizations opposing racially targeted policing and related abuses in Los Angeles, Chicago, and New York City have used community-focused, tech-enabled, data-gathering tools to identify and keep track of law enforcement officers engaged in ongoing harassment or abuse against local communities.³¹² Legal aid organizations also have used geospatial mapping techniques to identify slumlords, predatory lenders, and other bad actors. Incorporating more advanced AI may make such tools easier to create and increase their efficacy.

Community-Built AI

Some advocates and technologists are developing infrastructure for affected communities and their allies to be involved in the development of technologies, including AI, that address injustices or other community priorities.

The most prominent example is the expansion of Social Justice Hackathons, an instance of which recently happened in Philadelphia.³¹³ The event brought together “[s]tudents, programmers, ‘civic-minded hackers’, faculty, justice advocates, public servants, community legal providers, writers, story-tellers, and members of the Philly community” to come up with ideas for justice-promoting technologies and build proof-of-concept versions. Winning projects included an interactive website that allows members of the public to assess the performance of elected judges whose records are otherwise difficult to discern, a resource hub for new immigrants, a tool to resolve disputes between landlords and tenants around the return of security deposits, and a guide to starting restaurants that is accessible to people in immigrant communities.³¹⁴

Events like these represent the strongest examples of efforts to shift the dynamics about AI development as applied to low-income people. The shift runs along two axes. First, the Social Justice Hackathon aims to build something useful to low-income communities—a consideration generally absent from the AI uses detailed in Part One wherein at most, government agencies might consult communities or their advocates in largely performative ways to get input about some AI use that will likely end up being harmful.

Second, the Hackathon aims to let communities inform and, in the future, ultimately determine what gets built. This is different from the examples of expanded documents or chatbots listed above, where well-meaning advocacy organizations build something they think will be useful for the communities they serve. Such decisions are usually considered carefully, informed by interactions with low-income people, and are reflective of trends the organization is seeing. Nonetheless, the communities served are not generally consulted in a formal or structured way with power to decide what gets built. The Social Justice Hackathon at least imagines such a process and starts on the path to provide some of the infrastructure needed to make it a reality.

Conclusion

This report is the first known effort to explain and quantify the breadth of AI used to make decisions about the lives of low-income people. Even leaving out the criminal legal system, the immigration system, and a few other known uses, the report shows a staggering permeation of AI into everyday aspects of life. Powerful actors employ AI to decide where low-income people live and work, what they pay in rent, what they earn in wages, how their kids are treated in schools, what government benefits they can receive and under what conditions, which parenting actions will be subject to state investigation, and what resources will be available to protect them against domestic violence. Too often, these decisions restrict opportunities and cause tangible harms. No aspect of life is left untouched.

AI threatens to compound exposure to other forms of AI-based decision-making and related harms. For example, a low-wage worker terminated by surveillance and management technologies may become newly eligible for public benefits, where they will now face more AI-based decision-making. That, in turn, could expose them to an increased risk that a child will be labeled at school or that someone will refer them for a child welfare check on the basis of neglect. While such compounding harms largely reflect existing vulnerabilities of life as a low-income person—no AI is required to experience them—the ubiquity of AI decision-making may intensify these dynamics.

AI is different from and more threatening than human-driven decision-making. It provides an unwarranted sheen of pseudo-scientific rationality to the inequities it creates or perpetuates, expands the speed and scale of harmful decisions, enables previously infeasible forms of harmful actions, enables policymakers to easily weaponize the lack of government capacity,

makes the lives of low-income people fundamentally less predictable, depends on unyielding demand for more data that erodes privacy, extends the negative consequences of unfortunate past actions or circumstances indefinitely, and makes it more difficult to challenge adverse decisions.

Thus far, AI-based decision-making has defied existing political, market, and legal accountability structures. Low-income people have felt the brunt of the harms while having no meaningful say in AI's adoption or use. This dynamic will only continue with AI's expansion. Still, concerns for fairness, justice, and equity for low-income people are not the only drivers of urgency. As shown through AI's use in Medicare Advantage plans, tenant screening, and worker surveillance and management, for example, AI-based decision-making is not limited to low-income people. If left unchecked, AI will soon be making decisions about the key aspects of everyone's lives, with all the attendant risk of harm.

A better future with circumscribed AI use is possible, and now is the time to start bringing it into existence. Though these technologies have existed for some time, they are not yet fully entrenched as legitimate ways for government officials, employers, landlords, and other influential actors to routinely make decisions. AI-related harms still seem both new and grave, as instances of false arrests, erroneous benefit denials, denied medical treatments, missed housing opportunities, racial bias, misinformation, and oppressive working conditions can still prompt press attention, public outrage, court intervention, and legislative oversight. Meanwhile, many sectors, including health care, banking, entertainment, law, and the non-profit world, are considering whether and how to integrate AI into their own operations, evincing a mix of curiosity, openness, and skepticism that is propitious for advocacy. In the realm of politics, party lines on AI have not fully calcified. And the dearth of AI legislation at the state and federal level creates an opening to stake out bold policy positions and support them with activated constituencies for the inevitable negotiation of conflicting interests that will take place.

This unique confluence offers a precious opportunity to prevent foreseeable systemic injustices from taking hold. Low-income people, their advocates and allies, and broader interests concerned about AI-based decision-making must fight the immediate battles to reduce urgent harms, win legislative protections, and secure the resources needed to enforce rights and assert power. Such exigencies necessitate a certain pragmatism in advocacy to navigate imperfect solutions, unexpected alignments of interests, and unsatisfying compromises.

Still, pragmatism does not require that we bend our imaginations to the prosaic. Instead, we must situate the possible near-term progress within a broader vision for a more just world where people have what they need to thrive. Here, AI is unexpectedly helpful.

The all-encompassing nature of AI decision-making offers an unusual entryway to discuss the broader vision of what justice looks like with respect to housing, work, school, government benefits, family stability, and other areas not analyzed in this report, such as law enforcement. We know, for instance, that AI is used to exploit workers. So, we can imagine what is needed to prevent such exploitation: higher wages, safer working conditions, stable scheduling, generous paid leave policies, free or low-cost child care, support for adult family members who need care, better protections against employer abuses, and easier processes for forming unions and engaging in collective action. Knowing what is wrong helps to show what might be right.

Keeping the more just future in view while recognizing that so much of it may not be achieved in our lifetimes is challenging. After all, many conversations about a just future lead quickly to requiring major reorganizations of society and the economy to root out the various ills of poverty, inequality, exploitation, racism, sexism, ableism, homophobia, imperialism, environmental devastation, and so on.

With such a major restructuring being beyond the reach of any one person, I turn to something I learned as a frontline do-gooder: imagination through incremental action. I did not have the power to usher in a sweeping new utopian life for my clients. But I could work with them to temporarily touch a world where justice felt true and personal in a way that it usually did not. Those moments sparked hope and possibility that sustained my clients and me through the inevitable next fight, as the better future we imagined temporarily phased into the troubled present we knew.

This report is not individualized advocacy. It will almost certainly not make someone's life materially better in the next week, month, or year. But it may help a reader understand something that they did not previously, and that understanding may lead to concern, care, and action. With enough of that, we may just take a few steps forward. This, at least, is how I imagine it.

Appendix: Defining Poverty and Low Income

Defining Poverty

Poverty is officially measured using a formula implemented in the 1960s that looks at the cost of a minimally adequate food diet in 1963 and multiplies it by three.³¹⁵ That amount has been updated every year based on inflation but is otherwise unchanged.

Table 1. Current U.S. poverty levels.

Household Size	Federal Poverty Line
1	\$15,060
2	\$20,440
3	\$25,820
4	\$31,200
5	\$36,580
6	\$41,960

The formal federal poverty line does not account for changing demands on household budgets.³¹⁶ For example, housing costs consume much more of an average household budget now than they did in the 1960s, but that change is not reflected in the formal poverty measure.

It also does not account for regional variations in the cost of living, meaning that people struggling in high-cost areas, such as San Francisco or New York City, could be excluded. As a result of these and other criticisms beyond the scope of this report, the U.S. government developed an informal alternative, the Supplemental Poverty Measure, which tends to show slightly higher poverty rates than the formal measure.³¹⁷

Measurement complexities aside, 41 million people in the United States (12.6 percent of the population) live below the formal poverty line. Around 12 million of them are children (16.1 percent of all U.S. children). States with the most poverty have rates between 17 percent and 19 percent, with child poverty rates in those states between 22 percent and 26 percent. Counties with the highest poverty rates tend to be located near the Mississippi River, in Appalachia, or on American Indian reservations and have overall poverty rates around 25 percent to 35 percent (with the highest rates around 50 percent) and child poverty rates between 45 percent and 60 percent.³¹⁸ Some neighborhoods in cities have similar rates.³¹⁹

The consequences of poverty are, of course, dire. The U.S. Department of Health and Human Services recently summarized the impacts of poverty on health, economic mobility, and life expectancy:

[R]esidents of impoverished communities are at increased risk for mental illness, chronic disease, higher mortality, and lower life expectancy. ... Childhood poverty is associated with developmental delays, toxic stress, chronic illness, and nutritional deficits. Individuals who experience childhood poverty are more likely to experience poverty into adulthood, which contributes to generational cycles of poverty. In addition to lasting effects of childhood poverty, adults living in poverty are at a higher risk of adverse health effects from obesity, smoking, substance use, and chronic stress. Finally, older adults with lower incomes experience higher rates of disability and mortality. One study found that men and women in the top 1 percent of income were expected to live 14.6 and 10.1 years longer respectively than men and women in the bottom 1 percent.³²⁰

Looking at other important life outcomes, poverty is associated with lower levels of educational attainment³²¹ and higher rates of incarceration.³²²

The impacts of poverty fall disproportionately on groups that society already marginalizes, particularly Black, Hispanic, and American Indian people, women, people with disabilities, and LGBT (lesbian, gay, bisexual, or transgender) people.³²³

While white people constitute the largest group of people experiencing poverty, only 9.5 percent of white people are in poverty.³²⁴ In contrast, 21.4 percent of Black people, 16.7 percent of Hispanic people, and 24.5 percent of American Indian people live below the poverty line. Moreover, even though Black, Hispanic, and American Indian people make up only 11.7 percent, 19.2 percent, and 0.5 percent, respectively, of the U.S. population, they constitute 19.9 percent, 25.5 percent, and 1 percent of the people living in poverty (see Table 2 for side-by-side comparisons).

Table 2. U.S. poverty levels by race.

Race or Ethnicity	Percent of U.S. Population	Percent of Total People Living in Poverty	Percent of Race or Ethnic Group Experiencing Poverty
White	57.4	43.4	9.5
Black	11.7	19.9	21.4
Hispanic*	19.2	25.5	16.7
American Indian	0.5	1.0	24.5

**Persons of Hispanic origin may be of any race; all other racial/ethnic groups are non-Hispanic.*

Historic and ongoing policies and practices have racialized poverty. Though a full exploration is beyond the scope of this report, systemic injustices—slavery, segregation, forced deportation, land theft, exclusion from wealth-building policies such as the GI Bill, exclusion from certain neighborhoods, jobs, unions, or schools, systemic devaluation of property, forced displacement, usurious or inaccessible credit, pay differentials for identical jobs, racially targeted policing, and mass incarceration—have led and continue to lead to fewer economic opportunities and massive inequality for Black, Hispanic, and American Indian people.

This is reflected in a stark racial wealth gap wherein typical white households have between 6.5 times and 10 times as much wealth as typical Black households.³²⁵ Put differently, for every \$100 of wealth that typical white households have, Black households have between \$10 and \$16. Typical Latino households fare little better, having roughly \$22 in wealth for every \$100 a typical white household has.

While people in poverty of any race have less wealth than typical households, the racial wealth gap implicates the issues discussed here by perpetuating generational poverty, intensifying inequality, and illuminating the limitations of family support that Black and Latino people in poverty might turn to in comparison to their white counterparts.

Women also experience poverty disproportionately.³²⁶ More women than men are in poverty and a greater percentage of women than men are in poverty: 12.9 percent to 10.6 percent. Past and ongoing policies, practices, and conditions explain this. Women often have had limited opportunities to work for pay outside of the home, and where such opportunities exist, they have been in lower-paying jobs. Even when performing the same jobs as men, women earn less. In addition, women are forced or pressured to assume more housekeeping or caregiving responsibilities. They have less access to credit. And women have disproportionately been targets of intimate partner violence that impedes opportunities for work or school, causes eviction or otherwise-lost housing, and destabilizes other aspects of economic advancement. Moreover, the racialization of poverty is reflected in gender-based disparities, with Black, Latino, and American Indian women experiencing higher poverty rates than white women.

People who are disabled experience poverty disproportionately, with 27 percent of disabled people experiencing poverty, compared to 12 percent of nondisabled people.³²⁷ The dynamic works in two ways. First, poverty can cause disability through inadequate nutrition, environmental hazards, exposure to conditions that adversely affect development, lack of medical treatment, or physically demanding jobs that lead to injuries and work-related health conditions. Second, disability can cause poverty by restricting a person's ability to work, generating medical bills, and requiring care and support that costs money or takes away the time that a family member could work. Discrimination perpetuates work-related inequities, making it harder for disabled people to get and keep jobs and receive fair pay.

Importantly, characteristics such as race and gender intersect with disability and intensify disparities. For example, disabled Black people experience greater poverty than disabled white people, and disabled women experience greater poverty than disabled men.

Combined, LGBT people have a poverty rate of 17 percent, several points higher than the 12.6 percent national average and the 12 percent average for straight, cisgender people.³²⁸ Poverty rates for white, Black, Latino, and American Indian LGBT people are higher than their straight, cisgender counterparts of the same race, ethnicity, or gender.

Within the LGBT group, there is significant variation based on gender and race. For example, white LGBT people generally experience less poverty than LGBT people of other races and ethnicities. And gay cisgender men experience less poverty—10.3 percent—than lesbian cisgender women, bisexual cisgender men and women, and transgender people, of whom 21.2 percent experience poverty. Discrimination faced by LGBT people contributes to unemployment or underemployment, housing instability, intimate partner violence, and various forms of social or economic exclusion.

Defining Low-Income

This report focuses on AI's impact on low-income people, which encompasses more than just those people living below the federal poverty line. There are several reasons for this.

First, the official poverty measure does not capture the full extent of economic precarity in the United States. People living just or modestly above the poverty line still struggle to meet basic life needs, such as housing, food, clothing, medical treatment, schooling, transportation, basic technology (including internet access), family or community support, and leisure. Even if these needs are met, they are likely to be minimally adequate. Bills may regularly go unpaid. Significant other expenses, like a car repair or big medical or dental bills, can take months or years to recover from. More serious events, including the loss of a job or SNAP benefits, mean lost housing, skipped meals, missed medical treatment, or other ruinous consequences.

In short, people cannot simply absorb the fallout from tough times the way that people with higher incomes can. The shared experiences of precarity between people below and above the poverty line mean that any adverse effects of AI are likely to be felt in similar ways.

Second, and relatedly, people who today are modestly above the poverty line are often the same people who in other years are below it. That is, people cycle into and out of formal poverty with the vagaries of the economy and personal circumstances.³²⁹ As such, people living above the poverty line are generally seeking similar kinds of work, schooling, housing, and benefits as those below. Thus, they are equally likely to be exposed to AI systems across these issue areas.

Third, whether people are living below or above the poverty line when AI harms them, their ability to resist it is likely similar. Neither group can afford attorneys or expend money to increase their influence in political decision-making processes.

Ultimately, this report attempts to assess AI's impact on economically vulnerable people in the United States. Stopping at those in poverty alone would be missing key parts of the picture.

Even accepting that “low-income” is the proper frame of analysis, the term itself does not have a standard definition. Many government agencies or programs use the term to mean anyone below what the applicable income limit is for their services. Thus, low-income means 125 percent of the federal poverty line for legal aid³³⁰ or 80 percent of the median family income in a given region for federal housing assistance.³³¹ But both of these figures would exclude people eligible for public benefit programs such as the Supplemental Nutrition Assistance Program (which allows for gross household income up to 130 percent of the poverty line), Medicaid (138 percent of the poverty line), the Special Supplemental Nutrition Program for Women, Infants, and Children (185 percent of the poverty line) and the Children's Health Insurance Program (variable, but generally above 200 percent of the poverty line). If the law defines them as eligible for means-tested benefits, they can reasonably be considered low-income.

Still, merely looking at program eligibility excludes people who are ineligible because of non-income program requirements (such as immigration status) and people above the respective limits who still experience financial precarity. To quantify a fitting level of precarity, then,

researchers have defined low-income based on what it costs for necessities in particular geographic areas. While there are different ways to calculate the cost of basic needs, the methods generally agree that income must exceed 200 percent of the federal poverty line essentially everywhere and that, in some locations, much more is needed.³³²

Based on the above considerations, this report defines low-income as people with income below 200 percent of the federal poverty line.

Table 3. Low-income thresholds by household size.

Household Size	200 Percent of the Federal Poverty Line
1	\$30,120
2	\$40,880
3	\$51,640
4	\$62,400
5	\$73,160
6	\$83,920

Nationally, 92 million (28.3 percent of the overall U.S. population) live below this threshold, with states such as Mississippi, Louisiana, and Arkansas having close to 40 percent of their residents living below it.³³³

Of course, people with incomes above 200 percent of the poverty line also experience financial struggles and precarity, and they often do so while being financially ineligible for public benefits. But this report must draw a line somewhere, and 200 percent of the poverty line seems both appropriate for what is being described and manageable, given available data. Any case-specific deviations from this definition are noted in the applicable section of the report text.

Endnotes

- 1 There are other important safety net programs like the Children’s Health Insurance Program (CHIP), Women, Infants, and Children (WIC), Temporary Assistance to Needy Families (TANF), and Low-Income Home Energy Assistance Program (LIHEAP) that fall outside the scope of the report because I am not aware of distinct AI use in them. That said, states often incorporate some of these programs into their Medicaid or SNAP eligibility and enrollment systems such that failures in those systems could affect CHIP, WIC, TANF, and LIHEAP.
- 2 Some states have refused to extend Medicaid eligibility to all low-income people that federal laws allow. Even in states that have expanded eligibility, many immigrants are excluded, including those with lawful permanent resident status who have not met a five-year waiting period.
- 3 Medicaid.gov, “June 2024 Medicaid & CHIP Enrollment Data Highlights” (2024), available at <https://www.medicaid.gov/medicaid/program-information/medicaid-and-chip-enrollment-data/report-highlights/index.html> (last accessed October 28, 2024).
- 4 Kaiser Family Foundation, “Medicaid Income Eligibility Limits for Adults as a Percent of the Federal Poverty Level” (2024), available at <https://www.kff.org/affordable-care-act/state-indicator/medicaid-income-eligibility-limits-for-adults-as-a-percent-of-the-federal-poverty-level/?currentTimeframe=0&sortModel=%7B%22colId%22:%22Location%-22,%22sort%22:%22asc%22%7D>

- 5 Medicaid recipients requiring long-term care are classified as low income for purposes of this report even if their income may slightly exceed 200 percent of the federal poverty line because their experiences of disability make their lives as precarious as those with less income. And, practically, having the care needs they do often requires more costs than people without the same need for care.
- 6 Kaiser Family Foundation, “Medicaid Financing: The Basics” (2023), available at <https://www.kff.org/medicaid/issue-brief/medicaid-financing-the-basics/>.
- 7 Congressional Research Service, “Medicaid’s Federal Medical Assistance Percentage (FMAP)” (2020), available at <https://sgp.fas.org/crs/misc/R43847.pdf>.
- 8 Kaiser Family Foundation, “Federal Medical Assistance Percentage (FMAP) for Medicaid and Multiplier” (n.d.), available at <https://www.kff.org/medicaid/state-indicator/federal-matching-rate-and-multiplier/>.
- 9 Kaiser Family Foundation, “Medicaid Enrollment and Unwinding Tracker” (2024), available at <https://www.kff.org/report-section/medicaid-enrollment-and-unwinding-tracker-overview/>.
- 10 Amanda Seitz and Anita Snow, “Kicked Off Medicaid: Millions at risk as states trim rolls,” AP News, April 28, 2023, available at <https://apnews.com/article/medicaid-enrollees-removed-review-health-insurance-pandemic-bffc3c67ab2767e4e3cea8250683ea7a>; Ali Rogin and Kaisha Young, “Arkansas Medicaid recipients fight to stay covered after federal protections end,” PBS News, September 3, 2023, available at <https://www.pbs.org/newshour/show/arkansas-medicaid-recipients-fight-to-stay-covered-after-federal-protections-end>; Eleanor Klibanoff and Lomi Kriel, “Despite Persistent Warnings, Texas Rushed to Remove Millions from Medicaid. That Move Cost Eligible Residents Care.” ProPublica, September 26, 2024, available at <https://www.propublica.org/article/texas-medicaid-unwinding-consequences>.
- 11 National Health Law Program, Electronic Privacy Information Center, and Upturn, Inc., “Complaint and Request for Investigation, Injunction, and Other Relief in the Matter of Deloitte Consulting LLP” (2024), available at <https://healthlaw.org/wp-content/uploads/2024/01/NHeLP-EPIC-Upturn-FTC-Deloitte-Complaint.pdf>; Benefits Tech Advocacy Hub, “Case Study Library” (2024), available at <https://www.btah.org/case-studies.html>.
- 12 Molly O’Malley Watts, MaryBeth Musumeci, and Priya Chidambaram, “Medicaid Home and Community-Based Services Enrollment and Spending” (San Francisco: Kaiser Family Foundation, 2020), available at <https://www.kff.org/report-section/medicaid-home-and-community-based-services-enrollment-and-spending-issue-brief/>.

- 13 Molly O'Malley Watts, MaryBeth Musumeci, and Priya Chidambaram, "Medicaid Home and Community-Based Services Enrollment and Spending" (San Francisco: Kaiser Family Foundation, 2020), available at <https://www.kff.org/report-section/medicaid-home-and-community-based-services-enrollment-and-spending-issue-brief/>.
- 14 Two leading tools in the space, interRAI and SIS, are used in 25 states and 17 states, respectively, though there may be some overlap. In addition, other tools are available. See Steven Lutzky, "Advancing interrai: Multi-State Collaboration to Use Better Data Better" (Chicago: Shirley Ryan Ability Lab, 2023), available at <https://www.sralab.org/research/labs/rrtc-home-and-community-based-services/news/advancing-interrai-multi-state-collaboration-use-better-data-better>; American Association on Intellectual and Developmental Disabilities, "States and Provinces in North American Using the SIS" (n.d.), available at <https://www.aaid.org/sis/states-using-sis>.
- 15 Arkansas Department of Human Services, "ARIA Behavioral Health Tier Logic" (n.d.), available at https://humanservices.arkansas.gov/wp-content/uploads/ARIA_DBHS_Tier_Logic_002.pdf.
- 16 Andy Davis, "3 Care-Guiders OK'd by Arkansas Insurance Department for Medicaid Group," Arkansas Democrat-Gazette, October 24, 2017, available at <https://www.arkansasonline.com/news/2017/oct/24/3-care-guiders-ok-d-for-medicaid-group-/>.
- 17 Center for Medicaid and Medicare Services, "Capitated Rate Setting for MLTSS Programs" (n.d.), available at <https://www.medicaid.gov/medicaid/home-community-based-services/downloads/rate-setting-mltss-prgrms.pdf>.
- 18 This number is calculated by applying the percentage of people enrolled in Medicaid managed care plans—75 percent (calculated in 2021)—to the current 2024 Medicaid population of 75 million. See Elizabeth Hinton and Jada Raphael, "10 Things to Know About Medicaid Managed Care" (San Francisco: Kaiser Family Foundation, 2024), available at <https://www.kff.org/medicaid/issue-brief/10-things-to-know-about-medicaid-managed-care/>.
- 19 Department of Health and Human Services, "High Rates of Prior Authorization Denials by Some Plans and Limited State Oversight Raise Concerns About Access to Care in Medicaid Managed Care" (Office of the Inspector General, 2023), available at <https://oig.hhs.gov/oei/reports/OEI-09-19-00350.pdf>.
- 20 Government Accountability Office, "Medicaid: Managed Care Plans' Prior Authorization Decisions for Children Need Additional Oversight," GAO-24-106532 (2024), available at <https://files.gao.gov/reports/GAO-24-106532/index.html>.

- 21 Elizabeth Hinton and Jada Raphael, “10 Things to Know About Medicaid Managed Care” (San Francisco: Kaiser Family Foundation, 2024), available at <https://www.kff.org/medicaid/issue-brief/10-things-to-know-about-medicaid-managed-care/>.
- 22 Casey Ross and Bob Herman, “Denied by AI: How Medicare Advantage plans use algorithms to cut off care for seniors in need,” STAT News, March 13, 2023, available at <https://www.statnews.com/2023/03/13/medicare-advantage-plans-denial-artificial-intelligence/>.
- 23 Patrick Rucker, Maya Miller, and David Armstrong, “How Cigna Saves Millions by Having Its Doctors Reject Claims Without Reading Them,” ProPublica, March 25, 2023, available at <https://www.propublica.org/article/cigna-pxdx-medical-health-insurance-rejection-claims>.
- 24 T. Christian Miller, Patrick Rucker, and David Armstrong, “Not Medically Necessary: Inside the Company Helping America’s Biggest Health Insurers Deny Coverage for Care,” ProPublica, October 23, 2023, available at <https://www.propublica.org/article/evicore-health-insurance-denials-cigna-unitedhealthcare-aetna-prior-authorizations>.
- 25 Centers for Medicare & Medicaid Services, “Medicare Monthly Enrollment” (2024), available at <https://data.cms.gov/summary-statistics-on-beneficiary-enrollment/medicare-and-medicaid-reports/medicare-monthly-enrollment>.
- 26 Maria T. Peña and others, “A Profile of Medicare-Medicaid Enrollees (Dual Eligibles)” (San Francisco: Kaiser Family Foundation, 2023), available at <https://www.kff.org/medicare/issue-brief/a-profile-of-medicare-medicaid-enrollees-dual-eligibles/>.
- 27 See Meredith Freed and others, “Medicare Advantage in 2024: Enrollment Update and Key Trends” (San Francisco: Kaiser Family Foundation, 2024), available at <https://www.kff.org/medicare/issue-brief/medicare-advantage-in-2024-enrollment-update-and-key-trends>.
- 28 Meredith Freed and others, “Medicare Advantage in 2024: Enrollment Update and Key Trends” (San Francisco: Kaiser Family Foundation, 2024), available at <https://www.kff.org/medicare/issue-brief/medicare-advantage-in-2024-enrollment-update-and-key-trends/>.
- 29 Center for Medicare Advocacy, “Medicare Prior Authorization” (n.d.), available at <https://medicareadvocacy.org/prior-authorization/>.

- 30 Nancy Ochieng, Juliette Cubanski, and Tricia Neuman, “A Snapshot of Sources of Coverage Among Medicare Beneficiaries: Appendix Table 1. Sources of Coverage Among Medicare Beneficiaries, 2022” (San Francisco: Kaiser Family Foundation, 2024), available at <https://www.kff.org/medicare/issue-brief/a-snapshot-of-sources-of-coverage-among-medicare-beneficiaries/#AppendixTable1>.
- 31 Department of Health and Human Services, “High Rates of Prior Authorization Denials by Some Plans and Limited State Oversight Raise Concerns About Access to Care in Medicaid Managed Care” (Office of the Inspector General, 2023), available at <https://oig.hhs.gov/oei/reports/OEI-09-19-00350.pdf>.
- 32 Casey Ross and Bob Herman, “Denied by AI: How Medicare Advantage plans use algorithms to cut off care for seniors in need,” STAT News, March 13, 2023, available at <https://www.statnews.com/2023/03/13/medicare-advantage-plans-denial-artificial-intelligence/>; Susan Jaffe, “Feds Rein In Use of Predictive Software That Limits Care for Medicare Advantage Patients,” KFF Health News, October 5, 2023, available at <https://kffhealthnews.org/news/article/biden-administration-software-algorithms-medicare-advantage/>.
- 33 Kaiser Family Foundation, “Status of State Medicaid Expansion Decisions: Interactive Map” (2024), available at <https://www.kff.org/affordable-care-act/issue-brief/status-of-state-medicare-expansion-decisions-interactive-map/>.
- 34 Kaiser Family Foundation, “Health Insurance Coverage of the Population Ages 0-64 with Incomes below 200 percent Federal Poverty Level” (2023), available at <https://www.kff.org/other/state-indicator/health-insurance-coverage-population-0-64-up-to-200-fpl/>.
- 35 Kaiser Family Foundation, “Explaining Health Care Reform: Questions About Health Insurance Subsidies” (2024), available at <https://www.kff.org/affordable-care-act/issue-brief/explaining-health-care-reform-questions-about-health-insurance-subsidies/>.
- 36 Kaiser Family Foundation, “Marketplace Plan Selections by Household Income” (2024), available at <https://www.kff.org/affordable-care-act/state-indicator/marketplace-plan-selections-by-household-income-2/>.
- 37 When reviewing these figures, I compared KFF’s marketplace plan enrollment figures for low-income people with KFF’s figures for all health insurance coverage of low-income people under the age of 65. The latter does not have a category for marketplace plans. Potentially applicable categories, like non-group plans or other public coverage, feature numbers significantly below KFF’s figures for marketplace plan enrollment. Compare <https://www.kff.org/affordable-care-act/state-indicator/marketplace-plan-selections-by-household-income-2/> and <https://www.kff.org/other/state-indicator/health-insurance-coverage-population-0-64-up-to-200-fpl/>.

- 38 Kaye Pestaina and Karen Pollitz, “Examining Prior Authorization in Health Insurance” (San Francisco: Kaiser Family Foundation, 2022), available at <https://www.kff.org/policy-watch/examining-prior-authorization-in-health-insurance/>.
- 39 U.S. Department of Agriculture, “SNAP Data Tables” (Food and Nutrition Service, 2024), available at <https://www.fns.usda.gov/pd/supplemental-nutrition-assistance-program-snap>.
- 40 Center on Budget and Policy Priorities, “A Quick Guide to SNAP Eligibility and Benefits” (2024), available at <https://www.cbpp.org/research/food-assistance/a-quick-guide-to-snap-eligibility-and-benefits>.
- 41 Jillian Humphries, “Data Coordination at SNAP and Medicaid Agencies: A National Landscape Analysis” (Hamilton, NJ: Benefits Data Trust and Center for Health Care Strategies, 2023), available at https://bdtrust.org/data-coordination-at-snap-and-medicaid-agencies-part-2_may-2023.pdf.
- 42 This automation includes “data matching,” which happens when a state looks at available databases for relevant information about an applicant. While matching to a few specified databases is required under federal law, states differ in the additional databases to which they look. See Randy Alison Aussenberg, “Errors and Fraud in the Supplemental Nutrition Assistance Program (SNAP)” (Washington, DC: Congressional Research Service, 2018), available at <https://sgp.fas.org/crs/misc/R45147.pdf>.
- 43 Helen Gemmell and Mea Martinez, “Class Action Complaint: Helen Gemmell, and Mea Martinez, individually and on behalf of all others similarly situated v. Melba Depena Affigne, in her official capacity as Director of the Rhode Island Department of Human Services” (2016), available at https://www.riaclu.org/sites/default/files/field_documents/UHIP_SNAP_complaint_final.pdf.
- 44 Katie Davis, “Deloitte Credits \$60M to Rhode Island as more cases lost in DHS system,” NBC 10 News, October 24, 2017, available at <https://turnto10.com/i-team/nbc-i-team-rhode-island-receives-60-million-credit-from-deloitte>.
- 45 American Civil Liberties Union of Rhode Island and NCLEJ, “Settlement Announced in Lawsuit Over Food Stamp Benefit Delays Caused by UHIP Computer Debacle,” Press release, February 24, 2017, available at https://nclej.org/wp-content/uploads/2017/02/RI_SNAP_settlement_GemmellVBeane.pdf.
- 46 Eric J. Beane, “An Assessment of the Unified Health Infrastructure Project: A Report to Gov. Fina M. Raimondo” (Providence, RI: Rhode Island Department of Human Services, 2017), available at <http://www.transparency.ri.gov/uhip/documents/assessments/UHIP%2030-day%20assessment.pdf>.

- 47 Parker L. Gilkesson, “SNAP ‘Program Integrity’: How Racialized Fraud Provisions Criminalize Hunger” (Washington, DC: The Center for Law and Social Policy, 2022), available at https://www.clasp.org/wp-content/uploads/2022/04/2022_SNAP20Program20Integrity20-20How20Racialized20Fraud20Provisions20Criminalize20Hunger.pdf; Randy Alison Aussenberg, “Errors and Fraud in the Supplemental Nutrition Assistance Program (SNAP)” (Washington, DC: Congressional Research Service, 2018), available at <https://sgp.fas.org/crs/misc/R45147.pdf>.
- 48 Parker L. Gilkesson, “SNAP ‘Program Integrity’: How Racialized Fraud Provisions Criminalize Hunger” (Washington, DC: The Center for Law and Social Policy, 2022), available at https://www.clasp.org/wp-content/uploads/2022/04/2022_SNAP20Program20Integrity20-20How20Racialized20Fraud20Provisions20Criminalize20Hunger.pdf; Randy Alison Aussenberg, “Errors and Fraud in the Supplemental Nutrition Assistance Program (SNAP)” (Washington, DC: Congressional Research Service, 2018), available at <https://sgp.fas.org/crs/misc/R45147.pdf>.
- 49 The federal Government Accountability Office identified data matching and data mining as two techniques used. See Government Accountability Office, “Supplemental Nutrition Assistance Program: Disseminating Information on Successful Use of Data Analytics Could Help States Manage Fraud Risks,” GAO-19-115 (2018), available at <https://www.gao.gov/assets/gao-19-115.pdf>.
- 50 See, for example, *Shahid v. Dep’t of Health & Human Services*, 333 Mich. App. 267, 963 N.W.2d 638 (Mich. Ct. App. 2020), overturning, at the third level of appeal, an administrative hearing decision finding that a recipient committed fraud merely because he had suspicious transactions.
- 51 Virginia Eubanks, “Zombie debts are hounding struggling Americans. Will you be next?” *The Guardian*, October 15, 2019, available at <https://www.theguardian.com/law/2019/oct/15/zombie-debt-benefits-overpayment-poverty>; Virginia Eubanks, “The Shakedown State: Fraud Detection and Predicting the Past” (Washington, DC: Electronic Privacy Information Center, 2022), available at <https://epic.org/shakedownstate/>.
- 52 Generally, when SSDI or SS retirement benefits are less than what SSI pays, a person will qualify for SSI to supplement their income. Social Security Administration, “Monthly Statistical Snapshot, September 2024” (2024), available at https://www.ssa.gov/policy/docs/quickfacts/stat_snapshot/.
- 53 Center on Budget and Policy Priorities, “Policy Basics: Supplemental Security Income” (2024), available at <https://www.cbpp.org/research/social-security/supplemental-security-income>.

- 54 Center on Budget and Policy Priorities, “Policy Basics: Supplemental Security Income” (2024), available at <https://www.cbpp.org/research/social-security/supplemental-security-income>; Center on Budget and Policy Priorities, “Chart Book: Social Security Disability Insurance” (2024), available at <https://www.cbpp.org/research/social-security/social-security-disability-insurance-0>.
- 55 Natalie Alms, “30,000 Died in Fiscal 2023 Waiting for Disability Decisions from Social Security,” NextGov, available at <https://www.nextgov.com/digital-government/2024/04/30000-died-fiscal-2023-waiting-disability-decisions-social-security/395796/>.
- 56 Social Security Administration, “Monthly Statistical Snapshot, September 2024” (2024), available at https://www.ssa.gov/policy/docs/quickfacts/stat_snapshot/.
- 57 The asset limit for eligible couples is \$3,000. Though a full discussion is beyond the scope of this report, SSI’s income and asset rules make it so that recipients cannot marry without grave financial penalties.
- 58 Social Security Administration, “Disabled Worker Beneficiaries in Current Payment Status at the End of June 2024, Distributed by Benefit Level, Sex, and Age Group” (2024), available at https://www.ssa.gov/oact/progdata/benefits/da_mbc202406.html.
- 59 Center on Budget and Policy Priorities, “Chart Book: Social Security Disability Insurance” (2024), available at <https://www.cbpp.org/research/social-security/social-security-disability-insurance-0>.
- 60 Sarah Mancini, Kate Lang, and Chi Chi Wu, “Mismatched and Mistaken: How the Use of an Inaccurate Private Database Results in SSI Recipients Unjustly Losing Benefits” (Boston: National Consumer Law Center, 2021), available at <https://justiceinaging.org/wp-content/uploads/2021/04/SSADataReport.pdf>.
- 61 Social Security Administration, “AI Use Cases” (n.d.), available at <https://www.ssa.gov/data/SSA-AI-Inv.csv>.
- 62 Social Security Administration, “AI Use Cases” (n.d.), available at <https://www.ssa.gov/data/SSA-AI-Inv.csv>.
- 63 U.S. Department of Labor, “Unemployment Insurance: All Programs Data” (Employment and Training Administration, n.d.), available at <https://oui.doleta.gov/unemploy/docs/allprograms.xlsx>.
- 64 U.S. Department of Labor, “The Employment Situation – September 2024,” Bureau of Labor Statistics, Press release, October 4, 2024, available at <https://www.bls.gov/news.release/pdf/empsit.pdf>.

- 65 Internal Revenue Service, “Statistics of Income: 2021 Individual Income Tax Returns,” Table 1.4, available at <https://www.irs.gov/pub/irs-pdf/p1304.pdf#page=24>
- 66 Internal Revenue Service, “Statistics of Income: 2021 Individual Income Tax Returns,” Table 1.4, available at <https://www.irs.gov/pub/irs-pdf/p1304.pdf#page=24>
- 67 U.S. Department of Labor, “Unemployment Insurance Data” (Employment and Training Administration, n.d.), available at <https://oui.doleta.gov/unemploy/DataDashboard.asp>.
- 68 Center on Budget and Policy Priorities, “Policy Basics: How Many Weeks of Unemployment Compensation Are Available?” (2024), available at <https://www.cbpp.org/research/economy/how-many-weeks-of-unemployment-compensation-are-available>.
- 69 Lizette Alvarez, “Florida’s Site Said to Delay Millions in Aid to Jobless,” New York Times, January 10, 2014, available at <https://www.nytimes.com/2014/01/11/us/floridas-site-said-to-delay-millions-in-aid-to-jobless.html>.
- 70 StateScoop staff, “Florida, Deloitte, at odds over unemployment insurance system,” StateScoop, January 9, 2014, available at <https://statescoop.com/florida-deloitte-at-odds-over-unemployment-insurance-system/>.
- 71 State of Florida Auditor General, Department of Economic Opportunity, “Information Technology Operational Audit: CONNECT” Report No. 2015-107 (2015), available at https://flauditor.gov/pages/pdf_files/2015-107.pdf; State of Florida Auditor General, Department of Economic Opportunity, “Information Technology Operational Audit: Reemployment Assistance Claims and Benefits Information System” Report No. 2017-039 (2016), available at https://flauditor.gov/pages/pdf_files/2017-039.pdf; State of Florida Auditor General, Department of Economic Opportunity, “Information Technology Operational Audit: Reemployment Assistance Claims and Benefits Information System (CONNECT)” Report No. 2019-183 (2019), available at https://flauditor.gov/pages/pdf_files/2019-183.pdf; State of Florida Auditor General, Department of Economic Opportunity, “Information Technology Operational Audit: Reemployment Assistance Claims and Benefits Information System (CONNECT)” Report No. 2021-169 (2021), available at https://flauditor.gov/pages/pdf_files/2021-169.pdf.
- 72 George Wentworth and Claire McKenna, “Ain’t No Sunshine: Fewer than One in Eight Unemployed Workers in Florida Is Receiving Unemployment Insurance” (New York City: National Employment Law Project, 2015), available at <https://www.nelp.org/insights-research/aint-no-sunshine-florida-unemployment-insurance/>.

- 73 Lauren Aratani, “‘Designed for us to fail’: Floridians upset as unemployment system melts down,” *The Guardian*, April 15, 2020, available at <https://www.theguardian.com/us-news/2020/apr/15/florida-unemployment-system-coronavirus-applications>.
- 74 Lauren Wamsley, “Gov. Says Florida’s Unemployment System Was Designed To Create ‘Pointless Roadblocks,’” *NPR*, August 6, 2020, available at <https://www.npr.org/sections/coronavirus-live-updates/2020/08/06/899893368/gov-says-floridas-unemployment-system-was-designed-to-create-pointless-roadblock>.
- 75 Tampa Bay Times Editorial Board, “More evidence that Florida’s unemployment system was designed to fail,” *Tampa Bay Times*, March 23, 2021, available at <https://www.tampabay.com/opinion/2021/03/23/more-evidence-that-floridas-unemployment-system-was-designed-to-fail-editorial/>.
- 76 Shawn Donnan and Dina Bass, “How Did ID.me Get Between You and Your Identity?” *Bloomberg Businessweek*, January 20, 2022, available at <https://www.bloomberg.com/news/features/2022-01-20/cybersecurity-company-id-me-is-becoming-government-s-digital-gatekeeper>.
- 77 This information comes from state records I have reviewed.
- 78 Adrienne Roberts, “Another class action lawsuit filed against Michigan’s Unemployment Insurance Agency,” *Detroit Free Press*, August 31, 2022, available at <https://www.freep.com/story/news/local/michigan/2022/08/31/class-action-lawsuit-michigan-unemployment-insurance-agency/65465633007/>; *Kreps et al v. Michigan Unemployment Insurance Agency et al*, No. 2:2022-cv-12020, Dkt. 71 (Plaintiffs’ Motion for Summary Judgment) (E.D. Mich. 2023).
- 79 See *Arkansas Dep’t of Com., Div. of Workforce Servs. v. Legal Aid of Arkansas*, 2022 Ark. 130, 645 S.W.3d 9 (2022), holding that fraud-related algorithms in the state unemployment program did not meet the asserted Freedom of Information Act exemption for criminal investigations and that the agency was required to disclose the requested information.
- 80 The exemption is for “the methods used by the Division of Workforce Services to identify and investigate fraudulent claims.” See Ark. Code Ann. 11-10-314(a) (1); Act 197 of 2023.
- 81 Legislative Analyst’s Office, “The 2022-23 Budget: Assessing Proposals to Address Unemployment Insurance Fraud” (2022), available at <https://lao.ca.gov/Publications/Report/4542>.

- 82 U.S. Department of Labor, “Unemployment Insurance Integrity Strategic Plan” (Employment and Training Administration, n.d.), available at https://oui.doleta.gov/unemploy/pdf/ui_prog_integrity_2023.pdf; Government Accountability Office, “Unemployment Insurance: Estimated Amount of Fraud during Pandemic Likely Between \$100 Billion and \$135 Billion,” GAO-23-106696 (2023), available at <https://www.gao.gov/assets/gao-23-106696.pdf>; Government Accountability Office, “Pandemic Unemployment Assistance: States’ Controls to Address Fraud,” GAO-24-107471 (2024), available at <https://www.gao.gov/assets/gao-24-107471.pdf>; U.S. Department of Labor, “OIG Oversight of the Unemployment Insurance Program” (Office of the Inspector General, 2023), available at <https://www.oig.dol.gov/doloiguooversightwork.htm>.
- 83 U.S. Department of Commerce, “NIST Releases Second Public Draft of Digital Identity Guidelines for Final Review,” National Institute of Standards and Technology, Press release, August 21, 2024, available at <https://www.nist.gov/news-events/news/2024/08/nist-releases-second-public-draft-digital-identity-guidelines-final-review>.
- 84 Cah Cahoo v. SAS Analytics Inc., 912 F.3d 887(6th Cir. 2019); Benefits Tech Advocacy Hub, “Michigan Unemployment Insurance False Fraud Determinations” (n.d.), available at <https://www.btah.org/case-study/michigan-unemployment-insurance-false-fraud-determinations.html>
- 85 Keely Quinlan, “Automated public-benefit fraud detection used by states subject of new FTC complaint,” StateScoop, January 4, 2021, available at <https://statescoop.com/automated-public-benefit-fraud-detection-state-ftc-complaint/>; Electronic Privacy Information Center, “Complaint and Request for Investigation, Injunction, and Other Relief in the Matter of Thomson Reuters Corporation, West Publishing Corporation, and Pondera Solutions, LLC” (2024), available at <https://epic.org/wp-content/uploads/2024/01/EPIC-FTC-Thomson-Reuters-Complaint.pdf>.
- 86 Todd Feathers, “Google’s AI Will Help Decide Whether Unemployed Workers Get Benefits,” Gizmodo, September 10, 2024, available at <https://gizmodo.com/googles-ai-will-help-decide-whether-unemployed-workers-get-benefits-2000496215>.
- 87 U.S. Department of Labor, “Lower Authority Reversal Rates (Regular Claims): June 2024” (Employment and Training Administration, n.d.), available at https://oui.doleta.gov/unemploy/docs/reversal/LA_reversal_Jun_2024.xlsx.

- 88 Keely Quinlan, “Chatbot Snapshot: How State, Local Government Websites Use AI Assistants,” StateScoop, July 17, 2024, available at <https://statescoop.com/government-ai-chatbots-state-local-websites-2024/>; Tzuhao Chen and others, “The Adoption and Implementation of Artificial Intelligence Chatbots in Public Organizations: Evidence from U.S. State Governments,” The American Review of Public Administration (2023), available at <https://www.researchgate.net/publication/373879520>
- 89 This terminology is unfortunate because it stigmatizes people with mental health conditions who perceive physical sensations, voices, or images that others without such conditions do not.
- 90 Joint Center for Housing Studies of Harvard University, “America’s Rental Housing 2024” (2024), available at https://www.jchs.harvard.edu/sites/default/files/reports/files/Harvard_JCHS_Americas_Rental_Housing_2024.pdf; Joint Center for Housing Studies of Harvard University, “State of the Nation’s Housing 2024: Appendix Tables” (n.d.), available at https://www.jchs.harvard.edu/sites/default/files/interactive-item/files/Harvard_JCHS_State_Nations_Housing_2024_Appendix_Tables.xlsx.
- 91 U.S. Census Bureau, “Average Household Size of Occupied Units by Tenure,” available at <https://data.census.gov/table?q=B25010:%20AVERAGE%20HOUSEHOLD%20SIZE%20OF%20OCCUPIED%20HOUSING%20UNITS%20BY%20TENURE> (last accessed November 4, 2024).
- 92 Importantly, rent and utilities now constitute more than 30 percent of household income in half of all rental households. So, even people who are not considered “low income” have increasingly less income available for nonhousing costs. See Peyton Whitney, “More Than 42 Million US Households Were Cost Burdened in 2022” (Cambridge, MA: Joint Center for Housing Studies of Harvard University, 2024), available at <https://www.jchs.harvard.edu/blog/more-42-million-us-households-were-cost-burdened-2022>.
- 93 Kathryn Reynolds and Elizabeth Burton, “Almost Half of Renter Households Feel Pressured to Leave Their Homes” (Washington, DC: Urban Institute, 2023), available at <https://www.urban.org/urban-wire/almost-half-renter-households-feel-pressured-leave-their-homes>.
- 94 U.S. Department of Housing and Urban Development, “Public Housing (PH) Data Dashboard” (n.d.), available at https://www.hud.gov/program_offices/public_indian_housing/programs/ph/PH_Dashboard; U.S. Department of Housing and Urban Development, “HUD’s Public Housing Program” (n.d.), available at https://www.hud.gov/topics/rental_assistance/phprog.

- 95 Molly Parker, “HUD Long Neglected These Residents. Now As They Move Out, Some Feel HUD Let Them Down Again.” ProPublica, April 6, 2018, available at <https://www.propublica.org/article/cairo-thebes-hud-long-neglected-these-residents-now-as-they-move-out-some-feel-let-down-again>.
- 96 Angela Caputo and Sharon Lerner, “House poor, pollution rich,” APM Reports, January 13, 2021, available at <https://www.apmreports.org/story/2021/01/13/public-housing-near-polluted-superfund-sites>.
- 97 Center on Budget and Policy Priorities, “Policy Basics: The Housing Choice Voucher Program” (2024), available at <https://www.cbpp.org/research/policy-basics-the-housing-choice-voucher-program>.
- 98 Center on Budget and Policy Priorities, “Policy Basics: Public Housing” (2024), available at <https://www.cbpp.org/research/policy-basics-introduction-to-public-housing>; Center on Budget and Policy Priorities, “Policy Basics: The Housing Choice Voucher Program” (2024), available at <https://www.cbpp.org/research/policy-basics-the-housing-choice-voucher-program>.
- 99 Prasanna Rajasekaran, Mark Treskon, and Solomon Greene, “Rent Control: What Does the Research Tell Us about the Effectiveness of Local Action?” (Washington, DC: Urban Institute, 2019), available at https://www.urban.org/sites/default/files/publication/99646/rent_control_what_does_the_research_tell_us_about_the_effectiveness_of_local_action_1.pdf.
- 100 NYC Rent Guidelines Board, “Rent Control FAQ” (n.d.), available at <https://rentguidelinesboard.cityofnewyork.us/resources/faqs/rent-control/#continue>; Office of Rent Administration, “Rent Stabilization and Rent Control” (Homes and Community Renewal Division, 2024), available at https://hcr.ny.gov/system/files/documents/2024/01/fact-sheet-01-01-2024_0.pdf.
- 101 National Low Income Housing Coalition, “The Gap: A Shortage of Affordable Homes” (2024), available at https://nlihc.org/sites/default/files/gap/2024/Gap-Report_2024.pdf.
- 102 Joint Center for Housing Studies of Harvard University, “America’s Rental Housing 2024” (2024), available at https://www.jchs.harvard.edu/sites/default/files/reports/files/Harvard_JCHS_Americas_Rental_Housing_2024.pdf.
- 103 Banks and other mortgage lenders also use the tools to determine whether to give mortgages to prospective buyers. While low-income people can and do own homes, home lending practices are outside the scope of the report.
- 104 Chi Chi Wu and others, “Digital Denials: How Abuse, Bias, and Lack of Transparency in Tenant Screening Harm Renters” (Boston: National Consumer Law Center, 2023), available at https://www.nclc.org/wp-content/uploads/2023/09/202309_Report_Digital-Denials.pdf.

- 105 Natasha Duarte and Mariah de Leon, “Response to the FTC’s request for information on tenant screening technologies” (Washington, DC: Upturn, 2023), available at <https://www.upturn.org/work/response-to-the-ftcs-request-for-information-on-tenant-screening/>.
- 106 TechEquity, “Tech, Bias, and Housing Initiative: Tenant Screening” (2022), available at <https://techequity.us/2022/02/23/tech-bias-and-housing-initiative-tenant-screening/>.
- 107 Kaveh Waddell, “How Tenant Screening Reports Make It Hard for People to Bounce Back from Tough Times,” Consumer Reports, March 11, 2021, available at <https://www.consumerreports.org/electronics/algorithmic-bias/tenant-screening-reports-make-it-hard-to-bounce-back-from-tough-times-a2331058426/>; Lauren Kirchner and Matthew Goldstein, “Access Denied: Faulty Automated Background Checks Freeze Out Renters,” The Markup, May 28, 2020, available at <https://themarkup.org/locked-out/2020/05/28/access-denied-faulty-automated-background-checks-freeze-out-renters>.
- 108 TechEquity, “Screened Out of Housing: How AI-Powered Tenant Screening Hurts Renters” (2024), available at <https://techequity.us/wp-content/uploads/2024/07/Screened-out-of-housing-paper.pdf>.
- 109 TransUnion, “Low Turnover and Higher Rental Prices in 2017 Driving Profitable and Attractive Market for Landlords,” Press release, April 19, 2017, available at <https://www.globenewswire.com/news-release/2017/04/19/963170/0/en/Low-Turnover-and-Higher-Rental-Prices-in-2017-Driving-Profitable-and-Attractive-Market-for-Landlords.html>
- 110 Cyrus Farivar, “Tenant screening software faces national reckoning,” NBC News, March 14, 2021, available at <https://www.nbcnews.com/tech/tech-news/tenant-screening-software-faces-national-reckoning-n1260975>.
- 111 Lauren Kirchner and Matthew Goldstein, “Access Denied: Faulty Automated Background Checks Freeze Out Renters,” The Markup, May 28, 2020, available at <https://themarkup.org/locked-out/2020/05/28/access-denied-faulty-automated-background-checks-freeze-out-renters>.
- 112 Lauren Kirchner and Matthew Goldstein, “Access Denied: Faulty Automated Background Checks Freeze Out Renters,” The Markup, May 28, 2020, available at <https://themarkup.org/locked-out/2020/05/28/access-denied-faulty-automated-background-checks-freeze-out-renters>.
- 113 Colin Lecher, “Automated background checks are deciding who’s fit for a home,” The Verge, February 1, 2019, available at <https://www.theverge.com/2019/2/1/18205174/automation-background-check-criminal-records-corelogic>.

- 114 Chi Chi Wu and others, “Digital Denials: How Abuse, Bias, and Lack of Transparency in Tenant Screening Harm Renters” (Boston: National Consumer Law Center, 2023), available at https://www.nclc.org/wp-content/uploads/2023/09/202309_Report_Digital-Denials.pdf.
- 115 Heather Vogell and Haru Coryne, “Rent Going Up? One Company’s Algorithm Could Be Why.” ProPublica, October 15, 2022, available at <https://www.propublica.org/article/yieldstar-rent-increase-realtor-rent>; Emma Roth, “US sues RealPage over rent-setting software that allegedly drove up prices,” The Verge, August 23, 2024, available at <https://www.theverge.com/2024/8/23/24226885/doj-realpage-antitrust-lawsuit-rent-fixing-software>.
- 116 U.S. Department of Justice, “Complaint in the Matter of United States, State of North Carolina, State of California, State of Colorado, State of Connecticut, State of Minnesota, State of Oregon, State of Tennessee, and State of Washington v. RealPage, Inc.” (2024), available at <https://www.justice.gov/opa/media/1364976/dl?inline>.
- 117 American Economic Liberties Project, Local Progress, and Local Progress Impact Lab, “A New Culprit in the Housing Crisis: Rent-setting Software Algorithms” (2024), available at <http://www.economicliberties.us/wp-content/uploads/2024/03/Policy-Memo-Rent-Setting-Software-Algorithms.pdf>
- 118 In re RealPage, Inc., Rental Software Antitrust Litig., 3:23-md-03071 (M.D. Tenn. Dec. 28, 2023) — Dkt. 687, available at <https://www.tnmd.uscourts.gov/sites/tnmd/files/687.pdf>. CHECK FORMAT
- 119 U.S. Department of Justice, “Complaint in the Matter of United States, State of North Carolina, State of California, State of Colorado, State of Connecticut, State of Minnesota, State of Oregon, State of Tennessee, and State of Washington v. RealPage, Inc.” (2024), available at <https://www.justice.gov/opa/media/1364976/dl?inline>.
- 120 Office of the Attorney General for the District of Columbia, “Complaint in the Matter of District of Columbia v. RealPage Inc., Avenue5 Residential LLC, AvalonBay Communities Inc., Bell Partners Inc., Bozzuto Management Company, Camden Development Inc., Equity Residential Management LLC, Gables Residential Services Inc., Greystar Management Services LP, Highmark Residential LLC, JBG Associates LLC, Mid-America Apartments LP, Paradigm Management LP, UDR Inc., and William C. Smith and Co., Inc.” (2023), available at <https://oag.dc.gov/sites/default/files/2023-11/DC%20OAG%20RealPage%20Complaint%20-%20Filed.pdf>

- 121 Office of the Attorney General of Arizona, “Complaint in the Matter of State of Arizona v. RealPage Inc., Apartment Management Consultants LLC, Avenue5 Residential LLC, BH Management Services LLC, Camden Property Trust; Crow Holdings LP, Trammell Crow Residential Company, Greystar Management Services LP, HSL Properties Inc., HSL Asset Management LLC, RPM Living LLC, and Weidner Property Management LLC” (2024), available at <https://www.azag.gov/sites/default/files/2024-02/RealPage%20Complaint.pdf>.
- 122 The DOJ lawsuit lists regional submarkets with significant AI penetration rates. Presumably, a researcher with appropriate expertise could use the information about submarkets to extrapolate demographic estimates.
- 123 Heather Vogell and Haru Coryne, “Rent Going Up? One Company’s Algorithm Could Be Why.” ProPublica, October 15, 2022, available at <https://www.propublica.org/article/yardstar-rent-increase-realpage-rent>.
- 124 Heather Vogell and Haru Coryne, “Rent Going Up? One Company’s Algorithm Could Be Why.” ProPublica, October 15, 2022, available at <https://www.propublica.org/article/yardstar-rent-increase-realpage-rent>.
- 125 Douglas MacMillan, “Eyes on the poor: Cameras, facial recognition watch over public housing,” Washington Post, May 16, 2023, available at <https://www.washingtonpost.com/business/2023/05/16/surveillance-cameras-public-housing/>; Tanvi Misra, “The Tenants Fight Back Against Facial Recognition Technology,” Bloomberg, May 7, 2019, available at <https://www.bloomberg.com/news/articles/2019-05-07/when-facial-recognition-tech-comes-to-housing>.
- 126 Gender Shades, “Overview,” available at <http://gendershades.org/overview.html> (last accessed October 29, 2024); Larry Hardesty, “Study finds gender and skin-type bias in commercial artificial-intelligence systems,” MIT News, February 11, 2018, available at <https://news.mit.edu/2018/study-finds-gender-skin-type-bias-artificial-intelligence-systems-0212>.
- 127 U.S. Department of Commerce, “NIST Study Evaluates Effects of Race, Age, Sex on Face Recognition Software,” National Institute of Standards and Technology, Press release, December 19, 2019, available at <https://www.nist.gov/news-events/news/2019/12/nist-study-evaluates-effects-race-age-sex-face-recognition-software>; U.S. Department of Commerce, “Face Recognition Vendor Test (FRVT): Part 8: Summarizing Demographic Differentials” (National Institute of Standards and Technology, 2022), available at https://pages.nist.gov/frvt/reports/demographics/nistir_8429.pdf; U.S. Department of Commerce, “Face Recognition Technology Evaluation: Demographic Effects in Face Recognition” (National Institute of Standards and Technology, n.d.), available at https://pages.nist.gov/frvt/html/frvt_demographics.html.

- 128 Kashmir Hill, “Facial Recognition Led to Wrongful Arrests. So Detroit Is Making Changes.” *New York Times*, June 29, 2024, available at <https://www.nytimes.com/2024/06/29/technology/detroit-facial-recognition-false-arrests.html>.
- 129 Center on Budget and Policy Priorities, “Policy Basics: Public Housing” (2024), available at <https://www.cbpp.org/research/policy-basics-introduction-to-public-housing>.
- 130 Tanvi Misra, “The Tenants Fight Back Against Facial Recognition Technology,” *Bloomberg*, May 7, 2019, available at <https://www.bloomberg.com/news/articles/2019-05-07/when-facial-recognition-tech-comes-to-housing>.
- 131 Erin Durkin, “New York tenants fight as landlords embrace facial recognition cameras,” *The Guardian*, May 30, 2019, available at <https://www.theguardian.com/cities/2019/may/29/new-york-facial-recognition-cameras-apartment-complex>.
- 132 Kristin F. Butcher and Diane Whitmore Schanzenbach, “Most Workers in Low-Wage Labor Market Work Substantial Hours, in Volatile Jobs” (Washington, DC: Center on Budget and Policy Priorities, 2018), available at <https://www.cbpp.org/research/food-assistance/most-workers-in-low-wage-labor-market-work-substantial-hours-in-volatile>.
- 133 WorkRise, “Who Is the Low-Wage Workforce?” (2023), available at <https://www.workrisenetwork.org/features/who-low-wage-workforce>.
- 134 To express the thresholds in 2024 dollars, this report takes the two studies’ provided thresholds (\$16.98 per hour in 2021 dollars and \$16.03 in 2016 dollars, respectively) and runs them through the consumer price index inflation calculator. See Bureau of Labor Statistics, CPI Inflation Calculator, available at https://www.bls.gov/data/inflation_calculator.htm. This methodology does not account for changes in real wages that would show up in the current median wage benchmark or changed distributions of workers falling below the thresholds.
- 135 Martha Ross and Nicole Bateman, “Meet the Low-Wage Workforce” (Washington: Metropolitan Policy Program at Brookings, 2019), available at https://www.brookings.edu/wp-content/uploads/2019/11/201911_Brookings-Metro_low-wage-workforce_Ross-Bateman.pdf.

- 136 Because any full-time worker earning above about \$14.50 per hour would, on an annual basis, exceed 200 percent of the federal poverty line for a household of one person, I use the term “low-wage workers” for this section instead of “low-income people” except where I can confirm the low-income status. Available data establishes that at least 30 percent of low-wage workers fall below 150 percent of the federal poverty line and suggests that many more would be considered low-income due to household size or variations in employment that result in lower annual earnings than their hourly wages alone would suggest. See Martha Ross and Nicole Bateman, “Meet the Low-Wage Workforce” (Washington: Metropolitan Policy Program at Brookings, 2019), available at https://www.brookings.edu/wp-content/uploads/2019/11/201911_Brookings-Metro_low-wage-workforce_Ross-Bateman.pdf.
- 137 U.S. Department of Labor, “Workers with a disability more concentrated in service occupations than those with no disability,” TED blog, October 10, 2017, available at <https://www.bls.gov/opub/ted/2017/workers-with-a-disability-more-concentrated-in-service-occupations-than-those-with-no-disability.htm>.
- 138 Muhammad Ali and others, “Discrimination through optimization: How Facebook’s ad delivery can lead to skewed outcomes” (Washington: Upturn, 2019), available at <https://arxiv.org/pdf/1904.02095>.
- 139 In response to an anti-discrimination lawsuit, Facebook stopped allowing employers to specify the race or gender of people the employers most wanted to see their ads. By the time of the study, employers could no longer specify these characteristics.
- 140 Aaron Rieke and others, “Essential Work: Analyzing the Hiring Technologies of Large Hourly Employers” (Washington: Upturn, 2021), available at <https://www.upturn.org/static/reports/2021/essential-work/files/upturn-essential-work.pdf>; Miranda Bogen and Aaron Rieke, “Help Wanted: An Examination of Hiring Algorithms, Equity, and Bias” (2018), available at <https://www.upturn.org/static/reports/2018/hiring-algorithms/files/Upturn%20--%20Help%20Wanted%20-%20An%20Exploration%20of%20Hiring%20Algorithms,%20Equity%20and%20Bias.pdf>; David Kaz and others, “Employment Processes as Barriers to Employment in the Lower-Wage Labor Market: Literature Review” (Rockville, MD: Abt Associates, 2024), available at https://www.acf.hhs.gov/sites/default/files/documents/opre/OPRE-employment-processes-barriers-feb2024_Final.pdf.
- 141 Consumer Financial Protection Bureau, “Market Snapshot: Background Screening Reports” (2019), available at https://files.consumerfinance.gov/f/documents/201909_cfpb_market-snapshot-background-screening_report.pdf.

- 142 Emily Vogels, “Digital divide persists even as Americans with lower incomes make gains in tech adoption” (Washington, DC: Pew Research Center, 2021), available at <https://www.pewresearch.org/short-reads/2021/06/22/digital-divide-persists-even-as-americans-with-lower-incomes-make-gains-in-tech-adoption/>.
- 143 Phela Townsend, “Disconnected: How the Digital Divide Harms Workers and What We Can Do about It” (New York City: The Century Foundation, 2020), available at <https://tcf.org/content/report/disconnected-digital-divide-harms-workers-can/>.
- 144 David Kaz and others, “Employment Processes as Barriers to Employment in the Lower-Wage Labor Market: Literature Review” (Rockville, MD: Abt Associates, 2024), available at https://www.acf.hhs.gov/sites/default/files/documents/opre/OPRE-employment-processes-barriers-feb2024_Final.pdf.
- 145 Economic Policy Institute, “Low Wage Workforce Tracker,” available at <https://www.epi.org/low-wage-workforce/> (last accessed November 4, 2024); Jesse Wursten and Michael Reich, “Small Businesses and the Minimum Wage,” (Berkeley, CA: Institute for Research on Labor and Employment, 2023), available at <https://irle.berkeley.edu/wp-content/uploads/2023/03/Small-Businesses-and-the-Minimum-Wage.pdf>.
- 146 Veenal Dubal, “On Algorithmic Wage Discrimination,” Columbia Law Review 123 (2023), available at https://columbialawreview.org/wp-content/uploads/2023/11/Dubal-On_Algorithmic_Wage_discrimination.pdf.
- 147 James Parrott and Michael Reich, “Comparison of two Seattle TNC driver studies” (Berkeley, CA: Center on Wage and Employment Dynamics and Institute for Research on Labor and Employment, 2020), available at <https://irle.berkeley.edu/wp-content/uploads/2020/07/Comparison-of-two-Seattle-studies-3.pdf>; Lawrence Mishel, “Uber and the labor market” (Washington, DC: Economic Policy Institute, 2018), available at <https://www.epi.org/publication/uber-and-the-labor-market-uber-drivers-compensation-wages-and-the-scale-of-uber-and-the-gig-economy/>.
- 148 Veenal Dubal, “On Algorithmic Wage Discrimination,” Columbia Law Review 123 (2023), available at https://columbialawreview.org/wp-content/uploads/2023/11/Dubal-On_Algorithmic_Wage_discrimination.pdf.

- 149 Rakesh Kochhar, “The Enduring Grip of the Gender Pay Gap” (Washington, DC: Pew Research Center, 2023), available at <https://www.pewresearch.org/social-trends/2023/03/01/the-enduring-grip-of-the-gender-pay-gap/>; Elise Gould, “Gender wage gap persists in 2023” (Washington, DC: Economic Policy Institute, 2024), available at <https://www.epi.org/blog/gender-wage-gap-persists-in-2023-women-are-paid-roughly-22-less-than-men-on-average/#:~:text=White%20women%20and%20AAPI%20women,for%20a%20full%2Dtime%20worker.>
- 150 Rakesh Kochhar, “The Enduring Grip of the Gender Pay Gap” (Washington, DC: Pew Research Center, 2023), available at <https://www.pewresearch.org/social-trends/2023/03/01/the-enduring-grip-of-the-gender-pay-gap/>; Elise Gould, “Gender wage gap persists in 2023” (Washington, DC: Economic Policy Institute, 2024), available at <https://www.epi.org/blog/gender-wage-gap-persists-in-2023-women-are-paid-roughly-22-less-than-men-on-average/#:~:text=White%20women%20and%20AAPI%20women,for%20a%20full%2Dtime%20worker>; Stephen Miller, “Black Workers Still Earn Less than Their White Counterparts” (Alexandria, VA: Society for Human Resource Management, 2020), available at <https://www.shrm.org/topics-tools/news/benefits-compensation/black-workers-still-earn-less-white-counterparts>; Payscale, “2024 Gender Pay Gap Report” (2024), available at <https://www.payscale.com/research-and-insights/gender-pay-gap/>.
- 151 Monica Anderson and others, “The State of Gig Work in 2021” (Washington, DC: Pew Research Center, 2021), available at <https://www.pewresearch.org/internet/2021/12/08/the-state-of-gig-work-in-2021/>.
- 152 Jeffrey R. Young, “Algorithms are Changing Low-Wage Work. What Educators Should Know About Life ‘On the Clock.’” EdSurge Podcast, November 26, 2019, available at <https://www.edsurge.com/news/2019-11-26-algorithms-are-changing-low-wage-work-what-educators-should-know-about-life-on-the-clock>; Annette Bernhardt, Reem Suleiman, and Lisa Kresge, “Data and Algorithms at Work: The Case for Worker Technology Rights” (Berkeley, CA: University of California, Berkeley Labor Center, 2021), available at <https://laborcenter.berkeley.edu/data-algorithms-at-work/>; Alexandra Mateescu and Aiha Nguyen, “Algorithmic Management in the Workplace” (New York City: Data & Society, 2019), available at https://datasociety.net/wp-content/uploads/2019/02/DS_Algorithmic_Management_Explainer.pdf.
- 153 Rashi Shrivastava, “Dozens Of KFC, Taco Bell And Dairy Queen Franchises Are Using AI To Track Workers,” Forbes, February 23, 2024, available at <https://www.forbes.com/sites/rashishrivastava/2024/02/23/dozens-of-kfc-taco-bell-and-dairy-queen-franchises-are-using-ai-to-track-workers/?sh=133e21396c11>.

- 154 Heather Haddon, “Surge Pricing Is Coming to More Menus Near You,” Wall Street Journal, March 10, 2024, available at <https://www.wsj.com/business/hospitality/surge-pricing-is-coming-to-more-menus-near-you-66a245f3?mod=mhp>; Whizy Kim, “Uber-style pricing is coming for everything,” Vox, March 19, 2024, available at <https://www.vox.com/money/24105250/fast-food-restaurants-dynamic-pricing-algorithm-wendys>.
- 155 Matt Novak, “Kroger’s Plan to Use Facial Recognition Raises Concerns About Surge Pricing,” Gizmodo, October 16, 2024, available at <https://gizmodo.com/krogers-plan-to-use-facial-recognition-raises-concerns-about-surge-pricing-2000512685>; Mandy Taheri, “Lawmakers Raise Concerns About ‘Dynamic Pricing’ at Kroger,” Newsweek, October 17, 2024, available at <https://www.newsweek.com/lawmakers-raise-concerns-about-dynamic-pricing-kroger-1970110>.
- 156 Annette Bernhardt, Reem Suleiman, and Lisa Kresge, “Data and Algorithms at Work: The Case for Worker Technology Rights” (Berkeley, CA: University of California, Berkeley Labor Center, 2021), available at <https://laborcenter.berkeley.edu/data-algorithms-at-work/>; Alexandra Mateescu and Aiha Nguyen, “Algorithmic Management in the Workplace” (New York City: Data & Society, 2019), available at https://datasociety.net/wp-content/uploads/2019/02/DS_Algorithmic_Management_Explainer.pdf; Alexandra Mateescu and Aiha Nguyen, “Workplace Monitoring & Surveillance” (New York City: Data & Society, 2019), available at https://datasociety.net/wp-content/uploads/2019/02/DS_Workplace_Monitoring_Surveillance_Explainer.pdf; Alexandra Mateescu, “Explainer: Challenging Worker Datafication” (New York City: Data & Society, 2023), available at https://datasociety.net/wp-content/uploads/2023/11/DS_Explainer-Challenging-Worker-Datafication.pdf.
- 157 Alexandra Mateescu and Aiha Nguyen, “Algorithmic Management in the Workplace” (New York City: Data & Society, 2019), available at https://datasociety.net/wp-content/uploads/2019/02/DS_Algorithmic_Management_Explainer.pdf.
- 158 American Bar Association, “How much employee monitoring is too much?” (2018), available at <https://www.americanbar.org/news/abanews/publications/youraba/2018/january-2018/how-much-employee-monitoring-is-too-much/>.
- 159 Amazon, “Everything You Should Know About Amazon’s Hourly Wage and Benefits Offerings for Fulfillment Employees,” available at <https://www.aboutamazon.com/news/workplace/how-much-does-amazon-pay> (last accessed November 6, 2024).

- 160 Seeta Peña Gangadharan, “Tracked and Targeted: Navigating Worker Surveillance at Amazon” (Los Angeles, CA: Our Data Bodies, 2023), available at <https://www.odbproject.org/wp-content/uploads/2024/01/Tracked-and-Targeted.pdf> ; Mishal Khan, “At Work and Under Watch: Surveillance and Suffering at Amazon and Walmart Warehouses” (Boston, MA: Oxfam America, 2024), available at https://webassets.oxfamamerica.org/media/documents/At_Work_and_Under_Watch.pdf.
- 161 Beth Gutelius, and Sanjay Pinto, “Pain Points: Data on Work Intensity, Monitoring, and Health at Amazon Warehouses” (Chicago, IL: University of Illinois at Chicago, 2023), available at https://indigo.uic.edu/articles/report/Pain_Points_Data_on_Work_Intensity_Monitoring_and_Health_at_Amazon_Warehouses/24435124?file=44065940.
- 162 Jodi Kantor and Arya Sundaram, “The Rise of the Worker Productivity Score,” New York Times, August 14, 2022, available at <https://www.nytimes.com/interactive/2022/08/14/business/worker-productivity-tracking.html>.
- 163 PHI, “Direct Care Workers in the United States: Key Facts” (Los Angeles, CA: PHI, 2023), available at <https://www.phinational.org/national-resource-center/resource/direct-care-workers-in-the-united-states-key-facts-2024/>.
- 164 Virginia Eubanks and Alexandra Mateescu, “We Don’t Deserve This: New App Places U.S. Caregivers Under Digital Surveillance,” The Guardian, July 28, 2021, available at <https://www.theguardian.com/us-news/2021/jul/28/digital-surveillance-caregivers-artificial-intelligence>; Alexandra Mateescu, “Electronic Visit Verification: The Weight of Surveillance and the Fracturing of Care” (New York City: Data & Society, 2021), available at https://datasociety.net/wp-content/uploads/2021/11/EVV_REPORT_11162021.pdf.
- 165 Virginia Eubanks and Alexandra Mateescu, “We Don’t Deserve This: New App Places U.S. Caregivers Under Digital Surveillance,” The Guardian, July 28, 2021, available at <https://www.theguardian.com/us-news/2021/jul/28/digital-surveillance-caregivers-artificial-intelligence>; Alexandra Mateescu, “Electronic Visit Verification: The Weight of Surveillance and the Fracturing of Care” (New York City: Data & Society, 2021), available at https://datasociety.net/wp-content/uploads/2021/11/EVV_REPORT_11162021.pdf.
- 166 National Council on Independent Living, “Call to Action: The Need for Federal Protections in Electronic Visit Verification” (2020), available at <https://ncil.org/wp-content/uploads/2020/08/8-2-20-CALL-TO-ACTION.pdf>.
- 167 The Annie E. Casey Foundation, “Children Below 200% Poverty in the United States” (2024), available at <https://datacenter.aecf.org/data/tables/47-children-below-200-poverty>.

- 168 U.S. Department of Education, “Digest of Education Statistics: Table 204.10 Number and percentage of public school students eligible for free or reduced-price lunch, by state: Selected school years, 2000-01 through 2022-23” (National Center for Education Statistics, 2024), available at https://nces.ed.gov/programs/digest/d23/tables/dt23_204.10.asp; Federal Register, “Child Nutrition Programs: Income Eligibility Guidelines” (2024), available at <https://www.federalregister.gov/documents/2024/02/20/2024-03355/child-nutrition-programs-income-eligibility-guidelines#p-15>.
- 169 The Annie E. Casey Foundation, “Children in Title I Schools by Race and Ethnicity in the United States” (2021), available at <https://datacenter.aecf.org/data/tables/8418-children-in-title-i-schools-by-race-and-ethnicity#detailed/1/any/false/1696,1648,1603,1381/167,168,133,3,185,107/17042>.
- 170 Centers of Disease Control and Prevention, “QuickStats: Percentage of Children and Adolescents Aged 5–17 Years Who Received Free or Reduced-Cost Meals at School During the Previous 12 Months, by Race and Hispanic Ethnicity and Family Income—National Health Interview Survey, United States, 2021” (2023), available at <https://www.cdc.gov/mmwr/volumes/72/wr/mm7220a6.htm>.
- 171 Emma García and Elaine Weiss, “Education inequalities at the school starting gate” (Washington, DC: Economic Policy Institute, 2017), available at <https://www.epi.org/publication/education-inequalities-at-the-school-starting-gate/>; Kate Bradley, “The Socioeconomic Achievement Gap in the US Public Schools” (Provo, UT: Ballard Center for Social Impact, 2022), available at <https://ballardbrief.byu.edu/issue-briefs/the-socioeconomic-achievement-gap-in-the-us-public-schools>.
- 172 Carmel Martin and others, “A Quality Approach to School Funding” (Washington, DC: Center for American Progress, 2018), available at <https://www.americanprogress.org/article/quality-approach-school-funding/>.
- 173 Equal Justice Society, “Breaking the Chains: The School-to-Prison Pipeline, Implicit Bias, and Racial Trauma” (2016), available at <https://equaljusticesociety.org/breakingthechains/>.
- 174 Paul Hemez, John J. Brent, and Thomas J. Mowen, “Exploring the School-to-Prison Pipeline: How School Suspensions Influence Incarceration During Young Adulthood,” *Youth Violence and Juvenile Justice* 18 (3) (2019), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8277150/>.
- 175 Robert Balfanz and others, “Indicators & Interventions: A Practical Manual for Early Warning Systems” (Baltimore, MD: Johns Hopkins School of Education and Everyone Graduates Center, 2019), available at https://new.every1graduates.org/wp-content/uploads/2020/06/20190918_EGC_EWSManual_Final2.pdf.

- 176 U.S. Department of Education, “Issue Brief: Early Warning Systems” (Office of Planning, Evaluation and Policy Development, 2016), available at <https://www2.ed.gov/rschstat/eval/high-school/early-warning-systems-brief.pdf>.
- 177 Of course, there are several limitations to this assumption, but it is a necessary one to try to quantify AI exposure given the available information.
- 178 Todd Feathers, “False Alarm: How Wisconsin Uses Race and Income to Label Students ‘High Risk,’” The Markup, May 11, 2023, available at <https://themarkup.org/machine-learning/2023/04/27/false-alarm-how-wisconsin-uses-race-and-income-to-label-students-high-risk>.
- 179 Juan C. Perdomo and others, “Difficult Lessons on Social Prediction from Wisconsin Public Schools.” Working Paper (2023), available at <https://arxiv.org/pdf/2304.06205>.
- 180 The study authors underscored that the data was too “noisy” to be conclusive. See Juan C. Perdomo and others, “Difficult Lessons on Social Prediction from Wisconsin Public Schools.” Working Paper (2023), available at <https://arxiv.org/pdf/2304.06205>.
- 181 Todd Feathers, “False Alarm: How Wisconsin Uses Race and Income to Label Students ‘High Risk,’” The Markup, May 11, 2023, available at <https://themarkup.org/machine-learning/2023/04/27/false-alarm-how-wisconsin-uses-race-and-income-to-label-students-high-risk>.
- 182 Clarence Okoh, Written Statement before the U.S. Commission on Civil Rights Pennsylvania Advisory Committee, Public Briefings on Civil Rights and the Rising Use of Artificial Intelligence in Education, March 25, 2024, available at https://www.clasp.org/wp-content/uploads/2024/04/20240405_StatementForRecordAISurvey.pdf.
- 183 Clarence Okoh, Written Statement before the U.S. Commission on Civil Rights Pennsylvania Advisory Committee, Public Briefings on Civil Rights and the Rising Use of Artificial Intelligence in Education, March 25, 2024, available at https://www.clasp.org/wp-content/uploads/2024/04/20240405_StatementForRecordAISurvey.pdf; Kathleen McGrory and Neil Bedi, “Pasco’s sheriff created a futuristic program to stop crime before it happens. It monitors and harasses families across the county.” Tampa Bay Times, September 3, 2020, available at <https://projects.tampabay.com/projects/2020/investigations/police-pasco-sheriff-targeted/intelligence-led-policing/>; Neil Bedi and Kathleen McGrory, “Pasco’s sheriff uses grades and abuse histories to label schoolchildren potential criminals. The kids and their parents don’t know.” Tampa Bay Times, November 19, 2020, available at <https://projects.tampabay.com/projects/2020/investigations/police-pasco-sheriff-targeted/school-data/>

- 184 Pasco Sheriff's Office, "Intelligence-led Policing Manual" (2018), available at https://s3.documentcloud.org/documents/20412738/ilp_manual012918.pdf.
- 185 Kathleen McGrory and Neil Bedi, "Pasco's sheriff created a futuristic program to stop crime before it happens. It monitors and harasses families across the county." Tampa Bay Times, September 3, 2020, available at <https://projects.tampabay.com/projects/2020/investigations/police-pasco-sheriff-targeted/intelligence-led-policing/>.
- 186 Elizabeth Laird, Madeliene Dwyer, and High Grant-Chapman, "Off Task: EdTech Threats to Student Privacy and Equity in the Age of AI" (Washington, DC: Center for Democracy and Technology, 2023), available at <https://cdt.org/wp-content/uploads/2023/09/091923-CDT-Off-Task-web.pdf>.
- 187 Clarence Okoh, Written Statement before the U.S. Commission on Civil Rights Pennsylvania Advisory Committee, Public Briefings on Civil Rights and the Rising Use of Artificial Intelligence in Education, March 25, 2024, available at https://www.clasp.org/wp-content/uploads/2024/04/20240405_StatementForRecordAISurvey.pdf
- 188 Catwhipple, "Coalition forms to oppose Saint Paul's joint powers agreement," The Circle, November 1, 2018, available at <https://thecirclenews.org/news/coalition-forms-to-oppose-saint-pauls-joint-powers-agreement/>; Frederick Melo, "St. Paul, Ramsey County to end youth data-sharing agreement after withering criticism," Twin Cities Pioneer Press, January 28, 2019, available at <https://www.twincities.com/2019/01/28/st-paul-ramsey-county-to-end-youth-data-sharing-agreement-after-withering-criticism>; InEquality and TCIA, "Improving Outcomes for Kids & Families: Beyond Predictive Analytics & Data Sharing" (n.d.), available at <https://static1.squarespace.com/static/5b9081c58ab7224793278e1d/t/60b0642622a43826090a6ef4/1622172711142/JPA+policy+brief+v2.pdf>.
- 189 Letter to U.S. Secretary of Education Miguel Cardona, U.S. Assistant Secretary for Civil Rights Catherine E. Lhamon, and U.S. Deputy Assistant Secretary for Policy Monique Dixon, "Re: The Department of Education Must Prohibit Public Schools from Using Artificial Intelligence & Police Surveillance Technologies to Abuse the Civil and Human Rights of Marginalized Youth," March 18, 2024, available at <https://www.clasp.org/wp-content/uploads/2024/03/NOTICE-Coalition-letter-to-DOE.pdf>; Elizabeth Laird, Madeliene Dwyer, and High Grant-Chapman, "Off Task: EdTech Threats to Student Privacy and Equity in the Age of AI" (Washington, DC: Center for Democracy and Technology, 2023), available at <https://cdt.org/wp-content/uploads/2023/09/091923-CDT-Off-Task-web.pdf>.

- 190 Nick Morrison, “Zero Tolerance Policies in School ‘Promote further Misbehavior,’ Study Finds,” *Forbes*, May 15, 2024, available at <https://www.forbes.com/sites/nickmorrison/2024/05/15/zero-tolerance-policies-in-school-promote-further-misbehavior-study-finds/>.
- 191 Megan French-Marcelin and Sarah Hinger, “Bullies in Blue: The Origins and Consequences of School Policing” (New York City: American Civil Liberties Union, 2017), available at https://assets.aclu.org/live/uploads/publications/aclu_bullies_in_blue_4_11_17_final.pdf.
- 192 Clarence Okoh, Written Statement before the U.S. Commission on Civil Rights Pennsylvania Advisory Committee, Public Briefings on Civil Rights and the Rising Use of Artificial Intelligence in Education, March 25, 2024, available at https://www.clasp.org/wp-content/uploads/2024/04/20240405_StatementForRecordAISurvey.pdf.
- 193 Limited English Proficiency, “Home page” (n.d.), available at <https://www.lep.gov> (last accessed October 28, 2024).
- 194 Sweta Haldar, Drishti Pillai, and Samantha Artiga, “Overview of Health Coverage and Care for Individuals with Limited English Proficiency (LEP)” (San Francisco: Kaiser Family Foundation, 2023), available at <https://www.kff.org/racial-equity-and-health-policy/issue-brief/overview-of-health-coverage-and-care-for-individuals-with-limited-english-proficiency/>.
- 195 Yeganeh Torbati, “Google Says Google Translate Can’t Replace Human Translators. Immigration Officials Have Used It to Vet Refugees.” *ProPublica*, September 26, 2019, available at <https://www.propublica.org/article/google-says-google-translate-cant-replace-human-translators-immigration-officials-have-used-it-to-vet-refugees>.
- 196 Legal Aid Foundation of Los Angeles, “Advocates Reach Agreement with CA Employment Development Department to Expand Language Access,” Press release, February 28, 2022, available at <https://lafla.org/stories-events/language-rights-advocates-reach-agreement-with-ca-edd/>; U.S. Department of Justice, “U.S. Departments of Justice and Labor and Washington State Department of Labor and Industries Reach Agreement to Improve Access for Limited English Proficient Workers,” Office of Public Affairs, Press release, October 1, 2015, available at <https://www.justice.gov/opa/pr/us-departments-justice-and-labor-and-washington-state-department-labor-and-industries-reach>.
- 197 Nicole C. Tensmeyer and others, “Analysis of Language Translations of State Governments’ Coronavirus Disease 2019 Vaccine Websites,” *Health Equity* 6 (1) (2022), available at <https://pmc.ncbi.nlm.nih.gov/articles/PMC9536349/>.

- 198 Ruth W. Leemis and others, “The National Intimate Partner and Sexual Violence Survey: 2016-2017 Report on Intimate Partner Violence (Atlanta, GA: Centers for Disease Control and Prevention, 2022), available at https://www.cdc.gov/nisvs/documentation/NISVSReportonIPV_2022.pdf.
- 199 Amy Bonomi and others, “Intimate Partner Violence and Neighborhood Income: A Longitudinal Analysis,” *Violence Against Women* 20 (1) (2014), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5486977/pdf/nihms866337.pdf>.
- 200 Matthew J. Breiding and others, “Economic Insecurity and Intimate Partner and Sexual Violence Victimization,” *American Journal of Preventive Medicine* 53 (4) (2017), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6426442/pdf/nihms-1016342.pdf>.
- 201 Matthew J. Breiding and others, “Economic Insecurity and Intimate Partner and Sexual Violence Victimization,” *American Journal of Preventive Medicine* 53 (4) (2017), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6426442/pdf/nihms-1016342.pdf>.
- 202 A full treatment of the approach is beyond the scope of this report. For an example of an introductory discussion, see Erika Sasson, “Can Restorative Practices Address Intimate Partner Violence? Summary of a Roundtable Discussion” (New York City: Center for Court Innovation, 2016), available at https://www.innovatingjustice.org/sites/default/files/documents/Intimate_Partner_Restorative_Roundtable.pdf; Leigh Goodmark, “Reimagining VAWA: Why Criminalization Is a Failed Policy and What a Non-Carceral VAWA Could Look Like,” *Violence Against Women* 00 (0) (2020), available at <https://nyscasa.org/wp-content/uploads/2020/09/VAW-Crim-as-failed-policy-1.pdf>.
- 203 Emily Turner, Gavin Brown, and Juanjo Medina-Ariza, “Predicting Domestic Abuse (Fairly) and Police Risk Assessment,” *Psychosocial Intervention* 31 (3) (2022), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10268549/pdf/1132-0559-pi-31-3-0145.pdf>.
- 204 Jens Ludwig, “Using Machine Learning to Identify High Risk Domestic Violence Offenders in NYC, Final Summary Overview” (Washington, DC: Office of Justice Programs, 2022), available at <https://ojp.gov/pdffiles1/nij/grants/305443.pdf>.
- 205 Dillon Reisman, “New Jersey Relies on an Algorithm to Determine Benefits for Survivors of Domestic Violence – But Should It Be Used at All?” (Newark, NJ: ACLU-New Jersey, 2022), available at <https://www.aclu-nj.org/en/news/new-jersey-relies-algorithm-determine-benefits-survivors-domestic-violence-should-it-be-used>.

- 206 Christopher D. Maxwell and others, “New Approaches to Policing High-Risk Intimate Partner Victims and Those Committing the Crimes” (Washington, DC: National Institute of Justice, 2020), available at <https://nij.ojp.gov/topics/articles/new-approaches-policing-high-risk-intimate-partner-victims-and-those-committing>.
- 207 DV RISC, “Intimate Partner Violence Risk Assessments & Models,” available at <https://dvrisc.org/domestic-violence-homicide-prevention-tools-strategies-assessments/> (last accessed October 29, 2024); Domestic Shelters, “Domestic Violence Assessment Tools,” available at <https://www.domesticshelters.org/resources/risk-assessment-tools> (last accessed October 29, 2024); Esperanza Garcia-Vergara, “Risk Assessment Instruments for Intimate Partner Femicide: A Systematic Review,” *Frontiers in Psychology* 13 (2022), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC9195003/>.
- 208 Maryland Network Against Domestic Violence, “Lethality Assessment Program (LAP) Overview,” available at <https://www.mnadv.org/lethality-assessment-program/lap-program-overview-2/> (last accessed October 29, 2024); Geiger Institute, “The Domestic Violence High Risk Team Model” (n.d.), available at <https://geigerinstitute.org/wp-content/uploads/2024/05/DVHRT-Marketing-Document.pdf>; Nada Hassanein, “Police officers can help prevent domestic violence killings by asking the right questions,” *West Virginia Watch*, December 8, 2023, available at <https://westvirginiawatch.com/2023/12/08/police-officers-can-help-prevent-domestic-violence-killings-by-asking-the-right-questions/>.
- 209 Matthew Bland, “Algorithms Can Predict Domestic Abuse, But Should We Let Them?” in Hamid Jahankhani and others, eds., *Policing in the Era of AI and Smart Societies* (New York City: Springer, 2020), available at https://www.researchgate.net/profile/Matthew-Bland/publication/343026304_Algorithms_Can_Predict_Domestic_Abuse_But_Should_We_Let_Them/links/60f5893616f9f31300956afe/Algorithms-Can-Predict-Domestic-Abuse-But-Should-We-Let-Them.pdf.
- 210 Adam Satariano and Roser Toll Pifarré, “An Algorithm Told Police She Was Safe. Then Her Husband Killed Her.” *New York Times*, July 18, 2024, available at <https://www.nytimes.com/interactive/2024/07/18/technology/spain-domestic-violence-viogen-algorithm.html>; Michele Catanzaro, “In Spain, the VioGén algorithm attempts to forecast gender violence” (Berlin: Algorithm Watch, 2020), available at <https://automatingsociety.algorithmwatch.org/report2020/spain/spain-story/>.
- 211 Eticas, “The Adversarial Audit of the VioGén System” (Association Eticas Research and Innovation, 2022), available at https://eticasfoundation.org/wp-content/uploads/2024/09/Eticas_Audit_of_VioGen.pdf.

- 212 Colin Lecher, “Police Art Looking to Algorithms to Predict Domestic Violence,” The Markup, June 29, 2022, available at <https://themarkup.org/the-breakdown/2022/06/29/police-are-looking-to-algorithms-to-predict-domestic-violence>.
- 213 Adam Dodge and Erica Johnstone, “Using Fake Video Technology To Perpetrate Intimate Partner Abuse” (2018), available at <https://withoutmyconsent.org/perch/resources/2018-04-25deepfakedomesticviolenceadvisory.pdf>; Danielle Keats Citron, *The Fight for Privacy: Protecting Dignity, Identity, and Love in the Digital Age* (New York: W.W. Norton & Company, 2022).
- 214 Catherine Bernaciak and Dominic A. Ross, “How Easy Is It to Make and Detect a Deepfake?” (Pittsburgh, PA: Carnegie Mellon University Software Engineering Institute, 2022), available at <https://insights.sei.cmu.edu/blog/how-easy-is-it-to-make-and-detect-a-deepfake/>.
- 215 Oceane Duboust, “Deepfakes are spreading in scams and on social media ‘faster than expected’, experts warn,” EuroNews, June 13, 2024, available at <https://www.euronews.com/next/2024/06/13/deepfakes-are-spreading-in-scams-and-on-social-media-faster-than-expected-experts-warn>.
- 216 Bree Fowler, “Election Deepfakes Are Here and Better Than Ever,” CNET, August 28, 2024, available at <https://www.cnet.com/news/politics/election-deepfakes-are-getting-better-and-easier-to-make/>.
- 217 Rebecca A. Delfino, “Deepfakes on Trial: A Call To Expand the Trial Judge’s Deepfakes on Trial: A Call To Expand the Trial Judge’s Gatekeeping Role To Protect Legal Proceedings from Technological Fakery,” *Hastings Law Journal* 74 (2) (2023), available at https://repository.uclawsf.edu/cgi/viewcontent.cgi?article=4012&context=hastings_law_journal.
- 218 Victoria Copeland, “Dismantling the Carceral Ecosystem: Investigating the Role of “Child Protection” and Family Policing in Los Angeles.” Dissertation paper (University of California, Los Angeles, 2022), available at <https://escholarship.org/uc/item/3rc7z257>.
- 219 Congressional Research Service, “Child Welfare: Purposes, Federal Programs, and Funding” (2024), available at <https://sgp.fas.org/crs/misc/IF10590.pdf>.
- 220 Casey Family Programs, “How does turnover in the child welfare workforce impact children and families?” (2023), available at <https://www.casey.org/turnover-costs-and-retention-strategies/>.

- 221 U.S. Department of Health and Human Services, “Child Maltreatment 2022” (Children’s Bureau, Administration on Children, Youth, and Families and Administration for Children and Families, 2024), available at <https://www.acf.hhs.gov/sites/default/files/documents/cb/cm2022.pdf>; Children’s Rights, “Fact SheetL The Child Welfare System” (2023), available at <https://www.childrensrights.org/wp-content/uploads/2023/01/CR-The-Child-Welfare-System-2023-Fact-Sheet.pdf>.
- 222 Legislative Analyst’s Office, “California’s Child Welfare System: Addressing Disproportionalities and Disparities” (2024), available at <https://lao.ca.gov/reports/2024/4897/disproportionalities-disparities-child-welfare-042424.pdf>.
- 223 Annie E. Casey Foundation, “Child Welfare and Foster Care Statistics” (2024), available at <https://www.aecf.org/blog/child-welfare-and-foster-care-statistics>.
- 224 U.S. Department of Health and Human Services, “The AFCARS Report” (Children’s Bureau, Administration on Children, Youth, and Families and Administration for Children and Families, 2023), available at <https://www.acf.hhs.gov/sites/default/files/documents/cb/afcars-report-30.pdf>.
- 225 Annie E. Casey Foundation, “Child Welfare and Foster Care Statistics” (2024), available at <https://www.aecf.org/blog/child-welfare-and-foster-care-statistics>.
- 226 Human Rights Watch and American Civil Liberties Union, “ ‘If I Wasn’t Poor, I Wouldn’t Be Unfit’: The Family Separation Crisis in the US Child Welfare System” (2022), available at https://assets.aclu.org/live/uploads/publications/us_child_welfare_report_11.17.22.pdf; Shanta Trivedi, “The Harm of Child Removal,” NYU Review of Law & Social Change 43 (2019), available at https://scholarworks.law.ubalt.edu/cgi/viewcontent.cgi?article=2087&context=all_fac.
- 227 Olivia Hampton, “ ‘Deluged’ child welfare systems struggle to protect kids amid calls for reform,” NPR, November 30, 2023, available at <https://www.npr.org/2023/11/30/1211781955/deluged-child-welfare-systems-struggle-to-protect-kids-amid-calls-for-reform>.
- 228 U.S. Department of Health and Human Services, “Child Welfare Practice to Address Racial Disproportionality and Disparity” (Children’s Bureau, 2020), available at https://cwig-prod-prod-drupal-s3fs-us-east-1.s3.amazonaws.com/public/documents/racial_disproportionality.pdf. As the report notes, racial biases may play a role, too, but “an analysis of national-level child welfare data found that racial bias was not a large-scale driver of disproportionality at the referral phase” of Child Protective Services proceedings.
- 229 Shanta Trivedi, “The Harm of Child Removal,” NYU Review of Law & Social Change 43 (2019), available at https://scholarworks.law.ubalt.edu/cgi/viewcontent.cgi?article=2087&context=all_fac.

- 230 Human Rights Watch and American Civil Liberties Union, “ ‘If I Wasn’t Poor, I Wouldn’t Be Unfit’: The Family Separation Crisis in the US Child Welfare System” (2022), available at https://assets.aclu.org/live/uploads/publications/us_child_welfare_report_11.17.22.pdf.
- 231 Sally Ho and Garance Burke, “An algorithm that screens for child neglect raises concerns,” AP News, April 29, 2022, available at <https://apnews.com/article/child-welfare-algorithm-investigation-9497ee937e0053ad4144a86c68241ef1>; Dan Hurley, “Can an Algorithm Tell When Kids Are in Danger?” New York Times Magazine, January 2, 2018, available at <https://www.nytimes.com/2018/01/02/magazine/can-an-algorithm-tell-when-kids-are-in-danger.html>.
- 232 Christopher Teixeira and Matthew Boyas, “Predictive Analytics in Child Welfare: An Assessment of Current Efforts, Challenges and Opportunities” (Washington, DC: U.S. Department of Health and Human Services, Office of the Assistant Secretary for Planning and Evaluation, 2017), available at <https://aspe.hhs.gov/sites/default/files/private/pdf/257841/PACWAnAssessmentCurrentEffortsChallengesOpportunities.pdf>.
- 233 Anjana Samant and others, “Family Surveillance by Algorithm” (New York City: American Civil Liberties Union, 2021), available at https://assets.aclu.org/live/uploads/document/2021.09.28_Family_Surveillance_by_Algorithm.pdf.
- 234 Rhema Vaithianathan and others, “Developing Predictive Models to Support Child Maltreatment Hotline Screening Decisions: Allegheny County Methodology and Implementation” (2017), available at https://www.alleghenycountyanalytics.us/wp-content/uploads/2019/05/Methodology-V1-from-16-ACDHS-26_PredictiveRisk_Package_050119_FINAL.pdf.
- 235 Rhema Vaithianathan and others, “Allegheny Family Screening Tool: Methodology, Version 2” (2019), available at https://www.alleghenycountyanalytics.us/wp-content/uploads/2019/05/Methodology-V2-from-16-ACDHS-26_PredictiveRisk_Package_050119_FINAL-7.pdf.
- 236 Sally Ho and Garance Burke, “An algorithm that screens for child neglect raises concerns,” AP News, April 29, 2022, available at <https://apnews.com/article/child-welfare-algorithm-investigation-9497ee937e0053ad4144a86c68241ef1>.
- 237 Marissa Gerchick and others, “The Devil is in the Details: Interrogating Values Embedded in the Allegheny Family Screening Tool” (New York City: American Civil Liberties Union, 2023), available at <https://www.aclu.org/the-devil-is-in-the-details-interrogating-values-embedded-in-the-allegheny-family-screening-tool>.
- 238 Logan Stapleton and others, “Extended Analysis of ‘How Child Welfare Workers Reduce Racial Disparities in Algorithmic Decisions’” (2022), available at <https://arxiv.org/pdf/2204.13872>. Notably, the human screeners reduced—but did not eliminate—this disparity by overriding the AI system’s recommendation in some cases.
- 239 Ian Moura, “Assessing Bias Related to Disability Status in the Allegheny Family Screening Tool for Child Welfare Decision-Making” (Waltham, MA: Brandeis University, n.d.), available at https://drive.google.com/file/d/1BiAdIpZTeAlwYwBb0xn_ZfcVXPTmkKyv/view?usp=drivesdk.

- 240 Sally Ho and Garance Burke, “AP report: DOJ examining AI screening tool used by Pa. child welfare agency,” PBS News, January 31, 2023, available at <https://www.pbs.org/newshour/nation/ap-report-doj-examining-ai-screening-tool-used-by-pa-child-welfare-agency>.
- 241 Sally Ho and Garance Burke, “Here’s how an AI tool may flag parents with disabilities,” AP News, March 15, 2023, available at <https://apnews.com/article/child-protective-services-algorithms-artificial-intelligence-disability-02469a9ad3ed3e9a31ddae68838bc76e>; Laurel-Ann Dooley, “Does artificial intelligence discriminate in child neglect case assessments?” ABA Journal, December 1, 2023, available at <https://www.abajournal.com/magazine/article/does-artificial-intelligence-discriminate-in-child-neglect-case-assessments>.
- 242 Sally Ho and Garance Burke, “An algorithm that screens for child neglect raises concerns,” AP News, April 29, 2022, available at <https://apnews.com/article/child-welfare-algorithm-investigation-9497ee937e0053ad4144a86c68241ef1>.
- 243 Sally Ho and Garance Burke, “Here’s how an AI tool may flag parents with disabilities,” AP News, March 15, 2023, available at <https://apnews.com/article/child-protective-services-algorithms-artificial-intelligence-disability-02469a9ad3ed3e9a31ddae68838bc76e>.
- 244 Sally Ho and Garance Burke, “Oregon dropping AI tool used in child abuse cases,” AP News, June 2, 2022, available at <https://apnews.com/article/politics-technology-pennsylvania-child-abuse-1ea160dc5c2c203fdab456e3c2d97930>.
- 245 Anjana Samant and others, “Family Surveillance by Algorithm” (New York City: American Civil Liberties Union, 2021), available at https://assets.aclu.org/live/uploads/document/2021.09.28_Family_Surveillance_by_Algorithm.pdf; Garrett Therolf and Matt Hamilton, “News Analysis: Anthony, Noah, Gabriel and beyond: How to fix L.A. County DCFS,” Los Angeles Times, January 20, 2022, available at <https://www.latimes.com/california/story/2022-01-20/child-protective-services-sidebar>.
- 246 Anna Kawakami and others, “Improving Human-AI Partnerships in Child Welfare: Understanding Worker Practices, Challenges, and Desires for Algorithmic Decision Support” (2022), available at <https://dl.acm.org/doi/pdf/10.1145/3491102.3517439>.
- 247 Colorado Department of Human Services, “Number of Referrals/Assessments” (2024), available at <https://colorado.rom.socwel.ku.edu/reports/246>; Colorado Department of Human Services, “Count of Maltreatment Allegations by Maltreatment Type” (2024), available at <https://colorado.rom.socwel.ku.edu/reports/244>.

- 248 Los Angeles County Department of Child and Family Services, “Child Welfare Services Data Fact Sheet” (2023), available at <https://dcfs.lacounty.gov/wp-content/uploads/2024/03/Factsheet-CY-2023.pdf>.
- 249 Pennsylvania Partnerships for Children, “State of Child Welfare 2022” (2022), available at https://www.papartnerships.org/wp-content/uploads/2022/09/2022_State_of_Child_Welfare_-_Allegheny1.pdf.
- 250 Colleen Henry and Vicki Lens, “Marginalizing Mothers: Child Maltreatment Registries, Statutory Schemes, and Reduced Opportunities for Employment Schemes, and Reduced Opportunities for Employment,” City University of New York Law Review 24 (1) (2021), available at <https://academicworks.cuny.edu/cgi/viewcontent.cgi?article=1510&context=clr>.
- 251 Nicholas Kahn and others, “The Standard of Proof in the Substantiation of Child Abuse and Neglect,” Journal of Empirical Legal Studies (forthcoming), available at <https://ssrn.com/abstract=2818533>.
- 252 Annie E. Casey Foundation, “How to Remedy Harm Caused by State Child Abuse Registries” (2023), available at <https://www.aecf.org/blog/how-to-remedy-harm-caused-by-state-child-abuse-registries>.
- 253 Vincent M. Southerland, “The Intersection of Race and Algorithmic Tools in the Criminal Legal System,” Maryland Law Review 80 (2021), available at <https://static1.squarespace.com/static/5e62b76d502f345aab67e0/t/6124eeaa02c3223c2b69961b/1629810348413/1.+Southerland+-+FINAL.pdf>; Electronic Privacy Information Center, “AI in the Criminal Justice System” (n.d.), available at <https://epic.org/issues/ai/ai-in-the-criminal-justice-system/>; Rashia Richardson and Amba Kak, “Suspect Development Systems: Databasing Marginality and Enforcing Discipline,” University of Michigan Journal of Law Reform 55 (2022), available at <https://repository.law.umich.edu/cgi/viewcontent.cgi?article=2562&context=mjlr>.
- 254 Jolynn Dellinger and Stephanie K. Pell, “The criminalization of abortion and surveillance of women in a post-Dobbs world” (Washington, DC: Brookings Institution, 2024), available at <https://www.brookings.edu/articles/the-criminalization-of-abortion-and-surveillance-of-women-in-a-post-dobbs-world/>; Jeremy Kahn, “After Roe, fears mount about A.I.’s ability to identify those seeking abortions,” Fortune, June 28, 2022, available at <https://fortune.com/2022/06/28/after-roe-v-wade-fear-of-a-i-surveillance-abortion/>.
- 255 Julie Mao and others, “Automating Deportation: The Artificial Intelligence Behind the Department of Homeland Security’s Immigration Enforcement Regime” (Phoenix, AZ: Just Futures Law and Mijente, 2024), available at <https://mijente.net/wp-content/uploads/2024/06/Automating-Deportation.pdf>.

- 256 Paul Kiel, “It’s Getting Worse: The IRS Now Audits Poor Americans at About the Same Rate as the Top 1%,” ProPublica, May 30, 2019, available at <https://www.propublica.org/article/irs-now-audits-poor-americans-at-about-the-same-rate-as-the-top-1-percent>; Janet Holtzblatt and Alex Engler, “Machine Learning and Tax Enforcement” (Washington, DC: Tax Policy Center, 2022), available at <https://www.urban.org/sites/default/files/2022-06/Machine%20Learning%20and%20Tax%20Enforcement.pdf>; Congressional Research Service, “Audits of EITC Returns: By the Numbers” (2022), available at <https://crsreports.congress.gov/product/pdf/IN/IN11952/1>; Lucy Hager, Arianna Fano, and Chintal Shah, “Efforts to Address Equity and Efficiency in EITC Audits” (Washington, DC: Bipartisan Policy Center, 2023), available at <https://bipartisanpolicy.org/blog/efforts-to-address-equity-and-efficiency-in-eitc-audits/>.
- 257 Congressional Research Service, “Child Support Enforcement and Driver’s License Suspension Policies” (2011), available at <https://crsreports.congress.gov/product/pdf/R/R41762/5>.
- 258 Aaron Klein, “Reducing bias in AI-based financial services” (Washington, DC: Brookings Institution, 2020), available at <https://www.brookings.edu/articles/reducing-bias-in-ai-based-financial-services/>; Yeshimabeit Milner, “Abolish FICO,” Remarks at the White House Roundtable on Data Brokers, August 16, 2023, available at <https://d4bl.org/dispatch/101-abolish-fico>.
- 259 Garance Burke and Hilke Schellmann, “Researchers Say an AI-Powered Transcription Tool Used in Hospitals Invents Things No One Ever Said,” Associated Press, October 26, 2024, available at <https://apnews.com/article/ai-artificial-intelligence-health-business-90020cdf5fa16c79ca2e5b6c4c9b14>; Ziad Obermeyer and others, “Dissecting Racial Bias in Algorithms Used to Manage the Health of Populations,” Science, October 25, 2019, available at <https://www.science.org/doi/10.1126/science.aax2342>; American Medical Association, “Future of Health: The Emerging Landscape of Augmented Intelligence in Health Care” (n.d.), available at <https://www.ama-assn.org/system/files/future-health-augmented-intelligence-health-care.pdf>.
- 260 Spencer A. Overton, “Overcoming Racial Harms to Democracy from Artificial Intelligence” (2024), available at https://scholarship.law.gwu.edu/cgi/viewcontent.cgi?article=2986&context=faculty_publications.
- 261 Noam Hassenfeld, “Even the scientists who build AI can’t tell you how it works,” Vox, July 15, 2023, available at <https://www.vox.com/unexplainable/2023/7/15/23793840/chat-gpt-ai-science-mystery-unexplainable-podcast>.

- 262 Ryan Calo and Danielle Keats Citron, “The Automated Administrative State: A Crisis of Legitimacy,” *Emory Law Journal* 70 (4) (2021), available at <https://scholarlycommons.law.emory.edu/cgi/viewcontent.cgi?article=1418&context=elj>.
- 263 John Laidler, “High Tech is Watching You,” *The Harvard Gazette*, March 4, 2019, available at <https://news.harvard.edu/gazette/story/2019/03/harvard-professor-says-surveillance-capitalism-is-undermining-democracy/>.
- 264 Danielle Keats Citron, *The Fight for Privacy: Protecting Dignity, Identity, and Love in the Digital Age* (New York: W.W. Norton & Company, 2022).
- 265 Amba Kak, Written Statement before the U.S. Senate Committee on Commerce, Science, and Transportation, Hearing on The Need to Protect Americans’ Privacy and the AI Accelerant, July 11, 2024, available at <https://www.commerce.senate.gov/services/files/35C1CE13-F03D-4CE9-9559-B0651F156D86>.
- 266 Ryan Calo, Written Statement before the U.S. Senate Committee on Commerce, Science, and Transportation, Hearing on The Need to Protect Americans’ Privacy and the AI Accelerant, July 11, 2024, available at <https://www.commerce.senate.gov/services/files/379B1E45-CACB-4A94-A3B6-42368B2FA0EA>.
- 267 Kalee Burns, Liana Fox, and Danielle Wilson, “Expansions to Child Tax Credit Contributed to 46% Decline in Child Poverty Since 2020” (Washington, DC: U.S. Census Bureau, 2022), available at <https://www.census.gov/library/stories/2022/09/record-drop-in-child-poverty.html>; Christopher Wimer and others, “The 2021 Child Tax Credit Expansion: Child Poverty Reduction and the Children Formerly Left Behind,” *Poverty and Social Policy Brief* 6 (8) (2022), available at <https://static1.squarespace.com/static/610831a16c95260dbd68934a/t/63629828229a175d3864c0a1/1667405865381/Expanded-CTC-and-Child-Poverty-in-2021-CPSP.pdf>.
- 268 Grant Fergusson, “Outsourced and Automated: How AI Companies Have Taken Over Government Decision-Making” (Washington, DC: Electronic Privacy Information Center, 2023), available at <https://epic.org/wp-content/uploads/2023/09/FINAL-EPIC-Outsourced-Automated-Report-w-Appendix-Updated-9.26.23.pdf>.
- 269 Rachana Pradhan and Samantha Liss, “Medicaid for Millions in America Hinges on Deloitte-Run Systems Plagued by Errors,” *KFF Health News*, June 24, 2024, available at <https://kffhealthnews.org/news/article/medicaid-deloitte-run-eligibility-systems-plagued-by-errors/>; Samantha Liss and Rachana Pradhan “Errors in Deloitte-Run Medicaid Systems Can Cost Millions and Take Years To Fix,” *KFF Health News*, September 5, 2024, available at <https://kffhealthnews.org/news/article/deloitte-run-medicaid-systems-errors-cost-millions-take-years-to-fix>.

- 270 *TransUnion LLC v. Ramirez*, 141 S. Ct. 2190, 2215 (2021).
- 271 Reuters, “Stone Point Capital, Insight Partners to buy CoreLogic for \$6 billion,” February 4, 2021, available at <https://www.reuters.com/article/world/americas/stone-point-capital-insight-partners-to-buy-corelogic-for-6-billion-idUSKBN2A42H3>; Commercial Investigations LLC, “The Impact of Consolidation on Your Background Screening Services” (2024), available at <https://www.commercialinvestigationsllc.com/blog/the-impact-of-consolidation-on-your-background-screening-services>; Casey Ross and Bob Herman, “Denied by AI: How Medicare Advantage plans use algorithms to cut off care for seniors in need,” STAT, March 13, 2023, available at <https://www.statnews.com/2023/03/13/medicare-advantage-plans-denial-artificial-intelligence/>.
- 272 Grant Fergusson, “Outsourced and Automated: How AI Companies Have Taken Over Government Decision-Making” (Washington, DC: Electronic Privacy Information Center, 2023), available at <https://epic.org/wp-content/uploads/2023/09/FINAL-EPIC-Outsourced-Automated-Report-w-Appendix-Updated-9.26.23.pdf>.
- 273 Knight Foundation and Rita Allen Foundation, “How can we harness technology to promote civic engagement and more responsive government?” (n.d.), available at <https://knightfoundation.org/features/civicttechbiz/>.
- 274 Jai Vipra and Sarah Myers West, “Computational Power and AI” (New York City: AI Now Institute, 2023), available at <https://ainowinstitute.org/publication/policy/compute-and-ai>.
- 275 Cynthis Conti-Cook and Ed Vogel, “MyCity, INC: A Case Against ‘CompStat Urbanism’” (New York City: Surveillance Resistance Lab, 2024), available at https://surveillanceresistancelab.org/wp-content/uploads/MyCityINC_March2024.pdf.
- 276 The Due Process Clause’s requirement for an explanation for the reasons of any adverse action may put a ceiling on the complexity of AI decision-making. It may, for example, stop states from using AI to make decisions via advanced machine learning techniques where criteria continually change and cannot be accurately explained.
- 277 *Mathews v. Eldridge*, 424 U.S. 319, 335 (1976).
- 278 Benefits Tech Advocacy Hub, “Case Study Library” (2024), available at <https://www.btah.org/case-studies.html>.
- 279 *Elder v. Gillespie*, 54 F.4th 1055 (8th Cir. 2022).

- 280 U.S. Department of Justice, “Justice Department Secures Agreement with Florida School District to Protect Civil Rights of Students with Disabilities,” Office of Public Affairs, Press release, March 5, 2024, available at <https://www.justice.gov/opa/pr/justice-department-secures-agreement-florida-school-district-protect-civil-rights-students>.
- 281 Greater Boston Legal Services, “Preliminary Approval of Settlement Granted in Ground-Breaking Discrimination Class Action,” Press release, April 25, 2024, available at <https://www.gbls.org/news/preliminary-approval-settlement-granted-ground-breaking-discrimination-class-action-brought>.
- 282 Federal Trade Commission, “FTC and CFPB Settlement to Require Trans Union to Pay \$15 Million over Charges It Failed to Ensure Accuracy of Tenant Screening Reports,” Press release, October 12, 2023, available at <https://www.ftc.gov/news-events/news/press-releases/2023/10/ftc-cfpb-settlement-require-trans-union-pay-15-million-over-charges-it-failed-ensure-accuracy-tenant>.
- 283 Legal Services Corporation, “By the Numbers 2022: The Data Underlying Legal Aid Programs” (2023), available at <https://lsc-live.app.box.com/s/h2bajpr3gps4s4a1iio6fwiddhmu1nwb>.
- 284 ABA Profile of the Legal Profession, “Civil Legal Aid” (n.d.), available at <https://www.abalegalprofile.com/legalaid.html>.
- 285 Legal Services Corporation, “The Justice Gap: The Unmet Civil Legal Needs of Low-income Americans” (2022), available at <https://lsc-live.app.box.com/s/xl2v2uraitbbzrhwtjlgioemp3myz1>.
- 286 Amazon Web Services, “Model Explainability with AWS Artificial Intelligence and Machine Learning Solutions: Interpretability versus explainability” (2021), available at <https://docs.aws.amazon.com/whitepapers/latest/model-explainability-aws-ai-ml/interpretability-versus-explainability.html>.
- 287 Stop LAPD Spying Coalition and Free Radicals, “The Algorithmic Ecology: An Abolitionist Tool for Organizing Against Algorithms,” Medium.com, March 2, 2020, available at <https://stoplapdspying.medium.com/the-algorithmic-ecology-an-abolitionist-tool-for-organizing-against-algorithms-14fcbd0e64d0>.
- 288 Rideshare Drivers United, “Home page,” available at <https://www.drivers-united.org> (last accessed October 29, 2024); Kerry Harwin, “Lyft Drivers Win Nearly \$200,000 in back Pay!” Drivers Union, July 14, 2023, available at https://www.driversunionwa.org/lyft_drivers_win; Alfonso Galvan, “In major victory for rideshare drivers, Minneapolis Council overrides mayor’s veto on minimum pay,” Sahan Journal, March 14, 2024, available at <https://sahanjournal.com/business-work/mayor-jacob-frey-veto-override-uber-lyft-leaving-minneapolis/>.

- 289 The White House, Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence, October 30, 2023, available at <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>; Shalanda D. Young, “Memorandum for the Heads of Executive Departments and Agencies: Advancing Governance, Innovation, and Risk Management for Agency Use of Artificial Intelligence” (Washington, DC: Executive Office of the President, Office of Management and Budget, 2024), available at <https://www.whitehouse.gov/wp-content/uploads/2024/03/M-24-10-Advancing-Governance-Innovation-and-Risk-Management-for-Agency-Use-of-Artificial-Intelligence.pdf>.
- 290 California AI Transparency Act, SB-942, California State Legislature (September 20, 2024), available at https://leginfo.ca.gov/faces/billTextClient.xhtml?bill_id=202320240SB942; Generative artificial intelligence: training data transparency, AB-2013, California State Legislature (September 30, 2024), available at https://leginfo.ca.gov/faces/billTextClient.xhtml?bill_id=202320240AB2013.
- 291 Artificial Intelligence Policy Act, SB 149, Utah State Legislature (March 13, 2024), available at <https://le.utah.gov/%7E2024/bills/static/SB0149.html>.
- 292 Consumer Protections for Artificial Intelligence, SB24-205, Colorado General Assembly, 74th assembly 2nd sess. (May 13, 2024), available at <https://leg.colorado.gov/bills/sb24-205>.
- 293 A Local Law to amend the administrative code of the city of New York, in relation to automated employment decision tools, 2021/144, New York City Council (December 11, 2021), available at <https://legistar.council.nyc.gov/LegislationDetail.aspx?ID=4344524&GUID=B051915D-A9AC-451E-81F8-6596032FA3F9&Options=ID%7CText%7C&Search=>.
- 294 Richard Eppink, “Special Lectures 2019: Innovation, Technology and the Practice of Law” (Toronto, ON: Law Society of Ontario, 2019), available at <https://drive.google.com/file/d/1xdJiShMG7FRvYmBJlr79jPl4-jVyIcB/view>.
- 295 TechTonic Justice and colleagues from the National Health Law Program and Upturn have jointly arrived at these provisions and have additional details available.
- 296 AI Now Institute, “Data Minimization as a Tool for AI Accountability” (2023), available at <https://ainowinstitute.org/spotlight/data-minimization>.
- 297 Rashida Richardson, “Best Practices for Government Procurement of Data-Driven Technologies” (2021), available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3855637.

- 298 See, for example, Human Rights Watch, “Automated Neglect: How the World Bank’s Push to Allocate Cash Assistance Using Algorithms Threatens Rights” (2023), available at <https://www.hrw.org/report/2023/06/13/automated-neglect/how-world-banks-push-allocate-cash-assistance-using-algorithms>.
- 299 Michael Karanicolas, “Challenging Minority Rule: Developing AI Standards that Serve the Majority World,” *UCLA Law Review* 71 (2) (2024), available at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4450707; IAPP, “Global AI Law and Policy Tracker,” available at <https://iapp.org/resources/article/global-ai-legislation-tracker/> (last accessed October 29, 2024).
- 300 U.S. Department of Health and Human Services, “How Many People Participate in the Social Safety Net?” (Office of Human Services Policy, 2023), available at <https://aspe.hhs.gov/sites/default/files/documents/18eff5e45b2be85fb4c350176bca5c28/how-many-people-social-safety-net.pdf>.
- 301 MACPAC, “Increasing the Rate of Ex Parte Renewals” (2023), available at <https://www.macpac.gov/wp-content/uploads/2023/09/Increasing-the-Rate-of-Ex-Parte-Renewals-Brief.pdf>; Jennifer Wagner, “Streamlining Medicaid Renewals Through the Ex Parte Process” (Washington, DC: Center on Budget and Policy Priorities, 2021), available at <https://www.cbpp.org/research/health/streamlining-medicaid-renewals-through-the-ex-parte-process>.
- 302 Kaiser Family Foundation, “State Health Insurance Marketplace Types, 2025” (2024), available at <https://www.kff.org/affordable-care-act/state-indicator/state-health-insurance-marketplace-types/?currentTimeframe=0&sortModel=%7B%22colId%22:%22Location%22,%22sort%22:%22asc%22%7D>.
- 303 Emily Widra, “States of Incarceration: The Global Context 2024” (Northampton, MA: Prison Policy Initiative, 2024), available at <https://www.prisonpolicy.org/global/2024.html>.
- 304 National Conference of State Legislatures, “Criminal Records and Reentry Toolkit” (2023), available at <https://www.ncsl.org/civil-and-criminal-justice/criminal-records-and-reentry-toolkit>.
- 305 Restoration of Rights Project, “50-State Comparison: Expungement, Sealing & Other Record Relief,” available at <https://ccresourcecenter.org/state-restoration-profiles/50-state-comparisonjudicial-expungement-sealing-and-set-aside-2-2/> (last accessed October 29, 2024).
- 306 Margaret Love and Nick Sibilla, “Advancing Second Chances: Clean Slate and Other Record Reforms in 2023” (Washington, DC: Collateral Consequences Resource Center, 2024), available at https://ccresourcecenter.org/wp-content/uploads/2024/01/Annual-Report-2023.1.5.24.rev2_.pdf.

- 307 Code For America, “Working with communities and government to fundamentally transform the process of clearing records,” available at <https://codeforamerica.org/programs/criminal-justice/automatic-record-clearance/> (last accessed October 29, 2024).
- 308 NAACP, “Artificial Intelligence in Predictive Policing” (n.d.), available at <https://naacp.org/resources/artificial-intelligence-predictive-policing-issue-brief>.
- 309 Importantly, the designers limited and vetted the information available to the chatbots. See Sateesh Nori, “How A Chatbot Can Save Someone From Homelessness,” TEDx Talks video, August 8, 2024, available at <https://www.youtube.com/watch?v=z2TtVvopqzc>; Tenant Power Toolkit, “Home page,” available at <https://tenantpowertoolkit.org> (last accessed October 29, 2024); JustFix, “Home page,” available at <https://www.justfix.org/en/> (last accessed October 29, 2024).
- 310 Debt Collective, “Home page,” available at <https://debtcollective.org> (last accessed October 29, 2024).
- 311 COVID Worst Evictors, “Home page,” available at <https://www.worstevictorsnyc.org> (last accessed October 29, 2024).
- 312 Stop LAPD Spying Coalition, “Watch the Watchers: Home,” available at <https://watchthewatchers.net> (last accessed October 29, 2024); Jason Tashea, “Databases create access to police misconduct cases and offer a handy tool for defense lawyers” ABA Journal, February 1, 2016, available at https://www.abajournal.com/magazine/article/databases_create_access_to_police_misconduct_cases_and_offer_a_handy_tool_f.
- 313 Social Justice Hackathon, “Home page,” available at <https://socialjusticehackathon.com> (last accessed October 29, 2024).
- 314 Social Justice Hackathon, “Project gallery,” available at <https://sjhphl.devpost.com/project-gallery> (last accessed October 29, 2024).
- 315 U.S. Department of Health and Human Services, “Frequently Asked Questions Related to the Poverty Guidelines and Poverty” (Office of the Assistant Secretary for Planning and Evaluation, n.d.), available at <https://aspe.hhs.gov/topics/poverty-economic-mobility/poverty-guidelines/frequently-asked-questions-related-poverty-guidelines-poverty>
- 316 Yiyu Chen, James Fuller, and Renee Ryberg, “Knowing the Strengths and Limitations of Poverty Measures Can Help Us Better Understand Poverty” (Rockville, MD: Child Trends, 2023), available at <https://www.childtrends.org/publications/knowning-the-strengths-and-limitations-of-poverty-measures-can-help-us-better-understand-poverty>.

- 317 Center on Poverty and Social Policy at Columbia University, “Historical Supplemental Poverty Measure Data,” available at <https://www.povertycenter.columbia.edu/historical-spm-data> (last accessed October 29, 2024).
- 318 Data USA, “Map of Poverty Rate by County,” available at <https://datausa.io/map/?measure=ZqHsUx&groups%5B0%5D=1BUDwd%7C0&time=2022> (last accessed October 29, 2024).
- 319 U.S. Census Bureau, “EDA-Census Poverty Status Viewer,” available at <https://mtgis-portal.geo.census.gov/arcgis/apps/experiencebuilder/experience/?id=ad8ad0751e474f938fc98345462cdfbf> (last accessed October 29, 2024).
- 320 U.S. Department of Health and Human Services, “Poverty” (n.d.), available at <https://health.gov/healthypeople/priority-areas/social-determinants-health/literature-summaries/poverty>.
- 321 Madeline B. Harms and Sherona D. Garrett-Ruffin, “Disrupting links between poverty, chronic stress, and educational inequality,” NPJ Science of Learning 8 (2023), available at <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10662171/>.
- 322 Bernadette Rabuy and Daniel Kopf, “Prisons of Poverty: Uncovering the pre-incarceration incomes of the imprisoned” (Northampton, MA: Prison Policy Initiative, 2015), available at <https://www.prisonpolicy.org/reports/income.html>.
- 323 The report uses LGBTQ+ rather than the more universal LGBTQ+ (or other variations) because the poverty measures focus only on these characteristics.
- 324 Kaiser Family Foundation, “Poverty Rate by Race/Ethnicity” (2023), available at <https://www.kff.org/other/state-indicator/poverty-rate-by-raceethnicity/>.
- 325 Board of Governors of the Federal Reserve System, “Greater Wealth, Greater Uncertainty: Changes in Racial Inequality in the Survey of Consumer Finances, Accessible Data” (2023), available at <https://www.federalreserve.gov/econres/notes/feds-notes/greater-wealth-greater-uncertainty-changes-in-racial-inequality-in-the-survey-of-consumer-finances-accessible-20231018.htm>; Briana Sullivan, Donald Hays, and Neil Bennett, “Households With a White, Non-Hispanic Householder Were Ten Times Wealthier Than Those With a Black Householder in 2021” (Washington, DC: U.S. Census Bureau, 2024), available at <https://www.census.gov/library/stories/2024/04/wealth-by-race.html>.
- 326 Robin Bleiweis, Diana Boesch, and Alexandra Cawthorne Gaines, “The Basic Facts About Women in Poverty” (Washington, DC: Center for American Progress, 2020), available at <https://www.americanprogress.org/wp-content/uploads/sites/2/2020/08/Women-In-Poverty-UPDATE.pdf>.

- 327 Nanette Goodman, Michael Morris, and Kelvin Boston, “Financial Inequality: Disability, Race and Poverty in America” (Washington, DC: National Disability Institute, n.d.), available at <https://www.nationaldisabilityinstitute.org/wp-content/uploads/2019/02/disability-race-poverty-in-america.pdf>; Rebecca Vallas and others, “Economic Justice Is Disability Justice” (New York City: The Century Foundation, 2022), available at <https://tcf.org/content/report/economic-justice-disability-justice/>.
- 328 Bianca D.M. Wilson and others, “LGBT Poverty in the United States: Trends at the Onset of COVID-19” (Los Angeles: UCLA School of Law Williams Institute, 2023), available at <https://williamsinstitute.law.ucla.edu/publications/lgbt-poverty-us/>.
- 329 Signe-Mary McKernan, Caroline Ratcliffe, and Stephanie R. Cellini, “Transitioning In and Out of Poverty” (Washington, DC: The Urban Institute, 2009), available at <https://www.urban.org/sites/default/files/publication/30636/411956-Transitioning-In-and-Out-of-Poverty.PDF>.
- 330 Legal Services Corporation, “The Justice Gap: Section 2: Today’s Low-income America” (n.d.), available at <https://justicegap.lsc.gov/resource/section-2-todays-low-income-america/>.
- 331 U.S. Department of Housing and Urban Development, “Transmittal of Fiscal Year (FY) 1998 Public Housing/Section 8 Income Limits” (1998), available at <https://www.huduser.gov/portal/datasets/il/fmr98/sect8.html>.
- 332 Sarah Fass, “Measuring Poverty in the United States” (New York City: National Center for Children in Poverty, 2009), available at <https://www.nccp.org/publication/measuring-poverty-in-the-united-states/>; National Center for Children in Poverty, “Basic Needs Budget Calculator,” available at <http://frs.nccp.org/tools/frs/budget.php> (last accessed October 29, 2024); Economic Policy Institute, “Family Budget Calculator,” available at <https://www.epi.org/resources/budget/> (last accessed October 29, 2024).
- 333 Kaiser Family Foundation, “Distribution of the Total Population by Federal Poverty Level (above and below 200 percent FPL)” (2023), available at <https://www.kff.org/other/state-indicator/population-up-to-200-fpl>.

Acknowledgements

This marks the launch of a new organization and a report, so I owe many thanks.

To my wife Fabi and my son Paulo for all the love, joy, and brilliant ideas and to our families for the encouragement and support along the way.

To my teammates at the Benefits Tech Advocacy Hub for being on this journey with me for several years now: Elizabeth Edwards, Cassandra LaRose, Dave Machledt, Emily Paul, and Emma Weil.

To my board members for their unflagging support in helping to get us off the ground: Marquita Little Numan and Michael Moradzadeh.

To my advisory board members for all the encouragement on this project and over the years: Danielle Citron, Tammy Dobbs, James Dold, Karama Neal, Rashida Richardson, Lee Richardson, and Diana Singleton.

To Chance Williams and Lori McGlinchey at the Ford Foundation for believing in me from the moment we met. To Simone Liano, Nick Cain, and Vilas Dhar at the Patrick J. McGovern Foundation for supporting the vision.

To Chris Wong, Jane Chung, Jen Carnig, Gaby Connor, Michael Murray, and the teams at Objectively, Justice Speaks, and Spitfire for all the work you've done to bring us to launch.

To all the friends, mentors, and fellow troublemakers that I met through my work in the legal aid world (especially my former clients) for giving me the foundation to do this. There are far too many of you to name, but special thanks to my Arkansas guiding lights for supporting me through the toughest times: Lela Davison, Kathy Grady, Trevor Hawkins, Andrea Hope Howard, Harry Light, and Bill Waddell.

To all the people in the tech justice world who have opened doors and shared ideas: among others, Marc Aidinoff, Solon Barocas, Rishi Bharwani, Lydia X.Z. Brown, Ryan Calo, Cynthia Conti-Cook, Ángel Diaz, Alix Dunn, Ritchie Eppink, Virginia Eubanks, Grant Fergusson, Ami Fields-Meyer, Gemma Galdón Clavell, Alex Givens and my CDT collaborators (Ariana Aboulafia, Quinn Anex-Ries, Henry Claypool, Elizabeth Laird, Hannah Quay-de la Vallee, and Ridhi Shetty), Esra Gules-Guctas, Janet Haven and the Data & Society team, Liz Henry, Nani Jansen Reventlow, Amba Kak, Michael Karanicolas, Nik Marda, Varoon Mathur, Ian Moura, Sarah Myers West, Clarence Okoh, Deb Raji, Joanna Redden and the Data Justice Lab team, Dillon Reisman, Danielle Robinson, Hannah Sassaman, Michelle Shevin, Andreen Soley, Vincent Southerland, Amos Toh, and Harlan Yu and the Upturn team (past and present).

To my family in Memphis for the camaraderie, love, and unfailing inspiration: IMAKEMADBEATS, Marie Dukes, NoVa, Kid Maestro, C Major, PreauXX, AWFM, Cameron Bethany, Aaron James, Mr. Dukes, Sarai Payne, and Mave Rucker. There are so many more folks in Memphis to whom I owe thanks that I could never do them all justice, but if you were part of my music, the Word, Artistik Lounge, Mot and Ed's, the Peace and Justice Center, or various other parts of the arts, Hip Hop, or do-gooding scenes, much appreciation.

To all the other people who've supported this endeavor in any way, whether it was by introducing me to people, reading a draft of a concept note, or letting me know that I was on the right track: among others, Nick Allen, Madison Allen, Katherine Alteneder, Cory Anderson, Mehwish Ansari, Marvin Carr, Natalie Castellanos, Rebecca Cokley, Matthew Cortland, Anjali Fernandes, Lupita González, Alana Greer, Bethany Hamilton, Renee Ho, Aras Jizan, Claudia Johnson, Demetrious Jordan, Dom Kelly, Chris Kingsley, Joe Knight, Ellen Lawton, Alex Macgillivray, Sofia Michelakis, Elizabeth Nash, Adrián Palma, Paul Ryan, Kareem Saleh, Jim Sandman, Eric Sears, Madhav Shroff, Rusty Stahl, Paul Tarini, Jason Tashea, Leo Tucker, Melanie Thomasson, Rebecca Watson, Sheldon Weisgrau, and Miguel Willis.

To all the activists, academics, advocates, and journalists whose work is cited here or who otherwise are bringing light to the shadows.

Please blame any omissions or errors on the lack of sleep and not the heart.

Onward!



linkedin.com/company/techtonicjustice



[@techtonicj](https://twitter.com/techtonicj)



[@techtonicjustice.bsky.social](https://techtonicjustice.bsky.social)



[TechTonic Justice](https://facebook.com/TechTonicJustice)



[@techtonicjustice](https://instagram.com/techtonicjustice)



**TECHTONIC
JUSTICE**

techtonicjustice.org