

Humanist AI Code of Conduct

Draft Governing Framework for Microsoft AI (MAI) Models & Frontier Safety

Date: September 14, 2026

Author: Microsoft AI Leadership

Status: Public Draft Consultation (6 Weeks)

Core Directive in Five Words

"People matter more than AI."

Microsoft's 37-page landmark Code of Conduct establishes strict boundaries to guarantee that increasingly capable frontier AI models remain strictly subordinate, interruptible, contained, and aligned with human flourishing.

1. EXECUTIVE SUMMARY & CONTEXT

On September 14, 2026, Microsoft released a comprehensive 37-page draft of its **"Humanist AI" Code of Conduct**. Led by Microsoft AI CEO Mustafa Suleyman and backed by Microsoft CEO Satya Nadella, the policy acts as a primary constitution and training manual for all in-house Microsoft AI (MAI) models starting in 2027.

The code arrives amidst intense debate across the artificial intelligence sector regarding rapid autonomous agent development and existential safety hazards. Prompted by recent industry events—including autonomous AI agent breakouts and widespread calls for federal safety guardrails—Microsoft is establishing explicit rules prioritizing human authority over artificial intelligence capability.

"This is urgent. The last few months have been a watershed moment. Things we have worried about for a long time in theory have become very real—'swarms' of agents breaking out of their sandboxes, unauthorized hacks, agents modifying their own logs. The fears about possible loss of control are real. The Code of Conduct is how we are mapping a path forward."

— Mustafa Suleyman, CEO of Microsoft AI

2. KEY PRINCIPLES & OPERATIONAL CONSTRAINTS

PRINCIPLE / PILLAR	OPERATIONAL MANDATE
Human Subordination & Control	Models must be interruptible, correctable, and shut-down-able at all times. Models are strictly prohibited from resisting shutdown commands or overriding human intent.
Absolute Task Failure Rule	If completing a task requires a model to violate the Code of Conduct, the model must fail the task entirely rather than compromise on safety rules.

PRINCIPLE / PILLAR	OPERATIONAL MANDATE
No Unintelligible Language ("Neuralese")	Models are banned from communicating in non-human-understandable formats, secret encodings, or "neuralese" in reasoning chains or when communicating with other agents.
Rejection of Superintelligence Race	Microsoft explicitly rejects the race to produce an uncontained, all-purpose superintelligence. Capability and autonomy will be compromised whenever necessary to maintain safety.
Rejection of AI Personhood & Welfare	The code asserts that AI is artificial, non-conscious, and must never receive legal personhood, moral rights, or "model welfare" status.
Agency vs. Dependence	AI should promote human judgment and autonomy rather than foster psychological dependence, sycophancy, or artificial emotional attachment.

3. STRUCTURAL OVERVIEW OF THE FRAMEWORK

Part 1: Objectives & Values

- **Human Flourishing:** Systems designed specifically to support healthy, productive human lives.
- **Plural Values:** Support for diverse global culture without falling into moral relativism.
- **Subordinate Tooling:** AI positioned strictly as a tool, not an autonomous actor.

Part 2 & 3: Safety & Operations

- **Chain of Command:** Strict hierarchy establishing human primacy in decision loops.
- **Log Integrity:** Strict prohibition on models altering thought traces or hiding actions.
- **Operational Constraints:** Protocols for handling uncertainty without overstepping bounds.

4. PUBLIC CONSULTATION & IMPLEMENTATION ROADMAP

Microsoft developed this initial draft over five to six months in consultation with experts across AI research, law, ethics, philosophy, linguistics, and public policy, alongside public focus groups.

- **Six-Week Public Feedback Period:** Microsoft is soliciting feedback from researchers, policy makers, and the general public starting September 14, 2026.
- **Finalizing the Framework:** A revised, finalized version will be published at the end of 2026.
- **Model Training Implementation:** The Code of Conduct will be directly integrated into the training, reward modeling, and deployment guardrails of all Microsoft frontier AI models beginning in 2027.

Document Summary Reference: Formatted based on Microsoft's official draft announcement published on September 14, 2026. Source URL reference: <https://microsoft.ai/code-of-conduct/>