



Find community, meet experts, build skills, and discover the latest in AI. Join us at the Microsoft 365 Community Conference April 21-23. [Learn more >](#)

Microsoft 365 Copilot Blog

MICROSOFT 365 COPILOT BLOG 4 MIN READ

Introducing multi-model intelligence in Researcher



gauravanand  MICROSOFT

Mar 30, 2026

Today, Researcher—Microsoft 365 Copilot's deep research agent for work—takes a significant step forward. Designed to tackle complex research in the flow of work, Researcher now goes further with two new multi-model capabilities that raise the bar for




[Register](#) [Sign In](#)

and utilizes a combination of models from Frontier labs including Anthropic and OpenAI. One model leads the generation phase, planning the task, iterating through retrieval, and producing an initial draft, while a second model focuses on review and refinement, acting as an expert reviewer before the final report is produced. Our evaluations show that this architecture exceeds traditional single model approaches and delivers best in class deep research quality. This design provides clear optionality across generator and reviewer roles, with the ability to support and expand these roles over time as the system evolves.

DRACO Benchmark for Deep Research Quality







DRACO (Deep Research Accuracy, Completeness, and Objectivity) benchmark scores across 100 complex research tasks spanning 10 domains. All results are sourced from the original paper [Zhong et al., arXiv:2602.11685 (Feb. 2026)], except for Researcher with Critique. Researcher with Critique achieves a substantial improvement of +7.0 points (SEM ± 1.90) in the aggregated score, +13.88% over Perplexity Deep Research (Claude Opus 4.6 model), the best system reported in the paper.

Council brings multiple model responses side-by-side in the Researcher experience. Additionally, a cover letter provides valuable insights on where the models agree, where they diverge, and the unique insights each brings on the topic.

Critique & how it works

Many AI research workflows rely on a single model to handle planning, sourcing, synthesis, and writing, but Critique takes a different approach by dividing responsibilities between two AI partners—one optimized for deep exploration and structured synthesis, and a second focused on validating claims, improving presentation and strengthening the structure. By giving evaluation as much emphasis as generation, this architecture creates a powerful feedback loop that delivers higher-quality results across factual accuracy, analytical breadth, and presentation. Critique will be the default experience in Researcher, available when Auto is selected in the model picker.





Critique follows a review process similar to those conducted in academic and professional research settings. It's built around rubric-based evaluation—a structured review that focuses on strengthening the report without turning the reviewer into a second author. The reviewer examines the report from several angles and then generates an enhanced report, concentrating on the following dimensions:

- **Source Reliability Assessment.** The reviewer emphasizes the use of reputable, authoritative, and domain-appropriate sources, prioritizing evidence that is verifiable and suitable for your research context.
- **Report Completeness.** The reviewer evaluates if the final report satisfies the intent of your request comprehensively with relevant and unique insights.
- **Strict Evidence Grounding Enforcement.** The reviewer enforces a conservative grounding standard, requiring every key claim to be anchored to reliable sources with precise citations—strengthening factual accuracy, reliability, and trust in the final report.

Performance Validation on the DRACO Benchmark

We evaluated Critique on the DRACO benchmark (Deep Research Accuracy, Completeness, and Objectivity)—100 complex deep research tasks across 10 domains, introduced by researchers from Perplexity and academia in February 2026 [[Zhong et al.arXiv:2602.11685](#)]. These research tasks originate from anonymized real-world usage patterns executed in a large-scale research system. The system responses are graded against task-specific rubrics along four dimensions: factual accuracy, breadth and depth of analysis, presentation quality, and citation quality.

The DRACO results were evaluated using OpenAI's GPT-5.2 as the LLM judge - the strictest of the three judge models reported in the paper. We applied the same evaluation protocol and configuration published in the benchmark paper, helping ensure an apples-to-apples comparison. Across all measurements, results were computed by averaging scores over the full DRACO dataset, with each question evaluated across five independent runs.

To better understand the advantages of Critique, we compared the new architecture with the single-model Researcher (using the same GPT-5.2 judge) across the four evaluation axes that DRACO defines.

Evaluation Dimensions	Researcher with Critique (Multi-Model) vs. Researcher
	Improvement Contribution

	Score Improvement	Improvement Contribution to Final Score
Breadth and Depth of Analysis	+16.1 (+31.7%)	+3.33
Presentation Quality	+16.8 (+46.0%)	+3.04
Factual Accuracy	+4.1 (+7.6%)	+2.58
Citation Quality	+6.2 (+15.3%)	+0.55
Total	-	+9.5

We see the largest improvement in Breadth and Depth of Analysis (+3.33), followed by Presentation Quality (+3.04) and Factual Accuracy (+2.58). All dimensions show statistically significant improvements (paired t-test, $p < 0.0001$).

Critique pushes Researcher to identify missing analytical angles, close coverage gaps, sharpen formulations, and produce responses with stronger organization and clearer narrative flow. This accounts for the substantial improvements to the breadth and depth and presentation quality scores. The factual accuracy gain shows that Critique is challenging weak claims and enforcing higher precision. Citation quality improvements primarily reflect better use of existing sources, as the new layer emphasizes evidence selection and citation rather than increasing source coverage.


The DRACO query set spans 10 domains across medicine, technology, law and more. Researcher with Critique achieves higher scores than the single-model approach across domains, reinforcing its value as a horizontal quality layer for Researcher. At the domain level, statistically significant improvements are observed in 8 of the 10 domains (paired t-test, $p < 0.05$). The exceptions are Academic ($p=0.27$) and Needle-in-a-Haystack ($p=0.16$), both of which exhibit high variance.

Domain	Researcher with Critique (Multi-Model) vs. Researcher
	Score Improvement
Medicine	+21.8 (+54.5%)

General Knowledge	+13.1 (+30.0%)
Technology	+10.3 (+28.7%)
Needle in a Haystack	+12.9 (+24.5%)
Personalized Assistant	+10.5 (+22.1%)
Shopping/Product Comparison	+8.9 (+20.2%)
Finance	+7.3 (+17.7%)
Law	+11.6 (+16.2%)
UX Design	+5.8 (+12.4%)
Academic	+4.0 (+6.1%)

Council & how it works

Council is an alternative approach, designed for side-by-side comparison across multiple models. Available when Model Council is selected in the model picker in Researcher, Council runs an Anthropic and OpenAI model simultaneously, with each model producing a complete, standalone report—surfacing facts, citations, and analytical framings that the other may overlook or weigh differently. Once both reports are generated, a dedicated judge model evaluates the reports to create a distilled summary of key findings and highlights where the models meaningfully agree or diverge - including differences in magnitude, framing, or interpretation - and calls out unique contributions from each model.



Top Areas of Agreement		
Topic	GPT	Claude
Battery life priority	Cited by Marcus Webb and	8/10 customers mentioned

	Brian Kowalski; Phase 2 self-powered fiber directly addresses top pain point	battery anxiety; critical path validated by customer demand
App UX debt	3/10 customers cited clunky setup, Bluetooth drops, data-heavy reports; 2026 app investment confirmed	7/10 customers raised app friction; critical UX debt undermining hardware innovation
Healthcare certification gate	Phase 3 FDA 510(k) targeting 2H 2027; IRB studies starting late 2026 required for clinical credibility	Phase 3 medical certification primary unlock for B2B enterprise expansion
Predictive analytics ROI	Tomas Reyes lost \$400K to preventable injuries; 85%	Highest ROI feature for B2B segments; justifies premium

Enjoying the article? Sign in to share your thoughts.

Sign in

Get started with Critique and Council in Researcher

Share this page

Today, Critique and Council are broadly available in the [Frontier program](#). Learn how to power your [Frontier Transformation with Copilot and agents](#), and experience the power of multi-model intelligence.

What's new

Updated Mar 30, 2026 VERSION 1.0

Surface Laptop

Surface Laptop Studio 2

MICROSOFT 365 COPILOT

Copilot for personal use

AI in Windows

Explore Microsoft products



gauravanand



MICROSOFT

Joined March 24, 2025

[View Profile](#)

Microsoft Store

ofile



Microsoft 365 Copilot Blog

Follow this blog board to get notified when there's new activity

store support

Returns

Order tracking

Certified Refurbished

Microsoft Store Promise

Flexible Payments

Education

Microsoft in education

Devices for education

Microsoft Teams for Education

Like 16 Comment

Microsoft 365 Education

How to buy for your school

Educator training and development

Deals for students and parents

AI for education

Business

Microsoft AI

Microsoft Security

Dynamics 365

Microsoft 365

Microsoft Power Platform

Microsoft Teams

Microsoft 365 Copilot

Small Business

Developer & IT

Azure

Microsoft Developer

Microsoft Learn

Support for AI marketplace apps

Microsoft Tech Community

Microsoft Marketplace

Marketplace Rewards

Visual Studio

Company

Careers

About Microsoft

Company news

Privacy at Microsoft

Investors

Diversity and inclusion

Accessibility

Sustainability



Your Privacy Choices

Consumer Health Privacy

[Sitemap](#)

[Contact Microsoft](#)

[Privacy](#)

[Manage cookies](#)

[Terms of use](#)

[Trademarks](#)

[Safety & eco](#)

[Recycling](#)

[About our ads](#)

[© Microsoft 2026](#)