

Two Radical Ways to Share AI Ownership

Jack Henderson, Leon Erichsen, Matt Prewitt

June 30, 2026 · RadicalxChange

In the summer of 2021, months after Anthropic was founded, an internal memo titled “An Economic Model for Compensating Data Producers” was circulated among its leadership. Recently made public as an exhibit in the lawsuit *Bartz v. Anthropic*, it lays out a way to pay the people whose data is used to train AI. The memo opens with a diagnosis:

“distributed human labor is being used to train AI models by large, centralized actors, who concentrate the resulting profits while in the long run making the human labor obsolete.” It predicts three things if nothing changes: AI companies grow more unpopular but keep taking the data; creators get angry and try to wall their work off; and governments crack down, probably on AI generally rather than this issue in particular. In 2026, we can watch all three happening at once.

The memo’s alternative was to pay data producers a fraction of the profits from the model they helped train, to treat the trained model as a kind of corporate entity in which contributors hold equity in proportion to what they put in. The phrase it used was “a sort of ‘guaranteed basic capital’ for the machine age.”

We don’t know who wrote the memo, but the notes in its margins show that CEO Dario Amodei, co-founder Chris Olah, and researcher Andy Jones all took it seriously. Amodei saw why this would be hard. Things might settle into this arrangement on their own eventually, he noted, “but not without a bunch of hatred, threats, and mutual acrimony... The hope here would be to shortcut all that and just propose the equilibrium solution at the start.”

The logic of the bargain is simple. A universal basic income, or any transfer funded by AI profits, lasts only as long as it stays politically convenient. A stake makes people part-owners, on a footing that can’t easily be pulled out from under them. And a public that owns a piece of AI has at least one reason to support it and shape it instead of just resenting it; creators who get paid a share have a reason to make their work available instead of walling it off; and a deal struck in the open heads off the harsher crackdown that comes when nothing else has worked. The memo imagined companies granting this equity voluntarily at first; over time it would become a norm, and eventually a legal right.

Two settlements, we argue, follow from that logic. Data equity gives a stake to the people whose work trains the models; public equity gives one to society as a whole. This essay takes each in turn.

DATA EQUITY

Five years later, no AI company has built this, or even publicly proposed it. The idea is sound, but the market cannot produce it on its own. No single company can afford to pay for data, let alone hand out ownership in its models, while its competitors train on the same material for free. And the people who might collectively demand better terms aren't legally allowed to organize for them.

The first half of any grand settlement on AI has to fix the supply side of the training-data market, which has a built-in flaw that ordinary contracts can't repair. On one side: a few frontier labs with deep pockets and the ability to gather training data on a massive scale. On the other: hundreds of millions of writers, photographers, musicians, journalists, illustrators, and ordinary people whose work makes up the training material — each of them, individually, negotiating (if at all) with a company that already holds the very thing it's asking to license.

A market built this way does not produce prices that reflect what the work is worth. In a market like this, the price of information reflects bargaining power, not value. A person's data is worth little in their own hands; extracting value from it takes scale they don't have. So the people who produce the training material, with no leverage individually, see their work priced far below what it's worth to society.

The labs' first line of defense is legal: that training is fair use. The doctrine was built for human acts like reading and commentary that yield a new, inspectable work, not for a system that ingests nearly everything ever written, compresses it, and can reproduce large chunks on demand — and even where guardrails block that reproduction, the model's other capabilities still rest on its ability to do so “under the hood.” Several of the four factors cut against the labs (the use is commercial, it copies entire works wholesale, and at scale it can substitute for the originals), but no factor cleanly decides it. The real point isn't that the labs are wrong; it is that no clean ruling is possible. Call training fair use, and every future technology that strains the law gets a pass to ignore it. Call it infringement, and the damages run into the trillions, landing on a few companies the United States now treats as essential to national security. The rulings so far are holding actions, reasoning backward from that impossible result.

The fix is to let the scattered side of the market organize. But antitrust law currently forbids that, or at least makes people nervous enough that the organizing doesn't happen. What is needed is a narrow, explicit exception — a “safe harbor” — that lets qualifying groups of rights holders form joint negotiating bodies, bargain in good faith with the big AI developers, and, if talks stall, go to binding arbitration. This doesn't require Congress to write new copyright law, set a price, or resolve the fair-use fight now in the courts.

Antitrust law exists to stop concentrated power from squeezing scattered counterparties. But when the buyer is the concentrated side and the sellers are the scattered side, blocking the sellers from coordinating doesn't protect competition — it protects the buyer's monopsony. The U.S. recognized this in 1922, when the Capper-Volstead Act let farm cooperatives bargain together without breaking

antitrust law; similar carve-outs exist for forestry, fishing, and labor. The AI training market is that same pattern in unusually clear form, and the safe harbor simply restores the conditions for a real market to function.

RadicalxChange floated a “Data Freedom Act” in 2020, versions have surfaced in the EU’s data legislation, and the Trump administration’s own National AI Legislative Framework now urges Congress to let rights holders band together to negotiate with AI companies without running afoul of antitrust law. The very rule the memo said was missing is something the White House is now openly recommending. (Congress has flirted with a version of this for years. The Journalism Competition and Preservation Act, modeled on Australia’s News Media Bargaining Code, would do roughly this for news publishers facing dominant platforms; it has repeatedly cleared the Senate Judiciary Committee on bipartisan votes but never reached a floor vote.)

The labs’ second line of defense is that valuing data at scale is computationally impossible, because the methods for pricing any given datum, like Shapley values and influence functions, are expensive to run and would cost more than the data is worth. But a new paper by Glen Weyl and Raul Castro Fernandez points out that we don’t need an exact price for every datum, only workable estimates that divide the pie fairly across kinds of contributors. And the two ingredients for that fall out of training itself. The data mixture, the proportions in which data sources are combined, tells us how to divide the pie; and the scaling laws, the measured relationship between data and performance, tell us how big it is.

Because it is tied directly to the training process, the flow of payments concentrates on the data that measurably improves a model, rather than spreading into universal micropayments or pooling as a windfall to big publishers, and so doubles as a signal of what data is worth creating next. Paying data contributors in equity rather than cash keeps them invested in the models they go on to improve.

PUBLIC EQUITY

Data equity gives the people who supplied the training data a stake. The broader claim, that society as a whole should own a piece of AI, is gaining ground. Senator Bernie Sanders has introduced the American AI Sovereign Wealth Fund Act, a one-time tax, paid in stock, that would move half the equity of OpenAI, Anthropic, and other large developers into a sovereign wealth fund paying every American a dividend. OpenAI itself published a paper this spring calling for a public wealth fund seeded by AI companies. So it looks like the question is no longer whether the public gets a stake in AI, but how.

And how is a question of design. The radical-markets tradition this essay draws on holds that property at any scale is a form of monopoly: to own something is to hold the right to exclude others from it. Two troubles follow from that. The first is justice: socially generated value, the intelligence distilled from everyone’s work, is privately captured by a few firms. The second is efficiency: the

owner of a monopoly asset need not put it to its best use, and can sit on it, under-deploy it, or hold out, with nothing to push the asset toward whoever could do the most with it. Monopoly produces not only unfairness but waste. These two problems map onto the two halves of the argument: a narrow justice claim ("this particular work of mine was used without payment") grounds data equity, while the efficiency claim, joined to a broader justice claim that society underwrites AI's risks without sharing its upside, grounds public equity.

One proposal sounds like the most democratic option but should be resisted: handing shares directly to individuals. When 1990s Russia converted state enterprises to private ownership by distributing shares to every citizen, those shares were pried loose from ordinary people within weeks, bought up cheaply by the well-positioned, and funneled through corrupt auctions. The result was not broad ownership but the birth of an oligarchy. That is the case for holding the public's share in a collective, hard-to-alienate vehicle.

An institution that holds the asset on everyone's behalf settles who owns it and to what end, but not whether it ends up in the hands that can use it best.

One candidate design for this is partial common ownership (PCO). It works like this. A lab puts its own price on whatever model it currently runs above a certain capability level. It pays a recurring fee to a public fund based on that price. And it has to accept that, at the price it named, someone else can step in and take the model. Name a low price to dodge the fee, and you risk losing the model cheaply; name a high one, and you pay more.

What keeps the price honest is a credible threat to take the model at the named price, not the number of bidders who can make it. So restricting that right to the government alone need not weaken it. Taiwan ran such a system for land from 1954 to 1977, with the state as the only party able to force a sale: a single public taker, properly funded to act and bound to clear public rules, disciplines prices as well as a crowd of bidders.

But honest pricing is only half the point; the model also has to move to whoever can use it best. The fix is to separate who owns the model from who operates it. The model sits in a public trust that keeps ultimate control and never sells it, while the right to develop and deploy it is licensed to safety-approved operators on the same self-priced, takeable terms, so a more capable operator can take the license from a weaker one at the price the weaker one declared.

By moving the asset toward its best use even as it shares the gains, PCO is one of the rare designs that can lower inequality and raise growth at once.

The concept is borrowed from land, where the party willing to pay most for a parcel is usually the one who will put it to its most productive use. Productive use is roughly what society wants from land, and ability to pay is a fair proxy for best use. But with AI, that link is much less clear. The operator who bids highest for a model license is the one who can extract the most revenue from it, and maximal monetization of a frontier system is not obviously what the public wants; it may be

exactly what we want to restrain. This is one of the best arguments against applying PCO to frontier AI. But at least PCO creates space for democratic governance to set the boundary of permissible uses; the contestable market allocates only within it, moving the model to its best approved use rather than its most lucrative one.

This essay grew out of a RadicalxChange salon, and the shape of that conversation was itself a finding: the premise that AI is socially produced and that society should share in it was treated as consensus, and the energy focused on institutional design. Three tensions recurred. The collective vehicles meant to hold the public's stake grow less trustworthy, not more, as they scale; and a sovereign wealth fund invites capture. The international mismatch of a nationally bounded dividend for something built from global society. And the worry that a public paid to own AI might stop objecting to how it is used. The hard part, everyone agreed, is design and governance, not justification.

The clearest opening is the state-led AI programs now emerging among the “middle powers” that want their own stacks rather than dependence on the American or Chinese ones. Built with public money, they are a natural first venue for putting public equity into practice: the taxpayer is already the investor, so the case is strongest and the politics most tractable. But the same logic is increasingly on the table in the United States, where the government has shown it will take direct stakes in strategic industries, and where a public role in AI now has support across an unusually wide political coalition. Data equity mainly needs a change in the law, the safe harbor the White House now encourages. Any of these could be where a fairer and more efficient ownership arrangement is first put to the test.

The terms under which AI enters the world are still being set, and they turn on one choice being made now: who owns the value these systems create. It won't be settled by winning the argument and waiting for a statute. It will be settled when the first working pieces are built and the law and the market fall in behind them — the difference between reaching the equilibrium now and reaching it, in Amodei's words, after years of hatred, threats, and mutual acrimony.

With thanks to Glen Weyl and all the participants in the RadicalxChange salon for their feedback. This work was supported by The Rockefeller Foundation.

Source: RadicalxChange, “Two Radical Ways to Share AI Ownership,” June 30, 2026.
radicalxchange.org/updates/blog/two-radical-ways-to-share-ai-ownership/