AI and the Social Contract: From Redistribution to a Priced Data Commons

A positive perspective on an AI-cum-human settlement, and the open questions it leaves behind

Dr. Michael Jung

2026-06-13

Abstract

Modern social contracts tie income, taxation, and public provision to human work. If artificial intelligence substitutes for labour at scale, that link breaks, and the fiscal machinery of the welfare state breaks with it. This essay works through the argument step by step. It accepts the structural worry but resists its determinism: the best current evidence shows extraordinary disagreement about how far and how fast AI will displace labour, and the modern contract may bend rather than snap. It then surveys the redistribution proposals now in circulation, from taxes to universal basic income to sovereign AI funds, and the capture and market-distortion risks they carry. Against these it develops the outline's most original idea: treating data as a priced public good rather than a free commons. The proposal is attractive because it adds value instead of seizing it, and because it fits existing institutions. But it rests on a contestable premise, that data rather than compute is where AI rents accumulate, and it imports political and market risks of its own. The disagreements are collected in a closing agenda of open issues.

AI and the Social Contract

1. A crossroads, stated honestly

Public argument about artificial intelligence runs to extremes. One camp promises cures for most diseases and a compressed century of scientific progress. The other forecasts mass unemployment, concentrated power, or worse. Dario Amodei's *Machines of Loving Grace* is the most careful statement of the optimistic case, and it earns attention precisely because it refuses the singularity framing while still arguing that intelligence, aimed at the right bottlenecks, could transform biology, medicine, and development within a decade

[Amodei 2024]. The pessimistic literature is equally serious, and increasingly structural rather than apocalyptic [Aguirre 2025; Mishra 2026].

Both camps tend to skip the institution that sits between capability and outcome: the social contract. We do not experience technology directly. We experience it through the arrangements that decide who works, who pays, who is protected, and who decides. Those arrangements are already under visible strain. Defence, demography, climate, and debt are pressing against the fiscal limits of the post-war settlement before AI arrives at scale. The question this essay takes seriously is whether AI finishes a job that other pressures have started, and if so, what a deliberate redesign might look like.

The classical contract tradition is a useful anchor. Rousseau's problem was to "find a form of association that defends and protects the person and goods of each associate with the whole common force" while leaving each "as free as before" [Rousseau 1762]. The modern welfare state answered a narrower version of that problem through work. Wages fund consumption, employment funds taxation, and taxation funds the provision that holds the association together. AI's challenge is that it attacks the *mechanism* of the modern answer, paid human labour, rather than the underlying problem. The problem remains. The old solution may not.

What follows develops a positive perspective on an AI-cum-human settlement. It does not assume the worst. Where the original argument claims more certainty than the evidence supports, I say so and supply the counter-evidence, and those disagreements are gathered in Section 8 as an agenda for debate rather than buried.

2. The labour channel: a real risk, not a certainty

The core worry is a chain. AI reduces demand for human labour. Employment, work income, and work-based tax revenue shrink. Corporate revenues and the taxes on them shrink too, because they ultimately depend on consumer demand. Public income falls. Public spending on goods, redistribution, age, and care collapses. The social contract, starved of both private purchasing power and public funds, becomes untenable.

The chain is logically tight, and each link describes a genuine pressure already visible at the margin. It deserves to be taken seriously rather than waved away. But the argument's first premise, that AI will catalyse a *large-scale* reduction in the demand for human labour, and that the benign alternative is "hardly plausible," is the most contestable claim in the whole essay. The evidence does not support stating it with confidence.

A different view. Three bodies of evidence complicate the labour-collapse premise.

First, the historical mechanism of automation has usually been complementarity, not substitution. Autor's analysis of why there are "still so many jobs" after two centuries of mechanisation is the canonical statement. Automating some tasks raises the economic value of the tasks that remain, and those residual and newly created tasks have repeatedly absorbed displaced labour [Autor 2015]. The outline dismisses "an explosion of new products which perpetuate reliance on human labour" as implausible, but that is exactly what the last two centuries delivered, and the dismissal is asserted rather than argued. AI may differ. It targets cognition rather than muscle, and it may climb the task ladder faster than humans can move up it. But "may differ" is not "must differ."

Second, the macroeconomic estimates span an enormous range, and that range is itself the finding. Acemoglu's task-based accounting, which applies a Hulten-style theorem to the share of tasks AI can cheapen, yields total-factor-productivity gains of less than 0.55% over ten years: nontrivial but modest, and far from civilisation-altering [Acemoglu 2024]. At the other pole, Korinek and Suh entertain baseline GDP growth of 100% over a decade and "aggressive" AGI scenarios near 300%. Goldman Sachs sits in between at roughly 7% of global GDP. When credible economists differ by two orders of magnitude, the honest posture is conditional planning, not a single forecast of collapse.

Third, when you ask the experts directly, capability optimism and macroeconomic caution coexist. A large 2025 elicitation of economists, AI specialists, and superforecasters put a combined 61% probability on moderate-or-rapid AI progress by 2030. Yet the same economists expected GDP, productivity, and labour-force participation to stay near historical trend, with severe disruption confined to the lower-probability "rapid" branch [Forecasting Research Inst. 2025]. They may be wrong. Under-weighting a labour-substitution phase change is a real possibility, and the measurement work below suggests the change could be largely invisible until it is acute [Korinek & McKelvey 2026]. But the modal expert today does not endorse the collapse premise.

None of this refutes the worry. It reframes it. The defensible claim is conditional: if AI becomes a broad substitute for labour rather than a complement, then the fiscal cascade follows more or less as the outline describes. The collapse is a scenario to insure against, not a prophecy to organise around. That distinction matters for policy, because insuring against a tail risk and rebuilding the state around a certainty call for very different politics.

3. The fiscal cascade and the redesign proposals

Grant the conditional. If the labour channel does erode, the fiscal arithmetic is genuinely alarming, and for a structural reason the outline gets right. In a consumption-driven economy, corporate revenues and the taxes on profit, income, and value all rest on ultimate demand. Cut the wage base and you do not merely lose payroll taxes. You eventually lose the corporate and capital taxes that depend on households having money to spend. The revenue side and the spending side fail together.

This is why an unusual coalition has converged on the need for a new contract. Commercially minded AI firms and socially concerned reformers, "a funny alliance" as the outline puts it, share three motives: sustaining a life-world for people without work, sustaining the consumption-investment cycle that AI itself needs as a market, and smoothing the political path of AI expansion. This sits alongside the other legitimate AI worries, and should not crowd them out: misalignment, cybersecurity and bad-actor threats, deskilling, and cultural flattening [Aguirre 2025; Mishra 2026].

Current proposals cluster on the financing question, namely how to generate, distribute, and control funds to offset private and public income losses.

- **Generation.** Taxing AI-related revenues and profits; taxing AI-related income and capital gains; taxing AI-unrelated income and wealth; or public or private co-ownership of AI corporations.
- **Distribution.** Universal basic income under various eligibility rules; need-based basic income; AI-ownership dividends; and subsidised private or enlarged public goods.
- **Control.** Mostly public-agency variants, including sovereign wealth funds, new welfare regimes, and AI insurance schemes.

The policy literature is moving in this direction. The IPPR's "new politics of AI" argues that AI is qualitatively different enough to justify outcomes-focused rather than technology-neutral policy: set societal targets and steer AI investment toward them rather than waiting for markets to deliver [Jung/IPPR 2025]. Amodei, writing from inside a frontier lab, warns that even a benign AI transition could "lock in" power and calls for arrangements that prevent it [Amodei 2024]. The appetite for redesign is real, and it crosses ideological lines.

The trouble is that almost all of these proposals are, at root, redistributive. They take a claim on income or wealth that currently belongs to someone and reassign it. That is what makes them politically heavy, and it is the hinge on which the rest of the argument turns.

4. Caveats: capture, distortion, and the zero-sum trap

The proposals differ in incidence (who pays), in system impact (what they do to competition and governance), and in reach (whether they hit critical mass). Three risks recur, and the evidence sharpens each.

Political capture. A public AI wealth fund, or state co-ownership of AI corporations, concentrates an extraordinary asset in political hands. Whoever controls that fund controls a lever over the economy's most powerful general-purpose technology, and through it, potentially, over public discourse itself. The cloud-capitalism literature shows how readily AI infrastructure becomes a choke point. Control of compute confers leverage over every downstream industry, and that leverage has already begun fusing with state power into techno-nationalist alliances [Thelen 2025]. A sovereign AI fund does not dissolve that choke point. It moves the hand on the lever from a board to a ministry. That may be better or worse, but it is not obviously safer.

Private capture. Conversely, if the new machinery is built around private AI firms, with dividends routed through them and ownership stakes that entrench them, the effect could be to turbo-charge exactly the concentration the contract was meant to cushion. The structural-power dynamics here are not hypothetical. The "control inversion" and "gradual disempowerment" literatures argue that delegating authority to systems far more capable than their principals, and to the firms that run them, transfers power away from those principals through ordinary, locally rational steps, with no revolt required [Aguirre 2025; Krueger et al. 2026]. A contract that funnels rents through a few firms risks accelerating that transfer under the banner of fairness.

Market distortion. Every financing mechanism listed in Section 3 is non-neutral to the market's core functions of allocation, distribution, and innovation. Heavy AI-profit taxes blunt the incentive to deploy. Ownership mandates distort entry. Poorly designed dividends distort labour supply. The point is not that distortion is disqualifying, since all taxation distorts. The point is that these proposals distort the allocation mechanism at the moment society most needs it to work, when a general-purpose technology is reshaping which activities have value.

There is a deeper objection, and the outline names it. Introducing any of these proposals is close to a zero-sum game among powerful actors. Redistribution defines winners and losers in advance, which guarantees conflict, power-play, and the kind of degraded compromise that satisfies no one and is captured by whoever is strongest in the room. If the menu of options is intrinsically zero-sum, the political economy of adopting it is grim regardless of which option is "best" on paper. That is the real case for looking elsewhere.

5. Reconsidering options: data as a priced public good

The escape route the outline proposes is genuinely clever, and it is the part of the argument most worth developing. Instead of seizing existing income or wealth, and paying the zero-sum political price, invent a new category of resource and let new income streams flow from it. The candidate resource is data.

The logic runs as follows. Data are the indispensable input to AI, yet today they are treated mostly as a free commons: scraped, pooled, and monetised by whoever has the compute to exploit them, with little flowing back to the people and institutions that generate them. Establish data instead as a valuable public good with a deliberate market regime, one that weighs commercial value against privacy value, and data becomes a priced node in the AI value chain (data, then compute, then application), alongside the others. Income from that node can fund both private incomes and public provision without expropriating anyone's existing property.

The proposal carries three attractive promises.

- **Promise 1: value grows with AI.** If AI becomes, as expected, an all-encompassing envelope around commercial activity, the value of the globally available data it feeds on rises in step with AI expansion. The funding base grows precisely when the need does, which is the counter-cyclical property a labour-substituting transition requires.
- **Promise 2: institutional congruence.** "Data as a public good for sale" leaves private property, competitive dynamics, and meritocratic performance intact. It adds a market rather than abolishing one, which is what makes it positive-sum where redistribution is zero-sum.
- **Promise 3: recoverable market design.** Pricing data reopens classic market-design tools. Breaking the private data monopolies that currently lock value inside a few firms, or charging different prices to upstream and downstream users, could

spur competition and innovation rather than dampen it.

The intellectual lineage is respectable. The "unreasonable effectiveness of data" thesis established that, for many tasks, scale of data beat algorithmic cleverness, which is the empirical foundation for treating data as the strategic resource [Halevy et al. 2009]. The governance side has a serious precedent too. Prewitt's "data coalitions" argue that because data's value is collective, deriving from patterns across people rather than from any individual record, governance must be collective as well, through union-like bargaining entities rather than individual ownership [Prewitt 2021]. A public-good-for-sale regime and a data-coalition regime are two designs for the same insight: the value of data is social, so its proceeds should be distributed socially.

This is the positive core of the essay, and it is right to reframe the problem, from dividing a fixed pie to creating a new one. But it rests on a premise that deserves direct challenge.

6. A harder look at the data premise

A different view: is data really the binding constraint? The proposal's entire weight rests on the claim that data are *the* key resource of AI. The current evidence points, at least as strongly, at compute as the resource where AI's rents actually accumulate. That matters, because a scheme that prices the wrong layer funds nothing.

Three points cut against data-primacy.

First, the rent-bearing choke point today is infrastructure, not data. The cloud-capitalism analysis is explicit. Dominant firms shifted from asset-light data intermediation to asset-heavy compute ownership, and it is control of compute, not control of data, that confers structural leverage over downstream industries [Thelen 2025]. Frontier training data is increasingly abundant, synthetic, or substitutable. Frontier compute is scarce, capital-intensive, and concentrated in a handful of hands. If you want to tax where the surplus pools, the political economy points at the data centre, not the dataset.

Second, recent productivity gains have come from algorithms, not data volume. The PIIE measurement work finds that quality-adjusted AI output grew on the order of 2,000% or more per year, driven overwhelmingly by algorithmic progress, with inference prices falling roughly 94% per year at fixed performance, rather than by ingesting proportionally more data [Korinek & McKelvey 2026]. Even Halevy, Norvig, and Pereira's original "data wins" verdict was pre-deep-learning, and the retrospective judgement is that both data and architecture mattered. The slogan "just get more data" was always a misreading [Halevy et al. 2009]. Pricing data as the master resource may be fighting the last war.

Third, and most awkward for the proposal, a public data regime reintroduces the very risks Section 4 raised against the redistribution proposals. A state-run "data as a public good" market is a concentrated lever over the economy's key input, which is to say a political-capture risk wearing different clothes. Who sets the price? Who decides whose privacy is worth what? A single national or supranational data authority is a choke point by design, and the digital-sovereignty debate shows how quickly "sovereign control of data" slides into protectionism and state leverage over discourse [Thelen 2025]. The outline's own assessment question, whether "expensive data" are an impediment to benign progress, is the live one. Raising the price of the input that drives a positive-sum technology could slow the very expansion that was supposed to fund the contract.

This is not a refutation. It is a relocation of the burden of proof. The data-commons idea may still be the best available, for a reason the compute-primacy view actually strengthens. Compute is rivalrous, depreciating, and physically sited, so taxing it invites flight and obsolescence. Data is non-rivalrous and renewable, so a well-designed data market could be a more durable base even if it is not today's largest rent pool. But that is an argument to be made, with mechanism design and incidence analysis, not asserted. The honest version of the proposal reads: data could be established as a priced public good whose value would grow with AI; whether data, or compute, or some combination is the right base is an open empirical question that the measurement gap currently prevents us from answering [Korinek & McKelvey 2026].

There is also a constructive synthesis the outline under-explores. The choice is not public-good-for-sale *versus* data coalitions [Prewitt 2021]. A layered regime is conceivable: collective bargaining entities at the point of data generation, feeding a publicly chartered marketplace with transparent differential pricing, with proceeds split between the contributing communities and a public fund. That keeps property and competition intact (Promise 2), distributes the collective value of data to its collective producers, and dilutes the single-authority capture risk. It is harder to build than either pure design, but it is more robust to the objections above.

7. Challenges, and a substitute for the meaning of work

Suppose the data-commons route, in some layered form, is worth pursuing. The outline is candid that the challenges are large, and they are worth restating plainly.

A workable regime needs a fundamental accord. Data flow across sectors and borders, so it requires global, inter-sectoral agreement on governance, on pricing different uses, and on distributing the resulting income. This is a treaty-scale problem in a fragmenting, techno-nationalist world [Thelen 2025]. The deeper challenge is institutional: designing a settlement that performs better than today's on all three core dimensions, allocation, distribution, and innovation, rather than trading one off against the others. That is the bar a redesign must clear to be worth its disruption.

There is a dimension the financing debate keeps dropping. Work is not only income. The original chain treats the social contract as a flow of funds, but the welfare state also delivered meaning, dignity, and structure through paid work. Replace the income and you have still not replaced that. The most rigorous statement of the danger is the "extinction by success" argument: even if AI does everything right, optimisation-driven convergence can erode the role differentiation through which human meaning has historically been generated, leaving people materially provided for but causally inert [Mishra 2026]. The argument is contested. Internalist theories of meaning, and the historical record of humans inventing new vocations, both cut against its strongest form. But it identifies the gap that a financing-only redesign cannot fill.

The outline's most hopeful note is the right place to push back against that fatalism. If AI progress is open-ended, so is the space of AI-human skill combinations, and new paradigms of work and meaning may be discoverable rather than merely lost. A contract designed only to pay people not to work forecloses that search. A contract that funds participation, in care, deliberation, craft, and the domains where human judgement remains causally load-bearing, keeps it open. Pricing data to fund incomes is the easy half. Funding the institutions in which post-work meaning can be re-grown is the half that decides whether the settlement is humane.

8. Open issues for debate

The disagreements above are not loose ends to be tidied away. They are the agenda. A redesign of the social contract should be argued out, not assumed. The sharpest open questions:

1. **Is the labour-collapse premise sound, or merely possible?** Autor-style complementarity, Acemoglu's modest sub-0.55% TFP estimate, and the expert survey's near-trend expectations all sit against the collapse story. Korinek and Suh's 100-300% scenarios sit with it. The base case should be conditional planning across this range, not commitment to one branch [Autor 2015; Acemoglu 2024; Forecasting Research Inst. 2025; Korinek & McKelvey 2026].
2. **Data or compute: where do AI rents actually accrue?** If compute is the choke point and algorithms drive the gains, a data levy may tax the wrong layer. The counter-argument, that data is renewable and non-rivalrous and therefore a more durable base, needs formal incidence analysis [Thelen 2025; Korinek & McKelvey 2026; Halevy et al. 2009].
3. **Does a public data market just relocate capture risk?** A single data authority is itself a choke point over discourse and the economy. How is it insulated from the political-capture and digital-sovereignty pathologies it was meant to avoid [Thelen 2025]?
4. **Would "expensive data" slow benign progress?** Raising the price of AI's input could throttle the positive-sum expansion that funds the contract. What pricing regime funds redistribution without suppressing diffusion?
5. **Public-good-for-sale, data coalitions, or a layered hybrid?** Is the most robust design a chartered marketplace, union-style collective bargaining, or a layer that combines them [Prewitt 2021]?
6. **Can we even measure the base?** The AI economy is largely invisible in GDP statistics, and you cannot tax, or price, what you cannot measure. Satellite accounts may be a precondition for any of these schemes [Korinek & McKelvey 2026].
7. **What replaces work as a source of meaning?** Financing answers the income question and leaves dignity, voice, and structure unaddressed [Mishra 2026]. Which institutions keep human participation causally relevant?
8. **Is grand redesign politically reachable at all?** Are AI-induced needs, capacities, and anxieties sufficient to commit publics and politics to a redesign on this scale? Or does the zero-sum political economy of every financing option doom it to degraded compromise?

9. Feasibility, compatibility, simplicity, elegance

The test the outline closes on is the right one. A workable AI social contract should be feasible, meaning buildable inside today's fragmented institutions. It should be system-compatible, meaning it strengthens rather than disrupts allocation, distribution, and innovation. It should be simple enough to command consent, and elegant in the sense of positive-sum rather than a brokered carve-up.

By those tests, the redistribution menu scores poorly. It is feasible only through conflict, system-distorting by construction, complex, and zero-sum. The data-commons idea scores better on compatibility and elegance, because it adds value rather than seizing it. But its feasibility is unproven, and its simplicity is undercut by the treaty-scale coordination it needs and by the unresolved data-versus-compute question. The layered hybrid of Section 6 trades some simplicity for far more robustness, which on a problem this consequential is the right trade.

The most defensible conclusion is therefore modest in tone and ambitious in direction. The labour threat to the social contract is real but conditional, so we should insure against it rather than rebuild around it prematurely. The redistribution proposals treat the symptom, income, while importing capture and distortion risks that may be worse than the disease. The data-commons idea is the most promising move on the board, because it is positive-sum and institutionally congruent. But it has to survive the compute-primacy objection, the capture objection, and the measurement gap before it can carry the weight the argument places on it. Those are not reasons to abandon it. They are the work. A settlement that finishes that work, by pricing the right resource, distributing its collective value, insulating the mechanism from capture, and re-growing the institutions of human meaning, would be a contract worth the name. It would, in Rousseau's phrase, protect each associate with the common force while leaving them as free as before [Rousseau 1762].

References

- [**Acemoglu 2024**] Acemoglu, D. (2024). *The Simple Macroeconomics of AI.* (raw/papers/labour_market/The Simple Macroeconomics of AI.pdf)
- [**Aguirre 2025**] Aguirre, A. (2025). *Control Inversion: Why the Superintelligent AI Agents We Are Racing to Create Would Absorb Power, Not Grant It.* Future of

Life Institute. [[control-inversion-why-superintelligent-ai-agents]]

- [**Amodei 2024**] Amodei, D. (2024). *Machines of Loving Grace: How AI Could Transform the World for the Better.* [[machines-loving-grace-how-ai-could]]
- [**Autor 2015**] Autor, D. (2015). *Why Are There Still So Many Jobs? The History and Future of Workplace Automation.* Journal of Economic Perspectives. (raw/papers/labour_market/autor-why-still-so-many-jobs-jep2015.pdf)
- [**Forecasting Research Inst. 2025**] Forecasting Research Institute et al. (2025). *Forecasting the Economic Effects of AI: Predictions From Economists, AI Experts, and the Public.* [[forecasting-economic-effects-ai]]
- [**Halevy et al. 2009**] Halevy, A., Norvig, P., & Pereira, F. (2009). *The Unreasonable Effectiveness of Data.* IEEE Intelligent Systems. [[halevy-norvig-pereira-unreasonable-effectiveness-data]]
- [**Jung/IPPR 2025**] Jung, C. (2025). *The New Politics of AI: Why Fast Technological Change Requires Bold Policy Targets.* IPPR. [[new-politics-ai-fast-technological-change]]
- [**Korinek & McKelvey 2026**] Korinek, A., & McKelvey, M. (2026). *Where Is AI in GDP Statistics? Filling the Measurement Gap.* PIIE Policy Brief 26-7. [[where-ai-gdp-statistics-filling-measurement]]
- [**Krueger et al. 2026**] Krueger, D. et al. (2026). *Gradual Disempowerment.* [[gradual-disempowerment]]
- [**Mishra 2026**] Mishra (2026). *Obsolescence Without Hostility: Optimization, Uniformity, and the Erosion of Human Meaning in a Post-AI World.* AI & SOCIETY. [[obsolescence-without-hostility-optimization-uniformity-erosion]]
- [**Prewitt 2021**] Prewitt, M. (2021). *A View of the Future of Our Data.* Noema Magazine. [[view-future-data-noema-data-sovereignty]]
- [**Rousseau 1762**] Rousseau, J.-J. (1762). *The Social Contract.* (raw/papers/social_contract/rousseau-social-contract-1762.pdf)
- [**Thelen 2025**] Thelen, K. (2025). *Cloud Capitalism and the AI Transition.* Politics & Society. [[cloud-capitalism-ai-transition]]
- Keynes, J. M. (1930). *Economic Possibilities for our Grandchildren.* (raw/papers/social_contract/keynes-economic-possibilities-grandchildren-1930.pdf)